

Data Analytics in Finance

FINA 6333 for Spring 2025

Richard Herron

Table of contents

Week 1	11
McKinney Chapter 2 - Python Language Basics, IPython, and Jupyter Notebooks	12
Introduction	12
Language Semantics	12
Scalar Types	20
Control Flow	24
McKinney Chapter 2 - Practice - Blank	28
Announcements	28
Five-Minute Review	28
Practice	28
McKinney Chapter 2 - Practice - Sec 02	31
Announcements	31
Five-Minute Review	31
Practice	31
McKinney Chapter 2 - Practice - Sec 03	42
Announcements	42
Five-Minute Review	42
Practice	42
McKinney Chapter 2 - Practice - Sec 04	53
Announcements	53
Five-Minute Review	53
Practice	53
Week 2	65
McKinney Chapter 3 - Built-In Data Structures, Functions, and Files	66
Introduction	66
Data Structures and Sequences	66
List, Set, and Dict Comprehensions	75
Functions	77

Table of contents

McKinney Chapter 3 - Practice - Blank	81
Announcements	81
Five-Minute Review	81
Practice	81
McKinney Chapter 3 - Practice - Sec 02	85
Announcements	85
Five-Minute Review	85
Practice	88
McKinney Chapter 3 - Practice - Sec 03	100
Announcements	100
Five-Minute Review	100
Practice	103
McKinney Chapter 3 - Practice - Blank	115
Announcements	115
Five-Minute Review	115
Practice	118
Week 3	130
McKinney Chapter 4 - NumPy Basics: Arrays and Vectorized Computation	131
Introduction	131
The NumPy ndarray: A Multidimensional Array Object	132
Universal Functions: Fast Element-Wise Array Functions	144
Array-Oriented Programming with Arrays	146
McKinney Chapter 4 - Practice - Blank	151
Announcements	151
Five-Minute Review	151
Practice	151
McKinney Chapter 4 - Practice - Sec 02	154
Announcements	154
Five-Minute Review	154
Practice	158
McKinney Chapter 4 - Practice - Sec 03	171
Announcements	171
Five-Minute Review	171
Practice	175

Table of contents

McKinney Chapter 4 - Practice - Sec 04	187
Announcements	187
Five-Minute Review	187
Practice	191
 Week 4	 203
McKinney Chapter 5 - Getting Started with pandas	204
Introduction	204
Introduction to pandas Data Structures	205
Essential Functionality	215
Summarizing and Computing Descriptive Statistics	225
McKinney Chapter 5 - Practice - Blank	231
Announcements	231
Five-Minute Review	231
Practice	231
McKinney Chapter 5 - Practice - Sec 02	233
Announcements	233
Five-Minute Review	233
Practice	238
McKinney Chapter 5 - Practice - Sec 03	255
Announcements	255
Five-Minute Review	255
Practice	260
McKinney Chapter 5 - Practice - Sec 04	276
Announcements	276
Five-Minute Review	276
Practice	281
 Week 5	 297
McKinney Chapter 8 - Data Wrangling: Join, Combine, and Reshape	298
Introduction	298
Hierarchical Indexing	298
Combining and Merging Datasets	305
Reshaping and Pivoting	320

Table of contents

McKinney Chapter 8 - Practice - Blank	323
Announcements	323
Five-Minute Review	323
Practice	323
McKinney Chapter 8 - Practice - Sec 02	325
Announcements	325
Five-Minute Review	325
Practice	326
McKinney Chapter 8 - Practice - Sec 03	338
Announcements	338
Five-Minute Review	338
Practice	339
McKinney Chapter 8 - Practice - Sec 04	351
Announcements	351
Five-Minute Review	351
Practice	352
Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Blank	363
Log and Simple Returns	363
Portfolio Math	367
Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 02	373
Log and Simple Returns	373
Portfolio Math	377
Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 03	393
Log and Simple Returns	393
Portfolio Math	397
Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 04	413
Log and Simple Returns	413
Portfolio Math	417
Week 6	433
McKinney Chapter 10 - Data Aggregation and Group Operations	434
Introduction	434
GroupBy Mechanics	434
Data Aggregation	440
Apply: General split-apply-combine	444

Table of contents

Pivot Tables and Cross-Tabulation	446
McKinney Chapter 10 - Practice - Blank	449
Announcements	449
Five-Minute Review	449
Practice	449
Week 7	452
McKinney Chapter 11 - Time Series	453
Introduction	453
Time Series Basics	453
Date Ranges, Frequencies, and Shifting	462
Resampling and Frequency Conversion	468
Moving Window Functions	471
McKinney Chapter 11 - Practice - Blank	478
Announcements	478
Five-Minute Review	478
Practice	478
Week 8	480
Project 1	481
Purpose	481
Assignment	481
Criteria	482
Deliverables	482
Data	483
Quarto	485
Artificial Intelligence (AI)	486
Week 9	488
Herron Topic 2 - Trading Strategies Based on Technical Analysis	489
Introduction	489
What is technical analysis?	489
Why might trading strategies based on technical analysis work or not?	490
Implement a simple moving average (SMA) trading strategy	497

Table of contents

Herron Topic 2 - Practice	502
Announcements	502
Five-Minute Review	502
Practice	502
Herron Topic 2 - Practice - Sec 02	505
Announcements	505
Five-Minute Review	505
Practice	506
Herron Topic 2 - Practice - Sec 03	536
Announcements	536
Five-Minute Review	536
Practice	537
Herron Topic 2 - Practice - Sec 04	567
Announcements	567
Five-Minute Review	567
Practice	568
Week 10	598
Herron Topic 3 - Quantitative Value Investing	599
Introduction	599
Import the data from CRSP and Compustat	600
Create interim data frames	601
Performance of Portfolios Formed on B/M	607
Herron Topic 3 - Practice - Blank	611
Announcements	612
Five-Minute Review	612
Practice	612
Herron Topic 3 - Practice - Sec 02	613
Announcements	613
Five-Minute Review	614
Practice	614
Herron Topic 3 - Practice - Sec 03	635
Announcements	635
Five-Minute Review	636
Practice	636

Table of contents

Herron Topic 3 - Practice - Sec 04	657
Announcements	657
Five-Minute Review	658
Practice	658
 Week 11	 679
Project 2	680
Purpose	680
Assignment	680
Criteria	681
Deliverables	681
Data	682
Quarto	682
Artificial Intelligence (AI)	683
 Week 12	 685
Herron Topic 4 - Portfolio Optimization	686
Introduction	686
The 1/N Portfolio	686
SciPy's <code>minimize()</code> Function	690
The Minimum-Variance Portfolio	694
The Minimum-Variance Frontier	697
 Herron Topic 4 - Practice	 701
Announcements	701
Five-Minute Recap	701
Practice	701
 Herron Topic 4 - Practice - Sec 02	 703
Announcements	703
Five-Minute Recap	703
Practice	703
 Herron Topic 4 - Practice - Sec 03	 705
Announcements	705
Five-Minute Recap	705
Practice	705
 Herron Topic 4 - Practice - Sec 04	 707
Announcements	707

Table of contents

Five-Minute Recap	707
Practice	707
Week 13	709
Student's Choice 4	710
Week 14	711
MSFQ Assessment Exam	712
Week 15	713
Project 3	714

Welcome to FINA 6333 for Spring 2025 at the D'Amore-McKim School of Business at Northeastern University!

For each course topic, we will have one notebook for the pre-recorded lecture and one for the in-class practice. I will maintain these notebooks on here and everything else on Canvas. You have three choices to access these notebooks:

1. Download them from [OneDrive](#)
2. Download them from [GitHub](#)
3. Open them on [Google Colab](#)

Week 1

McKinney Chapter 2 - Python Language Basics, IPython, and Jupyter Notebooks

Introduction

We must understand the basics of Python before we can use it to analyze financial data. Chapter 2 of McKinney (2022) provides a crash course in Python's syntax, and Chapter 3 provides a crash course in Python's built-in data structures. This notebook focuses on the "Python Language Basics" in Section 2.3, which covers language semantics, scalar types, and control flow.

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Language Semantics

Indentation, not braces

Python uses whitespace (tabs or spaces) to structure code instead of using braces as in many other languages like R, C++, Java, and Perl.

Spaces are more than cosmetic in Python. Here is a `for` loop with an `if` statement that shows how Python uses indentation to separate code instead of parentheses and braces.

```
array = [1, 2, 3]
pivot = 2
less = []
greater = []

for x in array:
    if x < pivot:
        print(f'{x} is less than {pivot}')
        less.append(x)
    else:
```

```
print(f'{x} is NOT less than {pivot}')
greater.append(x)
```

```
1 is less than 2
2 is NOT less than 2
3 is NOT less than 2
```

```
less
```

```
[1]
```

```
greater
```

```
[2, 3]
```

Comments

Any text preceded by the hash mark (pound sign) # is ignored by the Python interpreter. This is often used to add comments to code. At times you may also want to exclude certain blocks of code without deleting them.

The Python interpreter ignores any code after a hash mark # on a given line. We can quickly comment/un-comment lines of code with the <Ctrl>-/ shortcut.

```
# We often use comments to leave notes for future us (or co-workers)
# 5 + 5
```

Function and object method calls

You call functions using parentheses and passing zero or more arguments, optionally assigning the returned value to a variable:

```
result = f(x, y, z)
g()
```

Almost every object in Python has attached functions, known as methods, that have access to the object's internal contents. You can call them using the following syntax:

```
obj.some_method(x, y, z)
```

Functions can take both positional and keyword arguments:

```
result = f(a, b, c, d=5, e='foo')
```

More on this later.

Here is a function named `add_numbers` that adds two numbers.

```
def add_numbers(a, b):
    return a + b
```

```
add_numbers(5, 5)
```

10

Here is a function named `add_strings` that adds or concatenates two strings separated by a space.

```
def add_strings(a, b):
    return a + ' ' + b
```

```
add_strings('5', '5')
```

'5 5'

What is the difference between `print()` and `return`?

- `print()` returns its argument to the console or “standard output”
- `return` returns its argument as an output we can assign to variables

Please see the following example.

```
def add_strings_2(a, b):
    string_to_print = a + ' ' + b + ' (this is from the print statement)'
    string_to_return = a + ' ' + b + ' (this is from the return statement)'
    print(string_to_print)
    return string_to_return
```

```
returned = add_strings_2('5', '5')
```

```
5 5 (this is from the print statement)
```

```
returned
```

```
'5 5 (this is from the return statement)'
```

Variables and argument passing

When assigning a variable (or name) in Python, you are creating a reference to the object on the righthand side of the equals sign.

```
a = [1, 2, 3]
b = a
```

If we assign `a` to a new variable `b`, both `a` and `b` refer to the *same* object, which is the list `[1, 2, 3]`.

```
a is b
```

```
True
```

If we modify a by appending 4, we also modify b because a and b refer to the same list.

```
a.append(4)
```

```
a
```

```
[1, 2, 3, 4]
```

```
b
```

```
[1, 2, 3, 4]
```

Likewise, if we modify b by appending 5, we also modify a.

```
b.append(5)
```

```
b
```

```
[1, 2, 3, 4, 5]
```

```
a
```

```
[1, 2, 3, 4, 5]
```

Dynamic references, strong types

In contrast with many compiled languages, such as Java and C++, object references in Python have no type associated with them.

Python has *dynamic references*. Therefore, we do not declare variable types, and we can change variable types. This behavior is because variables are names assigned to objects.

For example, above we assign `a` to a list, and below we can reassign it to an integer and then a string.

```
a
```

```
[1, 2, 3, 4, 5]
```

```
type(a)
```

```
list
```

```
a = 5  
type(a)
```

```
int
```

```
a = 'foo'  
type(a)
```

```
str
```

Python has *strong types*. Therefore, Python typically will not convert object types.

For example, '5' + 5 returns either '55' as a string or 10 as an integer in many programming languages. However, below '5' + 5 returns an error because Python will not implicitly convert the type of the string or integer.

```
# '5' + 5 #TypeError: can only concatenate str (not "int") to str
```

However, Python will implicitly convert integers to floats.

```
a = 4.5  
b = 2  
a / b
```

2.25

Attributes and methods

We can use tab completion to access attributes (characteristics stored inside objects) and methods (functions associated with objects). Tab completion is a feature of the IPython and Jupyter environments.

```
a = 'foo'  
  
a.capitalize()  
  
'Foo'  
  
a.upper().lower()  
  
'foo'  
  
a.count('o')
```

2

Binary operators and comparisons

Binary operators operate on two arguments.

5 - 7

-2

12 + 21.5

33.5

5 <= 2

False

Table 2-1 from McKinney (2022) summarizes the binary operators.

- `a + b` : Add a and b
- `a - b` : Subtract b from a
- `a * b` : Multiply a by b
- `a / b` : Divide a by b
- `a // b` : Floor-divide a by b, dropping any fractional remainder
- `a ** b` : Raise a to the b power
- `a & b` : True if both a and b are True; for integers, take the bitwise AND
- `a | b` : True if either a or b is True; for integers, take the bitwise OR
- `a ^ b` : For booleans, True if a or b is True , but not both; for integers, take the bitwise EXCLUSIVE-OR
- `a == b` : True if a equals b
- `a != b` : True if a is not equal to b
- `a <= b`, `a < b` : True if a is less than (less than or equal) to b
- `a > b`, `a >= b`: True if a is greater than (greater than or equal) to b
- `a is b` : True if a and b reference the same Python object
- `a is not b` : True if a and b reference different Python objects

Mutable and immutable objects

Most objects in Python, such as lists, dicts, NumPy arrays, and most user-defined types (classes), are mutable. This means that the object or values that they contain can be modified.

A list is a *mutable*, ordered collection of elements, which can be any data type. *Because lists are mutable, we can modify them.* Lists are defined using square brackets [] with elements separated by commas. Lists support indexing, slicing, and various methods for adding, removing, and modifying elements.

```
a_list = ['foo', 2, [4, 5]  
a_list
```

```
['foo', 2, [4, 5]]
```

Python is zero-indexed! The first element has a zero subscript [0]!

```
a_list[0]
```

```
'foo'
```

```
a_list[2]
```

```
[4, 5]
```

```
a_list[2][0]
```

```
4
```

```
a_list[2] = (3, 4)  
a_list
```

```
['foo', 2, (3, 4)]
```

A tuple is an *immutable*, ordered collection of elements, which can be any data type. *Because tuples are immutable, we cannot modify them.* Tuples are defined using optional but helpful parentheses (), with elements separated by commas.

```
a_tuple = (3, 5, (4, 5))  
a_tuple
```

```
(3, 5, (4, 5))
```

The Python interpreter returns an error if we try to modify `a_tuple` because tuples are immutable.

```
# a_tuple[1] = 'four' # TypeError: 'tuple' object does not support item assignment
```

The parentheses () are optional for tuples. However, parentheses () are helpful because they improve readability and remove ambiguity.

```
test = 1, 2, 3
type(test)
```

```
tuple
```

We will learn more about Python's built-in data structures in Chapter 3.

Scalar Types

Python along with its standard library has a small set of built-in types for handling numerical data, strings, boolean (True or False) values, and dates and time. These “single value” types are sometimes called scalar types and we refer to them in this book as scalars. See Table 2-4 for a list of the main scalar types. Date and time handling will be discussed separately, as these are provided by the datetime module in the standard library.

Table 2-2 from McKinney (2022) summarizes the standard scalar types.

- **None**: The Python “null” value (only one instance of the None object exists)
- **str**: String type; holds Unicode (UTF-8 encoded) strings
- **bytes**: Raw ASCII bytes (or Unicode encoded as bytes)
- **float**: Double-precision (64-bit) floating-point number (note there is no separate double type)
- **bool**: A True or False value
- **int**: Arbitrary precision signed integer

Numeric types

Integers are unbounded in Python. The ** binary operator raises the number on the left to the power on the right.

```
ival = 17239871
ival ** 6
```

```
26254519291092456596965462913230729701102721
```

Floats (decimal numbers) are 64-bit in Python.

```
fval = 7.243  
type(fval)
```

```
float
```

Dividing integers yields a float, if necessary.

```
3 / 2
```

```
1.5
```

We use `//` if we want integer division.

```
3 // 2
```

```
1
```

Booleans

The two Boolean values in Python are written as `True` and `False`. Comparisons and other conditional expressions evaluate to either `True` or `False`. Boolean values are combined with the `and` and `or` keywords.

We must type Booleans as `True` and `False` because Python is case sensitive.

```
True and True
```

```
True
```

```
(5 > 1) and (10 > 5)
```

```
True
```

```
False and True
```

```
False
```

```
False or True
```

```
True
```

```
(5 > 1) or (10 > 5)
```

```
True
```

We can substitute & for and and | for or.

```
True & True
```

```
True
```

```
False & True
```

```
False
```

```
False | True
```

```
True
```

Type casting

We can “recast” variables to change their types.

```
s = '3.14159'  
type(s)
```

```
str
```

```
1 + float(s)
```

4.14159

```
fval = float(s)
type(fval)
```

float

```
int(fval)
```

3

We can recast a string '5' to an integer or an integer 5 to a string to prevent the `5 + '5'` error above.

```
5 + int('5')
```

10

```
str(5) + '5'
```

'55'

None

`None` is null in Python. `None` is like #N/A or `=na()` in Excel.

```
a = None
a is None
```

True

```
b = 5
b is not None
```

True

```
type(None)
```

```
NoneType
```

Control Flow

Python has several built-in keywords for conditional logic, loops, and other standard control flow concepts found in other programming languages.

If you understand Excel's `if()`, then you understand Python's `if`, `elif`, and `else`.

if, elif, and else

```
x = -1
type(x)
```

```
int
```

```
if x < 0:
    print("It's negative")
```

```
It's negative
```

Single quotes and double quotes (' and ") are equivalent in Python. However, in the preceding code cell, we must use double quotes to differentiate between the enclosing quotes and the apostrophe in `It's`.

Python's `elif` avoids nested `if` statements. `elif` allows another `if` condition that is tested only if the preceding `if` and `elif` conditions were not `True`. An `else` runs if no other conditions are met.

```
x = 10
if x < 0:
    print("It's negative")
elif x == 0:
    print('Equal to zero')
elif 0 < x < 5:
```

```
    print('Positive but smaller than 5')
else:
    print('Positive and larger than or equal to 5')
```

Positive and larger than or equal to 5

We can combine comparisons with `and` and `or` (or `&` and `|`).

```
a = 5
b = 7
c = 8
d = 4
if (a < b) or (c > d):
    print('Made it')
```

Made it

for loops

We use `for` loops to loop over collections, like lists or tuples.

The `continue` keyword skips the remainder of the current iteration of the `for` loop, moving to the next iteration.

The `+=` operator adds and assigns values with one operator. That is, `a += 5` is an abbreviation for `a = a + 5`. There are equivalent operators for subtraction, multiplication, and division (i.e., `-=`, `*=`, and `/=`).

```
sequence = [1, 2, None, 4, None, 5, 'Alex']
total = 0
for value in sequence:
    if value is None or isinstance(value, str):
        continue
    total += value # the += operator is equivalent to "total = total + value"

total
```

12

The `break` keyword skips the remainder of the current and all remaining iterations of the `for` loop.

```
sequence = [1, 2, 0, 4, 6, 5, 2, 1]
total_until_5 = 0
for value in sequence:
    if value == 5:
        break
    total_until_5 += value

total_until_5
```

13

range

The `range` function returns an iterator that yields a sequence of evenly spaced integers.

The `range()` function quickly and efficiently generates iterators for `for` loops.

- With one argument, `range()` creates an iterator from 0 to that number *but excludes that number*, so `range(10)` is an iterator that starts at 0, stops at 9, with a length of 10
- With two arguments, the first argument is the *included* start value, and the second argument is the *excluded* stop value
- With three arguments, the third argument is the iterator step size

```
range(10)
```

```
range(0, 10)
```

We can cast a range to a list.

```
list(range(10))
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Python intervals are “closed” (included) on the left and “open” (excluded) on the right. The following is an empty list because we cannot count from 5 to 0 by steps of +1.

```
list(range(5, 0))
```

```
[]
```

However, we can count from 5 to 0 in steps of -1.

```
list(range(5, 0, -1))
```

```
[5, 4, 3, 2, 1]
```

Ternary expressions

We can complete simple comparisons on one line in Python.

```
x = -5
value = 'Non-negative' if x >= 0 else 'Negative'
value
```

```
'Negative'
```

McKinney Chapter 2 - Practice - Blank

Announcements

Five-Minute Review

Practice

Extract the year, month, and day from an 8-digit date (i.e., YYYYMMDD format) using // (integer division) and % (modulo division).

```
lb = 20080915
```

Write a function date that takes an 8-digit date argument and returns a year, month, and date tuple (e.g., return (year, month, day)).

Write a function date_2 that takes an 8-digit date as either integer or string.

Write a function date_3 that takes a list of 8-digit dates as integers or strings.

Write a for loop that prints the squares of integers from 1 to 10.

Write a for loop that prints the squares of even integers from 1 to 10.

Write a for loop that sums the squares of integers from 1 to 10.

Write a for loop that sums the squares of integers from 1 to 10 but stops before the sum exceeds 50.

FizzBuzz

Solve [FizzBuzz](#).

Use ternary expressions to make your FizzBuzz solution more compact.

Triangle

Write a function `triangle` that accepts a positive integer N and prints a numerical triangle of height $N - 1$. For example, `triangle(N=6)` should print:

```
1
22
333
4444
55555
```

Two Sum

Write a function `two_sum` that does the following.

Given a list of integers `nums` and an integer `target`, return the indices of the two numbers that add up to `target`.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Here are some examples:

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3]`, `target = 6`

Output: `[0,1]`

I saw this question on [LeetCode](#).

Best Time

Write a function `best_time` that solves the following.

You are given a list `prices` where `prices[i]` is the price of a given stock on the i^{th} day.

You want to maximize your profit by choosing a single day to buy one stock and choosing a different day in the future to sell that stock.

Return the maximum profit you can achieve from this transaction. If you cannot achieve any profit, return 0.

Here are some examples:

Example 1:

Input: `prices = [7,1,5,3,6,4]`

Output: 5

Explanation: Buy on day 2 (price = 1) and sell on day 5 (price = 6), profit = 6-1 = 5. Note that buying on day 2 and selling on day 1 is not allowed because you must buy before you sell.

Example 2:

Input: `prices = [7,6,4,3,1]`

Output: 0

Explanation: In this case, no transactions are done and the max profit = 0.

I saw this question on [LeetCode](#).

McKinney Chapter 2 - Practice - Sec 02

Announcements

1. Check your email inbox for an invitation to a free six-month subscription to DataCamp
 1. I added a few short courses to our course group
 2. These short courses are completely optional
 3. DataCamp has lots of resources to help you learn Python, R, SQL, Excel, etc.
2. Here are links to a few finance newsletters I strongly suggest:
 1. Matt Levine: <https://www.bloomberg.com/account/newsletters/money-stuff>
 2. Byrne Hobart: https://capitalgains.thediff.co/subscribe?ref=I0N1NGdmJq&_bhli_d=7fecfad9eb7fd8bcd529e945e11346b5897acdc
 3. Clifford Asness: <https://www.aqr.com/Insights/Perspectives>
 4. Owen Lamont: <https://www.acadian-asset.com/investment-insights/owenomics#>

Five-Minute Review

Practice

Extract the year, month, and day from an 8-digit date (i.e., YYYYMMDD format) using // (integer division) and % (modulo division).

```
lb = 20080915
```

```
lb
```

```
20080915
```

```
lb // 10_000 # // is integer division
```

2008

```
lb % 10_000 # % is modulo or remainder division
```

915

```
(lb % 10_000) // 100
```

9

```
lb % 100
```

15

What happened here?

- Floor or integer division `//` drops the digits on the right side (one digit per zero)
 - Modulo or remainder division `%` keeps the digits on the right side (one digit per zero)
-

Here is solution that approximates Excel's `LEFT()`, `MID()`, and `RIGHT()`. This works, but is not very Pythonic.

```
int(str(lb)[:4])
```

2008

```
int(str(lb)[4:6])
```

9

```
int(str(lb)[7:8])
```

5

Write a function date that takes an 8-digit date argument and returns a year, month, and date tuple (e.g., return (year, month, day)).

```
def date(ymd):
    year = ymd // 10_000 # // is integer division
    month = (ymd % 10_000) // 100
    day = ymd % 100
    return (year, month, day)
```

```
date(lb)
```

(2008, 9, 15)

```
date(20250110)
```

(2025, 1, 10)

Write a function date_2 that takes an 8-digit date as either integer or string.

```
def date_2(ymd):
    # if type(ymd) is str:
    #
    if isinstance(ymd, str):
        ymd = int(ymd)
    return date(ymd)
```

The `isinstance(ymd, str)` is better than `type(ymd) == str` because the `isinstance()` function also tests for all sub-classes of `str`. More here: <https://stackoverflow.com/questions/152580/whats-the-canonical-way-to-check-for-type-in-python/>.

```
date_2(str(lb))
```

(2008, 9, 15)

```
date_2('20250110')
```

(2025, 1, 10)

Write a function date_3 that takes a list of 8-digit dates as integers or strings.

```
dates_in = [20080915, 20250110]
```

```
def date_3(dates_in):
    dates_out = []
    for d in dates_in:
        # dates_out += [date_2(d)] # alternative
        dates_out.append(date_2(d))

    return dates_out
```

I have a slight preference for `.append()` over `+= []` because `.append()` modifies the list in place instead of making a copy. However, the speed differences in most cases in this course will be negligible. More here: <https://www.geeksforgeeks.org/difference-between-and-append-in-python/>.

```
date_3(dates_in)
```

```
[(2008, 9, 15), (2025, 1, 10)]
```

Write a for loop that prints the squares of integers from 1 to 10.

```
print('a', 'b', 'c', sep = '---')
```

```
a---b---c
```

```
for i in range(1, 11):
    print(i**2, end=' ')
```

```
1 4 9 16 25 36 49 64 81 100
```

Write a for loop that prints the squares of even integers from 1 to 10.

```
for i in range(1, 11):
    if i % 2 == 0:
        print(i**2, end=' ')
```

4 16 36 64 100

```
for i in range(2, 11, 2):
    print(i**2, end=' ')
```

4 16 36 64 100

Write a for loop that sums the squares of integers from 1 to 10.

```
total = 0
for i in range(1, 11):
    total += i**2
total
```

385

Write a for loop that sums the squares of integers from 1 to 10 but stops before the sum exceeds 50.

```
total = 0 # Initialize sum to zero

for i in range(1, 11): # Loop from 1 to 10
    # Check if adding the square of i would exceed 50
    if (total + i**2) > 50:
        # 'break' exits the loop completely, stopping further iterations
        # 'continue' would skip to the next iteration without executing further code in this
        break

    # Add the square of i to total
    total += i**2

# Print the final sum (implicit return in this case since it is the last line in the code cell)
total
```

30

FizzBuzzSolve [FizzBuzz](#).

Here is some pseudo code. The test for multiples of 3 and 5 must come first, otherwise it would never run!

```
# for i in range(1, 101):
#     # test for multiple of 3 & 5
#     #     print fizzbuzz
#     # test for multiple of 3
#     #     print fizz
#     # test for multiple of 5
#     #     print buzz
#     # otherwise print i
```

Here is my favorite FizzBuzz solution.

```
for i in range(1, 101):
    if (i % 3 == 0) & (i % 5 == 0):
        print('FizzBuzz', end=' ')
    elif (i % 3 == 0):
        print('Fizz', end=' ')
    elif (i % 5 == 0):
        print('Buzz', end=' ')
    else:
        print(i, end=' ')
```

1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz

Use ternary expressions to make your FizzBuzz solution more compact.

Here is a compact FizzBuzz solution. I consider the solution above easier to read and troubleshoot. The compact solution below uses the trick that we can multiply a string by `True` to return the string itself or by `False` to return an empty string.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) if (i%3==0) or (i%5==0) else i, end=' ')
```

```
1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz
```

Here is *an even more compact* FizzBuzz solution. The trick below is that Python's `or` returns its first truthy value. - If the concatenated string (`'Fizz'*(i%3==0) + 'Buzz'*(i%5==0)`) is not an empty string, which is falsy in Python, the `or` evaluates to that string. - If the string is empty, which means `i` is not divisible by 3 or 5, the `or` evaluates to `i`.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) or i, end=' ')
```

```
1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz
```

Triangle

Write a function `triangle` that accepts a positive integer N and prints a numerical triangle of height $N - 1$. For example, `triangle(N=6)` should print:

```
1
22
333
4444
55555
```

```
def triangle(N):
    for i in range(1, N):
        print(str(i) * i)
```

```
triangle(6)
```

```
1
22
333
4444
55555
```

The solution above works because multiplying a string by `i` concatenates `i` copies of that string.

```
'Test' + 'Test' + 'Test'
```

```
'TestTestTest'
```

```
'Test' * 3
```

```
'TestTestTest'
```

Two Sum

Write a function `two_sum` that does the following.

Given a list of integers `nums` and an integer `target`, return the indices of the two numbers that add up to target.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Here are some examples:

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3]`, `target = 6`

Output: `[0,1]`

I saw this question on [LeetCode](#).

```
def two_sum(nums, target):
    for i in range(1, len(nums)):
        for j in range(i):
            if nums[i] + nums[j] == target:
                return [j, i]
```

```
two_sum(nums = [2,7,11,15], target = 9)
```

[0, 1]

```
two_sum(nums = [3,2,4], target = 6)
```

[1, 2]

```
two_sum(nums = [3,3], target = 6)
```

[0, 1]

Best Time

Write a function `best_time` that solves the following.

You are given a list `prices` where `prices[i]` is the price of a given stock on the i^{th} day.

You want to maximize your profit by choosing a single day to buy one stock and choosing a different day in the future to sell that stock.

Return the maximum profit you can achieve from this transaction. If you cannot achieve any profit, return 0.

Here are some examples:

Example 1:

Input: `prices = [7,1,5,3,6,4]`

Output: 5

Explanation: Buy on day 2 (price = 1) and sell on day 5 (price = 6), profit = 6-1 = 5. Note that buying on day 2 and selling on day 1 is not allowed because you must buy before you sell.

Example 2:

Input: `prices = [7,6,4,3,1]`

Output: 0

Explanation: In this case, no transactions are done and the max profit = 0.

I saw this question on [LeetCode](#).

```
def max_profit(prices):
    # We start by assuming the first price is the lowest we've seen so far
    min_price = prices[0]

    # We initialize our maximum profit to zero, as no profit has been calculated yet
    max_profit = 0

    # Loop through each price in the list of prices
    for price in prices:
        # If the current price is lower than our lowest price seen, update min_price
        min_price = price if price < min_price else min_price

        # Calculate the profit if we were to sell at the current price
        profit = price - min_price

        # If this profit is better than our max profit so far, update max_profit
        max_profit = profit if profit > max_profit else max_profit

    # After checking all prices, return the maximum profit we've found
    return max_profit
```

```
max_profit(prices=[7,1,5,3,6,4])
```

5

```
max_profit(prices=[7,6,4,3,1])
```

0

We could replace the ternary statements with the `min()` and `max()` functions for a little more compact code.

```
def max_profit_2(prices):
    min_price = prices[0]
    max_profit = 0

    for price in prices:
        # Update min_price if current price is lower
        min_price = min(min_price, price)
```

```
# Calculate profit by selling at the current price  
current_profit = price - min_price  
  
# Update max_profit if the current_profit is higher  
max_profit = max(max_profit, current_profit)  
  
return max_profit
```

```
max_profit_2(prices=[7,1,5,3,6,4])
```

5

```
max_profit_2(prices=[7,6,4,3,1])
```

0

McKinney Chapter 2 - Practice - Sec 03

Announcements

1. Check your email inbox for an invitation to a free six-month subscription to DataCamp
 1. I added a few short courses to our course group
 2. These short courses are completely optional
 3. DataCamp has lots of resources to help you learn Python, R, SQL, Excel, etc.
2. Here are links to a few finance newsletters I strongly suggest:
 1. Matt Levine: <https://www.bloomberg.com/account/newsletters/money-stuff>
 2. Byrne Hobart: https://capitalgains.thediff.co/subscribe?ref=I0N1NGdmJq&_bhli_d=7fecfad9eb7fd8bcd529e945e11346b5897acdc
 3. Clifford Asness: <https://www.aqr.com/Insights/Perspectives>
 4. Owen Lamont: <https://www.acadian-asset.com/investment-insights/owenomics#>

Five-Minute Review

Practice

Extract the year, month, and day from an 8-digit date (i.e., YYYYMMDD format) using // (integer division) and % (modulo division).

```
lb = 20080915
```

```
lb
```

```
20080915
```

```
lb // 10_000 # // is integer division
```

2008

```
lb % 10_000 # % is modulo or remainder division
```

915

```
(lb % 10_000) // 100
```

9

```
lb % 100
```

15

What happened here?

- Floor or integer division `//` drops the digits on the right side (one digit per zero)
 - Modulo or remainder division `%` keeps the digits on the right side (one digit per zero)
-

Here is solution that approximates Excel's `LEFT()`, `MID()`, and `RIGHT()`. This works, but is not very Pythonic.

```
int(str(lb)[:4])
```

2008

```
int(str(lb)[4:6])
```

9

```
int(str(lb)[7:8])
```

5

Write a function date that takes an 8-digit date argument and returns a year, month, and date tuple (e.g., return (year, month, day)).

```
def date(x):
    year = x // 10_000 # // is integer division
    month = (x % 10_000) // 100
    day = x % 100
    return (year, month, day)
```

```
date(20250107)
```

(2025, 1, 7)

Write a function date_2 that takes an 8-digit date as either integer or string.

```
def date_2(x):
    # if type(x) is str:
    if isinstance(x, str):
        x = int(x)

    return date(x)
```

```
date_2(str(lb))
```

(2008, 9, 15)

```
date_2(lb)
```

(2008, 9, 15)

```
date_2('20250110')
```

(2025, 1, 10)

Write a function date_3 that takes a list of 8-digit dates as integers or strings.

```
ymds = [20080915, '20250110']
ymds
```

```
[20080915, '20250110']
```

This markdown cell is for *italicized* and **bold** text!

```
def date_3(ymds):
    ymds_out = []
    for ymd in ymds:
        ymds_out.append(date_2(ymd))

    return ymds_out
```

```
date_3(ymds)
```

```
[(2008, 9, 15), (2025, 1, 10)]
```

Write a for loop that prints the squares of integers from 1 to 10.

```
print('a', 'b', 'c', sep = '---')
```

```
a---b---c
```

```
for i in range(1, 11):
    print(i**2, end=' ')
```

```
1 4 9 16 25 36 49 64 81 100
```

Write a for loop that prints the squares of *even* integers from 1 to 10.

```
for i in range(1, 11):
    if i % 2 == 0:
        print(i**2, end=' ')
```

```
4 16 36 64 100
```

```
for i in range(2, 11, 2):
    print(i**2, end=' ')
```

```
4 16 36 64 100
```

Write a for loop that sums the squares of integers from 1 to 10.

```
total = 0
for i in range(1, 11):
    total += i**2

total
```

```
385
```

Write a for loop that sums the squares of integers from 1 to 10 but stops before the sum exceeds 50.

```
total = 0 # Initialize sum to zero

for i in range(1, 11): # Loop from 1 to 10
    # Check if adding the square of i would exceed 50
    if (total + i**2) > 50:
        # 'break' exits the loop completely, stopping further iterations
        # 'continue' would skip to the next iteration without executing further code in this
        break

    # Add the square of i to total
    total += i**2

# Print the final sum (implicit return in this case since it is the last line in the code cell)
total
```

```
30
```

FizzBuzz

Solve [FizzBuzz](#).

Here is some pseudo code. The test for multiples of 3 and 5 must come first, otherwise it would never run!

```
# for i in range(1, 101):
#     # test for multiple of 3 & 5
#     #     print fizzbuzz
#     # test for multiple of 3
#     #     print fizz
#     # test for multiple of 5
#     #     print buzz
#     # otherwise print i
```

Here is my favorite FizzBuzz solution.

```
for i in range(1, 101):
    if (i % 3 == 0) & (i % 5 == 0):
        print('FizzBuzz', end=' ')
    elif (i % 3 == 0):
        print('Fizz', end=' ')
    elif (i % 5 == 0):
        print('Buzz', end=' ')
    else:
        print(i, end=' ')
```

1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz

Use ternary expressions to make your FizzBuzz solution more compact.

Here is a compact FizzBuzz solution. I consider the solution above easier to read and troubleshoot. The compact solution below uses the trick that we can multiply a string by `True` to return the string itself or by `False` to return an empty string.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) if (i%3==0) or (i%5==0) else i, end=' ')
```

1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz

Here is *an even more compact* FizzBuzz solution. The trick below is that Python's `or` returns its first truthy value. - If the concatenated string (`'Fizz'*(i%3==0) + 'Buzz'*(i%5==0)`) is not an empty string, which is falsy in Python, the `or` evaluates to that string. - If the string is empty, which means `i` is not divisible by 3 or 5, the `or` evaluates to `i`.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) or i, end=' ')
```

```
1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz
```

Triangle

Write a function `triangle` that accepts a positive integer N and prints a numerical triangle of height $N - 1$. For example, `triangle(N=6)` should print:

```
1
22
333
4444
55555
```

```
def triangle(N):
    for i in range(1, N):
        print(str(i) * i)
```

```
triangle(6)
```

```
1
22
333
4444
55555
```

The solution above works because multiplying a string by `i` concatenates `i` copies of that string.

```
'Test' + 'Test' + 'Test'
```

```
'TestTestTest'
```

```
'Test' * 3
```

```
'TestTestTest'
```

Two Sum

Write a function `two_sum` that does the following.

Given a list of integers `nums` and an integer `target`, return the indices of the two numbers that add up to target.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Here are some examples:

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3]`, `target = 6`

Output: `[0,1]`

I saw this question on [LeetCode](#).

```
def two_sum(nums, target):
    for i in range(1, len(nums)):
        for j in range(i):
            if nums[i] + nums[j] == target:
                return [j, i]
```

```
two_sum(nums = [2,7,11,15], target = 9)
```

```
[0, 1]
```

```
two_sum(nums = [3,2,4], target = 6)
```

[1, 2]

```
two_sum(nums = [3,3], target = 6)
```

[0, 1]

Best Time

Write a function `best_time` that solves the following.

You are given a list `prices` where `prices[i]` is the price of a given stock on the i^{th} day.

You want to maximize your profit by choosing a single day to buy one stock and choosing a different day in the future to sell that stock.

Return the maximum profit you can achieve from this transaction. If you cannot achieve any profit, return 0.

Here are some examples:

Example 1:

Input: `prices = [7,1,5,3,6,4]`

Output: 5

Explanation: Buy on day 2 (price = 1) and sell on day 5 (price = 6), profit = 6-1 = 5. Note that buying on day 2 and selling on day 1 is not allowed because you must buy before you sell.

Example 2:

Input: `prices = [7,6,4,3,1]`

Output: 0

Explanation: In this case, no transactions are done and the max profit = 0.

I saw this question on [LeetCode](#).

```
def max_profit(prices):
    # We start by assuming the first price is the lowest we've seen so far
    min_price = prices[0]

    # We initialize our maximum profit to zero, as no profit has been calculated yet
    max_profit = 0
```

```
# Loop through each price in the list of prices
for price in prices:
    # If the current price is lower than our lowest price seen, update min_price
    min_price = price if price < min_price else min_price

    # Calculate the profit if we were to sell at the current price
    profit = price - min_price

    # If this profit is better than our max profit so far, update max_profit
    max_profit = profit if profit > max_profit else max_profit

# After checking all prices, return the maximum profit we've found
return max_profit
```

```
max_profit(prices=[7,1,5,3,6,4])
```

5

```
max_profit(prices=[7,6,4,3,1])
```

0

We could replace the ternary statements with the `min()` and `max()` functions for a little more compact code.

```
def max_profit_2(prices):
    min_price = prices[0]
    max_profit = 0

    for price in prices:
        # Update min_price if current price is lower
        min_price = min(min_price, price)

        # Calculate profit by selling at the current price
        current_profit = price - min_price

        # Update max_profit if the current_profit is higher
        max_profit = max(max_profit, current_profit)

    return max_profit
```

```
max_profit_2(prices=[7,1,5,3,6,4])
```

5

```
max_profit_2(prices=[7,6,4,3,1])
```

0

McKinney Chapter 2 - Practice - Sec 04

HELLO! NIRAMAY!

Announcements

1. Check your email inbox for an invitation to a free six-month subscription to DataCamp
 1. I added a few short courses to our course group
 2. These short courses are completely optional
 3. DataCamp has lots of resources to help you learn Python, R, SQL, Excel, etc.
2. Here are links to a few finance newsletters I strongly suggest:
 1. Matt Levine: <https://www.bloomberg.com/account/newsletters/money-stuff>
 2. Byrne Hobart: https://capitalgains.thediff.co/subscribe?ref=I0N1NGdmJq&_bhli_d=7fecfad9eb7fd8bcd529e945e11346b5897acdc
 3. Clifford Asness: <https://www.aqr.com/Insights/Perspectives>
 4. Owen Lamont: <https://www.acadian-asset.com/investment-insights/owenomics#>

Five-Minute Review

Practice

Extract the year, month, and day from an 8-digit date (i.e., YYYYMMDD format) using // (integer division) and % (modulo division).

```
lb = 20080915
```

```
lb
```

```
20080915
```

```
lb // 10_000 # // is integer division
```

2008

```
lb % 10_000 # % is modulo or remainder division
```

915

```
(lb % 10_000) // 100
```

9

```
lb % 100
```

15

What happened here?

- Floor or integer division `//` drops the digits on the right side (one digit per zero)
 - Modulo or remainder division `%` keeps the digits on the right side (one digit per zero)
-

Here is solution that approximates Excel's `LEFT()`, `MID()`, and `RIGHT()`. This works, but is not very Pythonic.

```
int(str(lb)[:4])
```

2008

```
int(str(lb)[4:6])
```

9

```
int(str(lb)[7:8])
```

5

Write a function date that takes an 8-digit date argument and returns a year, month, and date tuple (e.g., return (year, month, day)).

```
import this
```

The Zen of Python, by Tim Peters

Beautiful is better than ugly.
Explicit is better than implicit.
Simple is better than complex.
Complex is better than complicated.
Flat is better than nested.
Sparse is better than dense.
Readability counts.
Special cases aren't special enough to break the rules.
Although practicality beats purity.
Errors should never pass silently.
Unless explicitly silenced.
In the face of ambiguity, refuse the temptation to guess.
There should be one-- and preferably only one --obvious way to do it.
Although that way may not be obvious at first unless you're Dutch.
Now is better than never.
Although never is often better than *right* now.
If the implementation is hard to explain, it's a bad idea.
If the implementation is easy to explain, it may be a good idea.
Namespaces are one honking great idea -- let's do more of those!

```
def date(ymd):  
    year = ymd // 10_000 # // is integer division  
    month = (ymd % 10_000) // 100  
    day = ymd % 100  
    return (year, month, day)
```

```
lb
```

20080915

```
%who
```

```
date      lb  this
```

```
date(lb)
```

```
(2008, 9, 15)
```

```
date(20250110)
```

```
(2025, 1, 10)
```

```
type(date(20250110))
```

tuple

Write a function date_2 that takes an 8-digit date as either integer or string.

```
def date_2(ymd):
    # if type(ymd) is str:
    if isinstance(ymd, str):
        ymd = int(ymd)
    return date(ymd)
```

```
date_2(lb)
```

```
(2008, 9, 15)
```

```
date_2(str(lb))
```

```
(2008, 9, 15)
```

```
date_2(20250110)
```

```
(2025, 1, 10)
```

```
date_2('20250110')
```

```
(2025, 1, 10)
```

Write a function date_3 that takes a list of 8-digit dates as integers or strings.

```
ymds = [20080915, '20250110']  
ymds
```

```
[20080915, '20250110']
```

This markdown cell is for *italicized* and **bold** text!

```
def date_3(ymds):  
    ymds_out = []  
    for ymd in ymds:  
        ymds_out.append(date_2(ymd))  
  
    return ymds_out
```

```
date_3(ymds)
```

```
[(2008, 9, 15), (2025, 1, 10)]
```

Write a for loop that prints the squares of integers from 1 to 10.

```
print(1, 2, 3, sep='---')
```

```
1---2---3
```

```
for i in range(1, 11):  
    print(i**2, end=' ')
```

```
1 4 9 16 25 36 49 64 81 100
```

Write a for loop that prints the squares of even integers from 1 to 10.

```
for i in range(1, 11):
    if i % 2 == 0:
        print(i**2, end=' ')
```

4 16 36 64 100

```
for i in range(2, 11, 2):
    print(i**2, end=' ')
```

4 16 36 64 100

Write a for loop that sums the squares of integers from 1 to 10.

```
total = 0
for i in range(1, 11):
    total += i**2
total
```

385

Write a for loop that sums the squares of integers from 1 to 10 but stops before the sum exceeds 50.

```
total = 0 # Initialize sum to zero

for i in range(1, 11): # Loop from 1 to 10
    # Check if adding the square of i would exceed 50
    if (total + i**2) > 50:
        # 'break' exits the loop completely, stopping further iterations
        # 'continue' would skip to the next iteration without executing further code in this
        break

    # Add the square of i to total
    total += i**2

# Print the final sum (implicit return in this case since it is the last line in the code cell)
total
```

30

FizzBuzzSolve [FizzBuzz](#).

Here is some pseudo code. The test for multiples of 3 and 5 must come first, otherwise it would never run!

```
# for i in range(1, 101):
#     # test for multiple of 3 & 5
#     #     print fizzbuzz
#     # test for multiple of 3
#     #     print fizz
#     # test for multiple of 5
#     #     print buzz
#     # otherwise print i
```

Here is my favorite FizzBuzz solution.

```
for i in range(1, 101):
    if (i % 3 == 0) & (i % 5 == 0):
        print('FizzBuzz', end=' ')
    elif (i % 3 == 0):
        print('Fizz', end=' ')
    elif (i % 5 == 0):
        print('Buzz', end=' ')
    else:
        print(i, end=' ')
```

1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz

Use ternary expressions to make your FizzBuzz solution more compact.

Here is a compact FizzBuzz solution. I consider the solution above easier to read and troubleshoot. The compact solution below uses the trick that we can multiply a string by `True` to return the string itself or by `False` to return an empty string.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) if (i%3==0) or (i%5==0) else i, end=' ')
```

```
1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz
```

Here is *an even more compact* FizzBuzz solution. The trick below is that Python's `or` returns its first truthy value. - If the concatenated string (`'Fizz'*(i%3==0) + 'Buzz'*(i%5==0)`) is not an empty string, which is falsy in Python, the `or` evaluates to that string. - If the string is empty, which means `i` is not divisible by 3 or 5, the `or` evaluates to `i`.

```
for i in range(1, 101):
    print('Fizz'*(i%3==0) + 'Buzz'*(i%5==0) or i, end=' ')
```

```
1 2 Fizz 4 Buzz Fizz 7 8 Fizz Buzz 11 Fizz 13 14 FizzBuzz 16 17 Fizz 19 Buzz Fizz 22 23 Fizz
```

Triangle

Write a function `triangle` that accepts a positive integer N and prints a numerical triangle of height $N - 1$. For example, `triangle(N=6)` should print:

```
1
22
333
4444
55555
```

```
def triangle(N):
    for i in range(1, N):
        print(str(i) * i)
```

```
triangle(6)
```

```
1
22
333
4444
55555
```

The solution above works because multiplying a string by `i` concatenates `i` copies of that string.

```
'Test' + 'Test' + 'Test'
```

```
'TestTestTest'
```

```
'Test' * 3
```

```
'TestTestTest'
```

Two Sum

Write a function `two_sum` that does the following.

Given a list of integers `nums` and an integer `target`, return the indices of the two numbers that add up to target.

You may assume that each input would have exactly one solution, and you may not use the same element twice.

You can return the answer in any order.

Here are some examples:

Example 1:

Input: `nums = [2,7,11,15]`, `target = 9`

Output: `[0,1]`

Explanation: Because `nums[0] + nums[1] == 9`, we return `[0, 1]`.

Example 2:

Input: `nums = [3,2,4]`, `target = 6`

Output: `[1,2]`

Example 3:

Input: `nums = [3,3]`, `target = 6`

Output: `[0,1]`

I saw this question on [LeetCode](#).

```
def two_sum(nums, target):
    for i in range(1, len(nums)):
        for j in range(i):
            if nums[i] + nums[j] == target:
                return [j, i]
```

```
two_sum(nums = [2,7,11,15], target = 9)
```

[0, 1]

```
two_sum(nums = [3,2,4], target = 6)
```

[1, 2]

```
two_sum(nums = [3,3], target = 6)
```

[0, 1]

Best Time

Write a function `best_time` that solves the following.

You are given a list `prices` where `prices[i]` is the price of a given stock on the i^{th} day.

You want to maximize your profit by choosing a single day to buy one stock and choosing a different day in the future to sell that stock.

Return the maximum profit you can achieve from this transaction. If you cannot achieve any profit, return 0.

Here are some examples:

Example 1:

Input: `prices = [7,1,5,3,6,4]`

Output: 5

Explanation: Buy on day 2 (price = 1) and sell on day 5 (price = 6), profit = 6-1 = 5. Note that buying on day 2 and selling on day 1 is not allowed because you must buy before you sell.

Example 2:

Input: `prices = [7,6,4,3,1]`

Output: 0

Explanation: In this case, no transactions are done and the max profit = 0.

I saw this question on [LeetCode](#).

```
def max_profit(prices):
    # We start by assuming the first price is the lowest we've seen so far
    min_price = prices[0]

    # We initialize our maximum profit to zero, as no profit has been calculated yet
    max_profit = 0

    # Loop through each price in the list of prices
    for price in prices:
        # If the current price is lower than our lowest price seen, update min_price
        min_price = price if price < min_price else min_price

        # Calculate the profit if we were to sell at the current price
        profit = price - min_price

        # If this profit is better than our max profit so far, update max_profit
        max_profit = profit if profit > max_profit else max_profit

    # After checking all prices, return the maximum profit we've found
    return max_profit
```

```
max_profit(prices=[7,1,5,3,6,4])
```

5

```
max_profit(prices=[7,6,4,3,1])
```

0

We could replace the ternary statements with the `min()` and `max()` functions for a little more compact code.

```
def max_profit_2(prices):
    min_price = prices[0]
    max_profit = 0

    for price in prices:
        # Update min_price if current price is lower
        min_price = min(min_price, price)
```

```
# Calculate profit by selling at the current price  
current_profit = price - min_price  
  
# Update max_profit if the current_profit is higher  
max_profit = max(max_profit, current_profit)  
  
return max_profit
```

```
max_profit_2(prices=[7,1,5,3,6,4])
```

5

```
max_profit_2(prices=[7,6,4,3,1])
```

0

Week 2

McKinney Chapter 3 - Built-In Data Structures, Functions, and Files

Introduction

We must understand Python's core functionality to use NumPy and pandas. Chapter 3 of McKinney (2022) discusses Python's core functionality. We will focus on the following:

1. Data structures
 1. tuples
 2. lists
 3. dicts (also known as dictionaries)
2. List comprehensions
3. Functions
 1. Returning multiple values
 2. Using anonymous functions

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Data Structures and Sequences

Python's data structures are simple but powerful. Mastering their use is a critical part of becoming a proficient Python programmer.

Tuple

A tuple is a fixed-length, immutable sequence of Python objects.

We cannot change a tuple after we create it because tuples are immutable. A tuple is ordered, so we can subset or slice it with a numerical index. We will surround tuples with parentheses but they are not required.

```
tup = (4, 5, 6)
```

Python is zero-indexed, so zero accesses the first element in tup!

```
tup[0]
```

4

```
tup[1]
```

5

```
tup[2]
```

6

```
nested_tup = ((4, 5, 6), (7, 8))
```

```
nested_tup[0]
```

(4, 5, 6)

```
nested_tup[0][0]
```

4

```
tup = ('foo', [1, 2], True)
```

If an object inside a tuple is mutable, such as a list, you can modify it in-place.

```
# tup[2] = False # gives an error, because tuples are immutable (unchangeable)
```

```
tup[1].append(3)  
tup
```

('foo', [1, 2, 3], True)

You can concatenate tuples using the + operator to produce longer tuples:

Tuples are immutable, but we can combine two tuples into a new tuple.

```
(1, 2) + (1, 2)
```

```
(1, 2, 1, 2)
```

```
(4, None, 'foo') + (6, 0) + ('bar',)
```

```
(4, None, 'foo', 6, 0, 'bar')
```

Multiplying a tuple by an integer, as with lists, has the effect of concatenating together that many copies of the tuple:

This multiplication behavior is the logical extension of the addition behavior above. The output of `tup + tup` should be the same as that of `2 * tup`.

```
('foo', 'bar') + ('foo', 'bar')
```

```
('foo', 'bar', 'foo', 'bar')
```

```
('foo', 'bar') * 2
```

```
('foo', 'bar', 'foo', 'bar')
```

Unpacking tuples

If you try to assign to a tuple-like expression of variables, Python will attempt to unpack the value on the righthand side of the equals sign.

```
tup = (4, 5, 6)
a, b, c = tup
```

```
a
```

```
4
```

```
b
```

```
5
```

c

6

```
(d, e, f) = (7, 8, 9) # the parentheses are optional but helpful!
```

d

7

e

8

f

9

We can unpack nested tuples!

```
tup = 4, 5, (6, 7)
a, b, (c, d) = tup
```

Tuple methods

Since the size and contents of a tuple cannot be modified, it is very light on instance methods. A particularly useful one (also available on lists) is count, which counts the number of occurrences of a value.

```
a = (1, 2, 2, 2, 3, 4, 2)
a.count(2)
```

4

Python is zero-indexed!

```
a.index(2)
```

1

List

In contrast with tuples, lists are variable-length and their contents can be modified in-place. You can define them using square brackets [] or using the list type function.

```
a_list = [2, 3, 7, None]
tup = ('foo', 'bar', 'baz')
b_list = list(tup)
```

```
a_list
```

[2, 3, 7, None]

```
b_list
```

['foo', 'bar', 'baz']

Python is zero-indexed!

```
a_list[0]
```

2

Concatenating and combining lists

Similar to tuples, adding two lists together with + concatenates them.

```
[4, None, 'foo'] + [7, 8, (2, 3)]
```

[4, None, 'foo', 7, 8, (2, 3)]

The .append() method adds its argument as the last element in a list.

```
xx = [4, None, 'foo']
xx.append([7, 8, (2, 3)])
xx
```

```
[4, None, 'foo', [7, 8, (2, 3)]]
```

If you have a list already defined, you can append multiple elements to it using the extend method.

```
x = [4, None, 'foo']
x.extend([7, 8, (2, 3)])
x
```

```
[4, None, 'foo', 7, 8, (2, 3)]
```

Check your output! It will take you time to understand all these methods!

Slicing

Slicing is very important!

You can select sections of most sequence types by using slice notation, which in its basic form consists of start:stop passed to the indexing operator [].

Recall that Python is zero-indexed, so the first element has an index of 0. A consequence of zero-indexing is that start:stop is inclusive on the left edge (start) and exclusive on the right edge (stop).

```
seq = [7, 2, 3, 7, 5, 6, 0, 1]
seq
```

```
[7, 2, 3, 7, 5, 6, 0, 1]
```

```
seq[5]
```

6

Python is zero-indexed, so left edge of slide is included and right edge is excluded!

```
seq[1:5]
```

[2, 3, 7, 5]

Either the start or stop can be omitted, in which case they default to the start of the sequence and the end of the sequence, respectively.

```
seq[:5]
```

[7, 2, 3, 7, 5]

```
seq[3:]
```

[7, 5, 6, 0, 1]

Negative indices slice the sequence relative to the end.

```
seq[-1]
```

1

```
seq[-1:]
```

[1]

```
seq[-4:]
```

[5, 6, 0, 1]

```
seq[-4:-1]
```

[5, 6, 0]

```
seq[-6:-2]
```

[3, 7, 5, 6]

A step can also be used after a second colon to, say, take every other element.

```
seq
```

```
[7, 2, 3, 7, 5, 6, 0, 1]
```

```
seq[::2]
```

```
[7, 3, 5, 0]
```

```
seq[1::2]
```

```
[2, 7, 6, 1]
```

We can think of the trailing `:2` in the preceding code cells as “count by 2”. Therefore, the `1::2` slice:

- Starts at 1
- Stops at the end because of the first `:`
- Counts by 2 because of the trailing `:2`

A clever use of this is to pass `-1`, which has the useful effect of reversing a list or tuple.

```
seq[::-1]
```

```
[1, 0, 6, 5, 7, 3, 2, 7]
```

We will use slicing (subsetting) all semester, so we must understand the examples above.

dict

`dict` is likely the most important built-in Python data structure. A more common name for it is hash map or associative array. It is a flexibly sized collection of key-value pairs, where key and value are Python objects. One approach for creating one is to use curly braces `{}` and colons to separate keys and values.

Elements in dictionaries have named keys, while elements in tuples and lists have numerical indices. Dictionaries are handy for passing named arguments and returning named results.

```
empty_dict = {}
empty_dict
```

```
{}
```

A dictionary is a set of key-value pairs.

```
d1 = {'a': 'some value', 'b': [1, 2, 3, 4]}
d1
```

```
{'a': 'some value', 'b': [1, 2, 3, 4]}
```

```
d1['a']
```

```
'some value'
```

```
d1[7] = 'an integer'
d1
```

```
{'a': 'some value', 'b': [1, 2, 3, 4], 7: 'an integer'}
```

We access dictionary values by key names instead of key positions.

You can delete values either using the `del` keyword or the `pop` method (which simultaneously returns the value and deletes the key).

```
d1[5] = 'some value'
d1['dummy'] = 'another value'
d1
```

```
{'a': 'some value',
'b': [1, 2, 3, 4],
7: 'an integer',
5: 'some value',
'dummy': 'another value'}
```

```
del d1[5]
d1
```

```
{'a': 'some value',
'b': [1, 2, 3, 4],
7: 'an integer',
'dummy': 'another value'}
```

```
ret = d1.pop('dummy')
```

```
ret
```

```
'another value'
```

```
d1
```

```
{'a': 'some value', 'b': [1, 2, 3, 4], 7: 'an integer'}
```

The keys and values method give you iterators of the dict's keys and values, respectively. While the key-value pairs are not in any particular order, these functions output the keys and values in the same order.

```
d1.keys()
```

```
dict_keys(['a', 'b', 7])
```

```
d1.values()
```

```
dict_values(['some value', [1, 2, 3, 4], 'an integer'])
```

List, Set, and Dict Comprehensions

We will focus on list comprehensions, which are [Pythonic](#).

List comprehensions are one of the most-loved Python language features. They allow you to concisely form a new list by filtering the elements of a collection, transforming the elements passing the filter in one concise expression. They take the basic form:

```
[expr for val in collection if condition]
```

This is equivalent to the following for loop:

```
result = []
for val in collection:
    if condition:
        result.append(expr)
```

The filter condition can be omitted, leaving only the expression.

```
strings = ['a', 'as', 'bat', 'car', 'dove', 'python']
```

We could use a `for` loop over `strings` to keep only strings longer than two and then to capitalize them.

```
caps = []
for x in strings:
    if len(x) > 2:
        caps.append(x.upper())

caps
```

```
['BAT', 'CAR', 'DOVE', 'PYTHON']
```

A list comprehension is more Pythonic and replaces four lines of code with one. The general format of a list comprehension is [operation on `x` for `x` in list if condition]

```
[x.upper() for x in strings if len(x) > 2]
```

```
['BAT', 'CAR', 'DOVE', 'PYTHON']
```

Here is another example. The following code is a `for` loop and an equivalent list comprehension that squares the integers from 1 to 10.

```
squares = []
for i in range(1, 11):
    squares.append(i ** 2)

squares
```

```
[1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
```

```
[i**2 for i in range(1, 11)]
```

```
[1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
```

What if we wanted the squares of even numbers?

```
[i**2 for i in range(1, 11) if i%2==0]
```

```
[4, 16, 36, 64, 100]
```

```
[i**2 for i in range(2, 11, 2)]
```

```
[4, 16, 36, 64, 100]
```

Functions

Functions are the primary and most important method of code organization and reuse in Python. As a rule of thumb, if you anticipate needing to repeat the same or very similar code more than once, it may be worth writing a reusable function. Functions can also help make your code more readable by giving a name to a group of Python statements.

Functions are declared with the def keyword and returned from with the return keyword:

```
def my_function(x, y, z=1.5):
    if z > 1:
        return z * (x + y)
    else:
        return z / (x + y)
```

There is no issue with having multiple return statements. If Python reaches the end of a function without encountering a return statement, None is returned automatically.

Each function can have positional arguments and keyword arguments. Keyword arguments are most commonly used to specify default values or optional arguments. In the preceding function, x and y are positional arguments while z is a keyword argument. This means that the function can be called in any of these ways:

```
my_function(5, 6, z=0.7)
my_function(3.14, 7, 3.5)
my_function(10, 20)
```

The main restriction on function arguments is that the keyword arguments must follow the positional arguments (if any). You can specify keyword arguments in any order; this frees you from having to remember which order the function arguments were specified in and only what their names are.

Returning Multiple Values

We can write Python functions that return multiple objects. The function `f()` below returns one tuple that we can unpack to multiple objects.

```
def f():
    a = 5
    b = 6
    c = 7
    return (a, b, c)
```

```
f()
```

```
(5, 6, 7)
```

If we want to return multiple objects with names or labels, we can return a dictionary.

```
def f():
    a = 5
    b = 6
    c = 7
    return {'a' : a, 'b' : b, 'c' : c}
```

```
f()
```

```
{'a': 5, 'b': 6, 'c': 7}
```

```
f()['a']
```

5

Anonymous (Lambda) Functions

Python has support for so-called anonymous or lambda functions, which are a way of writing functions consisting of a single statement, the result of which is the return value. They are defined with the `lambda` keyword, which has no meaning other than “we are declaring an anonymous function.”

I usually refer to these as lambda functions in the rest of the book. They are especially convenient in data analysis because, as you’ll see, there are many cases where data transformation functions will take functions as arguments. It’s often less typing (and clearer) to pass a lambda function as opposed to writing a full-out function declaration or even assigning the lambda function to a local variable.

Lambda functions are Pythonic and let us to write simple functions on the fly.

```
strings = ['foo', 'card', 'bar', 'aaaa', 'abab']
```

```
strings.sort()  
strings
```

```
['aaaa', 'abab', 'bar', 'card', 'foo']
```

```
len(strings[0])
```

```
4
```

```
strings.sort(key=len)  
strings
```

```
['bar', 'foo', 'aaaa', 'abab', 'card']
```

For example, we could use a lambda function to sort `strings` by the last letter of each string.

```
strings.sort(key=lambda x: x[-1])  
strings
```

```
['aaaa', 'abab', 'card', 'foo', 'bar']
```

What if we want to sort by the *second* to last letter?

```
strings.sort(key=lambda x: x[-2])  
strings
```

```
['aaaa', 'abab', 'bar', 'foo', 'card']
```

McKinney Chapter 3 - Practice - Blank

Announcements

Five-Minute Review

Practice

Swap the values assigned to `a` and `b` using a third variable `c`.

```
a = 1
```

```
b = 2
```

Swap the values assigned to `a` and `b` *without* using a third variable `c`.

```
a = 1
```

```
b = 2
```

What is the output of the following code and why?

```
1, 1, 1 == (1, 1, 1)
```

(1, 1, False)

Create a list 11 of integers from 1 to 100.

Slice 11 to create a list 12 of integers from 60 to 50 (inclusive).

Create a list 13 of odd integers from 1 to 21.

Create a list 14 of the squares of integers from 1 to 100.

Create a list 15 that contains the squares of *odd* integers from 1 to 100.

Use a lambda function to sort strings by the last letter in each string.

```
strings = ['Clemson', 'Hinson', 'Pillsbury', 'Shubrick']
```

Given an integer array `nums` and an integer `k`, write a function to return the k^{th} largest element in the array.

Note that it is the k^{th} largest element in the sorted order, not the k^{th} distinct element.

Example 1:

Input: `nums` = [3,2,1,5,6,4], `k` = 2

Output: 5

Example 2:

Input: `nums` = [3,2,3,1,2,4,5,5,6], `k` = 4

Output: 4

I saw this question on [LeetCode](#).

```
nums = [3,2,3,1,2,4,5,5,6]
k = 4
```

Given an integer array `nums` and an integer `k`, write a function to return the `k` most frequent elements.

You may return the answer in any order.

Example 1:

Input: `nums` = [1,1,1,2,2,3], `k` = 2

Output: [1,2]

Example 2:

Input: `nums` = [1], `k` = 1

Output: [1]

I saw this question on [LeetCode](#).

```
nums = [1,1,1,2,2,3]  
k = 2
```

Test whether the given strings are palindromes.

Input: ["aba", "no"]

Output: [True, False]

```
tickers = ["AAPL", "GOOG", "XOX", "XOM"]
```

Write a function `calc_returns()` that accepts lists of prices and dividends and returns a list of returns.

```
prices = [100, 150, 100, 50, 100, 150, 100, 150]  
dividends = [1, 1, 1, 1, 2, 2, 2, 2]
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns lists of returns, capital gains yields, and dividend yields.

Write a function `rescale()` to rescale and shift numbers so that they cover the range [0, 1].

Input: [18.5, 17.0, 18.0, 19.0, 18.0]

Output: [0.75, 0.0, 0.5, 1.0, 0.5]

```
nums = [18.5, 17.0, 18.0, 19.0, 18.0]
```

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

```
data = {  
    "AAPL": (150.25, 10), # (price, shares)  
    "GOOGL": (2750.00, 2),  
    "MSFT": (300.75, 5)  
}
```

McKinney Chapter 3 - Practice - Sec 02

The `%precision` magic makes it easy to round all float on print to 4 decimal places.

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep forming groups on Canvas under *People* in the left sidebar
2. You must ask me for groups larger than four students

Five-Minute Review

List

A list is an ordered collection of objects that is changeable (mutable). You can create an empty list using either `[]` or `list()`.

```
my_list = [1, 2, 3, [1, 2, 3, [1, 2, 3]]]  
my_list
```

```
[1, 2, 3, [1, 2, 3, [1, 2, 3]]]
```

Python is zero-indexed!

```
my_list[0]
```

1

```
my_list[:3] # to get first 3 objects, :3
```

```
[1, 2, 3]
```

```
my_list[1:4] # to get next 3 objects from 1, go from 1 to 1+3 or 1:3
```

```
[2, 3, [1, 2, 3, [1, 2, 3]]]
```

Tuple

A tuple is similar to a list but un-changeable (immutable). You can create a tuple using parentheses () or the `tuple()` function.

```
my_tuple = (1, 2, 3, (1, 2, 3, (1, 2, 3)))  
my_tuple
```

```
(1, 2, 3, (1, 2, 3, (1, 2, 3)))
```

Python is zero-indexed!

```
my_tuple[0]
```

```
1
```

Tuples are immutable, so they cannot be changed!

```
# my_tuple[0] = 2_001  
# -----  
# TypeError Traceback (most recent call last)  
# Cell In[10], line 1  
# ----> 1 my_tuple[0] = 2_001  
  
# TypeError: 'tuple' object does not support item assignment
```

Dictionary

A dictionary is an ordered collection of key-value pairs that are changeable (mutable). You can create an empty dictionary using either {} or the `dict()` function.

```
my_dict = {'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}
```

- The *key* can be anything hashable (string, integer, tuple), but I (almost) always make the key a string
- The *value* can be any python object

```
my_dict['wb']
```

```
'Warren Buffett'
```

```
my_dict['pl'] = 'Peter Lynch'  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman', 'pl': 'Peter Lynch'}
```

List Comprehension

A list comprehension is a concise way of creating a new list by iterating over an existing list or other iterable object. It is more time and space-efficient than traditional for loops and offers a cleaner syntax. The basic syntax of a list comprehension is `new_list = [expression for item in iterable if condition]` where:

1. `expression` is the operation to be performed on each element of the iterable
2. `item` is the current element being processed
3. `iterable` is the list or other iterable object being iterated over
4. `condition` is an optional filter that only accepts items that evaluate to True.

For example, we can use the following list comprehension to create a new list of even numbers from 0 to 8: `even_numbers = [x for x in range(9) if x % 2 == 0]`

List comprehensions are a powerful tool in Python that can help you write more efficient and readable code (i.e., more Pythonic code).

What if we wanted multiples of 3 or 5 from 1 to 25?

```
threes_fives = [i for i in range(1, 26) if (i%3==0) | (i%5==0)]
```

```
threes_fives
```

```
[3, 5, 6, 9, 10, 12, 15, 18, 20, 21, 24, 25]
```

```
threes_fives_2 = [print(i) for i in range(1, 26) if (i%3==0) | (i%5==0)]
```

```
3  
5  
6  
9  
10  
12  
15  
18  
20  
21  
24  
25
```

```
threes_fives_2
```

```
[None, None, None]
```

Practice

Swap the values assigned to a and b using a third variable c.

```
a = 1
```

```
b = 2
```

```
c = a
```

```
a = b
```

```
b = c

print(f'a is {a} and b is {b}')


a is 2 and b is 1
```

More on f-strings!

F-strings offer a concise way to embed expressions inside string literals, using curly braces {}. Prefixed with f or F, these strings allow for easy formatting of variables, numbers, and expressions. For example:

```
name = "Alice"
print(f"Hello, {name}!")
```

This outputs “Hello, Alice!”. F-strings simplify complex formatting, making code more readable. For a deeper understanding and more examples: <https://realpython.com/python-f-strings/>

Swap the values assigned to a and b *without* using a third variable c.

```
a = 1
b = 2
b, a = a, b
print(f'a is {a} and b is {b}')


a is 2 and b is 1
```

```
a = 1
b = 2
a, b = b, a
print(f'a is {a} and b is {b}')


a is 2 and b is 1
```

What is the output of the following code and why?

```
1, 1, 1 == (1, 1, 1)
```

```
(1, 1, False)
```

Without parentheses (), Python reads the final element in the tuple as `1 == (1, 1, 1)`, which is `False`. We can use parentheses () to force Python to do what we want!

```
(1, 1, 1) == (1, 1, 1)
```

```
True
```

For this example, we must use parentheses () to be unambiguous!

Create a list 11 of integers from 1 to 100.

```
l1 = list(range(1, 101))
```

```
l1[:5]
```

```
[1, 2, 3, 4, 5]
```

```
l1[-5:]
```

```
[96, 97, 98, 99, 100]
```

Slice l1 to create a list 12 of integers from 60 to 50 (inclusive).

```
l2 = l1[59:48:-1]  
l2
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
11[48:59]
```

```
[49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59]
```

```
12_alt_1 = l1[49:60][::-1]  
12_alt_1
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
12_alt_2 = list(reversed(l1[49:60]))  
12_alt_2
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

Create a list 13 of odd integers from 1 to 21.

```
13 = list(range(1, 22, 2))  
13
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

Python != is the same as Excel's <>.

```
13_alt = [i for i in range(22) if i%2 != 0]  
13_alt
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_do_not_do_this = []  
for i in range(22):  
    if i%2 != 0:  
        13_alt_do_not_do_this.append(i)  
  
13_alt_do_not_do_this
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_do_not_do_this
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

Create a list 14 of the squares of integers from 1 to 100.

```
14 = [i**2 for i in range(1, 101)]  
14[:5]
```

```
[1, 4, 9, 16, 25]
```

Create a list 15 that contains the squares of *odd* integers from 1 to 100.

```
15 = [i**2 for i in range(1, 101) if i%2!=0]  
15[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15_alt = [i**2 for i in range(1, 101, 2)]  
15_alt[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15 == 15_alt
```

```
True
```

Which one is faster?!

```
%timeit [i**2 for i in range(1, 101) if i%2!=0]
```

```
12.3 s ± 2.96 s per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

```
%timeit [i**2 for i in range(1, 101, 2)]
```

```
4.74 s ± 956 ns per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

Premature optimization is the root of all evil

- Donald Knuth

Use a lambda function to sort strings by the last letter in each string.

```
strings = ['Pillsbury', 'Shubrick', 'Clemson', 'Hinson']
```

```
strings.sort()  
strings
```

```
['Clemson', 'Hinson', 'Pillsbury', 'Shubrick']
```

```
'Clemson'[-1]
```

```
'n'
```

```
strings.sort(key=lambda x: x[-1])  
strings
```

```
['Shubrick', 'Clemson', 'Hinson', 'Pillsbury']
```

```
strings.sort(key=len)  
strings
```

```
['Hinson', 'Clemson', 'Shubrick', 'Pillsbury']
```

Given an integer array `nums` and an integer `k`, write a function to return the k^{th} largest element in the array.

Note that it is the k^{th} largest element in the sorted order, not the k^{th} distinct element.

Example 1:

Input: `nums = [3,2,1,5,6,4]`, `k = 2`

Output: 5

Example 2:

Input: `nums = [3,2,3,1,2,4,5,5,6]`, `k = 4`

Output: 4

I saw this question on [LeetCode](#).

```
def get_klarge(nums=[3,2,3,1,2,4,5,5,6], k=4):
    return sorted(nums)[-k]
```

```
get_klarge(nums=[3,2,1,5,6,4], k=2)
```

5

The following code follows a bad practice and uses variables in the global scope instead of passing them as arguments or parameters of the functions.

```
nums = [3,2,3,1,2,4,5,5,6]
k = 4
```

```
def get_klarge_donotdo():
    return sorted(nums)[-k]
```

```
get_klarge_donotdo()
```

4

Here is an extreme example how this practice can lead to non-deterministic and confusing results that depend on how many times a function has been run.

```
x = [1, 2, 3, 4]
```

```
def confusing():
    x.append(1)
```

```
confusing()
confusing()
confusing()
confusing()
confusing()
```

```
x
```

```
[1, 2, 3, 4, 1, 1, 1, 1]
```

Given an integer array `nums` and an integer `k`, write a function to return the `k` most frequent elements.

You may return the answer in any order.

Example 1:

Input: `nums = [1,1,1,2,2,3], k = 2`
Output: `[1,2]`

Example 2:

Input: `nums = [1], k = 1`
Output: `[1]`

I saw this question on [LeetCode](#).

```
def get_kfreq(nums, k):
    counts = {}
    for n in nums:
        if n in counts:
            counts[n] += 1
        else:
            counts[n] = 1
    return [x[0] for x in sorted(counts.items(), key=lambda x: x[1], reverse=True)[:k]]
```

```
get_kfreq(nums=[1,1,1,2,2,3], k=2)
```

```
[1, 2]
```

Test whether the given strings are palindromes.

Input: ["aba", "no"]
Output: [True, False]

We can reverse a string with a `[::-1]` slice just like we reverse a string!

```
def is_palindrome(x):
    return [_x == _x[::-1] for _x in x]
```

```
is_palindrome(["aba", "no"])
```

[True, False]

```
tickers = ["AAPL", "GOOG", "XOX", "XOM"]
```

```
is_palindrome(tickers)
```

[False, True, True, False]

Write a function `calc_returns()` that accepts lists of prices and dividends and returns a list of returns.

```
prices = [100, 150, 100, 50, 100, 150, 100, 150]
dividends = [1, 1, 1, 1, 2, 2, 2, 2]
```

Although loop counters are un-Pythonic, this calculation is the rare case where I found loop counters more clear.

```
def calc_returns(p, d):
    r = []
    for i in range(1, len(p)):
        # uncomment this line to watch loop iterations
        # print(f'r is {r}, p[i] is {p[i]}, p[i-1] is {p[i-1]}, d[i] is {d[i]}')
        r.append((p[i] - p[i-1] + d[i]) / p[i-1])
    return r
```

```
calc_returns(p=prices, d=dividends)
```

```
[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]
```

We do not have to specify the argument names `p=` and `d=`, but they help me avoid errors.

```
calc_returns(prices, dividends)
```

```
[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]
```

We can do the same calculation without indexing! Instead, we can use `zip()` to simultaneously loop over prices, lagged prices, and dividends.

```
def calc_returns_zip(p, d):
    r = []
    for _p, _plag, _d in zip(p[1:], p[:-1], d[1:]):
        r.append((_p - _plag + _d) / _plag)
    return r
```

```
calc_returns_zip(p=prices, d=dividends)
```

```
[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]
```

```
calc_returns_zip(p=prices, d=dividends) == calc_returns_zip(p=prices, d=dividends)
```

```
True
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns lists of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):
    r, cg, dp = [], [], []
    for i in range(1, len(p)):
        r.append((p[i] + d[i] - p[i-1]) / p[i-1])
        cg.append((p[i] - p[i-1]) / p[i-1])
        dp.append(d[i] / p[i-1])

    return {'r':r, 'cg':cg, 'dp': dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': [0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200],  
'cg': [0.5000, -0.3333, -0.5000, 1.0000, 0.5000, -0.3333, 0.5000],  
'dp': [0.0100, 0.0067, 0.0100, 0.0400, 0.0200, 0.0133, 0.0200]}
```

```
calc_returns(p=prices, d=dividends) == calc_returns_2(p=prices, d=dividends)['r']
```

True

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1].

Input: [18.5, 17.0, 18.0, 19.0, 18.0]

Output: [0.75, 0.0, 0.5, 1.0, 0.5]

```
nums = [18.5, 17.0, 18.0, 19.0, 18.0]
```

```
def rescale(x):  
    x_min = min(x)  
    x_max = max(x)  
    return [(i - x_min) / (x_max - x_min) for i in x]
```

```
rescale(nums)
```

```
[0.7500, 0.0000, 0.5000, 1.0000, 0.5000]
```

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

```
data = {  
    "AAPL": (150.25, 10), # (price, shares)  
    "GOOGL": (2750.00, 2),  
    "MSFT": (300.75, 5)  
}
```

```
def calc_portval(data):
    portval = 0
    for p, n in data.values():
        portval += p * n
    return portval
```

```
calc_portval(data=data)
```

8506.2500

McKinney Chapter 3 - Practice - Sec 03

The `%precision` magic makes it easy to round all float on print to 4 decimal places.

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep forming groups on Canvas under *People* in the left sidebar
2. You must ask me for groups larger than four students

Five-Minute Review

List

A list is an ordered collection of objects that is changeable (mutable). You can create an empty list using either `[]` or `list()`.

```
my_list = [1, 2, 3, [1, 2, 3, [1, 2, 3]]]  
my_list
```

```
[1, 2, 3, [1, 2, 3, [1, 2, 3]]]
```

Python is zero-indexed!

```
my_list[0]
```

```
1
```

```
my_list[:3] # to get first 3 objects, :3
```

```
[1, 2, 3]
```

```
my_list[1:4] # to get next 3 objects from 1, go from 1 to 1+3 or 1:3
```

```
[2, 3, [1, 2, 3, [1, 2, 3]]]
```

Tuple

A tuple is similar to a list but un-changeable (immutable). You can create a tuple using parentheses () or the `tuple()` function.

```
my_tuple = (1, 2, 3, (1, 2, 3, (1, 2, 3)))  
my_tuple
```

```
(1, 2, 3, (1, 2, 3, (1, 2, 3)))
```

Python is zero-indexed!

```
my_tuple[0]
```

```
1
```

Tuples are immutable, so they cannot be changed!

```
# my_tuple[0] = 2_001  
# # -----  
# # TypeError Traceback (most recent call last)  
# # Cell In[10], line 1  
# # ----> 1 my_tuple[0] = 2_001  
# # TypeError: 'tuple' object does not support item assignment
```

Dictionary

A dictionary is an ordered collection of key-value pairs that are changeable (mutable). You can create an empty dictionary using either {} or the `dict()` function.

```
my_dict = {'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}
```

- The *key* can be anything hashable (string, integer, tuple), but I (almost) always make the key a string
- The *value* can be any python object

```
my_dict['wb']
```

```
'Warren Buffett'
```

```
my_dict['pl'] = 'Peter Lynch'  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman', 'pl': 'Peter Lynch'}
```

```
my_dict[(0, 1, 2)] = 'Risky! Do not do live demos!'  
my_dict
```

```
{'wb': 'Warren Buffett',  
'sk': 'Seth Klarman',  
'pl': 'Peter Lynch',  
(0, 1, 2): 'Risky! Do not do live demos!'}
```

List Comprehension

A list comprehension is a concise way of creating a new list by iterating over an existing list or other iterable object. It is more time and space-efficient than traditional for loops and offers a cleaner syntax. The basic syntax of a list comprehension is `new_list = [expression for item in iterable if condition]` where:

1. `expression` is the operation to be performed on each element of the iterable
2. `item` is the current element being processed
3. `iterable` is the list or other iterable object being iterated over
4. `condition` is an optional filter that only accepts items that evaluate to True.

For example, we can use the following list comprehension to create a new list of even numbers from 0 to 8: `even_numbers = [x for x in range(9) if x % 2 == 0]`

List comprehensions are a powerful tool in Python that can help you write more efficient and readable code (i.e., more Pythonic code).

What if we wanted multiples of 3 or 5 from 1 to 25?

```
threes_fives = [i for i in range(1, 26) if (i%3==0) | (i%5==0)]
```

```
threes_fives
```

```
[3, 5, 6, 9, 10, 12, 15, 18, 20, 21, 24, 25]
```

```
threes_fives_2 = [print(i) for i in range(1, 26) if (i%3==0) | (i%5==0)]
```

```
3  
5  
6  
9  
10  
12  
15  
18  
20  
21  
24  
25
```

```
threes_fives_2
```

```
[None, None, None]
```

Practice

Swap the values assigned to a and b using a third variable c.

```
a = 1  
b = 2  
c = [a, b]
```

```
a = c[1]
```

```
b = c[0]
```

```
print(f'a is {a} and b is {b}')
```

a is 2 and b is 1

Here is another way:

```
a = 1  
b = 2  
c = a  
a = b  
b = c
```

```
print(f'a is {a} and b is {b}')
```

a is 2 and b is 1

More on f-strings!

F-strings offer a concise way to embed expressions inside string literals, using curly braces {}. Prefixed with f or F, these strings allow for easy formatting of variables, numbers, and expressions. For example:

```
name = "Alice"  
print(f"Hello, {name}!")
```

This outputs “Hello, Alice!”. F-strings simplify complex formatting, making code more readable. For a deeper understanding and more examples: <https://realpython.com/python-f-strings/>

Swap the values assigned to a and b *without* using a third variable c.

```
a = 1
b = 2

b, a = a, b

print(f'a is {a} and b is {b}')
```

a is 2 and b is 1

What is the output of the following code and why?

```
1, 1, 1 == (1, 1, 1)
```

(1, 1, False)

Without parentheses (), Python reads the final element in the tuple as 1 == (1, 1, 1), which is **False**. We can use parentheses () to force Python to do what we want!

```
(1, 1, 1) == (1, 1, 1)
```

True

For this example, we must use parentheses () to be unambiguous!

Create a list l1 of integers from 1 to 100.

```
l1 = list(range(1, 101))
```

```
l1[:5]
```

[1, 2, 3, 4, 5]

```
11[-5:]
```

```
[96, 97, 98, 99, 100]
```

Slice 11 to create a list 12 of integers from 60 to 50 (inclusive).

```
11.index(60)
```

```
59
```

```
12 = 11[59:48:-1]
```

```
12_alt_1 = 11[49:60]
12_alt_1.reverse() # most list methods modify a list "in place"
12_alt_1
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
12_alt_2 = 11[49:60][::-1]
12_alt_2
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

Create a list 13 of odd integers from 1 to 21.

```
13 = list(range(1, 22, 2))
13
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_1 = []
for i in range(1, 22):
    if i%2 != 0:
        13_alt_1.append(i)

13_alt_1
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_2 = [i for i in range(1, 22) if i%2!=0]
13_alt_2
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

Create a list 14 of the squares of integers from 1 to 100.

```
14 = [i**2 for i in range(1, 101)]
14[:5]
```

```
[1, 4, 9, 16, 25]
```

Create a list 15 that contains the squares of *odd* integers from 1 to 100.

```
15 = [i**2 for i in range(1, 101) if i%2!=0]
15[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15_alt = [i**2 for i in range(1, 101, 2)]
15_alt[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15 == 15_alt
```

True

Which one is faster?!

```
%timeit [i**2 for i in range(1, 101) if i%2!=0]
```

```
12.2 s ± 4.94 s per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

```
%timeit [i**2 for i in range(1, 101, 2)]
```

```
5.96 s ± 998 ns per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

Premature optimization is the root of all evil

- Donald Knuth

Use a lambda function to sort strings by the last letter in each string.

```
strings = ['Pillsbury', 'Shubrick', 'Clemson', 'Hinson']
```

```
strings.sort()  
strings
```

```
['Clemson', 'Hinson', 'Pillsbury', 'Shubrick']
```

```
len('Pillsbury')
```

9

```
strings.sort(key=len)  
strings
```

```
['Hinson', 'Clemson', 'Shubrick', 'Pillsbury']
```

```
'Pillsbury'[-1]
```

'y'

```
strings.sort(key=lambda x: x[-1])  
strings
```

```
['Shubrick', 'Hinson', 'Clemson', 'Pillsbury']
```

Given an integer array `nums` and an integer `k`, write a function to return the k^{th} largest element in the array.

Note that it is the k^{th} largest element in the sorted order, not the k^{th} distinct element.

Example 1:

Input: `nums = [3,2,1,5,6,4]`, `k = 2`

Output: 5

Example 2:

Input: `nums = [3,2,3,1,2,4,5,5,6]`, `k = 4`

Output: 4

I saw this question on [LeetCode](#).

```
def get_klarge(nums=[3,2,3,1,2,4,5,5,6], k=4):
    return sorted(nums)[-k]
```

```
get_klarge(nums=[3,2,1,5,6,4], k=2)
```

5

The following code follows a bad practice and uses variables in the global scope instead of passing them as arguments or parameters of the functions.

```
nums = [3,2,3,1,2,4,5,5,6]
k = 4
```

```
def get_klarge_donotdo():
    return sorted(nums)[-k]
```

```
get_klarge_donotdo()
```

4

Here is an extreme example how this practice can lead to non-deterministic and confusing results that depend on how many times a function has been run.

```
x = [1, 2, 3, 4]
```

```
def confusing():
    x.append(1)
```

```
confusing()
confusing()
confusing()
confusing()
confusing()
```

```
x
```

```
[1, 2, 3, 4, 1, 1, 1, 1]
```

Given an integer array `nums` and an integer `k`, write a function to return the `k` most frequent elements.

You may return the answer in any order.

Example 1:

Input: `nums = [1,1,1,2,2,3], k = 2`
Output: `[1,2]`

Example 2:

Input: `nums = [1], k = 1`
Output: `[1]`

I saw this question on [LeetCode](#).

```
def get_kfreq(nums, k):
    counts = {}
    for n in nums:
        if n in counts:
            counts[n] += 1
        else:
            counts[n] = 1
    return [x[0] for x in sorted(counts.items(), key=lambda x: x[1], reverse=True)[:k]]
```

```
get_kfreq(nums=[1,1,1,2,2,3], k=2)
```

```
[1, 2]
```

Test whether the given strings are palindromes.

Input: ["aba", "no"]

Output: [True, False]

```
def is_palindrome(x):
    return [_x == _x[::-1] for _x in x]
```

```
is_palindrome(["aba", "no"])
```

[True, False]

```
tickers = ["AAPL", "GOOG", "XOX", "XOM"]
```

```
is_palindrome(tickers)
```

[False, True, True, False]

Write a function calc_returns() that accepts lists of prices and dividends and returns a list of returns.

```
prices = [100, 150, 100, 50, 100, 150, 100, 150]
dividends = [1, 1, 1, 1, 2, 2, 2, 2]
```

Although loop counters are un-Pythonic, this calculation is the rare case where I found loop counters more clear.

```
def calc_returns(p, d):
    r = []
    for i in range(1, len(p)):
        # the following print statement print r, p, and d values each iteration
        # print(f'r is {r}, p[i] is {p[i]}, p[i-1] is {p[i-1]}, d[i] is {d[i]},')
        r.append((p[i] + d[i] - p[i-1]) / p[i-1])
    return r
```

```
calc_returns(p=prices, d=dividends)
```

[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]

We can do the same calculation without indexing! Instead, we can use `zip()` to simultaneously loop over prices, lagged prices, and dividends.

```
def calc_returns_zip(p, d):
    r = []
    for _p, _plag, _d in zip(p[1:], p[:-1], d[1:]):
        # the following print statement print r, p, and d values each iteration
        # print(f'r is {r}, _p is {_p}, _plag is {_plag}, _d is {_d},')
        r.append((_p + _d - _plag) / _plag)
    return r
```

```
calc_returns_zip(p=prices, d=dividends)
```

```
[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]
```

```
calc_returns_zip(p=prices, d=dividends) == calc_returns_zip(p=prices, d=dividends)
```

```
True
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns lists of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):
    r = []
    cg = []
    dp = []
    # r, cg, dp = [], [], [] # tuple unpacking is very Pythonic!
    for i in range(1, len(p)):
        r.append((p[i] + d[i] - p[i-1]) / p[i-1])
        cg.append((p[i] - p[i-1]) / p[i-1])
        dp.append(d[i] / p[i-1])

    return {'r':r, 'cg':cg, 'dp':dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': [0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200],
 'cg': [0.5000, -0.3333, -0.5000, 1.0000, 0.5000, -0.3333, 0.5000],
 'dp': [0.0100, 0.0067, 0.0100, 0.0400, 0.0200, 0.0133, 0.0200]}
```

```
calc_returns(p=prices, d=dividends) == calc_returns_2(p=prices, d=dividends) ['r']
```

True

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1].

Input: [18.5, 17.0, 18.0, 19.0, 18.0]
Output: [0.75, 0.0, 0.5, 1.0, 0.5]

```
nums = [18.5, 17.0, 18.0, 19.0, 18.0]
```

```
def rescale(x):
    x_min = min(x)
    x_max = max(x)
    return [(i - x_min) / (x_max - x_min) for i in x]
```

```
rescale(nums)
```

[0.7500, 0.0000, 0.5000, 1.0000, 0.5000]

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

```
data = {
    "AAPL": (150.25, 10), # (price, shares)
    "GOOGL": (2750.00, 2),
    "MSFT": (300.75, 5)
}
```

```
def calc_portval(data):
    portval = 0
    for p, n in data.values():
        portval += p * n
    return portval
```

```
calc_portval(data=data)
```

8506.2500

McKinney Chapter 3 - Practice - Blank

The `%precision` magic makes it easy to round all float on print to 4 decimal places.

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep forming groups on Canvas under *People* in the left sidebar
2. You must ask me for groups larger than four students

Five-Minute Review

List

A list is an ordered collection of objects that is changeable (mutable). You can create an empty list using either `[]` or `list()`.

```
my_list = [1, 2, 3, [1, 2, 3, [1, 2, 3]]]  
my_list
```

```
[1, 2, 3, [1, 2, 3, [1, 2, 3]]]
```

Python is zero-indexed!

```
my_list[0]
```

```
1
```

```
my_list[:3] # to get first 3 objects, :3
```

```
[1, 2, 3]
```

```
my_list[1:4] # to get next 3 objects from 1, go from 1 to 1+3 or 1:3
```

```
[2, 3, [1, 2, 3, [1, 2, 3]]]
```

Tuple

A tuple is similar to a list but un-changeable (immutable). You can create a tuple using parentheses () or the `tuple()` function.

```
my_tuple = (1, 2, 3, (1, 2, 3, (1, 2, 3)))  
my_tuple
```

```
(1, 2, 3, (1, 2, 3, (1, 2, 3)))
```

Python is zero-indexed!

```
my_tuple[0]
```

```
1
```

Tuples are immutable, so they cannot be changed!

```
# my_tuple[0] = 2_001  
# # -----  
# # TypeError Traceback (most recent call last)  
# # Cell In[10], line 1  
# # ----> 1 my_tuple[0] = 2_001  
# # TypeError: 'tuple' object does not support item assignment
```

Dictionary

A dictionary is an ordered collection of key-value pairs that are changeable (mutable). You can create an empty dictionary using either {} or the `dict()` function.

```
my_dict = {'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman'}
```

- The *key* can be anything hashable (string, integer, tuple), but I (almost) always make the key a string
- The *value* can be any python object

```
my_dict['wb']
```

```
'Warren Buffett'
```

```
my_dict['pl'] = 'Peter Lynch'  
my_dict
```

```
{'wb': 'Warren Buffett', 'sk': 'Seth Klarman', 'pl': 'Peter Lynch'}
```

```
my_dict[(0, 1, 2)] = {'ad': 'Another dictionary!'}  
my_dict
```

```
{'wb': 'Warren Buffett',  
'sk': 'Seth Klarman',  
'pl': 'Peter Lynch',  
(0, 1, 2): {'ad': 'Another dictionary!'}}
```

List Comprehension

A list comprehension is a concise way of creating a new list by iterating over an existing list or other iterable object. It is more time and space-efficient than traditional for loops and offers a cleaner syntax. The basic syntax of a list comprehension is `new_list = [expression for item in iterable if condition]` where:

1. `expression` is the operation to be performed on each element of the iterable
2. `item` is the current element being processed
3. `iterable` is the list or other iterable object being iterated over
4. `condition` is an optional filter that only accepts items that evaluate to True.

For example, we can use the following list comprehension to create a new list of even numbers from 0 to 8: `even_numbers = [x for x in range(9) if x % 2 == 0]`

List comprehensions are a powerful tool in Python that can help you write more efficient and readable code (i.e., more Pythonic code).

What if we wanted multiples of 3 or 5 from 1 to 25?

```
threes_fives = [i for i in range(1, 26) if (i%3==0) | (i%5==0)]  
threes_fives
```

[3, 5, 6, 9, 10, 12, 15, 18, 20, 21, 24, 25]

```
threes_fives_2 = [print(i) for i in range(1, 26) if (i%3==0) | (i%5==0)]  
threes_fives_2
```

3
5
6
9
10
12
15
18
20
21
24
25

[None, None, None]

Practice

Swap the values assigned to a and b using a third variable c.

```
a = 1  
b = 2
```

```
c = a  
a = b  
b = c
```

```
print(f'a is {a} and b is {b}')
```

a is 2 and b is 1

More on f-strings!

F-strings offer a concise way to embed expressions inside string literals, using curly braces {}. Prefixed with f or F, these strings allow for easy formatting of variables, numbers, and expressions. For example:

```
name = "Alice"  
print(f"Hello, {name}!")
```

This outputs “Hello, Alice!”. F-strings simplify complex formatting, making code more readable. For a deeper understanding and more examples: <https://realpython.com/python-f-strings/>

Swap the values assigned to a and b *without* using a third variable c.

```
a = 1  
b = 2
```

```
b, a = a, b
```

```
print(f'a is {a} and b is {b}')
```

a is 2 and b is 1

What is the output of the following code and why?

```
1, 1, 1 == (1, 1, 1)
```

```
(1, 1, False)
```

Without parentheses (), Python reads the final element in the tuple as `1 == (1, 1, 1)`, which is `False`. We can use parentheses () to force Python to do what we want!

```
(1, 1, 1) == (1, 1, 1)
```

```
True
```

For this example, we must use parentheses () to be unambiguous!

Create a list 11 of integers from 1 to 100.

```
l1 = list(range(1, 101))
l1[:5]
```

```
[1, 2, 3, 4, 5]
```

Slice l1 to create a list 12 of integers from 60 to 50 (inclusive).

```
l1.index(60)
```

```
59
```

```
l2 = l1[59:48:-1]
l2
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
l2_alt_1 = l1[49:60]
l2_alt_1.reverse()
l2_alt_1
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
12_alt_2 = 11[49:60]
12_alt_2.sort(reverse=True)
12_alt_2
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

```
11[49:60][::-1]
```

```
[60, 59, 58, 57, 56, 55, 54, 53, 52, 51, 50]
```

Create a list 13 of odd integers from 1 to 21.

```
13 = list(range(1, 22, 2))
13
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_1 = []
for i in range(1, 22):
    if i%2 != 0:
        13_alt_1.append(i)

13_alt_1
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

```
13_alt_2 = [i for i in range(1, 22) if i%2!=0]
13_alt_2
```

```
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21]
```

Create a list 14 of the squares of integers from 1 to 100.

```
14 = [i**2 for i in range(1, 101)]
14[:5]
```

```
[1, 4, 9, 16, 25]
```

Create a list 15 that contains the squares of odd integers from 1 to 100.

```
15 = [i**2 for i in range(1, 101) if i%2!=0]
15[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15_alt = [i**2 for i in range(1, 101, 2)]
15_alt[:5]
```

```
[1, 9, 25, 49, 81]
```

```
15 == 15_alt
```

True

Which one is faster?!

```
%timeit [i**2 for i in range(1, 101) if i%2!=0]
```

```
13.1 s ± 4.81 s per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

```
%timeit [i**2 for i in range(1, 101, 2)]
```

```
6.51 s ± 771 ns per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

Premature optimization is the root of all evil

– Donald Knuth

Use a lambda function to sort strings by the last letter in each string.

```
strings = ['Pillsbury', 'Shubrick', 'Clemson', 'Hinson']
```

```
strings.sort()  
strings
```

```
['Clemson', 'Hinson', 'Pillsbury', 'Shubrick']
```

```
len('Pillsbury')
```

9

```
strings.sort(key=len)  
strings
```

```
['Hinson', 'Clemson', 'Shubrick', 'Pillsbury']
```

```
'Pillsbury' [-1]
```

'y'

```
strings.sort(key=lambda x: x[-1])  
strings
```

```
['Shubrick', 'Hinson', 'Clemson', 'Pillsbury']
```

Given an integer array `nums` and an integer `k`, write a function to return the k^{th} largest element in the array.

Note that it is the k^{th} largest element in the sorted order, not the k^{th} distinct element.

Example 1:

Input: `nums = [3,2,1,5,6,4]`, `k = 2`

Output: 5

Example 2:

Input: `nums = [3,2,3,1,2,4,5,5,6]`, `k = 4`

Output: 4

I saw this question on [LeetCode](#).

```
def get_klarge(nums, k):  
    return sorted(nums)[-k]
```

```
get_klarge(nums=[3,2,1,5,6,4], k=2)
```

5

```
get_klarge(nums=[3,2,3,1,2,4,5,5,6], k=4)
```

4

The following code follows a bad practice and uses variables in the global scope instead of passing them as arguments or parameters of the functions.

```
nums = [3,2,3,1,2,4,5,5,6]  
k = 4
```

```
def get_klarge_donotdo():  
    return sorted(nums)[-k]
```

```
get_klarge_donotdo()
```

4

Here is an extreme example how this practice can lead to non-deterministic and confusing results that depend on how many times a function has been run.

```
x = []
```

```
def confusing():  
    x.append(len(x))
```

```
confusing()  
confusing()  
confusing()
```

x

[0, 1, 2]

Given an integer array `nums` and an integer `k`, write a function to return the `k` most frequent elements.

You may return the answer in any order.

Example 1:

Input: `nums = [1,1,1,2,2,3]`, `k = 2`

Output: [1,2]

Example 2:

Input: `nums = [1]`, `k = 1`

Output: [1]

I saw this question on [LeetCode](#).

```
def get_kfreq(nums, k):
    counts = {}
    for n in nums:
        if n in counts:
            counts[n] += 1
        else:
            counts[n] = 1
    return [x[0] for x in sorted(counts.items(), key=lambda x: x[1], reverse=True)[:k]]
```

get_kfreq(`nums=[1,1,1,2,2,3]`, `k=2`)

[1, 2]

Test whether the given strings are palindromes.

Input: ["aba", "no"]

Output: [True, False]

```
def is_palindrome(x):
    return [_x == _x[::-1] for _x in x]
```

```
is_palindrome(["aba", "no"])
```

```
[True, False]
```

```
tickers = ["AAPL", "GOOG", "XOX", "XOM"]
```

```
is_palindrome(tickers)
```

```
[False, True, True, False]
```

Write a function calc_returns() that accepts lists of prices and dividends and returns a list of returns.

```
prices = [100, 150, 100, 50, 100, 150, 100, 150]
dividends = [1, 1, 1, 1, 2, 2, 2, 2]
```

Although loop counters are un-Pythonic, this calculation is the rare case where I found loop counters more clear.

```
def calc_returns(p, d):
    r = []
    for i in range(1, len(p)):
        # the following print statement print r, p, and d values each iteration
        # print(f'r is {r}, p[i] is {p[i]}, p[i-1] is {p[i-1]}, d[i] is {d[i]}')
        r.append((p[i] + d[i] - p[i-1]) / p[i-1])
    return r

calc_returns(p=prices, d=dividends)
```

```
[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]
```

We can do the same calculation without indexing! Instead, we can use `zip()` to simultaneously loop over prices, lagged prices, and dividends.

```
def calc_returns_zip(p, d):
    r = []
    for _p, _plag, _d in zip(p[1:], p[:-1], d[1:]):
        # the following print statement print r, p, and d values each iteration
        # print(f'r is {r}, _p is {_p}, _plag is {_plag}, _d is {_d},')
        r.append((_p + _d - _plag) / _plag)
    return r
```

```
calc_returns_zip(p=prices, d=dividends)
```

[0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200]

```
calc_returns_zip(p=prices, d=dividends) == calc_returns_zip(p=prices, d=dividends)
```

True

Rewrite the function calc_returns() as calc_returns_2() so it returns lists of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):
    r = []
    cg = []
    dp = []
    # r, cg, dp = [], [], [] # tuple unpacking is very Pythonic!
    for i in range(1, len(p)):
        r.append((p[i] + d[i] - p[i-1]) / p[i-1])
        cg.append((p[i] - p[i-1]) / p[i-1])
        dp.append(d[i] / p[i-1])

    return {'r':r, 'cg':cg, 'dp':dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': [0.5100, -0.3267, -0.4900, 1.0400, 0.5200, -0.3200, 0.5200],
 'cg': [0.5000, -0.3333, -0.5000, 1.0000, 0.5000, -0.3333, 0.5000],
 'dp': [0.0100, 0.0067, 0.0100, 0.0400, 0.0200, 0.0133, 0.0200]}
```

```
calc_returns(p=prices, d=dividends) == calc_returns_2(p=prices, d=dividends) ['r']
```

True

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1].

Input: [18.5, 17.0, 18.0, 19.0, 18.0]
Output: [0.75, 0.0, 0.5, 1.0, 0.5]

```
nums = [18.5, 17.0, 18.0, 19.0, 18.0]
```

```
def rescale(x):
    x_min = min(x)
    x_max = max(x)
    return [(i - x_min) / (x_max - x_min) for i in x]
```

```
rescale(nums)
```

[0.7500, 0.0000, 0.5000, 1.0000, 0.5000]

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

```
data = {
    "AAPL": (150.25, 10), # (price, shares)
    "GOOGL": (2750.00, 2),
    "MSFT": (300.75, 5)
}
```

```
def calc_portval(data):
    portval = 0
    for p, n in data.values():
        portval += p * n
    return portval
```

```
calc_portval(data=data)
```

8506.2500

Week 3

McKinney Chapter 4 - NumPy Basics: Arrays and Vectorized Computation

```
import numpy as np
```

```
%precision 4
```

```
'%.4f'
```

Introduction

Chapter 4 of McKinney (2022) discusses the NumPy package (an abbreviation of numerical Python), which is the foundation for numerical computing in Python, including pandas.

We will focus on:

1. Creating arrays
2. Slicing arrays
3. Applying functions and methods to arrays
4. Using conditional logic with arrays (i.e., `np.where()` and `np.select()`)

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Here is a simple example of NumPy's speed and syntax advantages relative to Python's built-in data structures. First, we create a list and a NumPy array with values from 0 to 999,999.

```
my_list = list(range(1_000_000))
my_arr = np.arange(1_000_000)
```

```
my_list[:5]
```

```
[0, 1, 2, 3, 4]
```

```
my_arr[:5]
```

```
array([0, 1, 2, 3, 4])
```

We must use a for loop or a list comprehension to double each value in `my_list`. We will comment this code because it prints a list with one million elements!

```
# [2 * x for x in my_list] # list comprehension to double each value
```

However, we can multiply `my_arr` by two because math “just works” with NumPy. Jupyter will pretty print the NumPy array, showing only the first and last few elements.

```
my_arr * 2
```

```
array([0, 2, 4, ..., 1999994, 1999996, 1999998],  
      shape=(1000000,))
```

We can use the “magic” function `%timeit` to time these two calculations.

```
%timeit [x * 2 for x in my_list]
```

```
72.5 ms ± 16.1 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
```

```
%timeit my_arr * 2
```

```
4.44 ms ± 867 s per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

The NumPy version is much faster than the list version! The NumPy version is also faster to type, read, and troubleshoot, and our time is more valuable than computer time!

The NumPy ndarray: A Multidimensional Array Object

One of the key features of NumPy is its N-dimensional array object, or `ndarray`, which is a fast, flexible container for large datasets in Python. Arrays enable you to perform mathematical operations on whole blocks of data using similar syntax to the equivalent operations between scalar elements.

```
np.random.seed(42) # makes random numbers repeatable
data = np.random.randn(2, 3)
data
```

```
array([[ 0.4967, -0.1383,  0.6477],
       [ 1.523 , -0.2342, -0.2341]])
```

Multiplying `data` by 10 multiplies each element in `data` by 10, and adding `data` to itself adds each element to itself (i.e., element-wise addition). NumPy arrays must contain homogeneous data types (e.g., all floats or integers) to achieve this common-sense behavior.

```
data * 10
```

```
array([[ 4.9671, -1.3826,  6.4769],
       [15.2303, -2.3415, -2.3414]])
```

```
data_2 = data + data
data_2
```

```
array([[ 0.9934, -0.2765,  1.2954],
       [ 3.0461, -0.4683, -0.4683]])
```

NumPy arrays have attributes. Recall that IPython and Jupyter provide tab completion.

```
data.ndim
```

```
2
```

```
data.shape
```

```
(2, 3)
```

```
data.dtype
```

```
dtype('float64')
```

We slice NumPy arrays using `[]`, the same as we slice lists and tuples.

```
data[0]
```

```
array([ 0.4967, -0.1383,  0.6477])
```

We chain []s with arrays, the same as we chain []s with lists and tuples.

```
data[0][0]
```

```
0.4967
```

However, with NumPy arrays, we can replace n chained []s with one pair of []s containing n indexes or slices, separated by commas. For example, [i] [j] becomes [i, j], and [i] [j] [k] becomes [i, j, k].

```
data[0, 0] # zero row, zero column
```

```
0.4967
```

```
data[0][0] == data[0, 0]
```

```
np.True_
```

Creating ndarrays

The easiest way to create an array is to use the array function. This accepts any sequence-like object (including other arrays) and produces a new NumPy array containing the passed data

```
data1 = [6, 7.5, 8, 0, 1]
arr1 = np.array(data1)
arr1
```

```
array([6. , 7.5, 8. , 0. , 1. ])
```

```
arr1.dtype
```

```
dtype('float64')
```

Here, `np.array()` implicitly casts the integers in `data1` to floats because NumPy arrays must have homogenous data types. We could explicitly cast all values to integers but would lose information.

```
np.array(data1, dtype=np.int64)
```

```
array([6, 7, 8, 0, 1])
```

We can cast a list of lists to a two-dimensional NumPy array.

```
data2 = [[1, 2, 3, 4], [5, 6, 7, 8]]  
arr2 = np.array(data2)  
arr2
```

```
array([[1, 2, 3, 4],  
       [5, 6, 7, 8]])
```

```
arr2.shape
```

```
(2, 4)
```

```
arr2.dtype
```

```
dtype('int64')
```

There are several other ways to create NumPy arrays.

```
np.zeros((3, 6))
```

```
array([[0., 0., 0., 0., 0., 0.],  
       [0., 0., 0., 0., 0., 0.],  
       [0., 0., 0., 0., 0., 0.]])
```

```
np.ones((3, 6))
```

```
array([[1., 1., 1., 1., 1., 1.],  
       [1., 1., 1., 1., 1., 1.],  
       [1., 1., 1., 1., 1., 1.]])
```

```
np.ones_like(arr2)
```

```
array([[1, 1, 1, 1],
       [1, 1, 1, 1]])
```

The `np.arange()` function is similar to Python's built-in `range()` but creates an array directly.

```
np.array(range(15))
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

```
np.arange(15)
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14])
```

Table 4-1 from McKinney (2022) summarizes NumPy array creation functions.

- `array`: Convert input data (list, tuple, array, or other sequence type) to an ndarray either by inferring a dtype or explicitly specifying a dtype; copies the input data by default
- `asarray`: Convert input to ndarray, but do not copy if the input is already an ndarray
- `arange`: Like the built-in range but returns an ndarray instead of a list
- `ones, ones_like`: Produce an array of all 1s with the given shape and dtype; `ones_like` takes another array and produces a `ones` array of the same shape and dtype
- `zeros, zeros_like`: Like `ones` and `ones_like` but producing arrays of 0s instead
- `empty, empty_like`: Create new arrays by allocating new memory, but do not populate with any values like ones and zeros
- `full, full_like`: Produce an array of the given shape and dtype with all values set to the indicated "fill value"
- `eye, identity`: Create a square N-by-N identity matrix (1s on the diagonal and 0s elsewhere)

Arithmetic with NumPy Arrays

Arrays are important because they enable you to express batch operations on data without writing any for loops. NumPy users call this vectorization. Any arithmetic operations between equal-size arrays applies the operation element-wise

```
arr = np.array([[1., 2., 3.], [4., 5., 6.]])  
arr
```

```
array([[1., 2., 3.],  
       [4., 5., 6.]])
```

NumPy array addition is elementwise.

```
arr + arr
```

```
array([[ 2.,  4.,  6.],  
       [ 8., 10., 12.]])
```

NumPy array multiplication is elementwise.

```
arr * arr
```

```
array([[ 1.,  4.,  9.],  
       [16., 25., 36.]])
```

NumPy array division is elementwise.

```
1 / arr
```

```
array([[1.      , 0.5     , 0.3333],  
       [0.25    , 0.2     , 0.1667]])
```

NumPy powers are elementwise, too.

```
arr ** 2
```

```
array([[ 1.,  4.,  9.],  
       [16., 25., 36.]])
```

We can also raise a single value to the elements in an array!

```
2 ** arr  
  
array([[ 2.,  4.,  8.],  
       [16., 32., 64.]])
```

Basic Indexing and Slicing

We index and slice one-dimensional arrays in the same way as lists and tuples.

```
arr = np.arange(10)  
arr  
  
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
arr[5]
```

```
np.int64(5)
```

```
arr[5:8]
```

```
array([5, 6, 7])
```

```
equiv_list = list(range(10))  
equiv_list
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
equiv_list[5:8]
```

```
[5, 6, 7]
```

We must jump through some hoops to replace elements 5, 6, and 7 with the value 12 in the list `equiv_list`.

```
# # TypeError: can only assign an iterable  
# equiv_list[5:8] = 12
```

```
equiv_list[5:8] = [12] * 3  
equiv_list
```

```
[0, 1, 2, 3, 4, 12, 12, 12, 8, 9]
```

However, this operation is easy with NumPy array `arr`!

```
arr[5:8] = 12  
arr
```

```
array([ 0,  1,  2,  3,  4, 12, 12, 12,  8,  9])
```

We call this behavior “broadcasting”.

As you can see, if you assign a scalar value to a slice, as in `arr[5:8] = 12`, the value is propagated (or broadcasted henceforth) to the entire selection. An important first distinction from Python’s built-in lists is that array slices are views on the original array. This means that the data is not copied, and any modifications to the view will be reflected in the source array.

```
arr_slice = arr[5:8]  
arr_slice
```

```
array([12, 12, 12])
```

```
arr_slice[1] = 12345  
arr_slice
```

```
array([ 12, 12345,     12])
```

```
arr
```

```
array([ 0,     1,     2,     3,     4,    12, 12345,    12,     8,  
         9])
```

The `:` slices every element in `arr_slice`.

```
arr_slice[:] = 64
arr_slice
```

```
array([64, 64, 64])
```

```
arr
```

```
array([ 0,  1,  2,  3,  4, 64, 64, 64,  8,  9])
```

If you want a copy of a slice of an ndarray instead of a view, you will need to explicitly copy the array—for example, `arr[5:8].copy()`.

```
arr_slice_2 = arr[5:8].copy()
arr_slice_2
```

```
array([64, 64, 64])
```

```
arr_slice_2[:] = 2_001
arr_slice_2
```

```
array([2001, 2001, 2001])
```

```
arr
```

```
array([ 0,  1,  2,  3,  4, 64, 64, 64,  8,  9])
```

Indexing with slices

We can slice across two or more dimensions and use the `[i, j]` notation.

```
arr2d = np.array([[1,2,3], [4,5,6], [7,8,9]])
arr2d
```

```
array([[1, 2, 3],
       [4, 5, 6],
       [7, 8, 9]])
```

```
arr2d[:,2]
```

```
array([[1, 2, 3],  
       [4, 5, 6]])
```

```
arr2d[:,2, 1:]
```

```
array([[2, 3],  
       [5, 6]])
```

A colon (`:`) by itself selects the entire dimension and is necessary to slice higher dimensions.

```
arr2d[:, :, 1]
```

```
array([[1],  
       [4],  
       [7]])
```

```
arr2d[:, 2, 1:] = 0  
arr2d
```

```
array([[1, 0, 0],  
       [4, 0, 0],  
       [7, 8, 9]])
```

Always check your output!

Boolean Indexing

We can use Booleans (i.e., `True` and `False`) to slice arrays, too. Boolean indexing in Python is like combining `index()` and `match()` in Excel.

```
names = np.array(['Bob', 'Joe', 'Will', 'Bob', 'Will', 'Joe', 'Joe'])  
np.random.seed(42)  
data = np.random.randn(7, 4)
```

```
names
```

```
array(['Bob', 'Joe', 'Will', 'Bob', 'Will', 'Joe', 'Joe'], dtype='<U4')
```

```
data
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [-0.2342, -0.2341,  1.5792,  0.7674],
       [-0.4695,  0.5426, -0.4634, -0.4657],
       [ 0.242 , -1.9133, -1.7249, -0.5623],
       [-1.0128,  0.3142, -0.908 , -1.4123],
       [ 1.4656, -0.2258,  0.0675, -1.4247],
       [-0.5444,  0.1109, -1.151 ,  0.3757]])
```

Here `names` provides seven names for the seven rows in `data`.

```
names == 'Bob'
```

```
array([ True, False, False,  True, False, False, False])
```

```
data[names == 'Bob']
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [ 0.242 , -1.9133, -1.7249, -0.5623]])
```

We can combine Boolean slicing with : slicing.

```
data[names == 'Bob', 2:]
```

```
array([[ 0.6477,  1.523 ],
       [-1.7249, -0.5623]])
```

We can use `~` to invert a Boolean.

```
cond = names == 'Bob'
data[~cond]
```

```
array([[-0.2342, -0.2341,  1.5792,  0.7674],  
      [-0.4695,  0.5426, -0.4634, -0.4657],  
      [-1.0128,  0.3142, -0.908 , -1.4123],  
      [ 1.4656, -0.2258,  0.0675, -1.4247],  
      [-0.5444,  0.1109, -1.151 ,  0.3757]])
```

For NumPy arrays, we must use `&` and `|` instead of `and` and `or`.

```
cond = (names == 'Bob') | (names == 'Will')  
data[cond]
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],  
      [-0.4695,  0.5426, -0.4634, -0.4657],  
      [ 0.242 , -1.9133, -1.7249, -0.5623],  
      [-1.0128,  0.3142, -0.908 , -1.4123]])
```

We can also create a Boolean for each element.

```
data
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],  
      [-0.2342, -0.2341,  1.5792,  0.7674],  
      [-0.4695,  0.5426, -0.4634, -0.4657],  
      [ 0.242 , -1.9133, -1.7249, -0.5623],  
      [-1.0128,  0.3142, -0.908 , -1.4123],  
      [ 1.4656, -0.2258,  0.0675, -1.4247],  
      [-0.5444,  0.1109, -1.151 ,  0.3757]])
```

```
data < 0
```

```
array([[False,  True, False, False],  
      [ True,  True, False, False],  
      [ True, False,  True,  True],  
      [False,  True,  True,  True],  
      [ True, False,  True,  True],  
      [False,  True, False,  True],  
      [ True, False,  True, False]])
```

```
data[data < 0] = 0
data

array([[0.4967, 0.      , 0.6477, 1.523 ],
       [0.      , 0.      , 1.5792, 0.7674],
       [0.      , 0.5426, 0.      , 0.      ],
       [0.242 , 0.      , 0.      , 0.      ],
       [0.      , 0.3142, 0.      , 0.      ],
       [1.4656, 0.      , 0.0675, 0.      ],
       [0.      , 0.1109, 0.      , 0.3757]])
```

Universal Functions: Fast Element-Wise Array Functions

A universal function, or ufunc, is a function that performs element-wise operations on data in ndarrays. You can think of them as fast vectorized wrappers for simple functions that take one or more scalar values and produce one or more scalar results.

```
arr = np.arange(10)
arr

array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

np.sqrt(arr)

array([0.      , 1.      , 1.4142, 1.7321, 2.      , 2.2361, 2.4495, 2.6458,
       2.8284, 3.      ])
```

Like above, we can raise a single value to a NumPy array of powers.

```
2**arr

array([ 1,  2,  4,  8, 16, 32, 64, 128, 256, 512])
```

`np.exp(x)` is e^x .

```
np.exp(arr)
```

```
array([1.0000e+00, 2.7183e+00, 7.3891e+00, 2.0086e+01, 5.4598e+01,
       1.4841e+02, 4.0343e+02, 1.0966e+03, 2.9810e+03, 8.1031e+03])
```

Table 4-4 from McKinney (2022) summarizes fast, element-wise unary functions:

- **abs, fabs**: Compute the absolute value element-wise for integer, floating-point, or complex values
- **sqrt**: Compute the square root of each element (equivalent to arr ** 0.5)
- **square**: Compute the square of each element (equivalent to arr ** 2)
- **exp**: Compute the exponent e^x of each element
- **log, log10, log2, log1p**: Natural logarithm (base e), log base 10, log base 2, and log(1 + x), respectively
- **sign**: Compute the sign of each element: 1 (positive), 0 (zero), or -1 (negative)
- **ceil**: Compute the ceiling of each element (i.e., the smallest integer greater than or equal to that number)
- **floor**: Compute the floor of each element (i.e., the largest integer less than or equal to each element)
- **rint**: Round elements to the nearest integer, preserving the dtype
- **modf**: Return fractional and integral parts of array as a separate array
- **isnan**: Return boolean array indicating whether each value is NaN (Not a Number)
- **isfinite, isnan**: Return boolean array indicating whether each element is finite (non-inf, non-NaN) or infinite, respectively
- **cos, cosh, sin, sinh, tan, tanh**: Regular and hyperbolic trigonometric functions
- **arccos, arccosh, arcsin, arcsinh, arctan, arctanh**: Inverse trigonometric functions
- **logical_not**: Compute truth value of not x element-wise (equivalent to ~arr).

These “unary” functions operate on one array and return a new array with the same shape. There are also “binary” functions that operate on two arrays and return one array.

```
np.random.seed(42)
x = np.random.randn(8)
y = np.random.randn(8)
```

```
x
```

```
array([ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342, -0.2341,  1.5792,
       0.7674])
```

```
y
```

```
array([-0.4695,  0.5426, -0.4634, -0.4657,  0.242 , -1.9133, -1.7249,
       -0.5623])
```

```
np.maximum(x, y)
```

```
array([ 0.4967,  0.5426,  0.6477,  1.523 ,  0.242 , -0.2341,  1.5792,
       0.7674])
```

Table 4-5 from McKinney (2022) summarizes fast, element-wise binary functions:

- `add`: Add corresponding elements in arrays
- `subtract`: Subtract elements in second array from first array
- `multiply`: Multiply array elements
- `divide`, `floor_divide`: Divide or floor divide (truncating the remainder)
- `power`: Raise elements in first array to powers indicated in second array
- `maximum`, `fmax`: Element-wise maximum; `fmax` ignores NaN
- `minimum`, `fmin`: Element-wise minimum; `fmin` ignores NaN
- `mod`: Element-wise modulus (remainder of division)
- `copysign`: Copy sign of values in second argument to values in first argument
- `greater`, `greater_equal`, `less`, `less_equal`, `equal`, `not_equal`: Perform element-wise comparison, yielding boolean array (equivalent to infix operators `>`, `>=`, `<`, `<=`, `==`, `!=`)
- `logical_and`, `logical_or`, `logical_xor`: Compute element-wise truth value of logical operation (equivalent to infix operators `&`, `|`, `^`)

Array-Oriented Programming with Arrays

Using NumPy arrays enables you to express many kinds of data processing tasks as concise array expressions that might otherwise require writing loops. This practice of replacing explicit loops with array expressions is commonly referred to as vectorization. In general, vectorized array operations will often be one or two (or more) orders of magnitude faster than their pure Python equivalents, with the biggest impact in any kind of numerical computations. Later, in Appendix A, I explain broadcasting, a powerful method for vectorizing computations.

Expressing Conditional Logic as Array Operations

The `numpy.where` function is a vectorized version of the ternary expression `x if condition else y`.

`np.where()` is an if-else statement, like Excel's `if()`.

```
xarr = np.array([1.1, 1.2, 1.3, 1.4, 1.5])
yarr = np.array([2.1, 2.2, 2.3, 2.4, 2.5])
cond = np.array([True, False, True, True, False])
```

```
np.where(cond, xarr, yarr)
```

```
array([1.1, 2.2, 1.3, 1.4, 2.5])
```

We could use a list comprehension instead, but it takes longer to type, read, and troubleshoot.

```
np.array([(x if c else y) for x, y, c in zip(xarr, yarr, cond)])
```

```
array([1.1, 2.2, 1.3, 1.4, 2.5])
```

`np.select()` lets us test more than one condition and has a default value if no condition is met.

```
np.select(
    condlist=[cond==True, cond==False],
    choicelist=[xarr, yarr]
)
```

```
array([1.1, 2.2, 1.3, 1.4, 2.5])
```

Mathematical and Statistical Methods

A set of mathematical functions that compute statistics about an entire array or about the data along an axis are accessible as methods of the array class. You can use aggregations (often called reductions) like `sum`, `mean`, and `std` (standard deviation) either by calling the array instance method or using the top-level NumPy function.

```
np.random.seed(42)
arr = np.random.randn(5, 4)
arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [-0.2342, -0.2341,  1.5792,  0.7674],
       [-0.4695,  0.5426, -0.4634, -0.4657],
       [ 0.242 , -1.9133, -1.7249, -0.5623],
       [-1.0128,  0.3142, -0.908 , -1.4123]])
```

```
arr.mean()
```

```
-0.1713
```

```
arr.sum()
```

```
-3.4260
```

The aggregation methods above aggregated the whole array. We can use the `axis` argument to aggregate columns (`axis=0`) and rows (`axis=1`).

```
arr.mean(axis=1)
```

```
array([ 0.6323,  0.4696, -0.214 , -0.9896, -0.7547])
```

```
arr[0].mean()
```

```
0.6323
```

```
arr[1].mean()
```

```
0.4696
```

```
arr.mean(axis=0)
```

```
array([-0.1956, -0.2858, -0.1739, -0.03 ])
```

```
arr[:, 0].mean()
```

-0.1956

```
arr[:, 1].mean()
```

-0.2858

The `.cumsum()` method returns the sum of all previous elements.

```
arr = np.array([0, 1, 2, 3, 4, 5, 6, 7]) # same output as np.arange(8)  
arr.cumsum()
```

```
array([ 0,  1,  3,  6, 10, 15, 21, 28])
```

We can also use the `.cumsum()` method along the axis of a multi-dimensional array.

```
arr = np.array([[0, 1, 2], [3, 4, 5], [6, 7, 8]])  
arr
```

```
array([[0, 1, 2],  
       [3, 4, 5],  
       [6, 7, 8]])
```

```
arr.cumsum(axis=0)
```

```
array([[ 0,  1,  2],  
       [ 3,  5,  7],  
       [ 9, 12, 15]])
```

```
arr.cumprod(axis=1)
```

```
array([[ 0,    0,    0],  
       [ 3,   12,   60],  
       [ 6,   42,  336]])
```

Table 4-6 from McKinney (2022) summarizes basic statistical methods:

- `sum`: Sum of all the elements in the array or along an axis; zero-length arrays have sum 0
- `mean`: Arithmetic mean; zero-length arrays have NaN mean
- `std, var`: Standard deviation and variance, respectively, with optional degrees of freedom adjustment (default denominator n)
- `min, max`: Minimum and maximum
- `argmin, argmax`: Indices of minimum and maximum elements, respectively
- `cumsum`: Cumulative sum of elements starting from 0
- `cumprod`: Cumulative product of elements starting from 1

McKinney Chapter 4 - Practice - Blank

```
import numpy as np
```

```
%precision 4
```

```
'%.4f'
```

Announcements

Five-Minute Review

Practice

Create a 1-dimensional array `a1` that counts from 0 to 24 by 1.

Create a 1-dimentional array `a2` that counts from 0 to 24 by 3.

Create a 1-dimentional array `a3` that counts from 0 to 100 by multiples of 3 or 5.

Create a 1-dimensional array `a4` that contains the squares of the even integers through 100,000.

Write a function `calc_pv()` that mimic Excel's PV function.

Excel's present value function is: `=PV(rate, nper, pmt, [fv], [type])`

The present value of an annuity payment is: $PV_{pmt} = \frac{pmt}{rate} \times \left(1 - \frac{1}{(1+rate)^{nper}}\right)$

The present value of a lump sum is: $PV_{fv} = \frac{fv}{(1+rate)^{nper}}$

Write a function calc_fv() that mimic Excel's FV function.

Excel's future value function is: =FV(rate, nper, pmt, [pv], [type])

Replace the negative values in data with -1 and positive values with +1.

```
np.random.seed(42)
data = np.random.randn(4, 4)
```

Write a function calc_n() that calculates the number of payments that generate x% of the present value of a perpetuity.

The present value of a growing perpetuity is $PV = \frac{C_1}{r-g}$, and the present value of a growing annuity is $PV = \frac{C_1}{r-g} \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

Write a function that calculates the internal rate of return of a NumPy array of cash flows.

Write a function calc_returns() that accepts NumPy arrays of prices and dividends and returns a NumPy array of returns.

```
prices = np.array([100, 150, 100, 50, 100, 150, 100, 150])
dividends = np.array([1, 1, 1, 1, 2, 2, 2, 2])
```

Rewrite the function calc_returns() as calc_returns_2() so it returns NumPy arrays of returns, capital gains yields, and dividend yields.

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1]

Input: np.array([18.5, 17.0, 18.0, 19.0, 18.0])
Output: np.array([0.75, 0.0, 0.5, 1.0, 0.5])

Write a function `calc_portval()` that accepts a dictionary of prices and share holdings and returns the portfolio value

First convert your dictionary to a NumPy array with one column for prices and another for shares.

```
data = {  
    "AAPL": (150.25, 10), # (price, shares)  
    "GOOGL": (2750.00, 2),  
    "MSFT": (300.75, 5)  
}
```

Write functions `calc_var()` and `calc_std()` that calculate variance and standard deviation.

NumPy's `.var()` and `.std()` methods return *population* statistics (i.e., denominators of n). The pandas equivalents return *sample* statistics (denominators of $n - 1$), which are more appropriate for financial data analysis where we have a sample instead of a population.

Your `calc_var()` and `calc_std()` functions should have a `sample` argument that is `True` by default so both functions return sample statistics by default.

Write a function `calc_ret()` to convert quantitative returns to qualitative returns

Returns within one standard deviation of the mean are “Medium”. Returns less than one standard deviation below the mean are “Low”, and returns greater than one standard deviation above the mean are “High”.

```
np.random.seed(42)  
returns = np.random.randn(100)
```

McKinney Chapter 4 - Practice - Sec 02

```
import numpy as np
```

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep signing up for project teams on [Canvas > People > Projects](#)
2. Claim a team, then ask to increase its limit if you want more than four teammates
3. Keep adding and voting for [students choice topics](#)

Five-Minute Review

NumPy is the foundation of scientific computing in Python! NumPy is also the foundation of pandas, so all these tools in this notebook are relevant to pandas, too!

Here are some random data for us to quickly review NumPy with. The `np.random.seed()` function makes our random number draws repeatable.

```
np.random.seed(42)
my_arr = np.random.randn(3, 5)
my_arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Slicing works the same as with lists and tuples. As well, we can replace chained slices, like `[i][j][k]`, with a Matlab notation, like `[i. j. k]`

```
my_arr[2][0]
```

-0.4634

```
my_arr[2, 0]
```

-0.4634

What if we want the *first two rows*?

```
my_arr[:2]
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426]])
```

What if we want the *first two columns*?

```
my_arr[:, :2]
```

```
array([[ 0.4967, -0.1383],
       [-0.2341,  1.5792],
       [-0.4634, -0.4657]])
```

What about linear algebra in NumPy?

```
A = np.array([[1, 2, 3], [4, 5, 6]])
B = np.array([[1, 2, 3], [4, 5, 6]])
print(f'A is:{\n {A}\n B is:{\n {B}}')
```

```
A is:
[[1 2 3]
 [4 5 6]]
B is:
[[1 2 3]
 [4 5 6]]
```

We can use the `.dot()` method to chain matrix multiplication.

```
A.dot(B.T)
```

```
array([[14, 32],  
       [32, 77]])
```

The `@` operator is matrix multiplication with NumPy arrays.

```
A @ B.T
```

```
array([[14, 32],  
       [32, 77]])
```

The `*` symbol is *elementwise* multiplication.

```
A * B
```

```
array([[ 1,  4,  9],  
       [16, 25, 36]])
```

The `np.where()` and `np.select()` functions are the NumPy equivalents of Excel's `if()`.

Say we want to replace values in `my_arr` greater than 0.5 with 0.5.

```
my_arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],  
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],  
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

```
np.where(my_arr > 0.5, 0.5, my_arr)
```

```
array([[ 0.4967, -0.1383,  0.5    ,  0.5    , -0.2342],  
       [-0.2341,  0.5    ,  0.5    , -0.4695,  0.5    ],  
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Or, we could use `np.minimum()` instead.

```
np.minimum(0.5, my_arr)
```

```
array([[ 0.4967, -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Now, say we want to replace values in `my_arr` greater than 0 with 0 and greater than 0.5 with 0.5.

```
np.where(
    my_arr>0.5, # first condition is most restrictive
    0.5, # if first condition True
    np.where( # otherwise
        my_arr>0, # second condition is less restrictive
        0, # if second condition True
        my_arr # otherwise
    )
)
```

```
array([[ 0.   , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.   , -1.9133, -1.7249]])
```

We have a better way to test more than one condition! `np.select()`!

```
np.select(
    condlist=[my_arr>0.5, my_arr>0],
    choicelist=[0.5, 0],
    default=my_arr
)

array([[ 0.   , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.   , -1.9133, -1.7249]])
```

Practice

Create a 1-dimensional array a1 that counts from 0 to 24 by 1.

```
np.array(range(25))
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

```
np.array(range(25))
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

```
a1 = np.arange(25)
a1
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

Create a 1-dimentional array a2 that counts from 0 to 24 by 3.

```
a2 = np.arange(0, 25, 3)
a2
```

```
array([ 0,  3,  6,  9, 12, 15, 18, 21, 24])
```

Create a 1-dimentional array a3 that counts from 0 to 100 by multiples of 3 or 5.

```
a3 = np.array([i for i in range(101) if (i%3==0) | (i%5==0)])
a3
```

```
array([ 0,  3,  5,  6,  9, 10, 12, 15, 18, 20, 21, 24, 25,
       27, 30, 33, 35, 36, 39, 40, 42, 45, 48, 50, 51, 54,
       55, 57, 60, 63, 65, 66, 69, 70, 72, 75, 78, 80, 81,
       84, 85, 87, 90, 93, 95, 96, 99, 100])
```

```
toy = np.arange(5)
toy[toy%2==0]
```

```
array([0, 2, 4])
```

```
my_arr_0_100 = np.arange(101)
a3_alt = my_arr_0_100[(my_arr_0_100%3==0) | (my_arr_0_100%5==0)]
a3_alt
```

```
array([ 0,   3,   5,   6,   9,  10,  12,  15,  18,  20,  21,  24,  25,
       27,  30,  33,  35,  36,  39,  40,  42,  45,  48,  50,  51,  54,
       55,  57,  60,  63,  65,  66,  69,  70,  72,  75,  78,  80,  81,
       84,  85,  87,  90,  93,  95,  96,  99, 100])
```

```
(a3 == a3_alt).all()
```

```
np.True_
```

```
np.allclose(a3, a3_alt)
```

```
True
```

Create a 1-dimensional array a4 that contains the squares of the even integers through 100,000.

```
%timeit np.arange(0, 100_001, 2) ** 2
a4 = np.arange(0, 100_001, 2) ** 2
a4
```

```
46.7 s ± 7.59 s per loop (mean ± std. dev. of 7 runs, 10,000 loops each)
```

```
array([ 0,          4,          16, ..., 9999200016,
       9999600004, 10000000000], shape=(50001,))
```

```
%timeit a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt
```

29.7 ms ± 2.85 ms per loop (mean ± std. dev. of 7 runs, 100 loops each)

```
array([          0,          4,         16, ..., 9999200016,
       9999600004, 10000000000], shape=(50001,))
```

```
(a4 == a4_alt).all()
```

np.True_

```
np.allclose(a4, a4_alt*1.0000001)
```

True

Write a function calc_pv() that mimic Excel's PV function.

Excel's present value function is: =PV(rate, nper, pmt, [fv], [type])

The present value of an annuity payment is: $PV_{pmt} = \frac{pmt}{rate} \times \left(1 - \frac{1}{(1+rate)^{nper}}\right)$

The present value of a lump sum is: $PV_{fv} = \frac{fv}{(1+rate)^{nper}}$

```
def calc_pv(rate, nper, pmt=None, fv=None, type='END'):
    # or we could set the pmt and fv defaults to 0
    if pmt is None:
        pmt = 0
    if fv is None:
        fv = 0
    if type not in ['BGN', 'END']:
        raise ValueError(f'type must be BGN or END; you provided {type}')

    pv_pmt = (pmt / rate) * (1 - (1 + rate) ** (-nper))
    pv_fv = fv * (1 + rate) ** (-nper)
    pv = pv_pmt + pv_fv

    if type == 'BGN':
```

```
pv *= 1 + rate  
  
return -pv  
  
calc_pv(rate=0.1, nper=10, pmt=100, fv=1_000)  
  
-1000.0000  
  
calc_pv(rate=0.1, nper=10, pmt=-100, fv=-1_000)  
  
1000.0000
```

Write a function calc_fv() that mimic Excel's FV function.

Excel's future value function is: =FV(rate, nper, pmt, [pv], [type])

```
def calc_fv(rate, nper, pmt=None, pv=None, type='END'):  
    if pmt is None:  
        pmt = 0  
    if pv is None:  
        pv = 0  
    if type not in ['BGN', 'END']:  
        raise ValueError('type must be BGN or END')  
  
    fv_pmt = (pmt / rate) * ((1 + rate)**nper - 1)  
    fv_pv = pv * (1 + rate)**nper  
    fv = fv_pmt + fv_pv  
  
    if type == 'BGN':  
        fv *= 1 + rate  
  
    return -fv
```

```
calc_fv(rate=0.1, nper=10, pmt=100, pv=-1_000)
```

1000.0000

The rule of 72!

```
calc_fv(rate=0.072, nper=10, pmt=0, pv=-1_000)
```

2004.2314

Replace the negative values in data with -1 and positive values with +1.

```
np.random.seed(42)
data = np.random.randn(4, 4)

data

array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [-0.2342, -0.2341,  1.5792,  0.7674],
       [-0.4695,  0.5426, -0.4634, -0.4657],
       [ 0.242 , -1.9133, -1.7249, -0.5623]])
```

We have at least three good solutions!

1. Slicing and broadcastng
2. np.where()
3. np.select()

First, here is the slicing and broadcasting solution.

```
data[data < 0] = -1
data[data > 0] = +1

data # here "1." and "-1." indicate that these values are floats
```

```
array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Second, here is the np.where() solution. NumPy's np.where() function has the same logic as Excel's if() function! I will recreate data so we start from the same point.

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.where(data < 0, -1, np.where(data > 0, +1, data))

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

The nested `np.where()` function calls can get confusing! An option is to insert white space!

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.where(
    data < 0, # condition
    -1, # result if True
    np.where(data > 0, +1, data) # result if False
)

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Third, here is the `np.select()` solution. NumPy's `np.select()` function lets us test *many* conditions! I will recreate `data` so we start from the same point.

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.select(
    condlist=[data<0, data>0],
    choicelist=[-1, +1],
    default=data
)

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Write a function `calc_n()` that calculates the number of payments that generate $x\%$ of the present value of a perpetuity.

The present value of a growing perpetuity is $PV = \frac{C_1}{r-g}$, and the present value of a growing annuity is $PV = \frac{C_1}{r-g} \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

Your `npmts()` should accept arguments c_1 , r , and g that represent C_1 , r , and g . The present value of a growing perpetuity is $PV = \frac{C_1}{r-g}$, and the present value of a growing annuity is $PV = \frac{C_1}{r-g} \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

We can use the growing annuity and perpetuity formulas to show: $x = \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

Then: $1 - x = \left(\frac{1+g}{1+r} \right)^t$.

Finally: $t = \frac{\log(1-x)}{\log\left(\frac{1+g}{1+r}\right)}$

We do not need to accept an argument c_1 because C_1 cancels out!

```
def npmts(x, r, g):
    return np.log(1-x) / np.log((1 + g) / (1 + r))
```

```
npmts(0.5, 0.1, 0.05)
```

14.9000

Write a function that calculates the internal rate of return of a NumPy array of cash flows.

Let us pick a set of cash flows c where we know the internal rate of return! For the following c , the IRR is 10%.

```
c = np.array([-100, +110])
r = 0.1
```

First we need an function to calculate net present value (NPV) from c and r ! The `npv()` function below uses NumPy arrays to calculate NPV as:

$$NPV = \sum_{t=0}^T \frac{c_t}{(1+r)^t}$$

```
def calc_npv(r, c):
    t = np.arange(len(c))
    return np.sum(c / (1 + r)**t)
```

```
calc_npv(r=r, c=c)
```

-0.0000

We can use a `for` loop to guess IRR values until we find an NPV close to zero (or reach the maximum number of iterations in `max_iter`). We can use the [Newton-Rapshon method](#) to make smarter guesses. If we have function $f(x)$ and guess x_t , our next guess should be $x_{t+1} = x_t - \frac{f(x_t)}{f'(x_t)}$. Here our $f(x)$ is $NPV(r)$, and we can approximate $f'(x_t)$ as $\frac{NPV(r+step) - NPV(r)}{step}$. We will make guess until $|NPV| < tol$.

```
def calc_irr(c, guess=0, step=1e-6, tol=1e-6, max_iter=1_000, verbose=False):
    irr = guess
    for i in range(max_iter):
        npv = calc_npv(r=irr, c=c)
        if abs(npv) < tol:
            if verbose:
                print(f'IRR: {irr:0.4f}\nNPV: {npv:0.4f}\nIterations: {i+1:d}')
            return irr

        deriv = (calc_npv(r=irr+step, c=c) - npv) / step
        irr -= npv / deriv

    raise ValueError(f'NPV did not converge to zero after {i+1} iterations.')
```

```
calc_irr(c=c, verbose=True)
```

IRR: 0.1000
NPV: 0.0000
Iterations: 4

0.1000

```
calc_irr(c=np.array([-100, 10, 10, 10, 10, 110]), verbose=True)
```

```
IRR: 0.1000  
NPV: 0.0000  
Iterations: 5
```

0.1000

```
# calc_irr(c=np.array([-2000, 1000, -500]))
```

Write a function `calc_returns()` that accepts NumPy arrays of prices and dividends and returns a NumPy array of returns.

```
prices = np.array([100, 150, 100, 50, 100, 150, 100, 150])  
dividends = np.array([1, 1, 1, 1, 2, 2, 2, 2])
```

We want to slice our arrays to “lag” or “shift” them! For example, we slice the `prices` array to calculate capital gains as follows.

```
prices[1:] - prices[:-1]
```

```
array([ 50, -50, -50,  50,  50, -50,  50])
```

```
def calc_returns(p, d):  
    return (p[1:] - p[:-1] + d[1:]) / p[:-1]
```

```
calc_returns(p=prices, d=dividends)
```

```
array([ 0.51 , -0.3267, -0.49 ,  1.04 ,  0.52 , -0.32 ,  0.52 ])
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns NumPy arrays of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):  
    cg = p[1:] / p[:-1] - 1  
    dp = d[1:] / p[:-1]  
    r = cg + dp  
    return {'r': r, 'cg': cg, 'dp': dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': array([ 0.51 , -0.3267, -0.49 ,  1.04 ,  0.52 , -0.32 ,  0.52 ]),
'cg': array([ 0.5 , -0.3333, -0.5 ,  1. ,  0.5 , -0.3333,  0.5 ]),
'dp': array([0.01 ,  0.0067,  0.01 ,  0.04 ,  0.02 ,  0.0133,  0.02 ])}
```

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1]

Input: np.array([18.5, 17.0, 18.0, 19.0, 18.0])

Output: np.array([0.75, 0.0, 0.5, 1.0, 0.5])

```
x = np.array([18.5, 17.0, 18.0, 19.0, 18.0])
x
```

array([18.5, 17. , 18. , 19. , 18.])

```
def rescale(x):
    return (x - x.min()) / (x.max() - x.min())
```

```
rescale(x=x)
```

array([0.75, 0. , 0.5 , 1. , 0.5])

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

First convert your dictionary to a NumPy array with one column for prices and another for shares.

```
data = {
    "AAPL": (150.25, 10), # (price, shares)
    "GOOGL": (2750.00, 2),
    "MSFT": (300.75, 5)
}
```

```
def calc_portval(data):
    x = np.array(list(data.values()))
    return x.prod(axis=1).sum()
```

```
calc_portval(data=data)
```

8506.2500

Write functions `calc_var()` and `calc_std()` that calculate variance and standard deviation.

NumPy's `.var()` and `.std()` methods return *population* statistics (i.e., denominators of n). The pandas equivalents return *sample* statistics (denominators of $n - 1$), which are more appropriate for financial data analysis where we have a sample instead of a population.

Your `calc_var()` and `calc_std()` functions should have a `sample` argument that is `True` by default so both functions return sample statistics by default.

We can use the methods `.mean()` and `.sum()` instead of the functions `np.mean()` and `np.sum()`. I find this format easier to read because it puts the data first.

```
def calc_var(x, sample=True):
    sq_err = (x - x.mean()) ** 2
    den = len(x)
    if sample:
        den -= 1
    return sq_err.sum() / den
```

We can re-use `calc_var()` in our `calc_std()` function.

```
def calc_std(x, sample=True):
    return calc_var(x=x, sample=sample) ** 0.5
```

```
np.random.seed(42)
arr = np.random.randn(1_000_000)
arr
```

```
array([ 0.4967, -0.1383,  0.6477, ..., -0.113 ,  1.4691,  0.4764],
      shape=(1000000,))
```

```
calc_var(arr)
```

1.0004

```
calc_std(arr)
```

```
1.0002
```

```
arr.var(ddof=1) == calc_var(arr)
```

```
np.True_
```

```
arr.std(ddof=1) == calc_std(arr)
```

```
np.True_
```

Write a function calc_ret() to convert quantitative returns to qualitative returns

Returns within one standard deviation of the mean are “Medium”. Returns less than one standard deviation below the mean are “Low”, and returns greater than one standard deviation above the mean are “High”.

```
np.random.seed(42)
returns = np.random.randn(10)
```

```
def calc_ret(r):
    mu = r.mean()
    sigma = r.std(ddof=1)
    return np.select(
        condlist=[
            r < (mu - sigma),
            r <= (mu + sigma),
            r > (mu + sigma)
        ],
        choicelist=[
            'Low',
            'Medium',
            'High'
        ],
        default=''
    )
```

```
calc_ret(returns)
```

```
array(['Medium', 'Medium', 'Medium', 'High', 'Medium', 'Medium', 'High',
       'Medium', 'Low', 'Medium'], dtype='<U6')
```

McKinney Chapter 4 - Practice - Sec 03

```
import numpy as np
```

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep signing up for project teams on [Canvas > People > Projects](#)
2. Claim a team, then ask to increase its limit if you want more than four teammates
3. Keep adding and voting for [students choice topics](#)

Five-Minute Review

NumPy is the foundation of scientific computing in Python! NumPy is also the foundation of pandas, so all these tools in this notebook are relevant to pandas, too!

Here are some random data for us to quickly review NumPy with. The `np.random.seed()` function makes our random number draws repeatable.

```
np.random.seed(42)
my_arr = np.random.randn(3, 5)
my_arr
```



```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Slicing works the same as with lists and tuples. As well, we can replace chained slices, like `[i][j][k]`, with a Matlab notation, like `[i. j. k]`

```
my_arr[2][4]
```

-1.7249

```
my_arr[2, 4]
```

-1.7249

What if we want the *first two rows*?

```
my_arr[:2]
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426]])
```

What if we want the *first two columns*?

```
my_arr[:, :2]
```

```
array([[ 0.4967, -0.1383],
       [-0.2341,  1.5792],
       [-0.4634, -0.4657]])
```

What about linear algebra in NumPy?

```
A = np.array([[1, 2, 3], [4, 5, 6]])
A
```

```
array([[1, 2, 3],
       [4, 5, 6]])
```

```
B = np.array([[1, 2, 3], [4, 5, 6]])
B
```

```
array([[1, 2, 3],  
       [4, 5, 6]])
```

We can use the `.dot()` method to chain matrix multiplication.

```
A.dot(B.T)
```

```
array([[14, 32],  
       [32, 77]])
```

The `@` operator is matrix multiplication with NumPy arrays.

```
A @ B.T
```

```
array([[14, 32],  
       [32, 77]])
```

The `*` symbol is *elementwise* multiplication.

```
A * B
```

```
array([[ 1,  4,  9],  
       [16, 25, 36]])
```

There is *another* way, too! The `np.matmul()` function is similar to the `.dot()` method!

```
np.matmul(A, B.T)
```

```
array([[14, 32],  
       [32, 77]])
```

The `np.where()` and `np.select()` functions are the NumPy equivalents of Excel's `if()`.

Say we want to replace values in `my_arr` greater than 0.5 with 0.5.

```
my_arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],  
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],  
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

```
np.where(my_arr>0.5, 0.5, my_arr)

array([[ 0.4967, -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Or, we could use `np.minimum()` instead.

```
np.minimum(0.5, my_arr)

array([[ 0.4967, -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Now, say we want to replace values in `my_arr` greater than 0 with 0 and greater than 0.5 with 0.5.

```
np.where(
    my_arr>0.5, # first condition is most restrictive
    0.5, # if first condition True
    np.where( # otherwise
        my_arr>0, # second condition is less restrictive
        0, # if second condition True
        my_arr # otherwise
    )
)

array([[ 0.      , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.      , -1.9133, -1.7249]])
```

We have a better way to test more than one condition! `np.select()`!

```
np.select(
    condlist=[my_arr>0.5, my_arr>0],
    choicelist=[0.5, 0],
    default=my_arr
)

array([[ 0.      , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.      , -1.9133, -1.7249]])
```

Practice

Create a 1-dimensional array a1 that counts from 0 to 24 by 1.

```
np.array(range(25))
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

```
a1 = np.arange(25)
a1
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

Create a 1-dimentional array a2 that counts from 0 to 24 by 3.

```
a2 = np.arange(0, 25, 3)
a2
```

```
array([ 0,  3,  6,  9, 12, 15, 18, 21, 24])
```

Create a 1-dimentional array a3 that counts from 0 to 100 by multiples of 3 or 5.

```
a3 = np.array([i for i in range(101) if (i%3==0) | (i%5==0)])
a3
```

```
array([ 0,  3,  5,  6,  9, 10, 12, 15, 18, 20, 21, 24, 25,
       27, 30, 33, 35, 36, 39, 40, 42, 45, 48, 50, 51, 54,
       55, 57, 60, 63, 65, 66, 69, 70, 72, 75, 78, 80, 81,
       84, 85, 87, 90, 93, 95, 96, 99, 100])
```

```
zero_to_100 = np.arange(101)
a3_alt = zero_to_100[(zero_to_100 % 3 == 0) | (zero_to_100 % 5 == 0)]
a3_alt
```

```
array([ 0,  3,  5,  6,  9, 10, 12, 15, 18, 20, 21, 24, 25,
       27, 30, 33, 35, 36, 39, 40, 42, 45, 48, 50, 51, 54,
       55, 57, 60, 63, 65, 66, 69, 70, 72, 75, 78, 80, 81,
       84, 85, 87, 90, 93, 95, 96, 99, 100])
```

```
(a3 == a3_alt).all()
```

```
np.True_
```

```
np.allclose(a3, a3_alt)
```

```
True
```

Create a 1-dimensional array a4 that contains the squares of the even integers through 100,000.

```
%timeit np.arange(0, 100_001, 2) ** 2
a4 = np.arange(0, 100_001, 2) ** 2
a4
```

```
43.3 s ± 10.3 s per loop (mean ± std. dev. of 7 runs, 10,000 loops each)
```

```
array([ 0,  4, 16, ..., 9999200016,
       9999600004, 10000000000], shape=(50001,))
```

```
%timeit a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt
```

```
29.5 ms ± 2.69 ms per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

```
array([ 0,  4, 16, ..., 9999200016,
       9999600004, 10000000000], shape=(50001,))
```

```
(a4 == a4_alt).all()
```

```
np.True_
```

```
np.allclose(a4, a4_alt*1.0000001)
```

True

Write a function calc_pv() that mimic Excel's PV function.

Excel's present value function is: =PV(rate, nper, pmt, [fv], [type])

The present value of an annuity payment is: $PV_{pmt} = \frac{pmt}{rate} \times \left(1 - \frac{1}{(1+rate)^{nper}}\right)$

The present value of a lump sum is: $PV_{fv} = \frac{fv}{(1+rate)^{nper}}$

```
def calc_pv(rate, nper, pmt=None, fv=None, type=0):
    # or we could set the pmt and fv defaults to 0
    if pmt is None:
        pmt = 0
    if fv is None:
        fv = 0
    if type not in [0, 1]:
        raise ValueError(f'type must be 0 or 1; you provided {type}')

    pv_pmt = (pmt / rate) * (1 - (1 + rate)**(-nper))
    pv_fv = fv * (1 + rate)**(-nper)
    pv = pv_pmt + pv_fv

    if type == 1:
        pv *= 1 + rate

    return -pv
```

```
calc_pv(rate=0.1, nper=10, pmt=10, fv=100)
```

-100.0000

```
calc_pv(rate=0.1, nper=10, pmt=10, fv=100, type=1)
```

-110.0000

Write a function calc_fv() that mimic Excel's FV function.

Excel's future value function is: =FV(rate, nper, pmt, [pv], [type])

```
def calc_fv(rate, nper, pmt=None, pv=None, type='END'):
    if pmt is None:
        pmt = 0
    if pv is None:
        pv = 0
    if type not in ['BGN', 'END']:
        raise ValueError('type must be BGN or END')

    fv_pmt = (pmt / rate) * ((1 + rate)**nper - 1)
    fv_pv = pv * (1 + rate)**nper
    fv = fv_pmt + fv_pv

    if type == 'BGN':
        fv *= 1 + rate

    return -fv
```

```
calc_fv(rate=0.1, nper=10, pmt=10, pv=-100)
```

100.0000

The rule of 72!

```
calc_fv(rate=0.072, nper=10, pmt=0, pv=-100)
```

200.4231

Replace the negative values in data with -1 and positive values with +1.

```
np.random.seed(42)
data = np.random.randn(4, 4)

data
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [-0.2342, -0.2341,  1.5792,  0.7674],
       [-0.4695,  0.5426, -0.4634, -0.4657],
       [ 0.242 , -1.9133, -1.7249, -0.5623]])
```

We have at least three good solutions!

1. Slicing and broadcasting
2. `np.where()`
3. `np.select()`

First, here is the slicing and broadcasting solution.

```
data[data < 0] = -1
data[data > 0] = +1

data # here "1." and "-1." indicate that these values are floats
```

```
array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Second, here is the `np.where()` solution. NumPy's `np.where()` function has the same logic as Excel's `if()` function! I will recreate `data` so we start from the same point.

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.where(data < 0, -1, np.where(data > 0, +1, data))
```

```
array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

The nested `np.where()` function calls can get confusing! An option is to insert white space!

```

np.random.seed(42)
data = np.random.randn(4, 4)

np.where(
    data < 0, # condition
    -1, # result if True
    np.where(data > 0, +1, data) # result if False
)

```

array([[1., -1., 1., 1.],
 [-1., -1., 1., 1.],
 [-1., 1., -1., -1.],
 [1., -1., -1., -1.]])

Third, here is the `np.select()` solution. NumPy's `np.select()` function lets us test *many* conditions! I will recreate `data` so we start from the same point.

```

np.random.seed(42)
data = np.random.randn(4, 4)

np.select(
    condlist=[data<0, data>0],
    choicelist=[-1, +1],
    default=data
)

```

array([[1., -1., 1., 1.],
 [-1., -1., 1., 1.],
 [-1., 1., -1., -1.],
 [1., -1., -1., -1.]])

Write a function `calc_n()` that calculates the number of payments that generate x% of the present value of a perpetuity.

The present value of a growing perpetuity is $PV = \frac{C_1}{r-g}$, and the present value of a growing annuity is $PV = \frac{C_1}{r-g} \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

We can use the growing annuity and perpetuity formulas to show: $x = \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

Then: $1 - x = \left(\frac{1+g}{1+r} \right)^t$.

Finally: $t = \frac{\log(1-x)}{\log\left(\frac{1+g}{1+r}\right)}$

We do not need to accept an argument c_1 because C_1 cancels out!

```
def npmts(x, r, g):
    return np.log(1-x) / np.log((1 + g) / (1 + r))

npmts(0.5, 0.1, 0.05)
```

14.9000

Write a function that calculates the internal rate of return of a NumPy array of cash flows.

Let us pick a set of cash flows c where we know the internal rate of return! For the following c , the IRR is 10%.

```
c = np.array([-100, 110])
r = 0.1
```

First we need a function to calculate net present value (NPV) from c and r ! The `npv()` function below uses NumPy arrays to calculate NPV as:

$$NPV = \sum_{t=0}^T \frac{c_t}{(1+r)^t}$$

```
def calc_npv(r, c):
    t = np.arange(len(c))
    return np.sum(c / (1 + r)**t)
```

```
calc_npv(r=r, c=c)
```

-0.0000

```
def calc_irr(c, guess=0, step=1e-6, tol=1e-6, max_iter=1_000, verbose=False):
    irr = guess
    for i in range(max_iter):
        npv = calc_npv(r=irr, c=c)
        if abs(npv) < tol:
```

```
if verbose:  
    print(f'IRR: {irr:0.4f}\nNPV: {npv:0.4f}\nIterations: {i+1:d}')  
return irr  
deriv = (calc_npv(r=irr+step, c=c) - npv) / step  
irr -= npv / deriv  
  
raise ValueError(f'NPV did not converge to zero after {i+1:d} iterations')  
  
calc_irr(c=c, verbose=True)
```

```
IRR: 0.1000  
NPV: 0.0000  
Iterations: 4
```

```
0.1000
```

```
calc_irr(c=np.array([-100, 10, 10, 10, 10, 110]), verbose=True)
```

```
IRR: 0.1000  
NPV: 0.0000  
Iterations: 5
```

```
0.1000
```

```
# calc_irr(c=np.array([-2000, 1000, -500]))
```

Write a function `calc_returns()` that accepts NumPy arrays of prices and dividends and returns a NumPy array of returns.

```
prices = np.array([100, 150, 100, 50, 100, 150, 100, 150])  
dividends = np.array([1, 1, 1, 2, 2, 2, 2])
```

We want to slice our arrays to “lag” or “shift” them! For example, we slice the `prices` array to calculate capital gains as follows.

```
prices[1:] - prices[:-1]
```

```
array([ 50, -50, -50,  50,  50, -50,  50])
```

```
def calc_returns(p, d):
    return (p[1:] - p[:-1] + d[1:]) / p[:-1]
```

```
calc_returns(p=prices, d=dividends)
```

```
array([ 0.51 , -0.3267, -0.49 ,  1.04 ,  0.52 , -0.32 ,  0.52 ])
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns *NumPy arrays* of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):
    cg = p[1:] / p[:-1] - 1
    dp = d[1:] / p[:-1]
    r = cg + dp
    return {'r': r, 'cg': cg, 'dp': dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': array([ 0.51 , -0.3267, -0.49 ,  1.04 ,  0.52 , -0.32 ,  0.52 ]),
 'cg': array([ 0.5   , -0.3333, -0.5   ,  1.     ,  0.5   , -0.3333,  0.5   ]),
 'dp': array([0.01 ,  0.0067,  0.01 ,  0.04 ,  0.02 ,  0.0133,  0.02 ])}
```

Write a function `rescale()` to rescale and shift numbers so that they cover the range [0, 1]

Input: `np.array([18.5, 17.0, 18.0, 19.0, 18.0])`

Output: `np.array([0.75, 0.0, 0.5, 1.0, 0.5])`

```
x = np.array([18.5, 17.0, 18.0, 19.0, 18.0])
x
```

```
array([18.5, 17. , 18. , 19. , 18. ])
```

```
def rescale(x):
    return (x - x.min()) / (x.max() - x.min())
```

```
rescale(x=x)
```

```
array([0.75, 0. , 0.5 , 1. , 0.5 ])
```

Write a function `calc_portval()` that accepts a dictionary of prices and share holdings and returns the portfolio value

First convert your dictionary to a NumPy array with one column for prices and another for shares.

```
data = {
    "AAPL": (150.25, 10), # (price, shares)
    "GOOGL": (2750.00, 2),
    "MSFT": (300.75, 5)
}
```

```
def calc_portval(data):
    x = np.array(list(data.values()))
    return x.prod(axis=1).sum()
```

```
calc_portval(data=data)
```

```
8506.2500
```

Write functions `calc_var()` and `calc_std()` that calculate variance and standard deviation.

NumPy's `.var()` and `.std()` methods return *population* statistics (i.e., denominators of n). The pandas equivalents return *sample* statistics (denominators of $n - 1$), which are more appropriate for financial data analysis where we have a sample instead of a population.

Your `calc_var()` and `calc_std()` functions should have a `sample` argument that is `True` by default so both functions return sample statistics by default.

```
def calc_var(x, sample=True):
    sq_err = (x - x.mean()) ** 2
    den = len(x)
    if sample:
        den -= 1
    return sq_err.sum() / den
```

We can re-use `calc_var()` in our `calc_std()` function.

```
def calc_std(x, sample=True):
    return calc_var(x=x, sample=sample) ** 0.5
```

```
np.random.seed(42)
arr = np.random.randn(1_000_000)
arr
```

```
array([ 0.4967, -0.1383,  0.6477, ..., -0.113 ,  1.4691,  0.4764],
      shape=(1000000,))
```

```
calc_var(arr)
```

1.0004

```
calc_std(arr)
```

1.0002

```
arr.var(ddof=1) == calc_var(arr)
```

```
np.True_
```

```
arr.std(ddof=1) == calc_std(arr)
```

```
np.True_
```

Write a function calc_ret() to convert quantitative returns to qualitative returns

Returns within one standard deviation of the mean are “Medium”. Returns less than one standard deviation below the mean are “Low”, and returns greater than one standard deviation above the mean are “High”.

```
np.random.seed(42)
returns = np.random.randn(10)
```

```
def calc_ret(r):
    mu = r.mean()
    sigma = r.std(ddof=1)
    return np.select(
        condlist=[
            r<(mu-sigma),
            r<=(mu+sigma),
            r>(mu+sigma)
        ],
        choicelist=[
            'Low',
            'Medium',
            'High'
        ],
        default=''
    )
```

```
calc_ret(returns)
```

```
array(['Medium', 'Medium', 'Medium', 'High', 'Medium', 'Medium', 'High',
       'Medium', 'Low', 'Medium'], dtype='<U6')
```

McKinney Chapter 4 - Practice - Sec 04

```
import numpy as np
```

```
%precision 4
```

```
'%.4f'
```

Announcements

1. Keep signing up for project teams on [Canvas > People > Projects](#)
2. Claim a team, then ask to increase its limit if you want more than four teammates
3. Keep adding and voting for [students choice topics](#)

Five-Minute Review

NumPy is the foundation of scientific computing in Python! NumPy is also the foundation of pandas, so all these tools in this notebook are relevant to pandas, too!

Here are some random data for us to quickly review NumPy with. The `np.random.seed()` function makes our random number draws repeatable.

```
np.random.seed(42)
my_arr = np.random.randn(3, 5)
my_arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Slicing works the same as with lists and tuples. As well, we can replace chained slices, like `[i][j][k]`, with a Matlab notation, like `[i. j. k]`

```
my_arr[2][3]
```

-1.9133

```
my_arr[2, 3]
```

-1.9133

What if we want the *first two rows*?

```
my_arr[:2]
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426]])
```

What if we want the *first two columns*?

```
my_arr[:, :2]
```

```
array([[ 0.4967, -0.1383],
       [-0.2341,  1.5792],
       [-0.4634, -0.4657]])
```

What about linear algebra in NumPy?

```
A = np.array([[1,2,3], [4,5,6]])
A
```

```
array([[1, 2, 3],
       [4, 5, 6]])
```

```
B = np.array([[1,2,3], [4,5,6]])
B
```

```
array([[1, 2, 3],  
       [4, 5, 6]])
```

We can use the `.dot()` method to chain matrix multiplication.

```
A.dot(B.T)
```

```
array([[14, 32],  
       [32, 77]])
```

The `@` operator is matrix multiplication with NumPy arrays.

```
A @ B.T
```

```
array([[14, 32],  
       [32, 77]])
```

The `*` operator is elementwise multiplication with NumPy arrays.

```
A * B
```

```
array([[ 1,  4,  9],  
       [16, 25, 36]])
```

The `np.where()` and `np.select()` functions are the NumPy equivalents of Excel's `if()`.

Say we want to replace values in `my_arr` greater than 0.5 with 0.5.

```
my_arr
```

```
array([[ 0.4967, -0.1383,  0.6477,  1.523 , -0.2342],  
       [-0.2341,  1.5792,  0.7674, -0.4695,  0.5426],  
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

```
np.where(my_arr>0.5, 0.5, my_arr)
```

```
array([[ 0.4967, -0.1383,  0.5   ,  0.5   , -0.2342],  
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],  
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Or, we could use `np.minimum()` instead.

```
np.minimum(my_arr, 0.5)
```

```
array([[ 0.4967, -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.242 , -1.9133, -1.7249]])
```

Now, say we want to replace values in `my_arr` greater than 0 with 0 and greater than 0.5 with 0.5.

```
np.where(
    my_arr>0.5, # first condition is most restrictive
    0.5, # if first condition True
    np.where( # otherwise
        my_arr>0, # second condition is less restrictive
        0, # if second condition True
        my_arr # otherwise
    )
)
```

```
array([[ 0.      , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.      , -1.9133, -1.7249]])
```

We have a better way to test more than one condition! `np.select()`!

```
np.select(
    condlist=[my_arr>0.5, my_arr>0],
    choicelist=[0.5, 0],
    default=my_arr
)

array([[ 0.      , -0.1383,  0.5   ,  0.5   , -0.2342],
       [-0.2341,  0.5   ,  0.5   , -0.4695,  0.5   ],
       [-0.4634, -0.4657,  0.      , -1.9133, -1.7249]])
```

Practice

Create a 1-dimensional array a1 that counts from 0 to 24 by 1.

```
a1 = np.arange(25)
a1
```



```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24])
```

Create a 1-dimentional array a2 that counts from 0 to 24 by 3.

```
a2 = np.arange(0, 25, 3)
a2
```



```
array([ 0,  3,  6,  9, 12, 15, 18, 21, 24])
```

Create a 1-dimentional array a3 that counts from 0 to 100 by multiples of 3 or 5.

```
a3 = np.array([i for i in range(101) if (i%3==0) | (i%5==0)])
a3
```



```
array([ 0,  3,  5,  6,  9, 10, 12, 15, 18, 20, 21, 24, 25,
       27, 30, 33, 35, 36, 39, 40, 42, 45, 48, 50, 51, 54,
       55, 57, 60, 63, 65, 66, 69, 70, 72, 75, 78, 80, 81,
       84, 85, 87, 90, 93, 95, 96, 99, 100])
```

```
zero_to_100 = np.arange(101)
a3_alt = zero_to_100[(zero_to_100%3==0) | (zero_to_100%5==0)]
a3_alt
```

```
array([ 0,  3,  5,  6,  9, 10, 12, 15, 18, 20, 21, 24, 25,
       27, 30, 33, 35, 36, 39, 40, 42, 45, 48, 50, 51, 54,
       55, 57, 60, 63, 65, 66, 69, 70, 72, 75, 78, 80, 81,
       84, 85, 87, 90, 93, 95, 96, 99, 100])
```

```
(a3 == a3_alt).all()
```

```
np.True_
```

```
np.allclose(a3, a3_alt * 1.000_001)
```

```
True
```

Create a 1-dimensional array a4 that contains the squares of the even integers through 100,000.

```
%timeit np.arange(0, 100_001, 2) ** 2
a4 = np.arange(0, 100_001, 2) ** 2
a4
```

```
30.1 s ± 6.69 s per loop (mean ± std. dev. of 7 runs, 10,000 loops each)
```

```
array([ 0, 4, 16, ..., 9999200016,
9999600004, 10000000000], shape=(50001,))
```

```
%timeit a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt = np.array([i**2 for i in range(100_001) if i%2==0])
a4_alt
```

```
22.1 ms ± 7.74 ms per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

```
array([ 0, 4, 16, ..., 9999200016,
9999600004, 10000000000], shape=(50001,))
```

```
(a4 == a4_alt).all()
```

```
np.True_
```

```
np.allclose(a4, a4_alt*1.0000001)
```

```
True
```

Write a function calc_pv() that mimic Excel's PV function.

Excel's present value function is: =PV(rate, nper, pmt, [fv], [type])

The present value of an annuity payment is: $PV_{pmt} = \frac{pmt}{rate} \times \left(1 - \frac{1}{(1+rate)^{nper}}\right)$

The present value of a lump sum is: $PV_{fv} = \frac{fv}{(1+rate)^{nper}}$

```
def calc_pv(rate, nper, pmt=None, fv=None, type=0):
    # or we could set the pmt and fv defaults to 0
    if pmt is None:
        pmt = 0
    if fv is None:
        fv = 0
    if type not in [0, 1]:
        raise ValueError(f'type must be 0 or 1; you provided {type}')

    pv_pmt = (pmt / rate) * (1 - (1 + rate)**(-nper))
    pv_fv = fv * (1 + rate)**(-nper)
    pv = pv_pmt + pv_fv

    if type == 1:
        pv *= 1 + rate

    return -pv
```

```
calc_pv(rate=0.1, nper=10, pmt=10, fv=100)
```

-100.0000

```
calc_pv(rate=0.1, nper=10, pmt=10, fv=100, type=1)
```

-110.0000

Write a function calc_fv() that mimic Excel's FV function.

Excel's future value function is: =FV(rate, nper, pmt, [pv], [type])

```
def calc_fv(rate, nper, pmt=None, pv=None, type='END'):
    if pmt is None:
        pmt = 0
    if pv is None:
        pv = 0
    if type not in ['BGN', 'END']:
        raise ValueError('type must be BGN or END')

    fv_pmt = (pmt / rate) * ((1 + rate)**nper - 1)
    fv_pv = pv * (1 + rate)**nper
    fv = fv_pmt + fv_pv

    if type == 'BGN':
        fv *= 1 + rate

    return -fv
```

```
calc_fv(rate=0.1, nper=10, pmt=10, pv=-100)
```

100.0000

The rule of 72!

```
calc_fv(rate=0.072, nper=10, pmt=0, pv=-100)
```

200.4231

Replace the negative values in data with -1 and positive values with +1.

```
np.random.seed(42)
data = np.random.randn(4, 4)

data

array([[ 0.4967, -0.1383,  0.6477,  1.523 ],
       [-0.2342, -0.2341,  1.5792,  0.7674],
       [-0.4695,  0.5426, -0.4634, -0.4657],
       [ 0.242 , -1.9133, -1.7249, -0.5623]])
```

We have at least three good solutions!

1. Slicing and broadcasting
2. `np.where()`
3. `np.select()`

First, here is the slicing and broadcasting solution.

```
data[data < 0] = -1
data[data > 0] = +1

data # here "1." and "-1." indicate that these values are floats

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Second, here is the `np.where()` solution. NumPy's `np.where()` function has the same logic as Excel's `if()` function! I will recreate `data` so we start from the same point.

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.where(data < 0, -1, np.where(data > 0, +1, data))

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

The nested `np.where()` function calls can get confusing! An option is to insert white space!

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.where(
    data < 0, # condition
    -1, # result if True
    np.where(data > 0, +1, data) # result if False
)
```

```
array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Third, here is the `np.select()` solution. NumPy's `np.select()` function lets us test *many* conditions! I will recreate `data` so we start from the same point.

```
np.random.seed(42)
data = np.random.randn(4, 4)

np.select(
    condlist=[data<0, data>0],
    choicelist=[-1, +1],
    default=data
)

array([[ 1., -1.,  1.,  1.],
       [-1., -1.,  1.,  1.],
       [-1.,  1., -1., -1.],
       [ 1., -1., -1., -1.]])
```

Write a function `calc_n()` that calculates the number of payments that generate x% of the present value of a perpetuity.

The present value of a growing perpetuity is $PV = \frac{C_1}{r-g}$, and the present value of a growing annuity is $PV = \frac{C_1}{r-g} \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

We can use the growing annuity and perpetuity formulas to show: $x = \left[1 - \left(\frac{1+g}{1+r} \right)^t \right]$.

Then: $1 - x = \left(\frac{1+g}{1+r} \right)^t$.

Finally: $t = \frac{\log(1-x)}{\log\left(\frac{1+g}{1+r}\right)}$

We do not need to accept an argument `c1` because C_1 cancels out!

```
def npmts(x, r, g):
    return np.log(1-x) / np.log((1 + g) / (1 + r))
```

```
npmts(0.5, 0.1, 0.05)
```

14.9000

Write a function that calculates the internal rate of return of a NumPy array of cash flows.

Let us pick a set of cash flows `c` where we know the internal rate of return! For the following `c`, the IRR is 10%.

```
c = np.array([-100, 110])
r = 0.1
```

First we need a function to calculate net present value (NPV) from `c` and `r`! The `npv()` function below uses NumPy arrays to calculate NPV as:

$$NPV = \sum_{t=0}^T \frac{c_t}{(1+r)^t}$$

```
def calc_npv(r, c):
    t = np.arange(len(c))
    return np.sum(c / (1 + r)**t)
```

```
calc_npv(r=r, c=c)
```

-0.0000

```
def calc_irr(c, guess=0, tol=1e-6, step=1e-6, max_iter=1_000, verbose=False):
    irr = guess
    for i in range(max_iter):
        npv = calc_npv(r=irr, c=c)
        if abs(npv) < tol:
            if verbose:
                print(f'IRR: {irr:0.4f}\nNPV: {npv:0.4f}\nIterations: {i+1:d}')
            return irr
        deriv = (calc_npv(r=irr+step, c=c) - npv) / step
        irr -= npv / deriv

    raise ValueError(f'NPV did not converge to zero after {i+1} iterations')
```

```
calc_irr(c)
```

```
0.1000
```

```
calc_irr(c=np.array([-100, 10, 10, 10, 10, 110]), verbose=True)
```

```
IRR: 0.1000
```

```
NPV: 0.0000
```

```
Iterations: 5
```

```
0.1000
```

```
# calc_irr(c=np.array([-2000, 1000, -500]))
```

Write a function `calc_returns()` that accepts NumPy arrays of prices and dividends and returns a NumPy array of returns.

```
prices = np.array([100, 150, 100, 50, 100, 150, 100, 150])
dividends = np.array([1, 1, 1, 1, 2, 2, 2, 2])
```

We want to slice our arrays to “lag” or “shift” them! For example, we slice the `prices` array to calculate capital gains as follows.

```
prices[1:] - prices[:-1]
```

```
array([ 50, -50, -50,  50,  50, -50,  50])
```

```
def calc_returns(p, d):
    return (p[1:] - p[:-1] + d[1:]) / p[:-1]
```

```
calc_returns(p=prices, d=dividends)
```

```
array([ 0.51 , -0.3267, -0.49 ,  1.04 ,  0.52 , -0.32 ,  0.52 ])
```

Rewrite the function `calc_returns()` as `calc_returns_2()` so it returns NumPy arrays of returns, capital gains yields, and dividend yields.

```
def calc_returns_2(p, d):
    cg = p[1:] / p[:-1] - 1
    dp = d[1:] / p[:-1]
    r = cg + dp
    return {'r': r, 'cg': cg, 'dp': dp}
```

```
calc_returns_2(p=prices, d=dividends)
```

```
{'r': array([ 0.51   , -0.3267, -0.49   ,  1.04   ,  0.52   , -0.32   ,  0.52   ]),
 'cg': array([ 0.5    , -0.3333, -0.5    ,  1.     ,  0.5    , -0.3333,  0.5    ]),
 'dp': array([0.01   ,  0.0067,  0.01   ,  0.04   ,  0.02   ,  0.0133,  0.02   ])}
```

Write a function rescale() to rescale and shift numbers so that they cover the range [0, 1]

Input: np.array([18.5, 17.0, 18.0, 19.0, 18.0])
Output: np.array([0.75, 0.0, 0.5, 1.0, 0.5])

```
x = np.array([18.5, 17.0, 18.0, 19.0, 18.0])
x
```

```
array([18.5, 17. , 18. , 19. , 18. ])
```

```
def rescale(x):
    return (x - x.min()) / (x.max() - x.min())
```

```
rescale(x=x)
```

```
array([0.75, 0. , 0.5 , 1. , 0.5 ])
```

Write a function calc_portval() that accepts a dictionary of prices and share holdings and returns the portfolio value

First convert your dictionary to a NumPy array with one column for prices and another for shares.

```
data = {
    "AAPL": (150.25, 10), # (price, shares)
    "GOOGL": (2750.00, 2),
    "MSFT": (300.75, 5)
}
```

```
def calc_portval(data):
    x = np.array(list(data.values()))
    return x.prod(axis=1).sum()
```

```
calc_portval(data=data)
```

8506.2500

Write functions `calc_var()` and `calc_std()` that calculate variance and standard deviation.

NumPy's `.var()` and `.std()` methods return *population* statistics (i.e., denominators of n). The pandas equivalents return *sample* statistics (denominators of $n - 1$), which are more appropriate for financial data analysis where we have a sample instead of a population.

Your `calc_var()` and `calc_std()` functions should have a `sample` argument that is `True` by default so both functions return sample statistics by default.

```
def calc_var(x, sample=True):
    sq_err = (x - x.mean()) ** 2
    den = len(x)
    if sample:
        den -= 1
    return sq_err.sum() / den
```

We can re-use `calc_var()` in our `calc_std()` function.

```
def calc_std(x, sample=True):
    return calc_var(x=x, sample=sample) ** 0.5
```

```
np.random.seed(42)
arr = np.random.randn(1_000_000)
arr
```

```
array([ 0.4967, -0.1383,  0.6477, ..., -0.113 ,  1.4691,  0.4764],  
      shape=(1000000,))
```

```
calc_var(arr)
```

1.0004

```
calc_std(arr)
```

1.0002

```
arr.var(ddof=1) == calc_var(arr)
```

```
np.True_
```

```
arr.std(ddof=1) == calc_std(arr)
```

```
np.True_
```

Write a function calc_ret() to convert quantitative returns to qualitative returns

Returns within one standard deviation of the mean are “Medium”. Returns less than one standard deviation below the mean are “Low”, and returns greater than one standard deviation above the mean are “High”.

```
np.random.seed(42)  
returns = np.random.randn(10)
```

```
def calc_ret(r):  
    mu = r.mean()  
    sigma = r.std(ddof=1)  
    return np.select(  
        condlist=[  
            r<(mu-sigma),  
            r<=(mu+sigma),  
            r>(mu+sigma)  
        ],  
        choicelist=[
```

```
'Low',
'Medium',
'High'
],
default=''
)
calc_ret(returns)

array(['Medium', 'Medium', 'Medium', 'High', 'Medium', 'Medium', 'High',
       'Medium', 'Low', 'Medium'], dtype='|<U6')
```

Week 4

McKinney Chapter 5 - Getting Started with pandas

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

Chapter 5 of McKinney (2022) discusses the fundamentals of pandas, which will be our main tool for the rest of the semester. pandas is an abbreviation of *panel data*. Panel data contain observations on multiple entities (e.g., individuals, firms, or countries) over multiple time periods. Panel data combine cross-sectional data (i.e., data across entities at a single point in time) with time-series data (i.e., data over time for a single entity). Panel data are widely used in finance and economics to analyze trends, relationships, and behaviors over time and across entities. We will use pandas every day for the rest of the course!

pandas will be a major tool of interest throughout much of the rest of the book. It contains data structures and data manipulation tools designed to make data cleaning and analysis fast and easy in Python. pandas is often used in tandem with numerical computing tools like NumPy and SciPy, analytical libraries like statsmodels and scikit-learn, and data visualization libraries like matplotlib. pandas adopts significant parts of NumPy's idiomatic style of array-based computing, especially array-based functions and a preference for data processing without for loops.

While pandas adopts many coding idioms from NumPy, the biggest difference is that pandas is designed for working with tabular or heterogeneous data. NumPy, by contrast, is best suited for working with homogeneous numerical array data.

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Introduction to pandas Data Structures

To get started with pandas, you will need to get comfortable with its two workhorse data structures: Series and DataFrame. While they are not a universal solution for every problem, they provide a solid, easy-to-use basis for most applications.

Series

A Series is a one-dimensional array-like object containing a sequence of values (of similar types to NumPy types) and an associated array of data labels, called its index. The simplest Series is formed from only an array of data.

The early examples use integer and string labels, but date-time labels are most useful.

```
obj = pd.Series([4, 7, -5, 3])
obj
```

```
0    4
1    7
2   -5
3    3
dtype: int64
```

Contrast obj with its NumPy array equivalent:

```
np.array([4, 7, -5, 3])
```

```
array([ 4,  7, -5,  3])
```

```
obj.values
```

```
array([ 4,  7, -5,  3])
```

```
obj.index # similar to range(4)
```

```
RangeIndex(start=0, stop=4, step=1)
```

We did not explicitly create an index for `obj`, so `obj` has an integer index that starts at 0. We can explicitly create an index with the `index=` argument.

```
obj2 = pd.Series([4, 7, -5, 3], index=['d', 'b', 'a', 'c'])  
obj2
```

```
d    4  
b    7  
a   -5  
c    3  
dtype: int64
```

```
obj2.index
```

```
Index(['d', 'b', 'a', 'c'], dtype='object')
```

```
obj2['a']
```

```
np.int64(-5)
```

```
obj2.loc['a']
```

```
np.int64(-5)
```

```
obj2.iloc[2]
```

```
np.int64(-5)
```

```
obj2['d'] = 6  
obj2
```

```
d    6  
b    7  
a   -5  
c    3  
dtype: int64
```

```
obj2[['c', 'a', 'd']]
```

```
c    3  
a   -5  
d    6  
dtype: int64
```

```
obj2.loc[['c', 'a', 'd']]
```

```
c    3  
a   -5  
d    6  
dtype: int64
```

A pandas series is like a NumPy array, and we can use Boolean filters and perform vectorized mathematical operations.

```
obj2 > 0
```

```
d    True  
b    True  
a   False  
c    True  
dtype: bool
```

```
obj2[obj2 > 0]
```

```
d    6  
b    7  
c    3  
dtype: int64
```

```
obj2.loc[obj2 > 0]
```

```
d      6  
b      7  
c      3  
dtype: int64
```

```
obj2 * 2
```

```
d     12  
b     14  
a    -10  
c      6  
dtype: int64
```

We can also create a pandas series from a dictionary. The dictionary keys become the series index.

```
sdata = {'Ohio': 35000, 'Texas': 71000, 'Oregon': 16000, 'Utah': 5000}  
obj3 = pd.Series(sdata)  
obj3
```

```
Ohio      35000  
Texas     71000  
Oregon    16000  
Utah      5000  
dtype: int64
```

If we also specify an index with list `states`, pandas will:

1. Respect the index order
2. Keep California because it was in the index
3. Drop Utah because it was not in the index

```
states = ['California', 'Ohio', 'Oregon', 'Texas']  
obj4 = pd.Series(sdata, index=states)  
obj4
```

```
California      NaN
Ohio          35000.0000
Oregon        16000.0000
Texas         71000.0000
dtype: float64
```

When we perform mathematical operations, pandas aligns series by their indexes. Here `NaN` is “not a number”, indicating missing values. `NaN` is a float, so the data type switches from `int64` to `float64`.

```
obj3 + obj4
```

```
California      NaN
Ohio          70000.0000
Oregon        32000.0000
Texas         142000.0000
Utah          NaN
dtype: float64
```

DataFrame

A pandas data frame is like a worksheet in an Excel workbook with row labels and column names that provide fast indexing.

A DataFrame represents a rectangular table of data and contains an ordered collection of columns, each of which can be a different value type (numeric, string, boolean, etc.). The DataFrame has both a row and column index; it can be thought of as a dict of Series all sharing the same index. Under the hood, the data is stored as one or more two-dimensional blocks rather than a list, dict, or some other collection of one-dimensional arrays. The exact details of DataFrame’s internals are outside the scope of this book.

There are many ways to construct a DataFrame, though one of the most common is from a dict of equal-length lists or NumPy arrays:

```
data = {
    'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
    'year': [2000, 2001, 2002, 2001, 2002, 2003],
    'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]
}
frame = pd.DataFrame(data)

frame
```

	state	year	pop
0	Ohio	2000	1.5000
1	Ohio	2001	1.7000
2	Ohio	2002	3.6000
3	Nevada	2001	2.4000
4	Nevada	2002	2.9000
5	Nevada	2003	3.2000

We did not specify an index, so `frame` has the default index of integers starting at 0.

```
frame2 = pd.DataFrame(
    data,
    columns=['year', 'state', 'pop', 'debt'],
    index=['one', 'two', 'three', 'four', 'five', 'six']
)

frame2
```

	year	state	pop	debt
one	2000	Ohio	1.5000	NaN
two	2001	Ohio	1.7000	NaN
three	2002	Ohio	3.6000	NaN
four	2001	Nevada	2.4000	NaN
five	2002	Nevada	2.9000	NaN
six	2003	Nevada	3.2000	NaN

If we extract one column with `df.column` or `df['column']`, we get a series. We can use the `df.colname` or `df['colname']` syntax to extract a column from a data frame as a series. **However, we must use the `df['colname']` syntax to add a column to a data frame.** Also, we must use the `df['colname']` syntax to extract or add a column whose name contains whitespace.

```
frame2['state']
```

one	Ohio
two	Ohio
three	Ohio
four	Nevada
five	Nevada

```
six      Nevada  
Name: state, dtype: object
```

```
frame2.state
```

```
one      Ohio  
two      Ohio  
three    Ohio  
four    Nevada  
five    Nevada  
six    Nevada  
Name: state, dtype: object
```

Data frames have two dimensions, so we must slice data frames more precisely than series.

1. The `.loc[]` method slices by row labels and column names
2. The `.iloc[]` method slices by *integer* row and label indexes

```
frame2.loc['three']
```

```
year      2002  
state    Ohio  
pop     3.6000  
debt      NaN  
Name: three, dtype: object
```

```
frame2.iloc[2]
```

```
year      2002  
state    Ohio  
pop     3.6000  
debt      NaN  
Name: three, dtype: object
```

We can use NumPy's `[row, column]` syntax within `.loc[]` and `.iloc[]`.

```
frame2.loc['three', 'state'] # row, column
```

```
'Ohio'
```

```
frame2.iloc[2, 1] # row, column
```

```
'Ohio'
```

```
frame2.loc['three', ['state', 'pop']] # row, column
```

```
state      Ohio
pop      3.6000
Name: three, dtype: object
```

```
frame2.iloc[2, [1, 2]] # row, column
```

```
state      Ohio
pop      3.6000
Name: three, dtype: object
```

We can assign either scalars or arrays to data frame columns.

1. Scalars will broadcast to every row in the data frame
2. Arrays must have the same length as the column

```
frame2['debt'] = 16.5
frame2
```

	year	state	pop	debt
one	2000	Ohio	1.5000	16.5000
two	2001	Ohio	1.7000	16.5000
three	2002	Ohio	3.6000	16.5000
four	2001	Nevada	2.4000	16.5000
five	2002	Nevada	2.9000	16.5000
six	2003	Nevada	3.2000	16.5000

```
frame2['debt'] = np.arange(6.)
frame2
```

	year	state	pop	debt
one	2000	Ohio	1.5000	0.0000
two	2001	Ohio	1.7000	1.0000
three	2002	Ohio	3.6000	2.0000
four	2001	Nevada	2.4000	3.0000
five	2002	Nevada	2.9000	4.0000
six	2003	Nevada	3.2000	5.0000

If we assign a series to a data frame column, pandas will use the index to align it with the data frame. Data frame rows not in the series will be `NaN`.

```
val = pd.Series([-1.2, -1.5, -1.7], index=['two', 'four', 'five'])
val
```

```
two    -1.2000
four   -1.5000
five   -1.7000
dtype: float64
```

```
frame2['debt'] = val
frame2
```

	year	state	pop	debt
one	2000	Ohio	1.5000	NaN
two	2001	Ohio	1.7000	-1.2000
three	2002	Ohio	3.6000	NaN
four	2001	Nevada	2.4000	-1.5000
five	2002	Nevada	2.9000	-1.7000
six	2003	Nevada	3.2000	NaN

We can add columns to our data frame, then delete them with `del`.

```
frame2['eastern'] = (frame2.state == 'Ohio')
frame2
```

	year	state	pop	debt	eastern
one	2000	Ohio	1.5000	NaN	True

	year	state	pop	debt	eastern
two	2001	Ohio	1.7000	-1.2000	True
three	2002	Ohio	3.6000	NaN	True
four	2001	Nevada	2.4000	-1.5000	False
five	2002	Nevada	2.9000	-1.7000	False
six	2003	Nevada	3.2000	NaN	False

```
del frame2['eastern']
frame2
```

	year	state	pop	debt
one	2000	Ohio	1.5000	NaN
two	2001	Ohio	1.7000	-1.2000
three	2002	Ohio	3.6000	NaN
four	2001	Nevada	2.4000	-1.5000
five	2002	Nevada	2.9000	-1.7000
six	2003	Nevada	3.2000	NaN

Index Objects

```
obj = pd.Series(range(3), index=['a', 'b', 'c'])
index = obj.index
index
```

```
Index(['a', 'b', 'c'], dtype='object')
```

Index objects are immutable!

```
# # TypeError: Index does not support mutable operations
# index[1] = 'd'
```

Indexes can contain duplicates, so an index does not guarantee that our data are duplicate-free.

```
dup_labels = pd.Index(['foo', 'foo', 'bar', 'bar'])
dup_labels
```

```
Index(['foo', 'foo', 'bar', 'bar'], dtype='object')
```

Essential Functionality

This section provides the most common pandas operations. It is difficult to provide an exhaustive reference, but this section introduces the most common operations.

Indexing, Selection, and Filtering

Indexing, selecting, and filtering will be among our most-used pandas features.

```
obj = pd.Series(np.arange(4.), index=['a', 'b', 'c', 'd'])  
obj
```

```
a    0.0000  
b    1.0000  
c    2.0000  
d    3.0000  
dtype: float64
```

```
obj['b']
```

```
1.0000
```

```
obj.loc['b']
```

```
1.0000
```

```
obj.iloc[1]
```

```
1.0000
```

```
obj.iloc[1:3]
```

```
b    1.0000  
c    2.0000  
dtype: float64
```

When we slice with labels, the left and right endpoints are inclusive.

```
obj.loc['b':'c']
```

```
b    1.0000
c    2.0000
dtype: float64
```

```
obj.loc['b':'c'] = 5
obj
```

```
a    0.0000
b    5.0000
c    5.0000
d    3.0000
dtype: float64
```

```
data = pd.DataFrame(
    data=np.arange(16).reshape((4, 4)),
    index=['Ohio', 'Colorado', 'Utah', 'New York'],
    columns=['one', 'two', 'three', 'four']
)
data
```

	one	two	three	four
Ohio	0	1	2	3
Colorado	4	5	6	7
Utah	8	9	10	11
New York	12	13	14	15

Indexing one column returns a series.

```
data['two']
```

```
Ohio      1
Colorado  5
Utah     9
New York 13
Name: two, dtype: int64
```

Indexing columns with a list returns a data frame.

```
data[['three']]
```

	three
Ohio	2
Colorado	6
Utah	10
New York	14

```
data[['three', 'one']]
```

	three	one
Ohio	2	0
Colorado	6	4
Utah	10	8
New York	14	12

Table 5-4 summarizes data frame indexing and slicing options:

- `df[val]`: Select single column or sequence of columns from the DataFrame; special case conveniences: boolean array (filter rows), slice (slice rows), or boolean DataFrame (set values based on some criterion)
- `df.loc[val]`: Selects single row or subset of rows from the DataFrame by label
- `df.loc[:, val]`: Selects single column or subset of columns by label
- `df.loc[val1, val2]`: Select both rows and columns by label
- `df.iloc[where]`: Selects single row or subset of rows from the DataFrame by integer position
- `df.iloc[:, where]`: Selects single column or subset of columns by integer position
- `df.iloc[where_i, where_j]`: Select both rows and columns by integer position
- `df.at[label_i, label_j]`: Select a single scalar value by row and column label
- `df.iat[i, j]`: Select a single scalar value by row and column position (integers) reindex method Select either rows or columns by labels
- `get_value, set_value` methods: Select single value by row and column label

pandas is powerful and these options can be overwhelming! We will typically use `df[val]` to select columns (here `val` is either a string or list of strings), `df.loc[val]` to select rows (here `val` is a row label), and `df.loc[val1, val2]` to select both rows and columns. The other options add flexibility, and we may occasionally use them. However, our data will be large enough that counting row and column number will be tedious, making `.iloc[]` impractical.

Arithmetic and Data Alignment

An important pandas feature for some applications is the behavior of arithmetic between objects with different indexes. When you are adding together objects, if any index pairs are not the same, the respective index in the result will be the union of the index pairs. For users with database experience, this is similar to an automatic outer join on the index labels.

```
s1 = pd.Series(
    data=[7.3, -2.5, 3.4, 1.5],
    index=['a', 'c', 'd', 'e']
)
s2 = pd.Series(
    data=[-2.1, 3.6, -1.5, 4, 3.1],
    index=['a', 'c', 'e', 'f', 'g']
)
```

```
s1
```

```
a    7.3000
c   -2.5000
d    3.4000
e    1.5000
dtype: float64
```

```
s2
```

```
a   -2.1000
c    3.6000
e   -1.5000
f    4.0000
g    3.1000
dtype: float64
```

```
s1 + s2
```

```
a    5.2000
c    1.1000
d      NaN
e    0.0000
f      NaN
```

```
g      NaN
dtype: float64
```

```
df1 = pd.DataFrame(
    data=np.arange(9.).reshape((3, 3)),
    columns=list('bcd'),
    index=['Ohio', 'Texas', 'Colorado']
)
df2 = pd.DataFrame(
    data=np.arange(12.).reshape((4, 3)),
    columns=list('bde'),
    index=['Utah', 'Ohio', 'Texas', 'Oregon']
)
```

```
df1
```

	b	c	d
Ohio	0.0000	1.0000	2.0000
Texas	3.0000	4.0000	5.0000
Colorado	6.0000	7.0000	8.0000

```
df2
```

	b	d	e
Utah	0.0000	1.0000	2.0000
Ohio	3.0000	4.0000	5.0000
Texas	6.0000	7.0000	8.0000
Oregon	9.0000	10.0000	11.0000

```
df1 + df2
```

	b	c	d	e
Colorado	NaN	NaN	NaN	NaN
Ohio	3.0000	NaN	6.0000	NaN
Oregon	NaN	NaN	NaN	NaN
Texas	9.0000	NaN	12.0000	NaN
Utah	NaN	NaN	NaN	NaN

Always check your output!

Arithmetic methods with fill values

```
df1 = pd.DataFrame(
    data=np.arange(12.).reshape((3, 4)),
    columns=list('abcd'))
)
df2 = pd.DataFrame(
    data=np.arange(20.).reshape((4, 5)),
    columns=list('abcde'))
)
df2.loc[1, 'b'] = np.nan
```

df1

	a	b	c	d
0	0.0000	1.0000	2.0000	3.0000
1	4.0000	5.0000	6.0000	7.0000
2	8.0000	9.0000	10.0000	11.0000

df2

	a	b	c	d	e
0	0.0000	1.0000	2.0000	3.0000	4.0000
1	5.0000	NaN	7.0000	8.0000	9.0000
2	10.0000	11.0000	12.0000	13.0000	14.0000
3	15.0000	16.0000	17.0000	18.0000	19.0000

df1 + df2

	a	b	c	d	e
0	0.0000	2.0000	4.0000	6.0000	NaN
1	9.0000	NaN	13.0000	15.0000	NaN
2	18.0000	20.0000	22.0000	24.0000	NaN
3	NaN	NaN	NaN	NaN	NaN

We can specify a fill value for `NaN` values. pandas fills would-be `NaN` values in each data frame *before* the arithmetic operation.

```
df1.add(df2, fill_value=0)
```

	a	b	c	d	e
0	0.0000	2.0000	4.0000	6.0000	4.0000
1	9.0000	5.0000	13.0000	15.0000	9.0000
2	18.0000	20.0000	22.0000	24.0000	14.0000
3	15.0000	16.0000	17.0000	18.0000	19.0000

Operations between DataFrame and Series

```
arr = np.arange(12.).reshape((3, 4))
arr
```

```
array([[ 0.,  1.,  2.,  3.],
       [ 4.,  5.,  6.,  7.],
       [ 8.,  9., 10., 11.]])
```

```
arr[0]
```

```
array([0., 1., 2., 3.])
```

```
arr - arr[0]
```

```
array([[ 0.,  0.,  0.,  0.],
       [ 4.,  4.,  4.,  4.],
       [ 8.,  8.,  8.,  8.]])
```

Arithmetic operations between series and data frames behave the same as in the example above.

```
frame = pd.DataFrame(
    data=np.arange(12.).reshape((4, 3)),
    columns=list('bde'),
    index=['Utah', 'Ohio', 'Texas', 'Oregon']
)

series = frame.iloc[0]
```

```
frame
```

	b	d	e
Utah	0.0000	1.0000	2.0000
Ohio	3.0000	4.0000	5.0000
Texas	6.0000	7.0000	8.0000
Oregon	9.0000	10.0000	11.0000

```
series
```

```
b    0.0000  
d    1.0000  
e    2.0000  
Name: Utah, dtype: float64
```

```
frame - series
```

	b	d	e
Utah	0.0000	0.0000	0.0000
Ohio	3.0000	3.0000	3.0000
Texas	6.0000	6.0000	6.0000
Oregon	9.0000	9.0000	9.0000

```
series2 = pd.Series(data=range(3), index=['b', 'e', 'f'])
```

```
frame
```

	b	d	e
Utah	0.0000	1.0000	2.0000
Ohio	3.0000	4.0000	5.0000
Texas	6.0000	7.0000	8.0000
Oregon	9.0000	10.0000	11.0000

```
series2
```

```
b    0  
e    1  
f    2  
dtype: int64
```

```
frame + series2
```

	b	d	e	f
Utah	0.0000	NaN	3.0000	NaN
Ohio	3.0000	NaN	6.0000	NaN
Texas	6.0000	NaN	9.0000	NaN
Oregon	9.0000	NaN	12.0000	NaN

pandas has a `.sub()` method that lets us chain operations, but we might need the `axis` argument to get the result we want!

```
series3 = frame['d']
```

```
frame
```

	b	d	e
Utah	0.0000	1.0000	2.0000
Ohio	3.0000	4.0000	5.0000
Texas	6.0000	7.0000	8.0000
Oregon	9.0000	10.0000	11.0000

```
series3
```

```
Utah      1.0000  
Ohio      4.0000  
Texas     7.0000  
Oregon   10.0000  
Name: d, dtype: float64
```

```
frame - series
```

	b	d	e
Utah	0.0000	0.0000	0.0000
Ohio	3.0000	3.0000	3.0000
Texas	6.0000	6.0000	6.0000
Oregon	9.0000	9.0000	9.0000

```
frame.sub(series3, axis=0)
```

	b	d	e
Utah	-1.0000	0.0000	1.0000
Ohio	-1.0000	0.0000	1.0000
Texas	-1.0000	0.0000	1.0000
Oregon	-1.0000	0.0000	1.0000

Function Application and Mapping

```
np.random.seed(42)
frame = pd.DataFrame(
    data=np.random.randn(4, 3),
    columns=list('bde'),
    index=['Utah', 'Ohio', 'Texas', 'Oregon']
)

frame
```

	b	d	e
Utah	0.4967	-0.1383	0.6477
Ohio	1.5230	-0.2342	-0.2341
Texas	1.5792	0.7674	-0.4695
Oregon	0.5426	-0.4634	-0.4657

```
frame.abs()
```

	b	d	e
Utah	0.4967	0.1383	0.6477
Ohio	1.5230	0.2342	0.2341
Texas	1.5792	0.7674	0.4695
Oregon	0.5426	0.4634	0.4657

Another frequent operation is applying a function on one-dimensional arrays to each column or row. DataFrame's apply method does exactly this:

```
frame.apply(lambda x: x.max() - x.min()) # implied axis=0
```

```
b    1.0825
d    1.2309
e    1.1172
dtype: float64
```

```
frame.apply(lambda x: x.max() - x.min(), axis=1) # explicit axis=1
```

```
Utah      0.7860
Ohio      1.7572
Texas     2.0487
Oregon    1.0083
dtype: float64
```

However, under the hood, the .apply() method is a `for` loop and slower than built-in methods.

```
%timeit frame['e'].abs()
```

```
44.9 s ± 10.7 s per loop (mean ± std. dev. of 7 runs, 100,000 loops each)
```

```
%timeit frame['e'].apply(np.abs)
```

```
96.9 s ± 16.6 s per loop (mean ± std. dev. of 7 runs, 10,000 loops each)
```

Summarizing and Computing Descriptive Statistics

```
df = pd.DataFrame(
    [[1.4, np.nan], [7.1, -4.5], [np.nan, np.nan], [0.75, -1.3]],
    index=['a', 'b', 'c', 'd'],
    columns=['one', 'two']
)

df
```

	one	two
a	1.4000	NaN
b	7.1000	-4.5000
c	NaN	NaN
d	0.7500	-1.3000

```
df.sum() # implied axis=0
```

```
one    9.2500
two   -5.8000
dtype: float64
```

```
df.sum(axis=1)
```

```
a    1.4000
b    2.6000
c    0.0000
d   -0.5500
dtype: float64
```

```
df.mean(axis=1, skipna=False)
```

```
a      NaN
b    1.3000
c      NaN
d   -0.2750
dtype: float64
```

The `.idxmax()` method returns the label for the maximum observation.

```
df
```

	one	two
a	1.4000	NaN
b	7.1000	-4.5000
c	NaN	NaN
d	0.7500	-1.3000

```
df.idxmax()
```

```
one      b
two      d
dtype: object
```

The `.describe()` returns summary statistics for each numerical column in a data frame.

```
df.describe()
```

	one	two
count	3.0000	2.0000
mean	3.0833	-2.9000
std	3.4937	2.2627
min	0.7500	-4.5000
25%	1.0750	-3.7000
50%	1.4000	-2.9000
75%	4.2500	-2.1000
max	7.1000	-1.3000

For non-numerical data, `.describe()` returns alternative summary statistics.

```
obj = pd.Series(['a', 'a', 'b', 'c'] * 4)
obj.describe()
```

```
count     16
unique      3
top        a
freq       8
dtype: object
```

Correlation and Covariance

 Note

Starting with version 0.2.51, the `yfinance` package changed the default behavior of the `auto_adjust` argument from `False` to `True`. By default, the `ya.download()` function now returns adjusted prices, without including the `Adj Close` column.

We prefer to work with raw data from Yahoo! Finance and explicitly calculate returns using the `Adj Close` column. Therefore, we will set `auto_adjust=False` in our `ya.download()` calls. See the [yfinance changelog](#) for release version 0.2.51.

Also, I will use the `progress=False` argument to improve the readability of the PDF and website I render from these notebooks.

```
data = ya.download(tickers='AAPL IBM MSFT GOOG', auto_adjust=False, progress=False)

data['Adj Close'].tail()
```

Ticker	AAPL	GOOG	IBM	MSFT
Date				
2025-02-24	247.1000	181.1900	261.8700	404.0000
2025-02-25	247.0400	177.3700	257.7500	397.9000
2025-02-26	240.3600	174.7000	255.8400	399.7300
2025-02-27	237.3000	170.2100	253.2300	392.5300
2025-02-28	241.8400	172.2200	252.4400	396.9900

Data frame `data` contains daily prices and volume for AAPL, IBM, MSFT, and GOOG. The `Adj Close` columns are reverse-engineered daily closing prices that account for dividends and stock splits (and reverse splits). As a result, the `.pct_change()` of `Adj Close` correctly considers dividends and price changes, so $r_t = \frac{(P_t + D_t) - P_{t-1}}{P_{t-1}} = \frac{\text{Adj Close}_t - \text{Adj Close}_{t-1}}{\text{Adj Close}_{t-1}}$.

```
returns = data['Adj Close'].pct_change().dropna()
returns
```

Ticker	AAPL	GOOG	IBM	MSFT
Date				
2004-08-20	0.0029	0.0794	0.0042	0.0029
2004-08-23	0.0091	0.0101	-0.0070	0.0044
2004-08-24	0.0280	-0.0414	0.0007	0.0000

Ticker	AAPL	GOOG	IBM	MSFT
Date				
2004-08-25	0.0344	0.0108	0.0043	0.0114
2004-08-26	0.0487	0.0180	-0.0045	-0.0040
...
2025-02-24	0.0063	-0.0021	0.0015	-0.0103
2025-02-25	-0.0002	-0.0211	-0.0157	-0.0151
2025-02-26	-0.0270	-0.0151	-0.0074	0.0046
2025-02-27	-0.0127	-0.0257	-0.0102	-0.0180
2025-02-28	0.0191	0.0118	-0.0031	0.0114

We multiply by 252 to annualize mean daily returns because means grow linearly with time and there are (about) 252 trading days per year.

```
returns.mean().mul(252)
```

```
Ticker
AAPL    0.3579
GOOG    0.2533
IBM     0.1112
MSFT    0.1905
dtype: float64
```

We multiply by $\sqrt{252}$ to annualize the volatility of daily returns because standard deviation is the square root of variance, variances grow linearly with time, and there are (about) 252 trading days per year. Ivo Welch explains this calculation at the bottom of Page 7 of Chapter 8 his [free corporate finance textbook](#).

```
returns.std().mul(np.sqrt(252))
```

```
Ticker
AAPL    0.3234
GOOG    0.3061
IBM     0.2281
MSFT    0.2693
dtype: float64
```

We can calculate pairwise correlations.

```
returns['MSFT'].corr(returns['IBM'])
```

0.4733

We can also calculate correlation matrices.

```
returns.corr()
```

Ticker	AAPL	GOOG	IBM	MSFT
Ticker				
AAPL	1.0000	0.5116	0.4121	0.5209
GOOG	0.5116	1.0000	0.3807	0.5603
IBM	0.4121	0.3807	1.0000	0.4733
MSFT	0.5209	0.5603	0.4733	1.0000

```
returns.corr().loc['MSFT', 'IBM']
```

0.4733

```
np.allclose(  
    a=returns['MSFT'].corr(returns['IBM']),  
    b=returns.corr().loc['MSFT', 'IBM']  
)
```

True

McKinney Chapter 5 - Practice - Blank

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

What are the mean daily returns for these four stocks?

```
tickers = 'AAPL IBM MSFT GOOG'
```

What are the standard deviations of daily returns for these four stocks?

What are the *annualized* means and standard deviations of daily returns for these four stocks?

Plot *annualized* means versus standard deviations of daily returns for these four stocks

Repeat the previous calculations and plot for the stocks in the Dow-Jones Industrial Index (DJIA)

We can find the current DJIA stocks on [Wikipedia](#). We must download new data, into `tickers_2`, `data_2`, and `returns_2`.

Calculate total returns for the stocks in the DJIA

Plot the distribution of total returns for the stocks in the DJIA

Which stocks have the minimum and maximum total returns?

Plot the cumulative returns for the stocks in the DJIA

Repeat the plot above with only the minimum and maximum total returns

McKinney Chapter 5 - Practice - Sec 02

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Please keep forming groups on Canvas > People > Projects. If you want a group with more than four students, please fill a group with four students, then email with the group number and size.
2. Please keep proposing and voting for students' choice topics [here](#).

Five-Minute Review

The pandas package makes it easy to manipulate panel data and we will use it all semester. Its name is an abbreviation of [panel data](#):

In statistics and econometrics, panel data and longitudinal data[1][2] are both multi-dimensional data involving measurements over time. Panel data is a subset of longitudinal data where observations are for the same subjects each time.

We will download some financial data from Yahoo! Finance to cover three important tools in pandas.

First, we can use the yfinance package to easily download stock data from Yahoo! Finance.

Note

Starting with version 0.2.51, the `yfinance` package changed the default behavior of the `auto_adjust` argument from `False` to `True`. By default, the `ya.download()` function now returns adjusted prices, without including the `Adj Close` column.

We prefer to work with raw data from Yahoo! Finance and explicitly calculate returns using the `Adj Close` column. Therefore, we will set `auto_adjust=False` in our `ya.download()` calls. See the [yfinance changelog](#) for release version 0.2.51.

Also, I will use the `progress=False` argument to improve the readability of the PDF and website I render from these notebooks.

```
df0 = ya.download(tickers='AAPL MSFT', auto_adjust=False, progress=False)
```

df0

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT	Low AAPL	Low MSFT	Open AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...
2025-02-24	247.1000	404.0000	247.1000	404.0000	248.8600	409.3700	244.4200	399.3200	244.9300
2025-02-25	247.0400	397.9000	247.0400	397.9000	250.0000	401.9200	244.9100	396.7000	248.0000
2025-02-26	240.3600	399.7300	240.3600	399.7300	244.9800	403.6000	239.1300	394.2500	244.3300
2025-02-27	237.3000	392.5300	237.3000	392.5300	242.4600	405.7400	237.0600	392.1700	239.4100
2025-02-28	241.8400	396.9900	241.8400	396.9900	242.0900	397.6300	230.2000	386.5700	236.9500

Second, we can slice rows and columns two ways: by integer locations with `.iloc[]` and by labels with `.loc[]`

```
df0.iloc[:6, :6]
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

```
df0.loc[:, '1980-12-19', : 'High']
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

i Note

To slice a DataFrame:

- Use `['Name']` to select specific columns by their names.
- Use `.loc[]` to slice rows, or rows and columns together, with labels or conditional expressions.

```
df0['High']
```

Ticker Date	AAPL	MSFT
1980-12-12	0.1289	NaN
1980-12-15	0.1222	NaN
1980-12-16	0.1133	NaN

Ticker	AAPL	MSFT
Date		
1980-12-17	0.1161	NaN
1980-12-18	0.1194	NaN
...
2025-02-24	248.8600	409.3700
2025-02-25	250.0000	401.9200
2025-02-26	244.9800	403.6000
2025-02-27	242.4600	405.7400
2025-02-28	242.0900	397.6300


```
# # KeyError: '1980-12-12'
# df0['1980-12-12']
```

Note, if we use string labels, like dates and words, pandas includes left and right edges! This string label behavior differs from the integer location behavior everywhere else in Python. However, it is easy to figure out the sequence of integer locations. It is difficult to figure our the sequence of string labels.

Third, there many methods we can apply to pandas objects (and chain)! At this point in the course, our most common methods will be:

1. `.pct_change()` to calculate simple returns from adjusted close prices
2. `.plot()` to quickly plot pandas objects
3. `.mean()`, `.std()`, `.describe()`, etc. to calculate summary statistics

```
( df0 # DataFrame containing AAPL and MSFT data from 1980-12-12 through today
  .loc['1990':'2024', 'Adj Close'] # Slice rows for 1990-2024 (inclusive) and the 'Adj Cl
  .pct_change() # Calculate daily percentage changes in 'Adj Close' (includes dividends an
  .add(1) # Prepare for compounding by adding 1 to daily returns
  .cumprod() # Compute cumulative product to get total return for each day since the start
  .sub(1) # Convert back to cumulative returns
  .plot() # Plot cumulative returns
)
```



The plot above is in decimal returns! pandas makes it easy to generate plots, but getting them beautiful and readable takes more work. The following code adds a title, labels, and formats the y axis.


```
from matplotlib.ticker import FuncFormatter

# Plot the data
ax = df0.loc['1990':'2024', 'Adj Close'].pct_change().add(1).cumprod().sub(1).plot()

# Format y-axis as percentages with comma separators
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:,.0f}'))

# Add labels and title if needed
plt.ylabel('Cumulative Return (%)')
plt.title('Cumulative Returns from 1990 to 2024')

# Show the plot and suppress text output (<Axes: xlabel='Date'> above)
plt.show()
```



Practice

What are the mean daily returns for these four stocks?

```
tickers = 'AAPL IBM MSFT GOOG'

returns = (
    yf.download(tickers=tickers, auto_adjust=False, progress=False)
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
)

returns
```

Ticker	AAPL	GOOG	IBM	MSFT
Date				
1962-01-02	NaN	NaN	NaN	NaN
1962-01-03	NaN	NaN	0.0087	NaN
1962-01-04	NaN	NaN	-0.0100	NaN
1962-01-05	NaN	NaN	-0.0197	NaN
1962-01-08	NaN	NaN	-0.0188	NaN
...
2025-02-21	-0.0011	-0.0271	-0.0123	-0.0190
2025-02-24	0.0063	-0.0021	0.0015	-0.0103
2025-02-25	-0.0002	-0.0211	-0.0157	-0.0151
2025-02-26	-0.0270	-0.0151	-0.0074	0.0046
2025-02-27	-0.0127	-0.0257	-0.0102	-0.0180

```

(
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .mean() # calculate mean of daily returns from GOOG IPO through yesterday
)

```

Ticker
 AAPL 0.0014
 GOOG 0.0010
 IBM 0.0004
 MSFT 0.0008
 dtype: float64

What are the standard deviations of daily returns for these four stocks?

```

(
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .std() # calculate standard deviation (volatility) of daily returns from GOOG IPO through yesterday
)

```

Ticker
 AAPL 0.0204

```

GOOG    0.0193
IBM     0.0144
MSFT    0.0170
dtype: float64

```

What are the *annualized* means and standard deviations of daily returns for these four stocks?

```

ann_means = (
    returns # daily returns from 1962 through today
    .dropna() # drop days with missing returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .mean() # calculate mean of daily returns from close of GOOG IPO through yesterday
    .mul(252) # means grow linearly with time, so multiply by 252
)

ann_means

```

```

Ticker
AAPL    0.3577
GOOG    0.2541
IBM     0.1119
MSFT    0.1909
dtype: float64

```

```

ann_stds = (
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .std() # calculate mean of daily returns from close of GOOG IPO through yesterday
    .mul(np.sqrt(252)) # variances grow linearly with time, so standard deviations grow sqrt
)

ann_stds

```

```


Ticker
AAPL    0.3234
GOOG    0.3060
IBM     0.2282
MSFT    0.2693
dtype: float64

```

Plot annualized means versus standard deviations of daily returns for these four stocks

Here is a crude plot!

```
plt.scatter(x=ann_stds, y=ann_means)
```



But we can do better than a crude plot! We will typically combine data into a data frame to make plotting easier. Because `ann_std` and `ann_means` are pandas' series, so we can use `pd.DataFrame` to combine them into a data frame.

```
df = pd.DataFrame({'Volatility': ann_stds, 'Mean Return': ann_means})  
df
```

	Volatility	Mean Return
Ticker		
AAPL	0.3234	0.3577

Ticker	Volatility	Mean Return
GOOG	0.3060	0.2541
IBM	0.2282	0.1119
MSFT	0.2693	0.1909

 Note

Below, we could use `enumerate()` instead of `.iterrows()`. However, `enumerate()` loops over *column names* instead of row indexes and contents. Therefore, with `enumerate()`, we would have to `.transpose()` our data frame, then use the tickers to slice the rows of our original data frame. Here `iterrows()` combines these several steps into one.


```
ax = df.plot(kind='scatter', x='Volatility', y='Mean Return')

for s, (v, mr) in df.iterrows():
    plt.text(s=s, x=v, y=mr)

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title('Return versus Risk for Tech Stocks')
plt.show()
```



Repeat the previous calculations and plot for the stocks in the Dow-Jones Industrial Index (DJIA)

We can find the current DJIA stocks on [Wikipedia](https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average). We must download new data, into `tickers_2`, `data_2`, and `returns_2`.

```
url_2 = 'https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average'  
wiki_2 = pd.read_html(io=url_2)
```

```
type(wiki_2)
```

```
list
```

```
tickers_2 = wiki_2[2]['Symbol'].to_list()
```

```
data_2 = yf.download(tickers=tickers_2, auto_adjust=False, progress=False)
```

```
returns_2 = (
    data_2
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
    .dropna()
)
```

```
df_2 = pd.DataFrame({
    'Volatility': returns_2.std().mul(np.sqrt(252)),
    'Mean Return': returns_2.mean().mul(252)
})
```

```
dates_2 = returns_2.index
```

```
dates_2[0]
```

```
Timestamp('2008-03-20 00:00:00')
```

```
dates_2[-1]
```

```
Timestamp('2025-02-27 00:00:00')
```

```
ax = df_2.plot(kind='scatter', x='Volatility', y='Mean Return')
```


```
for s, (v, mr) in df_2.iterrows():
    plt.text(s=s, x=v, y=mr)
```

```
dates_2 = returns_2.index
```

```
ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

```
plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')
```

```
plt.title(f'Return versus Risk for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



We can use the seaborn package to add a best-fit line! More on seaborn here: <https://seaborn.pydata.org/index.html>

```
import seaborn as sns

sns.regplot(
    data=df_2,
    x='Volatility',
    y='Mean Return'
)

for s, (v, mr) in df_2.iterrows():
    plt.text(s=s, x=v, y=mr)


ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

```

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title(f'Return versus Risk for DJIA Stocks\nfrom {dates_2[0]}\n{dates_2[-1]}')
plt.show()

```



NVDA is a real outlier! We can use the `.drop()` method to quickly drop NVDA. Theory predicts no relation between μ and σ for single stocks because single-stock risk is diversifiable. If we use a larger sample (or a different time period), we would see a flat or negatively sloped best-fit line.

```

sns.regplot(
    data=df_2.drop('NVDA'),
    x='Volatility',
    y='Mean Return'
)

```

```

for s, (v, mr) in df_2.drop('NVDA').iterrows():
    plt.text(s=s, x=v, y=mr)


dates_2 = returns_2.index

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title(f'Return versus Risk for DJIA Stocks (Minus NVDA)\nfrom {dates_2[0]}:{%B %Y} to {dates_2[-1]}:{%B %Y}')
plt.show()

```



Calculate total returns for the stocks in the DJIA

We can use the `.prod()` method to compound returns as $1 + R_T = \prod_{t=1}^T (1 + R_t)$. Technically, we should write R_T as $R_{0,T}$, but we typically omit the subscript 0.

In general, I prefer to do simple math on pandas objects (data frames and series) with methods instead of operators:

For example:

1. `.add(1)` instead of `+ 1`
2. `.sub(1)` instead of `- 1`
3. `.div(1)` instead of `/ 1`
4. `.mul(1)` instead of `* 1`

The advantage of methods over operators, is that we can easily chain methods without lots of parentheses.


```
total_returns_2 = returns_2.add(1).prod().sub(1)

ax = total_returns_2.sort_values().plot(kind='barh')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Total Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Plot the distribution of total returns for the stocks in the DJIA

We can plot a histogram, using either the `plt.hist()` function or the `.plot(kind='hist')` method.


A histogram is a great way to visualize data!

```
ax = total_returns_2.plot(kind='hist')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Distribution of Total Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Which stocks have the minimum and maximum total returns?

If we want the *values*, the `.min()` and `.max()` methods are the way to go!

```
total_returns_2.min()
```

2.1510

```
total_returns_2.max()
```

295.7450

The `.min()` and `.max()` methods give the values but not the tickers (or index). We use the `.idxmin()` and `.idxmax()` to get the tickers (or index).

```
total_returns_2.idxmin()
```

```
'VZ'
```

```
total_returns_2.idxmax()
```

```
'NVDA'
```

Here is what I would use to capture values and tickers!

```
total_returns_2.sort_values().iloc[[0, -1]]
```

```
Ticker
VZ      2.1510
NVDA    295.7450
dtype: float64
```

Not the exactly right tool here, but the `.nsmallest()` and `.nlargest()` methods are really useful!

```
total_returns_2.nsmallest(3)
```

```
Ticker
VZ      2.1510
BA      2.2134
CVX     2.7258
dtype: float64
```

```
total_returns_2.nlargest(3)
```

```
Ticker
NVDA    295.7450
AAPL    59.8113
AMZN    58.4955
dtype: float64
```

Plot the cumulative returns for the stocks in the DJIA


We can use the cumulative product method `.cumprod()` to calculate the right hand side of the formula above.

```
ax = (
    returns_2
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Cumulative Return (%)')

plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```




Repeat the plot above with only the minimum and maximum total returns

```
ax = (
    returns_2
    [total_returns_2.sort_values().iloc[[0, -1]].index] # slice min and max total return stock
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

plt.ylabel('Cumulative Return (%)')

```
plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}\nplt.show()
```



McKinney Chapter 5 - Practice - Sec 03

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Please keep forming groups on Canvas > People > Projects. If you want a group with more than four students, please fill a group with four students, then email with the group number and size.
2. Please keep proposing and voting for students' choice topics [here](#).

Five-Minute Review

The pandas package makes it easy to manipulate panel data and we will use it all semester. Its name is an abbreviation of [panel data](#):

In statistics and econometrics, panel data and longitudinal data[1][2] are both multi-dimensional data involving measurements over time. Panel data is a subset of longitudinal data where observations are for the same subjects each time.

We will download some financial data from Yahoo! Finance to cover three important tools in pandas.

First, we can use the yfinance package to easily download stock data from Yahoo! Finance.

Note

Starting with version 0.2.51, the `yfinance` package changed the default behavior of the `auto_adjust` argument from `False` to `True`. By default, the `ya.download()` function now returns adjusted prices, without including the `Adj Close` column.

We prefer to work with raw data from Yahoo! Finance and explicitly calculate returns using the `Adj Close` column. Therefore, we will set `auto_adjust=False` in our `ya.download()` calls. See the [yfinance changelog](#) for release version 0.2.51.

Also, I will use the `progress=False` argument to improve the readability of the PDF and website I render from these notebooks.

```
df0 = ya.download(tickers='AAPL MSFT', auto_adjust=False, progress=False)
```

df0

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT	Low AAPL	Low MSFT	Open AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...
2025-02-24	247.1000	404.0000	247.1000	404.0000	248.8600	409.3700	244.4200	399.3200	244.9300
2025-02-25	247.0400	397.9000	247.0400	397.9000	250.0000	401.9200	244.9100	396.7000	248.0000
2025-02-26	240.3600	399.7300	240.3600	399.7300	244.9800	403.6000	239.1300	394.2500	244.3300
2025-02-27	237.3000	392.5300	237.3000	392.5300	242.4600	405.7400	237.0600	392.1700	239.4100
2025-02-28	241.8400	396.9900	241.8400	396.9900	242.0900	397.6300	230.2000	386.5700	236.9500

Second, we can slice rows and columns two ways: by integer locations with `.iloc[]` and by labels with `.loc[]`

```
df0.iloc[:6, :6]
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

```
df0.loc[:, '1980-12-19', : 'High']
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

i Note

To slice a DataFrame:

- Use `['Name']` to select specific columns by their names.
- Use `.loc[]` to slice rows, or rows and columns together, with labels or conditional expressions.

```
df0['High']
```

Ticker Date	AAPL	MSFT
1980-12-12	0.1289	NaN
1980-12-15	0.1222	NaN
1980-12-16	0.1133	NaN

Ticker	AAPL	MSFT
Date		
1980-12-17	0.1161	NaN
1980-12-18	0.1194	NaN
...
2025-02-24	248.8600	409.3700
2025-02-25	250.0000	401.9200
2025-02-26	244.9800	403.6000
2025-02-27	242.4600	405.7400
2025-02-28	242.0900	397.6300


```
# # KeyError: '1980-12-12'
# df0['1980-12-12']
```

Note, if we use string labels, like dates and words, pandas includes left and right edges! This string label behavior differs from the integer location behavior everywhere else in Python. However, it is easy to figure out the sequence of integer locations. It is difficult to figure our the sequence of string labels.

Third, there many methods we can apply to pandas objects (and chain)! At this point in the course, our most common methods will be:

1. `.pct_change()` to calculate simple returns from adjusted close prices
2. `.plot()` to quickly plot pandas objects
3. `.mean()`, `.std()`, `.describe()`, etc. to calculate summary statistics

```
( df0  # DataFrame containing AAPL and MSFT data from 1980-12-12 through today
  .loc['1990':'2024', 'Adj Close']  # Slice rows for 1990-2024 (inclusive) and the 'Adj Cl
  .pct_change()  # Calculate daily percentage changes in 'Adj Close' (includes dividends an
  .add(1)  # Prepare for compounding by adding 1 to daily returns
  .cumprod()  # Compute cumulative product to get total return for each day since the start
  .sub(1)  # Convert back to cumulative returns
  .plot()  # Plot cumulative returns
)
```



The plot above is in decimal returns! pandas makes it easy to generate plots, but getting them beautiful and readable takes more work. The following code adds a title, labels, and formats the y axis.


```
from matplotlib.ticker import FuncFormatter

# Plot the data
ax = df0.loc['1990':'2024', 'Adj Close'].pct_change().add(1).cumprod().sub(1).plot()

# Format y-axis as percentages with comma separators
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:,.0f}'))

# Add labels and title if needed
plt.ylabel('Cumulative Return (%)')
plt.title('Cumulative Returns from 1990 to 2024')

# Show the plot and suppress text output (<Axes: xlabel='Date'> above)
plt.show()
```



Practice

What are the mean daily returns for these four stocks?

```
tickers = 'AAPL IBM MSFT GOOG'

returns = (
    yf.download(tickers=tickers, auto_adjust=False, progress=False)
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
)

(
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
```

```
.iloc[:-1] # drop today, which is likely a partial-day return  
.mean() # calculate mean of daily returns from GOOG IPO through yesterday  
)
```

```
Ticker  
AAPL    0.0014  
GOOG    0.0010  
IBM     0.0004  
MSFT    0.0008  
dtype: float64
```

What are the standard deviations of daily returns for these four stocks?

```
(  
    returns # daily returns from 1962 through today  
    .dropna() # drop days with incomplete returns  
    .iloc[:-1] # drop today, which is likely a partial-day return  
    .std() # calculate standard deviation (volatility) of daily returns from GOOG IPO through yesterday  
)
```

```
Ticker  
AAPL    0.0204  
GOOG    0.0193  
IBM     0.0144  
MSFT    0.0170  
dtype: float64
```

What are the *annualized* means and standard deviations of daily returns for these four stocks?

```
ann_means = (  
    returns # daily returns from 1962 through today  
    .dropna() # drop days with missing returns  
    .iloc[:-1] # drop today, which is likely a partial-day return  
    .mean() # calculate mean of daily returns from close of GOOG IPO through yesterday  
    .mul(252) # means grow linearly with time, so multiply by 252  
)
```

```
ann_means
```

```
Ticker
AAPL    0.3577
GOOG    0.2541
IBM     0.1119
MSFT    0.1909
dtype: float64
```

```
ann_stds = (
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .std() # calculate mean of daily returns from close of GOOG IPO through yesterday
    .mul(np.sqrt(252)) # variances grow linearly with time, so standard deviations grow sqrt
)


ann_stds
```

```
Ticker
AAPL    0.3234
GOOG    0.3060
IBM     0.2282
MSFT    0.2693
dtype: float64
```

Plot annualized means versus standard deviations of daily returns for these four stocks

Here is a crude plot!

```
plt.scatter(x=ann_stds, y=ann_means)
```



But we can do better than a crude plot! We will typically combine data into a data frame to make plotting easier. Because `ann_stds` and `ann_means` are pandas' series, so we can use `pd.DataFrame` to combine them into a data frame.

```
df = pd.DataFrame({'Volatility': ann_stds, 'Mean Return': ann_means})  
df
```

Ticker	Volatility	Mean Return
AAPL	0.3234	0.3577
GOOG	0.3060	0.2541
IBM	0.2282	0.1119
MSFT	0.2693	0.1909

i Note

Below, we could use `enumerate()` instead of `.iterrows()`. However, `enumerate()` loops over *column names* instead of row indexes and contents. Therefore, with `enumerate()`, we would have to `.transpose()` our data frame, then use the tickers to slice the rows of our original data frame. Here `iterrows()` combines these several steps into one.


```
ax = df.plot(kind='scatter', x='Volatility', y='Mean Return')

for i, (v, mr) in df.iterrows():
    plt.text(x=v, y=mr, s=i)

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title('Return versus Risk for Tech Stocks')
plt.show()
```



Repeat the previous calculations and plot for the stocks in the Dow-Jones Industrial Index (DJIA)

We can find the current DJIA stocks on [Wikipedia](https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average). We must download new data, into `tickers_2`, `data_2`, and `returns_2`.

```
url_2 = 'https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average'  
wiki_2 = pd.read_html(io=url_2)
```

```
type(wiki_2)
```

```
list
```

```
tickers_2 = wiki_2[2]['Symbol'].to_list()
```

```
returns_2 = (
    yf.download(tickers=tickers_2, auto_adjust=False, progress=False)
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
    .dropna()
)
```

```
df_2 = pd.DataFrame({
    'Volatility': returns_2.std().mul(np.sqrt(252)),
    'Mean Return': returns_2.mean().mul(252),
})
```

```
dates_2 = returns_2.index
```

```
dates_2[-1]
```

```
Timestamp('2008-03-20 00:00:00')
```

```
dates_2[0]
```

```
Timestamp('2025-02-27 00:00:00')
```


```
ax = df_2.plot(kind='scatter', x='Volatility', y='Mean Return')

for i, (v, mr) in df_2.iterrows():
    plt.text(x=v, y=mr, s=i, fontdict={'fontsize': 'x-small'})

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title(f'Return versus Risk for DJIA Stocks\n from {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



We can use the seaborn package to add a best-fit line! More on seaborn here: <https://seaborn.pydata.org/index.html>

```
import seaborn as sns

ax = sns.regplot(
    data=df_2,
    x='Volatility',
    y='Mean Return'
)

for i, (v, mr) in df_2.iterrows():
    plt.text(x=v, y=mr, s=i, fontdict={'fontsize': 'x-small'})


ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

```

plt.ylabel('Annualized Mean of Daily Returns (%)')
plt.xlabel('Annualized Volatility of Daily Returns (%)')

plt.title(f'Return versus Risk for DJIA Stocks\n from {dates_2[0] :%B %Y} to {dates_2[-1] :%B %Y}')
plt.show()

```



Calculate total returns for the stocks in the DJIA

We can use the `.prod()` method to compound returns as $1 + R_T = \prod_{t=1}^T (1 + R_t)$. Technically, we should write R_T as $R_{0,T}$, but we typically omit the subscript 0.

In general, I prefer to do simple math on pandas objects (data frames and series) with methods instead of operators:

For example:

1. `.add(1)` instead of `+ 1`
2. `.sub(1)` instead of `- 1`
3. `.div(1)` instead of `/ 1`
4. `.mul(1)` instead of `* 1`

The advantage of methods over operators, is that we can easily chain methods without lots of parentheses.


```
total_returns_2 = returns_2.add(1).prod().sub(1)

ax = total_returns_2.sort_values().plot(kind='barh')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Total Returns for DJIA Stocks\n from {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Plot the distribution of total returns for the stocks in the DJIA

We can plot a histogram, using either the `plt.hist()` function or the `.plot(kind='hist')` method.


A histogram is a great way to visualize data!

```
ax = total_returns_2.plot(kind='hist')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Distribution of Total Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Which stocks have the minimum and maximum total returns?

If we want the *values*, the `.min()` and `.max()` methods are the way to go!

```
total_returns_2.min()
```

2.1510

```
total_returns_2.max()
```

295.7450

The `.min()` and `.max()` methods give the values but not the tickers (or index). We use the `.idxmin()` and `.idxmax()` to get the tickers (or index).

```
total_returns_2.idxmin()
```

```
'VZ'
```

```
total_returns_2.idxmax()
```

```
'NVDA'
```

Here is what I would use to capture values and tickers!

```
total_returns_2.sort_values().iloc[[0, -1]]
```

```
Ticker
VZ      2.1510
NVDA    295.7450
dtype: float64
```

Not the exactly right tool here, but the `.nsmallest()` and `.nlargest()` methods are really useful!

```
total_returns_2.nsmallest(3)
```

```
Ticker
VZ      2.1510
BA      2.2134
CVX     2.7258
dtype: float64
```

```
total_returns_2.nlargest(3)
```

```
Ticker
NVDA    295.7450
AAPL    59.8113
AMZN    58.4955
dtype: float64
```

Plot the cumulative returns for the stocks in the DJIA


We can use the cumulative product method `.cumprod()` to calculate the right hand side of the formula above.

```
ax = (
    returns_2
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Cumulative Return (%)')

plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```




Repeat the plot above with only the minimum and maximum total returns

```
ax = (
    returns_2
    [total_returns_2.sort_values().iloc[[0, -1]].index] # slice min and max total return stock
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

plt.ylabel('Cumulative Return (%)')

```
plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}\nplt.show()
```



McKinney Chapter 5 - Practice - Sec 04

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Please keep forming groups on Canvas > People > Projects. If you want a group with more than four students, please fill a group with four students, then email with the group number and size.
2. Please keep proposing and voting for students' choice topics [here](#).

Five-Minute Review

The pandas package makes it easy to manipulate panel data and we will use it all semester. Its name is an abbreviation of [panel data](#):

In statistics and econometrics, panel data and longitudinal data[1][2] are both multi-dimensional data involving measurements over time. Panel data is a subset of longitudinal data where observations are for the same subjects each time.

We will download some financial data from Yahoo! Finance to cover three important tools in pandas.

First, we can use the yfinance package to easily download stock data from Yahoo! Finance.

Note

Starting with version 0.2.51, the `yfinance` package changed the default behavior of the `auto_adjust` argument from `False` to `True`. By default, the `ya.download()` function now returns adjusted prices, without including the `Adj Close` column.

We prefer to work with raw data from Yahoo! Finance and explicitly calculate returns using the `Adj Close` column. Therefore, we will set `auto_adjust=False` in our `ya.download()` calls. See the [yfinance changelog](#) for release version 0.2.51.

Also, I will use the `progress=False` argument to improve the readability of the PDF and website I render from these notebooks.

```
df0 = ya.download(tickers='AAPL MSFT', auto_adjust=False, progress=False)
```

```
df0
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT	Low AAPL	Low MSFT	Open AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...
2025-02-24	247.1000	404.0000	247.1000	404.0000	248.8600	409.3700	244.4200	399.3200	244.9300
2025-02-25	247.0400	397.9000	247.0400	397.9000	250.0000	401.9200	244.9100	396.7000	248.0000
2025-02-26	240.3600	399.7300	240.3600	399.7300	244.9800	403.6000	239.1300	394.2500	244.3300
2025-02-27	237.3000	392.5300	237.3000	392.5300	242.4600	405.7400	237.0600	392.1700	239.4100
2025-02-28	241.8400	396.9900	241.8400	396.9900	242.0900	397.6300	230.2000	386.5700	236.9500

Second, we can slice rows and columns two ways: by integer locations with `.iloc[]` and by labels with `.loc[]`

```
df0.iloc[:6, :6]
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

```
df0.loc[:, '1980-12-19', : 'High']
```

Price Ticker Date	Adj Close AAPL	Close MSFT	Close AAPL	Close MSFT	High AAPL	High MSFT
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN
1980-12-19	0.0970	NaN	0.1261	NaN	0.1267	NaN

i Note

To slice a DataFrame:

- Use `['Name']` to select specific columns by their names.
- Use `.loc[]` to slice rows, or rows and columns together, with labels or conditional expressions.

```
df0['High']
```

Ticker Date	AAPL	MSFT
1980-12-12	0.1289	NaN
1980-12-15	0.1222	NaN
1980-12-16	0.1133	NaN

Ticker	AAPL	MSFT
Date		
1980-12-17	0.1161	NaN
1980-12-18	0.1194	NaN
...
2025-02-24	248.8600	409.3700
2025-02-25	250.0000	401.9200
2025-02-26	244.9800	403.6000
2025-02-27	242.4600	405.7400
2025-02-28	242.0900	397.6300


```
# # KeyError: '1980-12-12'
# df0['1980-12-12']
```

Note, if we use string labels, like dates and words, pandas includes left and right edges! This string label behavior differs from the integer location behavior everywhere else in Python. However, it is easy to figure out the sequence of integer locations. It is difficult to figure our the sequence of string labels.

Third, there many methods we can apply to pandas objects (and chain)! At this point in the course, our most common methods will be:

1. `.pct_change()` to calculate simple returns from adjusted close prices
2. `.plot()` to quickly plot pandas objects
3. `.mean()`, `.std()`, `.describe()`, etc. to calculate summary statistics

```
( df0  # DataFrame containing AAPL and MSFT data from 1980-12-12 through today
  .loc['1990':'2024', 'Adj Close']  # Slice rows for 1990-2024 (inclusive) and the 'Adj Cl
  .pct_change()  # Calculate daily percentage changes in 'Adj Close' (includes dividends an
  .add(1)  # Prepare for compounding by adding 1 to daily returns
  .cumprod()  # Compute cumulative product to get total return for each day since the start
  .sub(1)  # Convert back to cumulative returns
  .plot()  # Plot cumulative returns
)
```



The plot above is in decimal returns! pandas makes it easy to generate plots, but getting them beautiful and readable takes more work. The following code adds a title, labels, and formats the y axis.


```
from matplotlib.ticker import FuncFormatter

# Plot the data
ax = df0.loc['1990':'2024', 'Adj Close'].pct_change().add(1).cumprod().sub(1).plot()

# Format y-axis as percentages with comma separators
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:,.0f}'))

# Add labels and title if needed
plt.ylabel('Cumulative Return (%)')
plt.title('Cumulative Returns from 1990 to 2024')

# Show the plot and suppress text output (<Axes: xlabel='Date'> above)
plt.show()
```



Practice

What are the mean daily returns for these four stocks?

```
tickers = 'AAPL IBM MSFT GOOG'
```

```
returns = (
    yf.download(tickers=tickers, auto_adjust=False, progress=False)
    ['Adj Close']
    .pct_change()
)
```

```
( 
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
```

```
.mean() # calculate mean of daily returns from GOOG IPO through yesterday  
)
```

```
Ticker  
AAPL    0.0014  
GOOG    0.0010  
IBM     0.0004  
MSFT    0.0008  
dtype: float64
```

What are the standard deviations of daily returns for these four stocks?

```
(  
    returns # daily returns from 1962 through today  
    .dropna() # drop days with incomplete returns  
    .iloc[:-1] # drop today, which is likely a partial-day return  
    .std() # calculate standard deviation (volatility) of daily returns from GOOG IPO through  
)
```

```
Ticker  
AAPL    0.0204  
GOOG    0.0193  
IBM     0.0144  
MSFT    0.0170  
dtype: float64
```

What are the *annualized* means and standard deviations of daily returns for these four stocks?

```
ann_means = (  
    returns # daily returns from 1962 through today  
    .dropna() # drop days with missing returns  
    .iloc[:-1] # drop today, which is likely a partial-day return  
    .mean() # calculate mean of daily returns from close of GOOG IPO through yesterday  
    .mul(252) # means grow linearly with time, so multiply by 252  
)  
  
ann_means
```

```
Ticker
AAPL    0.3570
GOOG    0.2528
IBM     0.1113
MSFT    0.1900
dtype: float64
```

```
ann_stds = (
    returns # daily returns from 1962 through today
    .dropna() # drop days with incomplete returns
    .iloc[:-1] # drop today, which is likely a partial-day return
    .std() # calculate mean of daily returns from close of GOOG IPO through yesterday
    .mul(np.sqrt(252)) # variances grow linearly with time, so standard deviations grow sqrt
)


ann_stds
```

```
Ticker
AAPL    0.3234
GOOG    0.3061
IBM     0.2282
MSFT    0.2693
dtype: float64
```

Plot annualized means versus standard deviations of daily returns for these four stocks

Here is a crude plot!

```
plt.scatter(x=ann_stds, y=ann_means)
```



But we can do better than a crude plot! We will typically combine data into a data frame to make plotting easier. Because `ann_std` and `ann_means` are pandas' series, so we can use `pd.DataFrame` to combine them into a data frame.

```
df = pd.DataFrame({'Volatility': ann_stds, 'Mean Return': ann_means})  
df
```

Ticker	Volatility	Mean Return
AAPL	0.3234	0.3570
GOOG	0.3061	0.2528
IBM	0.2282	0.1113
MSFT	0.2693	0.1900

i Note

Below, we could use `enumerate()` instead of `.iterrows()`. However, `enumerate()` loops over *column names* instead of row indexes and contents. Therefore, with `enumerate()`, we would have to `.transpose()` our data frame, then use the tickers to slice the rows of our original data frame. Here `iterrows()` combines these several steps into one.


```
ax = df.plot(kind='scatter', x='Volatility', y='Mean Return')

for i, (v, mr) in df.iterrows():
    plt.text(s=i, x=v, y=mr)

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Annualized Volatility of Daily Returns (%)')
plt.ylabel('Annualized Mean of Daily Returns (%)')

plt.title('Returns versus Risk for Tech Stocks')
plt.show()
```



Repeat the previous calculations and plot for the stocks in the Dow-Jones Industrial Index (DJIA)

We can find the current DJIA stocks on [Wikipedia](https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average). We must download new data, into `tickers_2`, `data_2`, and `returns_2`.

```
url_2 = 'https://en.wikipedia.org/wiki/Dow_Jones_Industrial_Average'  
wiki_2 = pd.read_html(io=url_2)
```

```
type(wiki_2)
```

```
list
```

```
tickers_2 = wiki_2[2]['Symbol'].to_list()
```

```
returns_2 = (
    yf.download(tickers=tickers_2, auto_adjust=False, progress=False)
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
    .dropna()
)
```

```
df_2 = pd.DataFrame({
    'Volatility': returns_2.std().mul(np.sqrt(252)),
    'Mean Return': returns_2.mean().mul(252)
})
```

```
dates_2 = returns_2.index
```

```
dates_2[0]
```

```
Timestamp('2008-03-20 00:00:00')
```

```
dates_2[-1]
```

```
Timestamp('2025-02-27 00:00:00')
```


```
ax = df_2.plot(kind='scatter', x='Volatility', y='Mean Return')

for i, (v, mr) in df_2.iterrows():
    plt.text(s=i, x=v, y=mr)

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Annualized Volatility of Daily Returns (%)')
plt.ylabel('Annualized Mean of Daily Returns (%)')

plt.title(f'Returns versus Risk for DJIA Stocks\n from {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



We can use the seaborn package to add a best-fit line! More on seaborn here: <https://seaborn.pydata.org/index.html>

```
import seaborn as sns


ax = sns.regplot(data=df_2, x='Volatility', y='Mean Return')

for i, (v, mr) in df_2.iterrows():
    plt.text(s=i, x=v, y=mr)

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Annualized Volatility of Daily Returns (%)')
plt.ylabel('Annualized Mean of Daily Returns (%)')
```

```
plt.title(f'Returns versus Risk for DJIA Stocks\n from {dates_2[0] :%B %Y} to {dates_2[-1] :%B %Y}')
plt.show()
```



Calculate total returns for the stocks in the DJIA

We can use the `.prod()` method to compound returns as $1 + R_T = \prod_{t=1}^T (1 + R_t)$. Technically, we should write R_T as $R_{0,T}$, but we typically omit the subscript 0.

In general, I prefer to do simple math on pandas objects (data frames and series) with methods instead of operators:

For example:

1. `.add(1)` instead of `+ 1`
2. `.sub(1)` instead of `- 1`
3. `.div(1)` instead of `/ 1`

4. `.mul(1)` instead of `* 1`

The advantage of methods over operators, is that we can easily chain methods without lots of parentheses.

We can use the `.clipboard()` method to quickly move data to Excel!

```
returns_2.to_clipboard()
```


```
total_returns_2 = returns_2.add(1).prod().sub(1)

ax = total_returns_2.sort_values().plot(kind='barh')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Total Returns for DJIA Stocks\n from {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Plot the distribution of total returns for the stocks in the DJIA

We can plot a histogram, using either the `plt.hist()` function or the `.plot(kind='hist')` method.


A histogram is a great way to visualize data!

```
ax = total_returns_2.plot(kind='hist')

ax.xaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.xlabel('Total Return (%)')

plt.title(f'Distribution of Total Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```



Which stocks have the minimum and maximum total returns?

If we want the *values*, the `.min()` and `.max()` methods are the way to go!

```
total_returns_2.min()
```

2.1510

```
total_returns_2.max()
```

295.7450

The `.min()` and `.max()` methods give the values but not the tickers (or index). We use the `.idxmin()` and `.idxmax()` to get the tickers (or index).

```
total_returns_2.idxmin()
```

```
'VZ'
```

```
total_returns_2.idxmax()
```

```
'NVDA'
```

Here is what I would use to capture values and tickers!

```
total_returns_2.sort_values().iloc[[0, -1]]
```

```
Ticker
VZ      2.1510
NVDA    295.7450
dtype: float64
```

Not the exactly right tool here, but the `.nsmallest()` and `.nlargest()` methods are really useful!

```
total_returns_2.nsmallest(3)
```

```
Ticker
VZ      2.1510
BA      2.2134
CVX     2.7258
dtype: float64
```

```
total_returns_2.nlargest(3)
```

```
Ticker
NVDA    295.7450
AAPL    59.8113
AMZN    58.4955
dtype: float64
```

Plot the cumulative returns for the stocks in the DJIA


We can use the cumulative product method `.cumprod()` to calculate the right hand side of the formula above.

```
ax = (
    returns_2
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))

plt.ylabel('Cumulative Return (%)')

plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}')
plt.show()
```




Repeat the plot above with only the minimum and maximum total returns

```
ax = (
    returns_2
    [total_returns_2.sort_values().iloc[[0, -1]].index] # slice min and max total return stock
    .add(1)
    .cumprod()
    .sub(1)
    .plot(legend=False) # with 30 stocks, this legend is too big to be useful
)

ax.yaxis.set_major_formatter(FuncFormatter(lambda x, _: f'{x*100:.0f}'))
```

plt.ylabel('Cumulative Return (%)')

```
plt.title(f'Cumulative Returns for DJIA Stocks\nfrom {dates_2[0]:%B %Y} to {dates_2[-1]:%B %Y}\nplt.show()
```



Week 5

McKinney Chapter 8 - Data Wrangling: Join, Combine, and Reshape

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

Chapter 8 of McKinney (2022) introduces a few important pandas concepts:

1. Joining or merging is combining 2+ data frames on 1+ indexes or columns into 1 data frame
2. Reshaping is rearranging a data frame so it has fewer columns and more rows (wide to long) or more columns and fewer rows (long to wide)

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Hierarchical Indexing

We must learn about hierarchical indexing before we learn about combining and reshaping data. A hierarchical index has two or more levels. For example, we could index rows by ticker and date. Or we could index columns by variable and ticker. Hierarchical indexing helps us work with high-dimensional data in a low-dimensional form.

```
np.random.seed(42)
data = pd.Series(
    data=np.random.randn(9),
    index=[['a', 'a', 'a', 'b', 'b', 'c', 'c', 'd', 'd'],
           [1, 2, 3, 1, 3, 1, 2, 2, 3]
      ]
)

data
```

```
a  1    0.4967
   2   -0.1383
   3    0.6477
b  1    1.5230
   3   -0.2342
c  1   -0.2341
   2    1.5792
d  2    0.7674
   3   -0.4695
dtype: float64
```

We can index this series to subset it.

```
data['b']
```

```
1    1.5230
3   -0.2342
dtype: float64
```

```
data.loc['b']
```

```
1    1.5230
3   -0.2342
dtype: float64
```

```
data['b':'c']
```

```
b    1    1.5230
     3   -0.2342
c    1   -0.2341
     2    1.5792
dtype: float64
```

```
data.loc['b':'c']
```

```
b    1    1.5230
     3   -0.2342
c    1   -0.2341
     2    1.5792
dtype: float64
```

We can subset on the index inner level, too. Here, the : slices all values in the outer index, and the 2 slices the three values with 2 indexes.

```
data.loc[:, 2]
```

```
a   -0.1383
c    1.5792
d    0.7674
dtype: float64
```

Here, `data` has a stacked or long format. We have multiple observations for each outer index (letters) with different inner indexes (numbers). We can un-stack `data` to convert the inner index level to columns. Now, we have an unstacked or wide format.

```
data.unstack()
```

	1	2	3
a	0.4967	-0.1383	0.6477
b	1.5230	NaN	-0.2342
c	-0.2341	1.5792	NaN
d	NaN	0.7674	-0.4695

We can create a data frame with hierarchical indexes or multi-indexes on rows *and* columns.

```
frame = pd.DataFrame(
    data=np.arange(12).reshape((4, 3)),
    index=[['a', 'a', 'b', 'b'], [1, 2, 1, 2]],
    columns[['Ohio', 'Ohio', 'Colorado'], ['Green', 'Red', 'Green']])
frame
```

		Ohio	Green	Red
	a	1	0	1
	a	2	3	4
	b	1	6	7
	b	2	9	10

We can name these multi-indexes, but index names are not required.

```
frame.index.names = ['key1', 'key2']
frame.columns.names = ['state', 'color']

frame
```

		state		Ohio
		color		Green
	key1	key2		
	a	1	0	
	a	2	3	
	b	1	6	
	b	2	9	

Recall that `df[val]` selects the `val` column. Here, `frame` has a multi-index for the columns, so `frame['Ohio']` selects all columns with Ohio as the outer index.

```
frame['Ohio']
```

		color	Green
key1		key2	
a		1	0
		2	3
b		1	6
		2	9

We can pass a tuple if we only want one column.

```
frame[['Ohio', 'Green']]
```

		state	Ohio
key1		color	Green
a		1	0
		2	3
b		1	6
		2	9

We must do more work to slice the inner level of the column index.

```
frame.loc[:, (slice(None), 'Green')]
```

		state	Ohio
key1		color	Green
a		1	0
		2	3
b		1	6
		2	9

We can use pd.IndexSlice[:, 'Green'] an alternative to (slice(None), 'Green').

```
frame.loc[:, pd.IndexSlice[:, 'Green']]
```

		state	Ohio
		color	Green
key1		key2	
a		1	0
		2	3
b		1	6
		2	9

Reordering and Sorting Levels

We can swap index levels with the `.swaplevel()` method. The default arguments are `i=-2` and `j=-1`, which swap the two innermost index levels.

```
frame
```

		state	Ohio
		color	Green
key1		key2	
a		1	0
		2	3
b		1	6
		2	9

```
frame.swaplevel().sort_index()
```

		state	Ohio
		color	Green
key2		key1	
1		a	0
		b	6
2		a	3
		b	9

We can use index *names*, too.

```
frame.swaplevel('key1', 'key2').sort_index()
```

		state	Ohio
	color	Green	
key2		key1	
1		a	0
		b	6
2		a	3
		b	9

Indexing with a DataFrame's columns

We can convert a column into an index and an index into a column with the `.set_index()` and `.reset_index()` methods.

```
frame = pd.DataFrame({
    'a': range(7),
    'b': range(7, 0, -1),
    'c': ['one', 'one', 'one', 'two', 'two', 'two', 'two'],
    'd': [0, 1, 2, 0, 1, 2, 3]
})

frame
```

	a	b	c	d
0	0	7	one	0
1	1	6	one	1
2	2	5	one	2
3	3	4	two	0
4	4	3	two	1
5	5	2	two	2
6	6	1	two	3

The `.set_index()` method converts columns to indexes and drops these columns by default.

```
frame2 = frame.set_index(['c', 'd'])

frame2
```

		a	b
c	d		
one		0 0 7	
		1 1 6	
		2 2 5	
		0 3 4	
two		1 4 3	
		2 5 2	
		3 6 1	

The `.reset_index()` method drops indexes, adds them as columns by default, and sets an integer index.

```
frame2.reset_index()
```

	c	d	a	b
0	one	0	0	7
1	one	1	1	6
2	one	2	2	5
3	two	0	3	4
4	two	1	4	3
5	two	2	5	2
6	two	3	6	1

Combining and Merging Datasets

pandas provides several methods and functions to combine and merge data. We can typically create the same output with several these methods or functions, but one may be more efficient.

When we want to combine data frames with similar indexes, we will tend to use the `.join()` method. The `.join()` can also combine three or more data frames.

Otherwise, we will use the `.merge()` method or `pd.merge()` function. The `pd.merge()` function is more flexible than the `.join()` method, so we will start with the `pd.merge()` function.

The [pandas website](#) provides helpful visualizations.

Database-Style DataFrame Joins

Merge or join operations combine datasets by linking rows using one or more keys. These operations are central to relational databases (e.g., SQL-based). The merge function in pandas is the main entry point for using these algorithms on your data.

```
df1 = pd.DataFrame({'key': ['b', 'b', 'a', 'c', 'a', 'a', 'b'], 'data1': range(7)})
df2 = pd.DataFrame({'key': ['a', 'b', 'd'], 'data2': range(3)})
```

df1

	key	data1
0	b	0
1	b	1
2	a	2
3	c	3
4	a	4
5	a	5
6	b	6

df2

	key	data2
0	a	0
1	b	1
2	d	2

```
pd.merge(df1, df2)
```

	key	data1	data2
0	b	0	1
1	b	1	1
2	a	2	0
3	a	4	0
4	a	5	0
5	b	6	1

The default is `how='inner'`, so `pd.merge()` inner joins left and right data frames by default, keeping only rows that appear in both. We can specify `how='outer'`, so `pd.merge()` outer joins left and right data frames, keeping all rows that appear in either.

```
pd.merge(df1, df2, how='outer')
```

	key	data1	data2
0	a	2.0000	0.0000
1	a	4.0000	0.0000
2	a	5.0000	0.0000
3	b	0.0000	1.0000
4	b	1.0000	1.0000
5	b	6.0000	1.0000
6	c	3.0000	NaN
7	d	NaN	2.0000

A `how='left'` merge keeps only rows that appear in the left data frame.

```
pd.merge(df1, df2, how='left')
```

	key	data1	data2
0	b	0	1.0000
1	b	1	1.0000
2	a	2	0.0000
3	c	3	NaN
4	a	4	0.0000
5	a	5	0.0000
6	b	6	1.0000

A `how='right'` merge keeps only rows that appear in the right data frame.

```
pd.merge(df1, df2, how='right')
```

	key	data1	data2
0	a	2.0000	0
1	a	4.0000	0
2	a	5.0000	0

	key	data1	data2
3	b	0.0000	1
4	b	1.0000	1
5	b	6.0000	1
6	d	NaN	2

By default, `pd.merge()` merges on any columns that appear in both data frames.

`on` : label or list Column or index level names to join on. These must be found in both DataFrames. If `on` is None and not merging on indexes then this defaults to the intersection of the columns in both DataFrames.

Here, `key` is the only common column between `df1` and `df2`. We *should* specify `on='key'` to avoid unexpected results.

```
pd.merge(df1, df2, on='key')
```

	key	data1	data2
0	b	0	1
1	b	1	1
2	a	2	0
3	a	4	0
4	a	5	0
5	b	6	1

We *must* specify `left_on` and `right_on` if our left and right data frames do not have a common column.

```
df3 = pd.DataFrame({'lkey': ['b', 'b', 'a', 'c', 'a', 'a', 'b'], 'data1': range(7)})
df4 = pd.DataFrame({'rkey': ['a', 'b', 'd'], 'data2': range(3)})
```

```
df3
```

	lkey	data1
0	b	0
1	b	1
2	a	2
3	c	3

	lkey	data1
4	a	4
5	a	5
6	b	6

df4

	rkey	data2
0	a	0
1	b	1
2	d	2

```
# pd.merge(df3, df4) # this code fails/errors because there are not common columns
# MergeError: No common columns to perform merge on. Merge options: left_on=None, right_on=None
```

```
pd.merge(df3, df4, left_on='lkey', right_on='rkey')
```

	lkey	data1	rkey	data2
0	b	0	b	1
1	b	1	b	1
2	a	2	a	0
3	a	4	a	0
4	a	5	a	0
5	b	6	b	1

Here, `pd.merge()` drops row c from `df3` and row d from `df4` because `pd.merge()` *inner* joins by default. An inner join keeps the intersection of the left and right data frame keys. If we want to keep rows c and d, we can *outer* join `df3` and `df4` with `how='outer'`.

```
pd.merge(df3, df4, left_on='lkey', right_on='rkey', how='outer')
```

	lkey	data1	rkey	data2
0	a	2.0000	a	0.0000
1	a	4.0000	a	0.0000
2	a	5.0000	a	0.0000

	lkey	data1	rkey	data2
3	b	0.0000	b	1.0000
4	b	1.0000	b	1.0000
5	b	6.0000	b	1.0000
6	c	3.0000	NaN	NaN
7	NaN	NaN	d	2.0000

Many-to-many merges have well-defined, though not necessarily intuitive, behavior.

```
df1 = pd.DataFrame({'key': ['b', 'b', 'a', 'c', 'a', 'b'], 'data1': range(6)})
df2 = pd.DataFrame({'key': ['a', 'b', 'a', 'b', 'd'], 'data2': range(5)})
```

df1

	key	data1
0	b	0
1	b	1
2	a	2
3	c	3
4	a	4
5	b	5

df2

	key	data2
0	a	0
1	b	1
2	a	2
3	b	3
4	d	4

```
pd.merge(df1, df2, on='key')
```

	key	data1	data2
0	b	0	1
1	b	0	3

	key	data1	data2
2	b	1	1
3	b	1	3
4	a	2	0
5	a	2	2
6	a	4	0
7	a	4	2
8	b	5	1
9	b	5	3

Many-to-many joins form the Cartesian product of the rows. Since there were three `b` rows in the left DataFrame and two in the right one, there are six `b` rows in the result. The join method only affects the distinct key values appearing in the result.

Be careful with many-to-many joins! In finance, we do not expect many-to-many joins because we expect at least one of the data frames to have unique observations. *pandas will not warn us if we accidentally perform a many-to-many join instead of a one-to-one or many-to-one join.*

```
# pd.merge(df1, df2, on='key', validate='1:1')
# MergeError: Merge keys are not unique in either left or right dataset; not a one-to-one me
```

We can merge on more than one key. For example, we can merge two data sets on ticker-date pairs or industry-date pairs.

```
left = pd.DataFrame({'key1': ['foo', 'foo', 'bar'],
                     'key2': ['one', 'two', 'one'],
                     'lval': [1, 2, 3]})

right = pd.DataFrame({'key1': ['foo', 'foo', 'bar', 'bar'],
                      'key2': ['one', 'one', 'one', 'two'],
                      'rval': [4, 5, 6, 7]})
```

```
left
```

	key1	key2	lval
0	foo	one	1
1	foo	two	2
2	bar	one	3

```
right
```

	key1	key2	rval
0	foo	one	4
1	foo	one	5
2	bar	one	6
3	bar	two	7

```
pd.merge(left, right, on=['key1', 'key2'], how='outer')
```

	key1	key2	lval	rval
0	bar	one	3.0000	6.0000
1	bar	two	NaN	7.0000
2	foo	one	1.0000	4.0000
3	foo	one	1.0000	5.0000
4	foo	two	2.0000	NaN

When column names overlap between the left and right data frames, `pd.merge()` appends `_x` and `_y` to the left and right versions of the overlapping column names.

```
pd.merge(left, right, on='key1')
```

	key1	key2_x	lval	key2_y	rval
0	foo	one	1	one	4
1	foo	one	1	one	5
2	foo	two	2	one	4
3	foo	two	2	one	5
4	bar	one	3	one	6
5	bar	one	3	two	7

I typically specify the `suffixes` argument to avoid confusion.

```
pd.merge(left, right, on='key1', suffixes=('_left', '_right'))
```

	key1	key2_left	lval	key2_right	rval
0	foo	one	1	one	4
1	foo	one	1	one	5
2	foo	two	2	one	4
3	foo	two	2	one	5
4	bar	one	3	one	6
5	bar	one	3	two	7

I read the `pd.merge()` docstring frequently! **Table 8-2** summarizes the commonly used arguments for `pd.merge()`.

- `left`: DataFrame to be merged on the left side.
- `right`: DataFrame to be merged on the right side.
- `how`: One of ‘inner’, ‘outer’, ‘left’, or ‘right’; defaults to ‘inner’.
- `on`: Column names to join on. Must be found in both DataFrame objects. If not specified and no other join keys given will use the intersection of the column names in left and right as the join keys.
- `left_on`: Columns in left DataFrame to use as join keys.
- `right_on`: Analogous to `left_on` for left DataFrame.
- `left_index`: Use row index in left as its join key (or keys, if a MultiIndex).
- `right_index`: Analogous to `left_index`.
- `sort`: Sort merged data lexicographically by join keys; True by default (disable to get better performance in some cases on large datasets).
- `suffixes`: Tuple of string values to append to column names in case of overlap; defaults to (`_x`, `_y`) (e.g., if ‘data’ in both DataFrame objects, would appear as ‘data_x’ and ‘data_y’ in result).
- `copy`: If False, avoid copying data into resulting data structure in some exceptional cases; by default always copies.
- `indicator`: Adds a special column `_merge` that indicates the source of each row; values will be ‘left_only’, ‘right_only’, or ‘both’ based on the origin of the joined data in each row.

Merging on Index

If we want to use `pd.merge()` to join on row indexes, we can use the `left_index` and `right_index` arguments.

```
left1 = pd.DataFrame({'key': ['a', 'b', 'a', 'a', 'b', 'c'], 'value': range(6)})
right1 = pd.DataFrame({'group_val': [3.5, 7]}, index=['a', 'b'])
```

```
left1
```

	key	value
0	a	0
1	b	1
2	a	2
3	a	3
4	b	4
5	c	5

```
right1
```

	group_val
a	3.5000
b	7.0000

```
pd.merge(left1, right1, left_on='key', right_index=True, how='outer')
```

	key	value	group_val
0	a	0	3.5000
2	a	2	3.5000
3	a	3	3.5000
1	b	1	7.0000
4	b	4	7.0000
5	c	5	NaN

The index arguments work for hierarchical indexes (multi indexes), too.

```
lefth = pd.DataFrame({'key1': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada'],
                      'key2': [2000, 2001, 2002, 2001, 2002],
                      'data': np.arange(5.)})
righth = pd.DataFrame(np.arange(12).reshape((6, 2)),
                      index=[[('Nevada', 'Nevada'), ('Ohio', 'Ohio', 'Ohio'),
                              [2001, 2000, 2000, 2001, 2002]],
                             columns=['event1', 'event2'])
```

```
pd.merge(lefth, righth, left_on=['key1', 'key2'], right_index=True, how='outer')
```

	key1	key2	data	event1	event2
4	Nevada	2000	NaN	2.0000	3.0000
3	Nevada	2001	3.0000	0.0000	1.0000
4	Nevada	2002	4.0000	NaN	NaN
0	Ohio	2000	0.0000	4.0000	5.0000
0	Ohio	2000	0.0000	6.0000	7.0000
1	Ohio	2001	1.0000	8.0000	9.0000
2	Ohio	2002	2.0000	10.0000	11.0000

```
left2 = pd.DataFrame([[1., 2.], [3., 4.], [5., 6.]],
                     index=['a', 'c', 'e'],
                     columns=['Ohio', 'Nevada'])
right2 = pd.DataFrame([[7., 8.], [9., 10.], [11., 12.], [13., 14.]],
                      index=['b', 'c', 'd', 'e'],
                      columns=['Missouri', 'Alabama'])
```

If we use both indexes, `pd.merge()` will keep the index.

```
pd.merge(left2, right2, how='outer', left_index=True, right_index=True)
```

	Ohio	Nevada	Missouri	Alabama
a	1.0000	2.0000	NaN	NaN
b	NaN	NaN	7.0000	8.0000
c	3.0000	4.0000	9.0000	10.0000
d	NaN	NaN	11.0000	12.0000
e	5.0000	6.0000	13.0000	14.0000

DataFrame has a convenient join instance for merging by index. It can also be used to combine together many DataFrame objects having the same or similar indexes but non-overlapping columns.

We can use the `.join()` method if both data frames have similar indexes.

```
left2
```

	Ohio	Nevada
a	1.0000	2.0000
c	3.0000	4.0000
e	5.0000	6.0000

```
right2
```

	Missouri	Alabama
b	7.0000	8.0000
c	9.0000	10.0000
d	11.0000	12.0000
e	13.0000	14.0000

```
left2.join(right2, how='outer')
```

	Ohio	Nevada	Missouri	Alabama
a	1.0000	2.0000	NaN	NaN
b	NaN	NaN	7.0000	8.0000
c	3.0000	4.0000	9.0000	10.0000
d	NaN	NaN	11.0000	12.0000
e	5.0000	6.0000	13.0000	14.0000

The `.join()` method left joins by default. Because the `.join()` method uses indexes, it requires fewer arguments than `.merge()`. The `.join()` method can also accept a list of data frames.

```
another = pd.DataFrame(
    data=[[7., 8.], [9., 10.], [11., 12.], [16., 17.]],
    index=['a', 'c', 'e', 'f'],
    columns=['New York', 'Oregon']
)

another
```

	New York	Oregon
a	7.0000	8.0000
c	9.0000	10.0000
e	11.0000	12.0000
f	16.0000	17.0000

```
left2.join([right2, another])
```

	Ohio	Nevada	Missouri	Alabama	New York	Oregon
a	1.0000	2.0000	NaN	NaN	7.0000	8.0000
c	3.0000	4.0000	9.0000	10.0000	9.0000	10.0000
e	5.0000	6.0000	13.0000	14.0000	11.0000	12.0000

```
left2.join([right2, another], how='outer')
```

	Ohio	Nevada	Missouri	Alabama	New York	Oregon
a	1.0000	2.0000	NaN	NaN	7.0000	8.0000
c	3.0000	4.0000	9.0000	10.0000	9.0000	10.0000
e	5.0000	6.0000	13.0000	14.0000	11.0000	12.0000
b	NaN	NaN	7.0000	8.0000	NaN	NaN
d	NaN	NaN	11.0000	12.0000	NaN	NaN
f	NaN	NaN	NaN	NaN	16.0000	17.0000

Concatenating Along an Axis

The `pd.concat()` function provides a flexible way to combine data frames and series along an axis. I typically use `pd.concat()` to combine:

1. A list of data frames with similar layouts
2. A list of series because series do not have `.join()` or `.merge()` methods

```
s1 = pd.Series([0, 1], index=['a', 'b'])
s2 = pd.Series([2, 3, 4], index=['c', 'd', 'e'])
s3 = pd.Series([5, 6], index=['f', 'g'])
```

```
s1
```

```
a    0  
b    1  
dtype: int64
```

```
s2
```

```
c    2  
d    3  
e    4  
dtype: int64
```

```
s3
```

```
f    5  
g    6  
dtype: int64
```

```
pd.concat([s1, s2, s3]) # implicit axis=0
```

```
a    0  
b    1  
c    2  
d    3  
e    4  
f    5  
g    6  
dtype: int64
```

```
pd.concat([s1, s2, s3], axis=1) # explicit axis=1
```

	0	1	2
a	0.0000	NaN	NaN
b	1.0000	NaN	NaN
c	NaN	2.0000	NaN
d	NaN	3.0000	NaN
e	NaN	4.0000	NaN

	0	1	2
f	NaN	NaN	5.0000
g	NaN	NaN	6.0000

```
result = pd.concat([s1, s2, s3], keys=['one', 'two', 'three']) # implicit axis=0
result
```

```
one    a    0
      b    1
two    c    2
      d    3
      e    4
three   f    5
      g    6
dtype: int64
```

```
result.unstack(level=0)
```

	one	two	three
a	0.0000	NaN	NaN
b	1.0000	NaN	NaN
c	NaN	2.0000	NaN
d	NaN	3.0000	NaN
e	NaN	4.0000	NaN
f	NaN	NaN	5.0000
g	NaN	NaN	6.0000

```
pd.concat([s1, s2, s3], axis=1, keys=['one', 'two', 'three']) # explicit axis=1
```

	one	two	three
a	0.0000	NaN	NaN
b	1.0000	NaN	NaN
c	NaN	2.0000	NaN
d	NaN	3.0000	NaN
e	NaN	4.0000	NaN
f	NaN	NaN	5.0000
g	NaN	NaN	6.0000

```
df1 = pd.DataFrame(
    data=np.arange(6).reshape(3, 2),
    index=['a', 'b', 'c'],
    columns=['one', 'two']
)
df2 = pd.DataFrame(
    data=5 + np.arange(4).reshape(2, 2),
    index=['a', 'c'],
    columns=['three', 'four']
)

pd.concat([df1, df2], axis=1, keys=['level1', 'level2'])
```

	level1		level2	
	one	two	three	four
a	0	1	5.0000	6.0000
b	2	3	NaN	NaN
c	4	5	7.0000	8.0000

```
pd.concat([df1, df2], axis=1, keys=['level1', 'level2'], names=['upper', 'lower'])
```

	upper	level1		level2	
	lower	one	two	three	four
a	0	1	5.0000	6.0000	
b	2	3	NaN	NaN	
c	4	5	7.0000	8.0000	

Reshaping and Pivoting

Above, we briefly explore reshaping data with `.stack()` and `.unstack()`. Here, we more deeply explore reshaping data.

Reshaping with Hierarchical Indexing

Hierarchical indexes (multi-indexes) help reshape data.

There are two primary actions:

- stack: This “rotates” or pivots from the columns in the data to the rows
- unstack: This pivots from the rows into the columns

```
data = pd.DataFrame(np.arange(6).reshape((2, 3)),
                    index=pd.Index(['Ohio', 'Colorado'], name='state'),
                    columns=pd.Index(['one', 'two', 'three'],
                                   name='number'))
```

data

	number	one	two	three
state				
Ohio	0	1	2	
Colorado	3	4	5	

```
result = data.stack()
```

result

```
state      number
Ohio        one      0
              two      1
              three     2
Colorado    one      3
              two      4
              three     5
dtype: int64
```

```
result.unstack()
```

	number	one	two	three
state				
Ohio	0	1	2	
Colorado	3	4	5	

```
s1 = pd.Series([0, 1, 2, 3], index=['a', 'b', 'c', 'd'])
s2 = pd.Series([4, 5, 6], index=['c', 'd', 'e'])
data2 = pd.concat([s1, s2], keys=['one', 'two'])

data2
```

```
one  a    0  
     b    1  
     c    2  
     d    3  
two  c    4  
     d    5  
     e    6  
dtype: int64
```

Un-stacking may introduce missing values because data frames are rectangular.

```
data2.unstack()
```

	a	b	c	d	e
one	0.0000	1.0000	2.0000	3.0000	NaN
two	NaN	NaN	4.0000	5.0000	6.0000

Stacking drops these missing values by default. However, this behavior may change soon, so check your output!

```
data2.unstack().stack()
```

```
one  a    0.0000  
     b    1.0000  
     c    2.0000  
     d    3.0000  
two  c    4.0000  
     d    5.0000  
     e    6.0000  
dtype: float64
```

McKinney provides two more subsections on reshaping data with the `.pivot()` and `.melt()` methods. Unlike, the stacking methods, the pivoting methods can aggregate data and do not require an index. We will skip these additional aggregation methods for now.

McKinney Chapter 8 - Practice - Blank

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

Download data from Yahoo! Finance for BAC, C, GS, JPM, MS, and PNC and assign to data frame stocks_wide.

Reshape stocks_wide from wide to long with dates and tickers as row indexes and assign to data frame stocks_long.

Add daily returns to both stocks_wide and stocks_long under the name Returns.

Hint: Use pd.MultiIndex() to create a multi index for the wide data frame stocks_wide.

Download the daily benchmark return factors from Ken French's data library.

Hint: Use the `DataReader()` function in the `pandas-datareader` package. We imported this package above with the `pdr.` prefix.

Add the daily benchmark return factors to `stocks_wide` and `stocks_long`.

Write a function `download()` that accepts tickers and returns a wide data frame of returns with the daily benchmark return factors.

We can even add a `shape` argument to return a wide or long data frame!

Download earnings per share for the stocks in `stocks_long` and combine to a long data frame `earnings`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Combine `earnings` with the returns from `stocks_long`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Plot the relation between daily returns and earnings surprises

Repeat the earnings exercise with the S&P 100 stocks

With more data, we can more clearly see the positive relation between earnings surprises and returns!

McKinney Chapter 8 - Practice - Sec 02

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. *The deadline for forming project groups is Tuesday, 2/11.* That evening, I will create random project groups from the unassigned students.
2. *The deadline for proposing (and voting on) students' choice topics is Tuesday, 2/25.* That evening, I will finalize our schedule for the second half of the semester.

Five-Minute Review

Chapter 8 of McKinney covers 3 important topics.

1. **Hierarchical Indexing:** Hierarchical indexes (or multi-indexes) organize data at multiple levels instead of just a flat, two-dimensional structure. They help us work with high-dimensional data in a low-dimensional form. For example, we can index rows by multiple levels like `Ticker` and `Date`, or columns by `Variable` and `Ticker`.
2. **Combining Data:** We will use three functions and methods to combine datasets on one or more keys. All three offer `inner`, `outer`, `left`, or `right` combinations.
 1. The `pd.merge()` function (or the `.merge()` method) provides the most flexible way to perform database-style joins on data frames.
 2. The `.join()` method combines data frames with similar indexes.

3. The `pd.concat()` function combines similarly-shaped series and data frames.
3. **Reshaping Data:** We can reshape data to change its structure, such as pivoting from wide to long format or vice versa. We will most often use the `.stack()` and `.unstack()` methods, which pivot columns to rows and rows to columns, respectively. Later in the course we will learn about the `.pivot()` method for aggregating data and the `.melt()` method for more advanced reshaping.

Practice

Download data from Yahoo! Finance for BAC, C, GS, JPM, MS, and PNC and assign to data frame `stocks_wide`.

```
stocks_wide = yf.download(tickers='BAC, C, GS, JPM, MS, PNC', auto_adjust=False, progress=False)
```

```
stocks_wide.tail()
```

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	261.3400
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	257.4000
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	258.7900
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	259.0500
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	264.6500

Side Note: I do not know of a good variable inspector for JupyterLab. When I need to visually interact with a data frame, I use the `.to_clipboard()` method to copy it to my clipboard, and then I paste it into Excel.

```
stocks_wide.to_clipboard()
```

I can quickly move from Excel to pandas with the `pd.read_clipboard()` function.

Reshape stocks_wide from wide to long with dates and tickers as row indexes and assign to data frame stocks_long.

We use the `.stack()` method to go from wider to longer, and the `.unstack()` method to go from long longer to wider. Note that we set `future_stack=True` to accept the future default arguments for `.stack()` and suppress the FutureWarning. A FutureWarning is not an error, just a warning about some expected change that could cause an error in the future.

```
stocks_long = stocks_wide.stack(future_stack=True)
```

```
stocks_long.tail()
```

Date	Price	Adj C
	Ticker	
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

Add daily returns to both stocks_wide and stocks_long under the name Returns.

Hint: Use `pd.MultiIndex()` to create a multi index for the wide data frame `stocks_wide`.

We can use create a multi-index from all the combinations of `['Returns']` and the tickers in `stocks_wide['Adj Close']`! This approach is much easier than manually creating all the variable-ticker pairs we would need, even with only size tickers!

We assign this multi-index to the variable `_` (underscore). We only assign and use `_` in the same code cell, so we do not have to worry about accidentally using it elsewhere in the notebook.

```
_ = pd.MultiIndex.from_product([['Returns'], stocks_wide['Adj Close'].columns])

stocks_wide[_] = (
    stocks_wide
    ['Adj Close']
    .iloc[:-1] # do not use mid-day Adj Close for returns calculation
    .pct_change()
)
```

```
stocks_wide.tail()
```

Price Ticker Date	Adj Close						Close					
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM		
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	261.3400		
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	257.4000		
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	258.7900		
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	259.0500		
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	264.6500		

To add returns to `stocks_long` we have two options. I prefer the first option, but I will present the second option to show an application of the `.join()` method. I will assign the results of these two options to `stocks_long_1` and `stocks_long_2` so we can keep the original `stocks_long` as-is.

Option 1: Re-use `stocks_wide`!

```
stocks_long_1 = stocks_wide.stack(future_stack=True)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_1.tail(12)
```

Date	Price Ticker	Adj C
	BAC	44.1200
	C	78.8700
	GS	605.0000
	JPM	259.0500
	MS	129.2400
	PNC	188.5600
2025-02-27	BAC	46.1000
	C	79.9500
	GS	622.2900
	JPM	264.6500
2025-02-28	MS	133.1100
	PNC	191.9200

Date	Price	Adj C
	Ticker	

Option 2: Use `.join()`!

Here we will use `_` as a temporary variable for our long data frame with daily returns.

```
_ = (
    stocks_wide
    ['Adj Close']
    .iloc[:-1]
    .pct_change()
    .stack(future_stack=True)
    .to_frame('Returns')
)

stocks_long_2 = stocks_long.join(_)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
	BAC	44.120
	C	78.870
	GS	605.00
2025-02-27	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11
	PNC	191.92

We can test the equality of `stocks_long_1` and `stocks_long_2` most easily with the `.equals()` method.

```
stocks_long_1.equals(stocks_long_2)
```

True

Download the daily benchmark return factors from Ken French's data library.

Hint: Use the `DataReader()` function in the `pandas-datareader` package. We imported this package above with the `pdr.` prefix.

I often cannot remember the exact name for the daily factors. We can use the `pdr.famafrench.get_available_datasets()` to list all the data in Kenneth French's data library.

```
pdr.famafrench.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

I set `start=1900` to keep all available data. Otherwise, `pdr.DataReader()` keeps only the most recent five years in any data set.

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors_daily',
    data_source='famafrench',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_14108\875599436.py:1: FutureWarning: The argument
```

```
ff = pdr.DataReader(
```

```
type(ff)
```

```
dict
```

The daily factors only have one data frame (in the 0 key) and the data set description (in the `DESCR` key).

```
ff.keys()
```

```
dict_keys([0, 'DESCR'])
```

i Note

Data from the Kenneth French data library are *percent* returns instead of *decimal* returns!

```
ff[0]
```

Date	Mkt-RF	SMB	HML	RF
1926-07-01	0.1000	-0.2500	-0.2700	0.0090
1926-07-02	0.4500	-0.3300	-0.0600	0.0090
1926-07-06	0.1700	0.3000	-0.3900	0.0090
1926-07-07	0.0900	-0.5800	0.0200	0.0090
1926-07-08	0.2100	-0.3800	0.1900	0.0090
...
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
print(ff['DESCR'])
```

```
F-F Research Data Factors daily
```

```
This file was created by CMPT_ME_BEME_RET_DAILY using the 202412 CRSP database. The Tbill r
```

```
0 : (25901 rows x 4 cols)
```

Add the daily benchmark return factors to stocks_wide and stocks_long.

Since both `ff[0]` and `stocks_long_2` have date indexes, we can easily combine them with the `.join()` method.

```
ff[0].tail()
```

	Mkt-RF	SMB	HML	RF
Date				
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
	BAC	44.120
	C	78.870
	GS	605.00
2025-02-27	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11
	PNC	191.92

We can quickly combine `stocks_long_2` and `ff[0]` because both have indexes with daily dates named `Date`. Two notes:

1. The `.join()` method left joins by default, so the combined output has only dates in `stocks_long_2`
2. Kenneth French provides *percent* returns, so we divide them by 100 to convert them to *decimal* returns to match our Yahoo! Finance data

```
stocks_long_2.join(ff[0].div(100))
```

Date	Ticker	Adj C
1973-02-21	BAC	1.5426
	C	NaN
	GS	NaN
	JPM	NaN
	MS	NaN
...
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

We could instead convert the Yahoo! Finance *decimal* returns to *percent* returns. I do not have a strong preference on all decimal returns or all percent returns, but all returns should have the same form.

```
(stocks_long_2
    .assign(Returns=lambda x: 100 * x['Returns'])
    .join(ff[0])
)
```

Date	Ticker	Adj C
1973-02-21	BAC	1.5426
	C	NaN
	GS	NaN
	JPM	NaN
	MS	NaN
...
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

With `stocks_wide`, we have to do a little more work because of its column multi-index! We will use the `pd.MultiIndex.from_product()` trick from above.

```
_ = pd.MultiIndex.from_product([['Factors'], ff[0].columns])
stocks_wide[_] = ff[0].div(100)

stocks_wide.loc[:'2024'].tail()
```

Price Ticker Date	Adj Close					Close				
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
2024-12-24	44.3800	70.5117	579.9144	241.0650	126.2201	192.4933	44.3800	71.0000	582.7900	24
2024-12-26	44.5500	70.8593	578.3621	241.8907	127.1837	193.1777	44.5500	71.3500	581.2300	24
2024-12-27	44.3400	70.5117	573.3370	239.9308	125.9221	191.6999	44.3400	71.0000	576.1800	24
2024-12-30	43.9100	69.9059	570.7200	238.0904	124.9188	190.9560	43.9100	70.3900	573.5500	23
2024-12-31	43.9500	69.9059	569.7946	238.4783	124.8890	191.2734	43.9500	70.3900	572.6200	23

Write a function `download()` that accepts tickers and returns a wide data frame of returns with the daily benchmark return factors.

We can even add a `shape` argument to return a wide or long data frame!

```
import warnings

def download(tickers, shape='wide'):
    """
    Download stock price data and Fama-French factors, returning in either 'wide' or 'long' format.

    Parameters:
    - tickers (str or list of str): Stock ticker(s) to download.
    - shape (str): Output format, either 'wide' (default) or 'long'.

    Returns:
    - pd.DataFrame: A DataFrame containing stock prices, returns, and Fama-French factors.
    """

    # shape must be wide or long
    if shape not in ['wide', 'long']:
```

```

raise ValueError('Invalid shape: must be "wide" or "long".')

# Download stock data
stocks = yf.download(tickers=tickers, auto_adjust=False, progress=False)

# Download Fama-French factors
# (suppressing FutureWarning for 'date_parser')
with warnings.catch_warnings():
    warnings.simplefilter('ignore', category=FutureWarning)
    factors = pdr.DataReader(
        name='F-F_Research_Data_Factors_daily',
        data_source='famafrench',
        start='1900'
    )[0].div(100) # Convert percentages to decimals

# Multi-index case
if isinstance(stocks.columns, pd.MultiIndex):
    # Compute daily returns
    _ = pd.MultiIndex.from_product([['Returns'], stocks['Adj Close'].columns])
    stocks[_] = stocks['Adj Close'].pct_change()

    if shape == 'wide':
        # Add factors with multi-index
        _ = pd.MultiIndex.from_product([['Factors'], factors.columns])
        stocks[_] = factors
        return stocks

    # Convert to long format then add factors
else:
    return stocks.stack(future_stack=True).join(factors)

# Single index case
# (redundant with recent versions of yfinance that always return a multi-index)
stocks['Returns'] = stocks['Adj Close'].pct_change()
return stocks.join(factors)

download(tickers='AAPL TSLA')

```

Price Ticker Date	Adj Close AAPL	Close TSLA	Close AAPL	Close TSLA	High AAPL	High TSLA	Low AAPL	Low TSLA	Open AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...
2025-02-24	247.1000	330.5300	247.1000	330.5300	248.8600	342.4000	244.4200	324.7000	244.9300
2025-02-25	247.0400	302.8000	247.0400	302.8000	250.0000	328.8900	244.9100	297.2500	248.0000
2025-02-26	240.3600	290.8000	240.3600	290.8000	244.9800	309.0000	239.1300	288.0400	244.3300
2025-02-27	237.3000	281.9500	237.3000	281.9500	242.4600	297.2300	237.0600	280.8800	239.4100
2025-02-28	241.8400	292.9800	241.8400	292.9800	242.0900	293.8800	230.2000	273.6000	236.9500

i Note

The `yfinance` package is a powerful tool for downloading market data, financial statements, and analyst estimates from Yahoo! Finance.

However, because `yfinance` relies on Yahoo! Finance's API, changes to the API can disrupt its functionality.

Recently, Yahoo! Finance changed API access to earnings forecasts and announcement dates, so we **cannot complete the earnings announcement exercise I had planned.**

Instead, I will prepare an alternative set of exercises for us to work on in class on Friday. Thank you for your flexibility!

Download earnings per share for the stocks in `stocks_long` and combine to a long data frame `earnings`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Combine earnings with the returns from `stocks_long`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`.

Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Plot the relation between daily returns and earnings surprises

Repeat the earnings exercise with the S&P 100 stocks

With more data, we can more clearly see the positive relation between earnings surprises and returns!

McKinney Chapter 8 - Practice - Sec 03

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. *The deadline for forming project groups is Tuesday, 2/11.* That evening, I will create random project groups from the unassigned students.
2. *The deadline for proposing (and voting on) students' choice topics is Tuesday, 2/25.* That evening, I will finalize our schedule for the second half of the semester.

Five-Minute Review

Chapter 8 of McKinney covers 3 important topics.

1. **Hierarchical Indexing:** Hierarchical indexes (or multi-indexes) organize data at multiple levels instead of just a flat, two-dimensional structure. They help us work with high-dimensional data in a low-dimensional form. For example, we can index rows by multiple levels like `Ticker` and `Date`, or columns by `Variable` and `Ticker`.
2. **Combining Data:** We will use three functions and methods to combine datasets on one or more keys. All three offer `inner`, `outer`, `left`, or `right` combinations.
 1. The `pd.merge()` function (or the `.merge()` method) provides the most flexible way to perform database-style joins on data frames.
 2. The `.join()` method combines data frames with similar indexes.

3. The `pd.concat()` function combines similarly-shaped series and data frames.
3. **Reshaping Data:** We can reshape data to change its structure, such as pivoting from wide to long format or vice versa. We will most often use the `.stack()` and `.unstack()` methods, which pivot columns to rows and rows to columns, respectively. Later in the course we will learn about the `.pivot()` method for aggregating data and the `.melt()` method for more advanced reshaping.

Practice

Download data from Yahoo! Finance for BAC, C, GS, JPM, MS, and PNC and assign to data frame `stocks_wide`.

```
stocks_wide = yf.download(tickers='BAC, C, GS, JPM, MS, PNC', auto_adjust=False, progress=False)
```

```
stocks_wide.tail()
```

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	261.3400
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	257.4000
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	258.7900
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	259.0500
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	264.6500

Reshape `stocks_wide` from wide to long with dates and tickers as row indexes and assign to data frame `stocks_long`.

We use the `.stack()` method to go from wider to longer, and the `.unstack()` method to go from long longer to wider. Note that we set `future_stack=True` to accept the future default arguments for `.stack()` and suppress the `FutureWarning`. A `FutureWarning` is not an error, just a warning about some expected change that could cause an error in the future.

```
stocks_long = stocks_wide.stack(future_stack=True)
```

```
stocks_long.tail()
```

Date	Price	Adj C
	Ticker	
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

i Note

The `.melt()` methods can reshape data frames from wide to long. However, our data frame has a column multi-index, which makes `.melt()` difficult to use and `.stack()` a better option.

Add daily returns to both `stocks_wide` and `stocks_long` under the name `Returns`.

Hint: Use `pd.MultiIndex()` to create a multi index for the wide data frame `stocks_wide`.

```
stocks_wide['Adj Close'].columns
```

```
Index(['BAC', 'C', 'GS', 'JPM', 'MS', 'PNC'], dtype='object', name='Ticker')
```

```
_ = pd.MultiIndex.from_product([['Returns'], stocks_wide['Adj Close'].columns])

stocks_wide[_] = (
    stocks_wide
    ['Adj Close']
    .iloc[:-1] # do not use mid-day Adj Close for returns calculation
    .pct_change()
)
```

```
stocks_wide.tail()
```

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JP
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	26
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	25
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	25
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	25
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	26

To add returns to `stocks_long` we have two options. I prefer the first option, but I will present the second option to show an application of the `.join()` method. I will assign the results of these two options to `stocks_long_1` and `stocks_long_2` so we can keep the original `stocks_long` as-is.

Option 1: Make `stocks_wide` long!

```
stocks_long_1 = stocks_wide.stack(future_stack=True)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_1.tail(12)
```

Date	Price Ticker	Adj C
	BAC	44.1200
	C	78.8700
	GS	605.00
2025-02-27	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.1000
	C	79.9500
2025-02-28	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

Option 2: Calculate returns from `stocks_wide`, make them long, then `.join()` them to `stocks_long`!

```
_ = stocks_wide['Adj Close'].iloc[:-1].pct_change().stack().to_frame('Returns')

stocks_long_2 = stocks_long.join(_)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
2025-02-27	BAC	44.120
	C	78.870
	GS	605.00
	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
2025-02-28	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

We can test the equality of `stocks_long_1` and `stocks_long_2` most easily with the `.equals()` method.

```
stocks_long_1.equals(stocks_long_2)
```

True

Download the daily benchmark return factors from Ken French's data library.

Hint: Use the `DataReader()` function in the `pandas-datareader` package. We imported this package above with the `pdr.` prefix.

I often cannot remember the exact name for the daily factors. We can use the `pdr.famafrench.get_available_datasets()` to list all the data in Kenneth French's data library.

```
pdr.famafrance.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors_daily',
    data_source='famafrance',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_15448\875599436.py:1: FutureWarning: The argument
```

```
    ff = pdr.DataReader(
```

```
type(ff)
```

```
dict
```

The daily factors only have one data frame (in the 0 key) and the data set description (in the DESCR key).

```
ff.keys()
```

```
dict_keys([0, 'DESCR'])
```

 Note

Data from the Kenneth French data library are *percent* returns instead of *decimal* returns!

```
ff[0]
```

	Mkt-RF	SMB	HML	RF
Date				
1926-07-01	0.1000	-0.2500	-0.2700	0.0090
1926-07-02	0.4500	-0.3300	-0.0600	0.0090
1926-07-06	0.1700	0.3000	-0.3900	0.0090
1926-07-07	0.0900	-0.5800	0.0200	0.0090
1926-07-08	0.2100	-0.3800	0.1900	0.0090
...
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
print(ff['DESCR'])
```

F-F Research Data Factors daily

This file was created by CMPT_ME_BEME_RET_DAILY using the 202412 CRSP database. The Tbill re

0 : (25901 rows x 4 cols)

Add the daily benchmark return factors to stocks_wide and stocks_long.

Since both ff[0] and stocks_long_2 have date indexes, we can easily combine them with the .join() method.

```
ff[0].tail()
```

	Mkt-RF	SMB	HML	RF
Date				
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
	BAC	44.120
	C	78.870
	GS	605.00
2025-02-27	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
2025-02-28	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

We can quickly combine `stocks_long_2` and `ff[0]` because both have indexes with daily dates named `Date`. Two notes:

1. The `.join()` method left joins by default, so the combined output has only dates in `stocks_long_2`
2. Kenneth French provides *percent* returns, so we divide them by 100 to convert them to *decimal* returns to match our Yahoo! Finance data

```
stocks_long_2.join(ff[0].div(100))
```

Date	Ticker	Adj C
	BAC	1.5426
	C	NaN
1973-02-21	GS	NaN
	JPM	NaN
	MS	NaN
...
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11

Date	Ticker	Adj C
	PNC	191.92

We could instead convert the Yahoo! Finance *decimal* returns to *percent* returns. I do not have a strong preference on all decimal returns or all percent returns, but all returns should have the same form.

```
(  
    stocks_long_2  
    .assign(Returns=lambda x: 100 * x['Returns'])  
    .join(ff[0])  
)
```

Date	Ticker	Adj C
	BAC	1.5426
	C	NaN
1973-02-21	GS	NaN
	JPM	NaN
	MS	NaN
...
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11
	PNC	191.92

With `stocks_wide`, we have to do a little more work because of its column multi-index! We will use the `pd.MultiIndex.from_product()` trick from above.

```
_ = pd.MultiIndex.from_product([['Factors'], ff[0].columns])  
stocks_wide[_] = ff[0].div(100)
```

```
stocks_wide.loc[:'2024'].tail()
```

Price Ticker Date	Adj Close						Close					
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JP		
2024-12-24	44.3800	70.5117	579.9144	241.0650	126.2201	192.4933	44.3800	71.0000	582.7900	24		
2024-12-26	44.5500	70.8593	578.3621	241.8907	127.1837	193.1777	44.5500	71.3500	581.2300	24		
2024-12-27	44.3400	70.5117	573.3370	239.9308	125.9221	191.6999	44.3400	71.0000	576.1800	24		
2024-12-30	43.9100	69.9059	570.7200	238.0904	124.9188	190.9560	43.9100	70.3900	573.5500	23		
2024-12-31	43.9500	69.9059	569.7946	238.4783	124.8890	191.2734	43.9500	70.3900	572.6200	23		

Write a function download() that accepts tickers and returns a wide data frame of returns with the daily benchmark return factors.

We can even add a `shape` argument to return a wide or long data frame!

We can even add a `shape` argument to return a wide or long data frame!

```
import warnings

def download(tickers, shape='wide'):
    """
    Download stock price data and Fama-French factors, returning in either 'wide' or 'long' format.

    Parameters:
    - tickers (str or list of str): Stock ticker(s) to download.
    - shape (str): Output format, either 'wide' (default) or 'long'.

    Returns:
    - pd.DataFrame: A DataFrame containing stock prices, returns, and Fama-French factors.
    """

    # shape must be wide or long
    if shape not in ['wide', 'long']:
        raise ValueError('Invalid shape: must be "wide" or "long".')
    
    # Download stock data
    stocks = yf.download(tickers=tickers, auto_adjust=False, progress=False)

    # Download Fama-French factors
```

```

# (suppressing FutureWarning for 'date_parser')
with warnings.catch_warnings():
    warnings.simplefilter('ignore', category=FutureWarning)
factors = pdr.DataReader(
    name='F-F_Research_Data_Factors_daily',
    data_source='famafrench',
    start='1900'
)[0].div(100) # Convert percentages to decimals

# Multi-index case
if isinstance(stocks.columns, pd.MultiIndex):
    # Compute daily returns
    _ = pd.MultiIndex.from_product([['Returns'], stocks['Adj Close'].columns])
    stocks[_] = stocks['Adj Close'].pct_change()

    if shape == 'wide':
        # Add factors with multi-index
        _ = pd.MultiIndex.from_product([['Factors'], factors.columns])
        stocks[_] = factors
        return stocks

    # Convert to long format then add factors
else:
    return stocks.stack(future_stack=True).join(factors)

# Single index case
# (redundant with recent versions of yfinance that always return a multi-index)
stocks['Returns'] = stocks['Adj Close'].pct_change()
return stocks.join(factors)

```

```
download(tickers='AAPL TSLA')
```

Price Ticker Date	Adj Close		Close		High		Low		Open
	AAPL	TSLA	AAPL	TSLA	AAPL	TSLA	AAPL	TSLA	AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...

Price Ticker Date	Adj Close AAPL	Close TSLA	Close AAPL	Close TSLA	High AAPL	High TSLA	Low AAPL	Low TSLA	Open AAPL
2025-02-24	247.1000	330.5300	247.1000	330.5300	248.8600	342.4000	244.4200	324.7000	244.9300
2025-02-25	247.0400	302.8000	247.0400	302.8000	250.0000	328.8900	244.9100	297.2500	248.0000
2025-02-26	240.3600	290.8000	240.3600	290.8000	244.9800	309.0000	239.1300	288.0400	244.3300
2025-02-27	237.3000	281.9500	237.3000	281.9500	242.4600	297.2300	237.0600	280.8800	239.4100
2025-02-28	241.8400	292.9800	241.8400	292.9800	242.0900	293.8800	230.2000	273.6000	236.9500

i Note

The `yfinance` package is a powerful tool for downloading market data, financial statements, and analyst estimates from Yahoo! Finance.

However, because `yfinance` relies on Yahoo! Finance's API, changes to the API can disrupt its functionality.

Recently, Yahoo! Finance changed API access to earnings forecasts and announcement dates, so we **cannot complete the earnings announcement exercise I had planned.**

Instead, I will prepare an alternative set of exercises for us to work on in class on Friday. Thank you for your flexibility!

Download earnings per share for the stocks in `stocks_long` and combine to a long data frame `earnings`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Combine `earnings` with the returns from `stocks_long`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Plot the relation between daily returns and earnings surprises

Repeat the earnings exercise with the S&P 100 stocks

With more data, we can more clearly see the positive relation between earnings surprises and returns!

McKinney Chapter 8 - Practice - Sec 04

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. *The deadline for forming project groups is Tuesday, 2/11.* That evening, I will create random project groups from the unassigned students.
2. *The deadline for proposing (and voting on) students' choice topics is Tuesday, 2/25.* That evening, I will finalize our schedule for the second half of the semester.

Five-Minute Review

Chapter 8 of McKinney covers 3 important topics.

1. **Hierarchical Indexing:** Hierarchical indexes (or multi-indexes) organize data at multiple levels instead of just a flat, two-dimensional structure. They help us work with high-dimensional data in a low-dimensional form. For example, we can index rows by multiple levels like `Ticker` and `Date`, or columns by `Variable` and `Ticker`.
2. **Combining Data:** We will use three functions and methods to combine datasets on one or more keys. All three offer `inner`, `outer`, `left`, or `right` combinations.
 1. The `pd.merge()` function (or the `.merge()` method) provides the most flexible way to perform database-style joins on data frames.
 2. The `.join()` method combines data frames with similar indexes.

3. The `pd.concat()` function combines similarly-shaped series and data frames.
3. **Reshaping Data:** We can reshape data to change its structure, such as pivoting from wide to long format or vice versa. We will most often use the `.stack()` and `.unstack()` methods, which pivot columns to rows and rows to columns, respectively. Later in the course we will learn about the `.pivot()` method for aggregating data and the `.melt()` method for more advanced reshaping.

Practice

Download data from Yahoo! Finance for BAC, C, GS, JPM, MS, and PNC and assign to data frame stocks_wide.

```
stocks_wide = yf.download(tickers='BAC, C, GS, JPM, MS, PNC', auto_adjust=False, progress=False)
```

```
stocks_wide.tail()
```

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	261.3400
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	257.4000
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	258.7900
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	259.0500
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	264.6500

Reshape stocks_wide from wide to long with dates and tickers as row indexes and assign to data frame stocks_long.

We use the `.stack()` method to go from wider to longer, and the `.unstack()` method to go from long longer to wider. Note that we set `future_stack=True` to accept the future default arguments for `.stack()` and suppress the `FutureWarning`. A `FutureWarning` is not an error, just a warning about some expected change that could cause an error in the future.

```
stocks_long = stocks_wide.stack(future_stack=True)
```

```
stocks_long.tail()
```

Date	Price	Adj C
	Ticker	
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

Add daily returns to both stocks_wide and stocks_long under the name Returns.

Hint: Use pd.MultiIndex() to create a multi index for the wide data frame stocks_wide.

```
stocks_wide['Adj Close'].columns
```

```
Index(['BAC', 'C', 'GS', 'JPM', 'MS', 'PNC'], dtype='object', name='Ticker')
```

```
_ = pd.MultiIndex.from_product([['Returns'], stocks_wide['Adj Close'].columns])

stocks_wide[_] = (
    stocks_wide
    ['Adj Close']
    .iloc[:-1] # do not use mid-day Adj Close for returns calculation
    .pct_change()
)
```

```
stocks_wide.tail()
```

Price	Adj Close						Close			
Ticker	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
Date										
2025-02-24	44.4600	78.5400	623.0505	261.3400	129.9700	186.9600	44.4600	78.5400	626.1400	261.3400
2025-02-25	43.9400	78.1400	611.8759	257.4000	129.6000	186.4700	43.9400	78.1400	614.9100	257.4000
2025-02-26	43.9400	79.0700	614.7218	258.7900	131.0500	187.0200	43.9400	79.0700	617.7700	258.7900
2025-02-27	44.1200	78.8700	605.0000	259.0500	129.2400	188.5600	44.1200	78.8700	608.0000	259.0500

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JP
2025-02-28	46.1000	79.9500	622.2900	264.6500	133.1100	191.9200	46.1000	79.9500	622.2900	26

To add returns to `stocks_long` we have two options. I prefer the first option, but I will present the second option to show an application of the `.join()` method. I will assign the results of these two options to `stocks_long_1` and `stocks_long_2` so we can keep the original `stocks_long` as-is.

Option 1: Stack `stocks_wide`!

```
stocks_long_1 = stocks_wide.stack(future_stack=True)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_1.tail(12)
```

Date	Price Ticker	Adj C
2025-02-27	BAC	44.120
	C	78.870
	GS	605.00
	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

Option 2: Calculate returns, make them long, then join them!

```
_ = stocks_wide['Adj Close'].iloc[:-1].pct_change().stack(future_stack=True).to_frame('Returns')
stocks_long_2 = stocks_long.join(_)
```

Recall, we omitted returns for the most recent trading day, which could include a partial data return.

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
	BAC	44.120
	C	78.870
	GS	605.00
2025-02-27	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11
	PNC	191.92

We can test the equality of `stocks_long_1` and `stocks_long_2` most easily with the `.equals()` method.

```
stocks_long_1.equals(stocks_long_2)
```

True

Download the daily benchmark return factors from Ken French's data library.

Hint: Use the `DataReader()` function in the `pandas-datareader` package. We imported this package above with the `pdr.` prefix.

I often cannot remember the exact name for the daily factors. We can use the `pdr.famafrench.get_available_datasets()` to list all the data in Kenneth French's data library.

```
pdr.famafrench.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors_daily',
    data_source='famafrench',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_28432\875599436.py:1: FutureWarning: The argument
```

```
    ff = pdr.DataReader(
```

```
type(ff)
```

```
dict
```

The daily factors only have one data frame (in the 0 key) and the data set description (in the DESCR key).

```
ff.keys()
```

```
dict_keys([0, 'DESCR'])
```

i Note

Data from the Kenneth French data library are *percent* returns instead of *decimal* returns!

```
ff[0]
```

Date	Mkt-RF	SMB	HML	RF
1926-07-01	0.1000	-0.2500	-0.2700	0.0090
1926-07-02	0.4500	-0.3300	-0.0600	0.0090
1926-07-06	0.1700	0.3000	-0.3900	0.0090

	Mkt-RF	SMB	HML	RF
Date				
1926-07-07	0.0900	-0.5800	0.0200	0.0090
1926-07-08	0.2100	-0.3800	0.1900	0.0090
...
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
print(ff['DESCR'])
```

F-F Research Data Factors daily

This file was created by CMPT_ME_BEME_RET_DAILY using the 202412 CRSP database. The Tbill r

0 : (25901 rows x 4 cols)

Add the daily benchmark return factors to stocks_wide and stocks_long.

Since both ff[0] and stocks_long_2 have date indexes, we can easily combine them with the .join() method.

```
ff[0].tail()
```

	Mkt-RF	SMB	HML	RF
Date				
2024-12-24	1.1100	-0.0900	-0.0500	0.0170
2024-12-26	0.0200	1.0400	-0.1900	0.0170
2024-12-27	-1.1700	-0.6600	0.5600	0.0170
2024-12-30	-1.0900	0.1200	0.7400	0.0170
2024-12-31	-0.4600	0.0000	0.7100	0.0170

```
stocks_long_2.tail(12)
```

Date	Ticker	Adj C
2025-02-27	BAC	44.120
	C	78.870
	GS	605.00
	JPM	259.05
	MS	129.24
	PNC	188.56
	BAC	46.100
	C	79.950
2025-02-28	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

We can quickly combine `stocks_long_2` and `ff[0]` because both have indexes with daily dates named `Date`. Two notes:

1. The `.join()` method left joins by default, so the combined output has only dates in `stocks_long_2`
2. Kenneth French provides *percent* returns, so we divide them by 100 to convert them to *decimal* returns to match our Yahoo! Finance data

```
stocks_long_2.join(ff[0].div(100))
```

Date	Ticker	Adj C
1973-02-21	BAC	1.5426
	C	NaN
	GS	NaN
	JPM	NaN
	MS	NaN
...
2025-02-28	C	79.950
	GS	622.29
	JPM	264.65
	MS	133.11
	PNC	191.92

We could instead convert the Yahoo! Finance *decimal* returns to *percent* returns. I do not have a strong preference on all decimal returns or all percent returns, but all returns should have the same form.

```
(  
    stocks_long_2  
    .assign(Returns=lambda x: 100 * x['Returns'])  
    .join(ff[0])  
)
```

Date	Ticker	Adj C
1973-02-21	BAC	1.5426
	C	NaN
	GS	NaN
	JPM	NaN
	MS	NaN
...
	C	79.950
	GS	622.29
2025-02-28	JPM	264.65
	MS	133.11
	PNC	191.92

With `stocks_wide`, we have to do a little more work because of its column multi-index! We will use the `pd.MultiIndex.from_product()` trick from above.

```
_ = pd.MultiIndex.from_product([['Factors'], ff[0].columns])  
stocks_wide[_] = ff[0].div(100)
```

```
stocks_wide.loc[:'2024'].tail()
```

Price	Adj Close					Close				
Ticker	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JPM
Date										
2024-12-24	44.3800	70.5117	579.9144	241.0650	126.2201	192.4933	44.3800	71.0000	582.7900	24

Price Ticker Date	Adj Close						Close			
	BAC	C	GS	JPM	MS	PNC	BAC	C	GS	JP
2024-12-26	44.5500	70.8593	578.3621	241.8907	127.1837	193.1777	44.5500	71.3500	581.2300	24
2024-12-27	44.3400	70.5117	573.3370	239.9308	125.9221	191.6999	44.3400	71.0000	576.1800	24
2024-12-30	43.9100	69.9059	570.7200	238.0904	124.9188	190.9560	43.9100	70.3900	573.5500	23
2024-12-31	43.9500	69.9059	569.7946	238.4783	124.8890	191.2734	43.9500	70.3900	572.6200	23

Write a function download() that accepts tickers and returns a wide data frame of returns with the daily benchmark return factors.

We can even add a `shape` argument to return a wide or long data frame!

We can even add a `shape` argument to return a wide or long data frame!

```
import warnings

def download(tickers, shape='wide'):
    """
    Download stock price data and Fama-French factors, returning in either 'wide' or 'long' format.

    Parameters:
    - tickers (str or list of str): Stock ticker(s) to download.
    - shape (str): Output format, either 'wide' (default) or 'long'.

    Returns:
    - pd.DataFrame: A DataFrame containing stock prices, returns, and Fama-French factors.
    """
    # shape must be wide or long
    if shape not in ['wide', 'long']:
        raise ValueError('Invalid shape: must be "wide" or "long".')
    # Download stock data
    stocks = yf.download(tickers=tickers, auto_adjust=False, progress=False)

    # Download Fama-French factors
    # (suppressing FutureWarning for 'date_parser')
    with warnings.catch_warnings():
        warnings.simplefilter('ignore', category=FutureWarning)
        factors = pdr.DataReader(

```

```

        name='F-F_Research_Data_Factors_daily',
        data_source='famafrench',
        start='1900'
)[0].div(100) # Convert percentages to decimals

# Multi-index case
if isinstance(stocks.columns, pd.MultiIndex):
    # Compute daily returns
    _ = pd.MultiIndex.from_product([['Returns'], stocks['Adj Close'].columns])
    stocks[_] = stocks['Adj Close'].pct_change()

    if shape == 'wide':
        # Add factors with multi-index
        _ = pd.MultiIndex.from_product([['Factors'], factors.columns])
        stocks[_] = factors
        return stocks

    # Convert to long format then add factors
else:
    return stocks.stack(future_stack=True).join(factors)

# Single index case
# (redundant with recent versions of yfinance that always return a multi-index)
stocks['Returns'] = stocks['Adj Close'].pct_change()
return stocks.join(factors)

```

```
download(tickers='AAPL TSLA')
```

Price Ticker Date	Adj Close		Close		High		Low		Open
	AAPL	TSLA	AAPL	TSLA	AAPL	TSLA	AAPL	TSLA	AAPL
1980-12-12	0.0987	NaN	0.1283	NaN	0.1289	NaN	0.1283	NaN	0.1283
1980-12-15	0.0936	NaN	0.1217	NaN	0.1222	NaN	0.1217	NaN	0.1222
1980-12-16	0.0867	NaN	0.1127	NaN	0.1133	NaN	0.1127	NaN	0.1133
1980-12-17	0.0889	NaN	0.1155	NaN	0.1161	NaN	0.1155	NaN	0.1155
1980-12-18	0.0914	NaN	0.1189	NaN	0.1194	NaN	0.1189	NaN	0.1189
...
2025-02-24	247.1000	330.5300	247.1000	330.5300	248.8600	342.4000	244.4200	324.7000	244.9300
2025-02-25	247.0400	302.8000	247.0400	302.8000	250.0000	328.8900	244.9100	297.2500	248.0000
2025-02-26	240.3600	290.8000	240.3600	290.8000	244.9800	309.0000	239.1300	288.0400	244.3300
2025-02-27	237.3000	281.9500	237.3000	281.9500	242.4600	297.2300	237.0600	280.8800	239.4100

Price Ticker Date	Adj Close AAPL	Close TSLA	Close AAPL	Close TSLA	High AAPL	Low TSLA	Open TSLA	Open AAPL
2025-02-28	241.8400	292.9800	241.8400	292.9800	242.0900	293.8800	230.2000	273.6000

i Note

The `yfinance` package is a powerful tool for downloading market data, financial statements, and analyst estimates from Yahoo! Finance.

However, because `yfinance` relies on Yahoo! Finance's API, changes to the API can disrupt its functionality.

Recently, Yahoo! Finance changed API access to earnings forecasts and announcement dates, so we **cannot complete the earnings announcement exercise I had planned.**

Instead, I will prepare an alternative set of exercises for us to work on in class on Friday. Thank you for your flexibility!

Download earnings per share for the stocks in `stocks_long` and combine to a long data frame `earnings`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Combine earnings with the returns from `stocks_long`.

Use the `.earnings_dates` method described [here](#). Use `pd.concat()` to combine the result of each the `.earnings_date` data frames and assign them to a new data frame `earnings`. Name the row indexes `Ticker` and `Date` and swap to match the order of the row index in `stocks_long`.

Plot the relation between daily returns and earnings surprises

Repeat the earnings exercise with the S&P 100 stocks

With more data, we can more clearly see the positive relation between earnings surprises and returns!

Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Blank

This notebook covers two topics:

1. Log and simple returns
2. Portfolio returns, plus two applications of portfolio returns

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Log and Simple Returns

We will typically use *simple* returns, calculated as $r_{simple,t} = \frac{p_t + d_t - p_{t-1}}{p_{t-1}}$. This simple return is the return that investors earn on their investments. We can calculate simple returns from Yahoo Finance data with the `.pct_change()` method on the adjusted close column (i.e., `Adj Close`), which adjusts for dividends and splits. The adjusted close column is a reverse-engineered close price (i.e., end-of-trading-day price) that incorporates dividends and splits, making simple return calculations easy.

However, we may see *log* returns elsewhere, which are the (natural) log of one plus simple returns: $r_{log,t} = \log(1 + r_{simple,t})$. Therefore, we calculate log returns as either the log of one plus simple returns or the difference of the logs of the adjusted close column. Log returns are also known as *continuously-compounded* returns.


This section explains the differences between simple and log returns and where each is appropriate.

Simple and Log Returns are Similar for Small Returns

$x \approx \log(1 + x)$ for small values of x , so simple returns and log returns are similar for small returns. Returns are typically small at daily and monthly horizons, so the difference between simple and log returns is small at daily and monthly horizons. The following figure shows that $r_{simple,t} \approx r_{log,t}$ for small values of r .

```
simpler = np.linspace(-0.75, 0.75, 100)
logr = np.log1p(simpler)
```

```
plt.plot(simpler, logr)
plt.plot([-1, 1], [-1, 1])
plt.xlabel('Simple Return')
plt.ylabel('Log Return')
plt.title('Log Versus Simple Returns')
plt.legend(['Actual Relation', 'Relation If Log = Simple'])
plt.show()
```



Simple Return Advantage: Portfolio Calculations

For a portfolio of N assets with portfolio weights w_i , the portfolio return r_p is the weighted average of the returns of its assets: $r_p = \sum_{i=1}^N w_i r_i$. For example, for an equal-weighted portfolio with two stocks, $r_p = 0.5r_1 + 0.5r_2 = \frac{r_1+r_2}{2}$. Therefore, we cannot calculate portfolio returns with log returns because the sum of logs is the log of products. That is $\log(1 + r_i) + \log(1 + r_j) = \log((1 + r_i) \times (1 + r_j))$, which is not what we want to measure! **We cannot perform portfolio calculations with log returns!**

Log Return Advantage: Log Returns are Additive

We compound simple returns with multiplication, *but we compound log returns with addition*. This additive property of log returns makes code simple, computations fast, and proofs easy when we must compound returns.

We compound returns from $t = 0$ to $t = T$ as follows:

$$1 + r_{0,T} = (1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)$$

Next, we take the log of both sides of the previous equation and use the property that the log of products is the sum of logs:

$$\log(1 + r_{0,T}) = \log((1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)) = \log(1 + r_1) + \log(1 + r_2) + \cdots + \log(1 + r_T) = \sum_{t=1}^T \log(1 + r_t)$$

Next, we exponentiate both sides of the previous equation:

$$e^{\log(1 + r_{0,T})} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Next, we use the property that $e^{\log(x)} = x$ to simplify the previous equation:

$$1 + r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Finally, we subtract 1 from both sides:

$$r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)} - 1$$

So, the return $r_{0,T}$ from $t = 0$ to $t = T$ is the exponentiated sum of log returns. The pandas developers assume users understand the math above and focus on optimizing sums!

```
np.random.seed(42)
df = pd.DataFrame(data={'r': np.exp(np.random.randn(10_000)) - 1})

df.describe()
```

	r
count	10000.0000
mean	0.6529
std	2.1918
min	-0.9802
25%	-0.4896
50%	-0.0026
75%	0.9564
max	49.7158

We can time the calculation of 10-observation rolling returns. We use `.apply()` for the simple return version because `.rolling()` does not have a product method. We find that `.rolling()` is slower with `.apply()` than with `.sum()` by a factor of about 1,000. *We will learn about `.rolling()` and `.apply()` in a few weeks, but they provide the best example of when to use log returns.*

```
%%timeit
df['r10_via_prod'] = (
    df['r']
    .add(1)
    .rolling(10)
    .apply(lambda x: x.prod())
    .sub(1)
)
```

936 ms ± 119 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

```
%%timeit
df['r10_via_sum'] = (
    df['r']
    .add(1)
    .pipe(np.log)
    .rolling(10)
    .sum()
    .pipe(np.exp)
    .sub(1)
)
```

1.15 ms ± 159 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
df.head(15)
```

	r	r10_via_prod	r10_via_sum
0	0.6433	NaN	NaN
1	-0.1291	NaN	NaN
2	0.9111	NaN	NaN
3	3.5861	NaN	NaN
4	-0.2088	NaN	NaN
5	-0.2087	NaN	NaN
6	3.8511	NaN	NaN
7	1.1542	NaN	NaN
8	-0.3747	NaN	NaN
9	0.7204	87.2886	87.2886
10	-0.3709	32.8006	32.8006
11	-0.3723	23.3617	23.3617
12	0.2737	15.2369	15.2369
13	-0.8524	-0.4774	-0.4774
14	-0.8218	-0.8823	-0.8823

```
np.allclose(df['r10_via_prod'], df['r10_via_sum'], equal_nan=True)
```

True

These two approaches calculate the same return series, but the simple-return approach using `.prod()` is about 1,000 times slower than the log-return approach using `.sum()!` **We can use log returns to calculate total returns very quickly!**

Portfolio Math

Portfolio return r_p is the weighted average of its asset returns, so $r_p = \sum_{i=1}^N w_i r_i$. Here N is the number of assets, w_i is the weight on asset i , and $\sum_{i=1}^N w_i = 1$.

The 1/N Portfolio

The $\frac{1}{N}$ portfolio equally weights portfolio assets, so $w_1 = w_2 = \dots = w_N = \frac{1}{N}$. If $w_i = \frac{1}{N}$, then $r_p = \sum_{i=1}^N \frac{1}{N} r_i = \frac{\sum_{i=1}^N r_i}{N} = \bar{r}$. Therefore, we can use `.mean(axis=1)` to calculate $\frac{1}{N}$ portfolio returns!

```
df2 = yf.download(tickers='AAPL AMZN GOOG MSFT NVDA TSLA', auto_adjust=False, progress=False)
returns2 = df2['Adj Close'].pct_change().dropna()

returns2.describe()
```

Ticker	AAPL	AMZN	GOOG	MSFT	NVDA	TSLA
count	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000
mean	0.0011	0.0012	0.0009	0.0010	0.0021	0.0021
std	0.0175	0.0205	0.0173	0.0161	0.0288	0.0362
min	-0.1286	-0.1405	-0.1110	-0.1474	-0.1876	-0.2106
25%	-0.0074	-0.0089	-0.0071	-0.0070	-0.0122	-0.0164
50%	0.0010	0.0010	0.0009	0.0007	0.0017	0.0012
75%	0.0102	0.0119	0.0093	0.0093	0.0161	0.0194
max	0.1198	0.1575	0.1605	0.1422	0.2981	0.2440

```
returns2.mean() # implied axis=0
```

```
Ticker
AAPL    0.0011
AMZN    0.0012
GOOG    0.0009
MSFT    0.0010
NVDA    0.0021
TSLA    0.0021
dtype: float64
```

```
rp2_via_mean = returns2.mean(axis=1)

rp2_via_mean
```

```
Date
2010-06-30   -0.0123
2010-07-01   -0.0107
2010-07-02   -0.0271
2010-07-06   -0.0223
2010-07-07    0.0254
...
2025-02-21   -0.0272
2025-02-24   -0.0127
```

```
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Note that when we apply the same portfolio weights every period, we rebalance at the same frequency as the returns data. If we have daily data, rebalance daily. If we have monthly data, we rebalance monthly, and so on.

A More General Solution

If we combine portfolio weights into vector w and the time series of asset returns into matrix \mathbf{R} , then we can calculate the time series of portfolio returns as $r_p = w^T \mathbf{R}$. The pandas version of this calculation is `R.dot(w)`, where `R` is a data frame of asset returns and `w` is a series or an array of portfolio weights. We can use this approach to calculate $\frac{1}{N}$ portfolio returns, too.

```
weights2 = np.ones(returns2.shape[1]) / returns2.shape[1]
```

```
weights2
```

```
array([0.1667, 0.1667, 0.1667, 0.1667, 0.1667, 0.1667])
```

```
rp2_via_dot = returns2.dot(weights2)
```

```
rp2_via_dot
```

```
Date
2010-06-30    -0.0123
2010-07-01    -0.0107
2010-07-02    -0.0271
2010-07-06    -0.0223
2010-07-07    0.0254
...
2025-02-21    -0.0272
2025-02-24    -0.0127
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Both approaches give the same answer!

```
np.allclose(rp2_via_mean, rp2_via_dot, equal_nan=True)
```

True

Portfolio Math Application 1: All stocks half the time or half stocks all the time?

Are you better off investing:

1. 100% in stocks 50% of the time and the riskless asset the other 50% of the time *or*
2. 50% in stocks and 50% in the riskless asset 100% of the time?

Here is a roadmap for convincing yourself with data!

Download *annual* market and risk-free asset returns from Kenneth French's data library

Convert these factors to decimal returns and calculate the market return series

Add a portfolio return series that is half stocks all the time

You might call this portfolio return series Balanced

Add a portfolio return series that switches between stocks and bills every year with stocks in odd years

You might call this portfolio return series Switching Stocks Odd

Add a portfolio return series that switches between stocks and bills every year with stocks in even years

You might call this portfolio return series Switching Stocks Even

Plot the cumulative returns calculate the summary statistics for the Balanced and Switching series

Use the `.describe()` method to calcualte summary statistics.

Which strategy do you prefer?

Why? How sure are you?

Use the simulate() function to simulate 10,000 different outcomes for the U.S. market

simulate() calculates one Switching return series because the randomization also randomizes the odd-year and even-year choice.

```
def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced and switching

    Parameters:
    df (pd.DataFrame): The input DataFrame.
    cols (list): List of column names to sample.
    n_iter (int): Number of iterations for simulation.

    Returns:
    pd.DataFrame: A concatenated DataFrame with simulation results.
    """
    return pd.concat(
        objs=[(
            df[cols]
            .sample(frac=1, ignore_index=True, random_state=i)
            .assign(
                Balanced=lambda x: x[cols].mean(axis=1),
                Switching=lambda x: np.where(x.index % 2 == 0, x[cols[0]], x[cols[1]])
            )
        ) for i in range(n_iter)],
        keys=range(n_iter),
        names=['Simulation', 'Year']
    )
```

Calculate the summary statistics for these new Balanced and Switching series

Which strategy do you prefer?

Why? How sure are you?

Portfolio Math Application 2: What are the benefits of diversification?

Use random portfolios of S&P 100 stocks of various portfolio sizes to show that portfolio volatility falls quickly, then slowly, then not at all as we increase portfolio size.

Download daily data for the stocks in the S&P 100

Wikipedia provides tickers for the stocks in the [S&P 100](#). Use a list comprehension to replace . in tickers with - for compatibility with Yahoo! Finance.

Calculate the past five years of daily returns for these stocks

Calculate the volatilities of 20 equal-weighted random portfolios of various portfolio sizes

Random portfolios should have portfolio sizes of 1, 2, 4, 6, 8, 10, 20, 30, 40, or 50 stocks each.

You can combine the `.sample(n=?, axis=1, random_state=?), .mean(axis=1)`, and `.std()` to calculate the volatilities of equal-weighted portfolios. You can collect these volatilities in a list of lists built with two `for` loops or list comprehensions. Replace the ?s in `.sample()` with loop counters. The inner loop will calculate a portfolio volatility for each portfolio size, and the outer loop will collect 20 versions of each portfolio. Using the outer loop counter for `random_state=` makes your analysis repeatable!

Combine this list of list into a data frame

Calculate the mean volatility for each portfolio size and replicate the plot above

Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 02

This notebook covers two topics:

1. Log and simple returns
2. Portfolio returns, plus two applications of portfolio returns

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Log and Simple Returns

We will typically use *simple* returns, calculated as $r_{simple,t} = \frac{p_t + d_t - p_{t-1}}{p_{t-1}}$. This simple return is the return that investors earn on their investments. We can calculate simple returns from Yahoo Finance data with the `.pct_change()` method on the adjusted close column (i.e., `Adj Close`), which adjusts for dividends and splits. The adjusted close column is a reverse-engineered close price (i.e., end-of-trading-day price) that incorporates dividends and splits, making simple return calculations easy.

However, we may see *log* returns elsewhere, which are the (natural) log of one plus simple returns: $r_{log,t} = \log(1 + r_{simple,t})$. Therefore, we calculate log returns as either the log of one plus simple returns or the difference of the logs of the adjusted close column. Log returns are also known as *continuously-compounded* returns.


This section explains the differences between simple and log returns and where each is appropriate.

Simple and Log Returns are Similar for Small Returns

$x \approx \log(1 + x)$ for small values of x , so simple returns and log returns are similar for small returns. Returns are typically small at daily and monthly horizons, so the difference between simple and log returns is small at daily and monthly horizons. The following figure shows that $r_{simple,t} \approx r_{log,t}$ for small values of r .

```
simpler = np.linspace(-0.75, 0.75, 100)
logr = np.log1p(simpler)
```

```
plt.plot(simpler, logr)
plt.plot([-1, 1], [-1, 1])
plt.xlabel('Simple Return')
plt.ylabel('Log Return')
plt.title('Log Versus Simple Returns')
plt.legend(['Actual Relation', 'Relation If Log = Simple'])
plt.show()
```



Simple Return Advantage: Portfolio Calculations

For a portfolio of N assets with portfolio weights w_i , the portfolio return r_p is the weighted average of the returns of its assets: $r_p = \sum_{i=1}^N w_i r_i$. For example, for an equal-weighted portfolio with two stocks, $r_p = 0.5r_1 + 0.5r_2 = \frac{r_1+r_2}{2}$. Therefore, we cannot calculate portfolio returns with log returns because the sum of logs is the log of products. That is $\log(1 + r_i) + \log(1 + r_j) = \log((1 + r_i) \times (1 + r_j))$, which is not what we want to measure! **We cannot perform portfolio calculations with log returns!**

Log Return Advantage: Log Returns are Additive

We compound simple returns with multiplication, *but we compound log returns with addition*. This additive property of log returns makes code simple, computations fast, and proofs easy when we must compound returns.

We compound returns from $t = 0$ to $t = T$ as follows:

$$1 + r_{0,T} = (1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)$$

Next, we take the log of both sides of the previous equation and use the property that the log of products is the sum of logs:

$$\log(1 + r_{0,T}) = \log((1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)) = \log(1 + r_1) + \log(1 + r_2) + \cdots + \log(1 + r_T) = \sum_{t=1}^T \log(1 + r_t)$$

Next, we exponentiate both sides of the previous equation:

$$e^{\log(1 + r_{0,T})} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Next, we use the property that $e^{\log(x)} = x$ to simplify the previous equation:

$$1 + r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Finally, we subtract 1 from both sides:

$$r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)} - 1$$

So, the return $r_{0,T}$ from $t = 0$ to $t = T$ is the exponentiated sum of log returns. The pandas developers assume users understand the math above and focus on optimizing sums!

```
np.random.seed(42)
df = pd.DataFrame(data={'r': np.exp(np.random.randn(10_000)) - 1})

df.describe()
```

	r
count	10000.0000
mean	0.6529
std	2.1918
min	-0.9802
25%	-0.4896
50%	-0.0026
75%	0.9564
max	49.7158

We can time the calculation of 10-observation rolling returns. We use `.apply()` for the simple return version because `.rolling()` does not have a product method. We find that `.rolling()` is slower with `.apply()` than with `.sum()` by a factor of about 1,000. *We will learn about `.rolling()` and `.apply()` in a few weeks, but they provide the best example of when to use log returns.*

```
%%timeit
df['r10_via_prod'] = (
    df['r']
    .add(1)
    .rolling(10)
    .apply(lambda x: x.prod())
    .sub(1)
)
```

233 ms ± 84.1 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
%%timeit
df['r10_via_sum'] = (
    df['r']
    .add(1)
    .pipe(np.log)
    .rolling(10)
    .sum()
    .pipe(np.exp)
    .sub(1)
)
```

1.71 ms ± 297 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
df.head(15)
```

	r	r10_via_prod	r10_via_sum
0	0.6433	NaN	NaN
1	-0.1291	NaN	NaN
2	0.9111	NaN	NaN
3	3.5861	NaN	NaN
4	-0.2088	NaN	NaN
5	-0.2087	NaN	NaN
6	3.8511	NaN	NaN
7	1.1542	NaN	NaN
8	-0.3747	NaN	NaN
9	0.7204	87.2886	87.2886
10	-0.3709	32.8006	32.8006
11	-0.3723	23.3617	23.3617
12	0.2737	15.2369	15.2369
13	-0.8524	-0.4774	-0.4774
14	-0.8218	-0.8823	-0.8823

```
np.allclose(df['r10_via_prod'], df['r10_via_sum'], equal_nan=True)
```

True

These two approaches calculate the same return series, but the simple-return approach using `.prod()` is about 1,000 times slower than the log-return approach using `.sum()!` **We can use log returns to calculate total returns very quickly!**

Portfolio Math

Portfolio return r_p is the weighted average of its asset returns, so $r_p = \sum_{i=1}^N w_i r_i$. Here N is the number of assets, w_i is the weight on asset i , and $\sum_{i=1}^N w_i = 1$.

The 1/N Portfolio

The $\frac{1}{N}$ portfolio equally weights portfolio assets, so $w_1 = w_2 = \dots = w_N = \frac{1}{N}$. If $w_i = \frac{1}{N}$, then $r_p = \sum_{i=1}^N \frac{1}{N} r_i = \frac{\sum_{i=1}^N r_i}{N} = \bar{r}$. Therefore, we can use `.mean(axis=1)` to calculate $\frac{1}{N}$ portfolio returns!

```
df2 = yf.download(tickers='AAPL AMZN GOOG MSFT NVDA TSLA', auto_adjust=False, progress=False)
returns2 = df2['Adj Close'].pct_change().dropna()

returns2.describe()
```

Ticker	AAPL	AMZN	GOOG	MSFT	NVDA	TSLA
count	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000
mean	0.0011	0.0012	0.0009	0.0010	0.0021	0.0021
std	0.0175	0.0205	0.0173	0.0161	0.0288	0.0362
min	-0.1286	-0.1405	-0.1110	-0.1474	-0.1876	-0.2106
25%	-0.0074	-0.0089	-0.0071	-0.0070	-0.0122	-0.0164
50%	0.0010	0.0010	0.0009	0.0007	0.0017	0.0012
75%	0.0102	0.0119	0.0093	0.0093	0.0161	0.0194
max	0.1198	0.1575	0.1605	0.1422	0.2981	0.2440

```
returns2.mean() # implied axis=0
```

```
Ticker
AAPL    0.0011
AMZN    0.0012
GOOG    0.0009
MSFT    0.0010
NVDA    0.0021
TSLA    0.0021
dtype: float64
```

```
rp2_via_mean = returns2.mean(axis=1)

rp2_via_mean
```

```
Date
2010-06-30   -0.0123
2010-07-01   -0.0107
2010-07-02   -0.0271
2010-07-06   -0.0223
2010-07-07    0.0254
...
2025-02-21   -0.0272
2025-02-24   -0.0127
```

```
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Note that when we apply the same portfolio weights every period, we rebalance at the same frequency as the returns data. If we have daily data, rebalance daily. If we have monthly data, we rebalance monthly, and so on.

A More General Solution

If we combine portfolio weights into vector w and the time series of asset returns into matrix \mathbf{R} , then we can calculate the time series of portfolio returns as $r_p = w^T \mathbf{R}$. The pandas version of this calculation is `R.dot(w)`, where `R` is a data frame of asset returns and `w` is a series or an array of portfolio weights. We can use this approach to calculate $\frac{1}{N}$ portfolio returns, too.

```
weights2 = np.ones(returns2.shape[1]) / returns2.shape[1]

weights2
```

```
array([0.1667, 0.1667, 0.1667, 0.1667, 0.1667, 0.1667])
```

```
rp2_via_dot = returns2.dot(weights2)

rp2_via_dot
```

```
Date
2010-06-30    -0.0123
2010-07-01    -0.0107
2010-07-02    -0.0271
2010-07-06    -0.0223
2010-07-07    0.0254
...
2025-02-21    -0.0272
2025-02-24    -0.0127
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Both approaches give the same answer!

```
np.allclose(rp2_via_mean, rp2_via_dot, equal_nan=True)
```

```
True
```

Portfolio Math Application 1: All stocks half the time or half stocks all the time?

Are you better off investing:

1. 100% in stocks 50% of the time and the riskless asset the other 50% of the time *or*
2. 50% in stocks and 50% in the riskless asset 100% of the time?

Here is a roadmap for convincing yourself with data!

Please see Kritzman (2000, Chapter 5) for a more detailed solution!

Download annual market and risk-free asset returns from Kenneth French's data library

```
pdr.famafrench.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors',
    data_source='famafrench',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_2148\304025689.py:1: FutureWarning: The argument 'data_source' was passed to 'read_csv'. This argument is deprecated, and will be removed in a future version of pandas. You can still pass 'data_source' as a keyword argument to 'read_csv', but it will be ignored.
  ff = pdr.DataReader(
C:\Users\r.herron\AppData\Local\Temp\ipykernel_2148\304025689.py:1: FutureWarning: The argument 'data_source' was passed to 'read_csv'. This argument is deprecated, and will be removed in a future version of pandas. You can still pass 'data_source' as a keyword argument to 'read_csv', but it will be ignored.
  ff = pdr.DataReader(
```

```
print(ff['DESCR'])
```

F-F Research Data Factors

This file was created by CMPT_ME_BEME_RET using the 202412 CRSP database. The 1-month TBill

```
0 : (1182 rows x 4 cols)
1 : Annual Factors: January-December (98 rows x 4 cols)
```

Convert these factors to decimal returns and calculate the market return series

```
df3 = ff[1].div(100)
```

```
df3['Mkt'] = df3['Mkt-RF'] + df3['RF']
```

Add a portfolio return series that is half stocks all the time

You might call this portfolio return series Balanced

```
df3['Balanced'] = df3[['Mkt', 'RF']].mean(axis=1)
```

Add a portfolio return series that switches between stocks and bills every year with stocks in odd years

You might call this portfolio return series Switching Stocks Odd

```
df3['Switching Stocks Odd'] = np.where(df3.index.year % 2 == 1, df3['Mkt'], df3['RF'])
```

Add a portfolio return series that switches between stocks and bills every year with stocks in even years

You might call this portfolio return series Switching Stocks Even

```
df3['Switching Stocks Even'] = np.where(df3.index.year % 2 == 0, df3['Mkt'], df3['RF'])
```

```
df3.head()
```

Date	Mkt-RF	SMB	HML	RF	Mkt	Balanced	Switching Stocks Odd	Switching Stocks Even
1927	0.2947	-0.0204	-0.0454	0.0312	0.3259	0.1785	0.3259	0.0312
1928	0.3539	0.0451	-0.0617	0.0356	0.3895	0.2126	0.0356	0.3895
1929	-0.1954	-0.3070	0.1167	0.0475	-0.1479	-0.0502	-0.1479	0.0475
1930	-0.3123	-0.0517	-0.1154	0.0241	-0.2882	-0.1321	0.0241	-0.2882
1931	-0.4511	0.0370	-0.1395	0.0107	-0.4404	-0.2149	-0.4404	0.0107


Plot the cumulative returns on a \$1 investment and calculate the summary statistics for the Balanced and Switching series

Use the `.describe()` method to calculate summary statistics.

```
portfolios = ['Balanced', 'Switching Stocks Odd', 'Switching Stocks Even']
```

```
import matplotlib.ticker as ticker
```

```
(  
    df3[portfolios]  
    .add(1)  
    .cumprod()  
    .plot()  
)  
  
plt.yscale('log')  
plt.ylabel('Value of $1 Investment')  
plt.title(f'Value of $1 in Investment in Three Strategies\nfrom {df3.index.year[0]} through  
plt.show()
```



Which strategy do you prefer?

Why? How sure are you?

The **Switching Stocks Odd** strategy *seems* best in this example! However, this apparent superiority is largely a matter of **luck**, specifically this combination of the sample period and starting year. In this sample, the number of even and odd years is equal, but the market happened to perform better in odd years. This outcome is entirely driven by **this particular historical sequence**.

To draw broader conclusions, we must break free from the “luck of the draw” tied to this single realization of history. Below, we will explore two approaches to determine which strategy is better:

1. **Theory:** We will analyze the expected returns and variances mathematically to assess the underlying return-risk tradeoff for each strategy.

2. **Simulation:** To rely on data instead of theory, we will simulate thousands of alternative historical sequences.

By shuffling the 98 years of market data repeatedly, we can generate 10,000 random samples and evaluate the performance of each strategy across these simulations.

This approach allows us to:

- Confirm that **both strategies have the same arithmetic average return.**
- Show that the **balanced strategy consistently delivers a better return-risk tradeoff (Sharpe ratio)** when evaluated across many possible outcomes.

We will do the simulation first.

Use the `simulate()` function to simulate 10,000 different outcomes for the U.S. market

`simulate()` calculates one **Switching** return series because the randomization also randomizes the odd-year and even-year choice.

```
df3[['Mkt', 'RF']].sample(frac=1, ignore_index=True, random_state=42).head()
```

	Mkt	RF
0	0.2886	0.0837
1	0.2870	0.0421
2	0.2361	0.0004
3	0.3871	0.0033
4	-0.3674	0.0160

```
def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced and switching

    Parameters:
    df (pd.DataFrame): The input DataFrame.
    cols (list): List of column names to sample.
    n_iter (int): Number of iterations for simulation.

    Returns:
    pd.DataFrame: A concatenated DataFrame with simulation results.
    """
    return pd.concat(
        objs=[(
```

```

df[cols]
    .sample(frac=1, ignore_index=True, random_state=i)
    .assign(
        Balanced=lambda x: x[cols].mean(axis=1),
        Switching=lambda x: np.where(x.index % 2 == 0, x[cols[0]], x[cols[1]])
    )
) for i in range(n_iter)],
keys=range(n_iter),
names=['Simulation', 'Year']
)

```

```

df4 = simulate(df3)

```

Calculate the summary statistics for these new Balanced and Switching series

```
df4[['Balanced', 'Switching']].agg(['mean', 'std'])
```

	Balanced	Switching
mean	0.0772	0.0774
std	0.1001	0.1487

We see that **Balanced** and **Switching** have the same mean return, but **Balanced** has much lower volatility than **Switching**! A risk-averse investor prefers **Balanced** because it has a higher return/risk ratio than **Switching**. We can quantify this ratio as the Sharpe ratio, which is the mean excess return divided by the volatility of excess returns: $S_i = \frac{r_i - r_f}{\sigma(r_i - r_f)}$. We can calculate excess returns and Sharpe ratios in one code snippet.

```

(
    df4
    [['Balanced', 'Switching']]
    .sub(df4['RF'], axis=0)
    .agg(lambda x: x.mean() / x.std())
    .to_frame('Sharpe ratio')
)

```

Sharpe ratio	
Balanced	0.4368
Switching	0.2961

We have to do a little more work if we want to combine mean and volatility of *raw* returns with the Sharpe ratio of *excess* returns.

```
pd.concat(
    objs=[
        df4[['Balanced', 'Switching']].agg(['mean', 'std']).transpose(),
        df4[['Balanced', 'Switching']].sub(df4['RF'], axis=0).agg(lambda x: x.mean() / x.std)
    ],
    axis=1
)
```

	mean	std	Sharpe ratio
Balanced	0.0772	0.1001	0.4368
Switching	0.0774	0.1487	0.2961

Which strategy do you prefer?

Why? How sure are you?

We prefer **Balanced** because it has a Sharpe ratio about 50% greater than **Switching**! We can see this in the data above, and here is the theory.

For **Balanced**, $\sigma_p^2 = w_m^2 \sigma_m^2 + w_f^2 \sigma_f^2 + 2w_m w_f \sigma_m \sigma_f \rho_{m,f}$. Because $\sigma_f^2 \approx 0$ and $\rho_{m,f} \approx 0$, $\sigma_p^2 \approx w_m^2 \sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx w_m \sigma_m = \frac{1}{2} \sigma_m$. We see this in the data!

```
0.5 * df3['Mkt'].std()
```

0.0998

```
df4['Balanced'].std()
```

0.1001

For **Switching**, $\sigma_p^2 = w_m\sigma_m^2 + w_f\sigma_f^2 + w_m w_f(\mu_m - \mu_f)^2$. We have a different formula because **Switching** is diversified *over time* instead at a point in time! Because $\sigma_f^2 \approx 0$ and $(\mu_m - \mu_f)^2 \approx 0$, $\sigma_p^2 \approx w_m\sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx \sqrt{w_m}\sigma_m = \sqrt{\frac{1}{2}}\sigma_m$. We see this in the data!

```
np.sqrt(0.5) * df3['Mkt'].std()
```

0.1411

```
df4['Switching'].std()
```

0.1487

The $(\mu_m - \mu_f)^2$ is close to zero but not exactly zero. We can get even closer to the observed **Switching** portfolio volatility if we consider this term!

```
np.sqrt(0.5 * df3['Mkt'].var() + 0.5 * 0.5 * (df3['Mkt'].mean() - df3['RF'].mean())**2)
```

0.1478

Please see Kritzman (2000, Chapter 5) for a more detailed solution!

Here are two after-class additions to this question.

Can we add a progress bar?

Yes! We have to rewrite the `simulate()` function and replace the list comprehension with a for loop. However, this addition requires the `tqdm` package. To avoid any conflicts, I will not install this package and provide the code as markdown instead of executable code.

```
from tqdm import tqdm

def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced
    and switching portfolio returns.

```

```
Parameters:
-----
df : pd.DataFrame
    The input DataFrame.
cols : list
    List of column names to sample.
n_iter : int
    Number of iterations for simulation.

Returns:
-----
pd.DataFrame
    A concatenated DataFrame with simulation results.
"""
all_resamples = []

# Use tqdm to visualize progress
for i in tqdm(range(n_iter), desc="Simulating"):
    # Random re-sampling
    resampled = (
        df[cols]
        .sample(frac=1, ignore_index=True, random_state=i)
        .assign(
            Balanced=lambda x: x[cols].mean(axis=1),
            Switching=lambda x: np.where(
                x.index % 2 == 0,
                x[cols[0]],
                x[cols[1]]
            )
        )
    )
    all_resamples.append(resampled)

# Concatenate final results
return pd.concat(all_resamples, keys=range(n_iter), names=['Simulation', 'Year'])
```

Ideally, we calculate the mean and volatility of each simulation, then take the average

We can do this easily with either `.groupby()` or `.pivot_table()!`

```
(  
    df4  
    .groupby(level='Simulation')  
    [['Balanced', 'Switching']]  
    .agg(['mean', 'std'])  
    .mean()  
    .unstack()  
)
```

	mean	std
Balanced	0.0772	0.1006
Switching	0.0774	0.1488

These statistics are similar, but not identical!

Portfolio Math Application 2: What are the benefits of diversification?

Use random portfolios of S&P 100 stocks of various portfolio sizes to show that portfolio volatility falls quickly, then slowly, then not at all as we increase portfolio size.

Download daily data for the stocks in the S&P 100

Wikipedia provides tickers for the stocks in the [S&P 100](#). Use a list comprehension to replace . in tickers with - for compatibility with Yahoo! Finance.

```
wiki = pd.read_html('https://en.wikipedia.org/wiki/S%26P_100')  
tickers = [i.replace('.', '-') for i in wiki[2]['Symbol']]  
data = yf.download(tickers=tickers, auto_adjust=False, progress=False).iloc[:-1]
```

Calculate the past five years of daily returns for these stocks

```
returns = (
    data
    ['Adj Close']
    .dropna(axis=1, how='all')
    .pct_change()
    .iloc[-5*252:]
)
```

Calculate the volatilities of 20 equal-weighted random portfolios of various portfolio sizes

Random portfolios should have portfolio sizes of 1, 2, 4, 6, 8, 10, 20, 30, 40, or 50 stocks each.

You can combine the `.sample(n=?, axis=1, random_state=?), .mean(axis=1),` and `.std()` to calculate the volatilities of equal-weighted portfolios. You can collect these volatilities in a list of lists built with two `for` loops or list comprehensions. Replace the `?`s in `.sample()` with loop counters. The inner loop will calculate a portfolio volatility for each portfolio size, and the outer loop will collect 20 versions of each portfolio. Using the outer loop counter for `random_state=` makes your analysis repeatable!

```
portfolio_size = [1, 2, 4, 6, 8, 10, 20, 30, 40, 50]
portfolio_number = 20

list_of_volatilities = []
for i in range(portfolio_number):
    list_of_volatilities.append([returns.sample(n=j, axis=1, random_state=i).mean(axis=1).std() for j in portfolio_size])

list_of_volatilities[0]
```

```
[0.0153,
 0.0127,
 0.0121,
 0.0118,
 0.0134,
 0.0129,
 0.0128,
 0.0129,
 0.0129,
 0.0129]
```

Combine this list of lists into a data frame


```
volatilities = (
    pd.DataFrame(
        data=list_of_volatilities,
        index=range(1, 1+portfolio_number),
        columns=portfolio_size
    )
    .rename_axis(index='Portfolio Number', columns='Portfolio Size')
)
```

volatilities

Portfolio Size	1	2	4	6	8	10	20	30	40	50
Portfolio Number										
1	0.0153	0.0127	0.0121	0.0118	0.0134	0.0129	0.0128	0.0129	0.0129	0.0129
2	0.0193	0.0144	0.0127	0.0145	0.0144	0.0145	0.0137	0.0136	0.0136	0.0133
3	0.0160	0.0165	0.0161	0.0151	0.0142	0.0144	0.0132	0.0134	0.0133	0.0134
4	0.0193	0.0177	0.0172	0.0171	0.0157	0.0156	0.0141	0.0141	0.0142	0.0139
5	0.0205	0.0199	0.0151	0.0164	0.0141	0.0133	0.0127	0.0131	0.0128	0.0131
6	0.0182	0.0157	0.0123	0.0136	0.0136	0.0139	0.0137	0.0143	0.0137	0.0136
7	0.0199	0.0174	0.0171	0.0160	0.0148	0.0142	0.0124	0.0133	0.0130	0.0128
8	0.0280	0.0178	0.0142	0.0149	0.0144	0.0131	0.0141	0.0138	0.0134	0.0134
9	0.0150	0.0136	0.0115	0.0116	0.0117	0.0119	0.0131	0.0132	0.0131	0.0133
10	0.0141	0.0171	0.0164	0.0150	0.0146	0.0147	0.0138	0.0131	0.0128	0.0126
11	0.0285	0.0210	0.0185	0.0170	0.0163	0.0167	0.0143	0.0142	0.0136	0.0137
12	0.0172	0.0193	0.0148	0.0154	0.0137	0.0140	0.0137	0.0133	0.0125	0.0128
13	0.0150	0.0131	0.0124	0.0126	0.0142	0.0134	0.0138	0.0133	0.0131	0.0135
14	0.0280	0.0198	0.0160	0.0147	0.0147	0.0137	0.0137	0.0131	0.0133	0.0134
15	0.0282	0.0197	0.0152	0.0142	0.0148	0.0142	0.0143	0.0140	0.0133	0.0134
16	0.0188	0.0150	0.0148	0.0136	0.0138	0.0140	0.0138	0.0141	0.0137	0.0136
17	0.0165	0.0132	0.0149	0.0141	0.0133	0.0128	0.0130	0.0133	0.0133	0.0131
18	0.0213	0.0198	0.0177	0.0166	0.0161	0.0156	0.0145	0.0142	0.0135	0.0139
19	0.0132	0.0136	0.0143	0.0140	0.0150	0.0145	0.0139	0.0134	0.0131	0.0127
20	0.0205	0.0209	0.0169	0.0154	0.0151	0.0143	0.0137	0.0134	0.0132	0.0132

Calculate the mean volatility for each portfolio size and replicate the plot above

```
volatilities.mul(100 * np.sqrt(252)).mean().plot()
plt.xlabel('Portfolio Size')
plt.ylabel('Mean of Annualized Volatility (%)')
plt.title('Diminishing Returns to Diversification (Elton and Gruber, 1977)')
plt.show()
```



Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 03

This notebook covers two topics:

1. Log and simple returns
2. Portfolio returns, plus two applications of portfolio returns

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Log and Simple Returns

We will typically use *simple* returns, calculated as $r_{simple,t} = \frac{p_t + d_t - p_{t-1}}{p_{t-1}}$. This simple return is the return that investors earn on their investments. We can calculate simple returns from Yahoo Finance data with the `.pct_change()` method on the adjusted close column (i.e., `Adj Close`), which adjusts for dividends and splits. The adjusted close column is a reverse-engineered close price (i.e., end-of-trading-day price) that incorporates dividends and splits, making simple return calculations easy.

However, we may see *log* returns elsewhere, which are the (natural) log of one plus simple returns: $r_{log,t} = \log(1 + r_{simple,t})$. Therefore, we calculate log returns as either the log of one plus simple returns or the difference of the logs of the adjusted close column. Log returns are also known as *continuously-compounded* returns.


This section explains the differences between simple and log returns and where each is appropriate.

Simple and Log Returns are Similar for Small Returns

$x \approx \log(1 + x)$ for small values of x , so simple returns and log returns are similar for small returns. Returns are typically small at daily and monthly horizons, so the difference between simple and log returns is small at daily and monthly horizons. The following figure shows that $r_{simple,t} \approx r_{log,t}$ for small values of r .

```
simpler = np.linspace(-0.75, 0.75, 100)
logr = np.log1p(simpler)
```

```
plt.plot(simpler, logr)
plt.plot([-1, 1], [-1, 1])
plt.xlabel('Simple Return')
plt.ylabel('Log Return')
plt.title('Log Versus Simple Returns')
plt.legend(['Actual Relation', 'Relation If Log = Simple'])
plt.show()
```



Simple Return Advantage: Portfolio Calculations

For a portfolio of N assets with portfolio weights w_i , the portfolio return r_p is the weighted average of the returns of its assets: $r_p = \sum_{i=1}^N w_i r_i$. For example, for an equal-weighted portfolio with two stocks, $r_p = 0.5r_1 + 0.5r_2 = \frac{r_1+r_2}{2}$. Therefore, we cannot calculate portfolio returns with log returns because the sum of logs is the log of products. That is $\log(1 + r_i) + \log(1 + r_j) = \log((1 + r_i) \times (1 + r_j))$, which is not what we want to measure! **We cannot perform portfolio calculations with log returns!**

Log Return Advantage: Log Returns are Additive

We compound simple returns with multiplication, *but we compound log returns with addition*. This additive property of log returns makes code simple, computations fast, and proofs easy when we must compound returns.

We compound returns from $t = 0$ to $t = T$ as follows:

$$1 + r_{0,T} = (1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)$$

Next, we take the log of both sides of the previous equation and use the property that the log of products is the sum of logs:

$$\log(1 + r_{0,T}) = \log((1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)) = \log(1 + r_1) + \log(1 + r_2) + \cdots + \log(1 + r_T) = \sum_{t=1}^T \log(1 + r_t)$$

Next, we exponentiate both sides of the previous equation:

$$e^{\log(1 + r_{0,T})} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Next, we use the property that $e^{\log(x)} = x$ to simplify the previous equation:

$$1 + r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Finally, we subtract 1 from both sides:

$$r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)} - 1$$

So, the return $r_{0,T}$ from $t = 0$ to $t = T$ is the exponentiated sum of log returns. The pandas developers assume users understand the math above and focus on optimizing sums!

```
np.random.seed(42)
df = pd.DataFrame(data={'r': np.exp(np.random.randn(10_000)) - 1})

df.describe()
```

	r
count	10000.0000
mean	0.6529
std	2.1918
min	-0.9802
25%	-0.4896
50%	-0.0026
75%	0.9564
max	49.7158

We can time the calculation of 10-observation rolling returns. We use `.apply()` for the simple return version because `.rolling()` does not have a product method. We find that `.rolling()` is slower with `.apply()` than with `.sum()` by a factor of about 1,000. *We will learn about `.rolling()` and `.apply()` in a few weeks, but they provide the best example of when to use log returns.*

```
%%timeit
df['r10_via_prod'] = (
    df['r']
    .add(1)
    .rolling(10)
    .apply(lambda x: x.prod())
    .sub(1)
)
```

980 ms ± 90 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

```
%%timeit
df['r10_via_sum'] = (
    df['r']
    .add(1)
    .pipe(np.log)
    .rolling(10)
    .sum()
    .pipe(np.exp)
    .sub(1)
)
```

1.38 ms ± 69.6 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
df.head(15)
```

	r	r10_via_prod	r10_via_sum
0	0.6433	NaN	NaN
1	-0.1291	NaN	NaN
2	0.9111	NaN	NaN
3	3.5861	NaN	NaN
4	-0.2088	NaN	NaN
5	-0.2087	NaN	NaN
6	3.8511	NaN	NaN
7	1.1542	NaN	NaN
8	-0.3747	NaN	NaN
9	0.7204	87.2886	87.2886
10	-0.3709	32.8006	32.8006
11	-0.3723	23.3617	23.3617
12	0.2737	15.2369	15.2369
13	-0.8524	-0.4774	-0.4774
14	-0.8218	-0.8823	-0.8823

```
np.allclose(df['r10_via_prod'], df['r10_via_sum'], equal_nan=True)
```

True

These two approaches calculate the same return series, but the simple-return approach using `.prod()` is about 1,000 times slower than the log-return approach using `.sum()!` **We can use log returns to calculate total returns very quickly!**

Portfolio Math

Portfolio return r_p is the weighted average of its asset returns, so $r_p = \sum_{i=1}^N w_i r_i$. Here N is the number of assets, w_i is the weight on asset i , and $\sum_{i=1}^N w_i = 1$.

The 1/N Portfolio

The $\frac{1}{N}$ portfolio equally weights portfolio assets, so $w_1 = w_2 = \dots = w_N = \frac{1}{N}$. If $w_i = \frac{1}{N}$, then $r_p = \sum_{i=1}^N \frac{1}{N} r_i = \frac{\sum_{i=1}^N r_i}{N} = \bar{r}$. Therefore, we can use `.mean(axis=1)` to calculate $\frac{1}{N}$ portfolio returns!

```
df2 = yf.download(tickers='AAPL AMZN GOOG MSFT NVDA TSLA', auto_adjust=False, progress=False)
returns2 = df2['Adj Close'].pct_change().dropna()

returns2.describe()
```

Ticker	AAPL	AMZN	GOOG	MSFT	NVDA	TSLA
count	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000
mean	0.0011	0.0012	0.0009	0.0010	0.0021	0.0021
std	0.0175	0.0205	0.0173	0.0161	0.0288	0.0362
min	-0.1286	-0.1405	-0.1110	-0.1474	-0.1876	-0.2106
25%	-0.0074	-0.0089	-0.0071	-0.0070	-0.0122	-0.0164
50%	0.0010	0.0010	0.0009	0.0007	0.0017	0.0012
75%	0.0102	0.0119	0.0093	0.0093	0.0161	0.0194
max	0.1198	0.1575	0.1605	0.1422	0.2981	0.2440

```
returns2.mean() # implied axis=0
```

```
Ticker
AAPL    0.0011
AMZN    0.0012
GOOG    0.0009
MSFT    0.0010
NVDA    0.0021
TSLA    0.0021
dtype: float64
```

```
rp2_via_mean = returns2.mean(axis=1)

rp2_via_mean
```

```
Date
2010-06-30   -0.0123
2010-07-01   -0.0107
2010-07-02   -0.0271
2010-07-06   -0.0223
2010-07-07    0.0254
...
2025-02-21   -0.0272
2025-02-24   -0.0127
```

```
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Note that when we apply the same portfolio weights every period, we rebalance at the same frequency as the returns data. If we have daily data, rebalance daily. If we have monthly data, we rebalance monthly, and so on.

A More General Solution

If we combine portfolio weights into vector w and the time series of asset returns into matrix \mathbf{R} , then we can calculate the time series of portfolio returns as $r_p = w^T \mathbf{R}$. The pandas version of this calculation is `R.dot(w)`, where `R` is a data frame of asset returns and `w` is a series or an array of portfolio weights. We can use this approach to calculate $\frac{1}{N}$ portfolio returns, too.

```
weights2 = np.ones(returns2.shape[1]) / returns2.shape[1]

weights2
```

```
array([0.1667, 0.1667, 0.1667, 0.1667, 0.1667, 0.1667])
```

```
rp2_via_dot = returns2.dot(weights2)

rp2_via_dot
```

```
Date
2010-06-30    -0.0123
2010-07-01    -0.0107
2010-07-02    -0.0271
2010-07-06    -0.0223
2010-07-07    0.0254
...
2025-02-21    -0.0272
2025-02-24    -0.0127
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Both approaches give the same answer!

```
np.allclose(rp2_via_mean, rp2_via_dot, equal_nan=True)
```

```
True
```

Portfolio Math Application 1: All stocks half the time or half stocks all the time?

Are you better off investing:

1. 100% in stocks 50% of the time and the riskless asset the other 50% of the time *or*
2. 50% in stocks and 50% in the riskless asset 100% of the time?

Here is a roadmap for convincing yourself with data!

Please see Kritzman (2000, Chapter 5) for a more detailed solution!

Download annual market and risk-free asset returns from Kenneth French's data library

```
pdr.famafrench.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors',
    data_source='famafrench',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_17392\304025689.py:1: FutureWarning: The argument 'data_source' is deprecated and will be removed in a future version. Use 'source' instead.
  ff = pdr.DataReader(
C:\Users\r.herron\AppData\Local\Temp\ipykernel_17392\304025689.py:1: FutureWarning: The argument 'data_source' is deprecated and will be removed in a future version. Use 'source' instead.
  ff = pdr.DataReader(
```

```
print(ff['DESCR'])
```

F-F Research Data Factors

This file was created by CMPT_ME_BEME_RETs using the 202412 CRSP database. The 1-month TBill

```
0 : (1182 rows x 4 cols)
1 : Annual Factors: January-December (98 rows x 4 cols)
```

Convert these factors to decimal returns and calculate the market return series

```
df3 = ff[1].div(100)
```

```
df3['Mkt'] = df3['Mkt-RF'] + df3['RF']
```

Add a portfolio return series that is half stocks all the time

You might call this portfolio return series Balanced

```
df3['Balanced'] = df3[['Mkt', 'RF']].mean(axis=1)
```

Add a portfolio return series that switches between stocks and bills every year with stocks in odd years

You might call this portfolio return series Switching Stocks Odd

```
df3['Switching Stocks Odd'] = np.where(df3.index.year % 2 == 1, df3['Mkt'], df3['RF'])
```

Add a portfolio return series that switches between stocks and bills every year with stocks in even years

You might call this portfolio return series Switching Stocks Even

```
df3['Switching Stocks Even'] = np.where(df3.index.year % 2 == 0, df3['Mkt'], df3['RF'])
```

```
df3.head()
```

Date	Mkt-RF	SMB	HML	RF	Mkt	Balanced	Switching Stocks Odd	Switching Stocks Even
1927	0.2947	-0.0204	-0.0454	0.0312	0.3259	0.1785	0.3259	0.0312
1928	0.3539	0.0451	-0.0617	0.0356	0.3895	0.2126	0.0356	0.3895
1929	-0.1954	-0.3070	0.1167	0.0475	-0.1479	-0.0502	-0.1479	0.0475
1930	-0.3123	-0.0517	-0.1154	0.0241	-0.2882	-0.1321	0.0241	-0.2882
1931	-0.4511	0.0370	-0.1395	0.0107	-0.4404	-0.2149	-0.4404	0.0107


Plot the cumulative returns on a \$1 investment and calculate the summary statistics for the Balanced and Switching series

Use the `.describe()` method to calcualte summary statistics.

```
portfolios = ['Balanced', 'Switching Stocks Odd', 'Switching Stocks Even']
```

```
import matplotlib.ticker as ticker
```

```
(  
    df3[portfolios]  
    .add(1)  
    .cumprod()  
    .plot()  
)  
  
plt.yscale('log')  
plt.ylabel('Value of $1 Investment')  
plt.title(f'Value of $1 in Investment in Three Strategies\nfrom {df3.index.year[0]} through  
plt.show()
```



Which strategy do you prefer?

Why? How sure are you?

The **Switching Stocks Odd** strategy *seems* best in this example! However, this apparent superiority is largely a matter of **luck**, specifically this combination of the sample period and starting year. In this sample, the number of even and odd years is equal, but the market happened to perform better in odd years. This outcome is entirely driven by **this particular historical sequence**.

To draw broader conclusions, we must break free from the “luck of the draw” tied to this single realization of history. Below, we will explore two approaches to determine which strategy is better:

- Theory:** We will analyze the expected returns and variances mathematically to assess the underlying return-risk tradeoff for each strategy.

2. **Simulation:** To rely on data instead of theory, we will simulate thousands of alternative historical sequences.

By shuffling the 98 years of market data repeatedly, we can generate 10,000 random samples and evaluate the performance of each strategy across these simulations.

This approach allows us to:

- Confirm that **both strategies have the same arithmetic average return.**
- Show that the **balanced strategy consistently delivers a better return-risk tradeoff (Sharpe ratio)** when evaluated across many possible outcomes.

We will do the simulation first.

Use the `simulate()` function to simulate 10,000 different outcomes for the U.S. market

`simulate()` calculates one **Switching** return series because the randomization also randomizes the odd-year and even-year choice.

```
df3[['Mkt', 'RF']].sample(frac=1, ignore_index=True, random_state=42).head()
```

	Mkt	RF
0	0.2886	0.0837
1	0.2870	0.0421
2	0.2361	0.0004
3	0.3871	0.0033
4	-0.3674	0.0160

```
def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced and switching

    Parameters:
    df (pd.DataFrame): The input DataFrame.
    cols (list): List of column names to sample.
    n_iter (int): Number of iterations for simulation.

    Returns:
    pd.DataFrame: A concatenated DataFrame with simulation results.
    """
    return pd.concat(
        objs=[(
```

```

df[cols]
    .sample(frac=1, ignore_index=True, random_state=i)
    .assign(
        Balanced=lambda x: x[cols].mean(axis=1),
        Switching=lambda x: np.where(x.index % 2 == 0, x[cols[0]], x[cols[1]])
    )
) for i in range(n_iter)],
keys=range(n_iter),
names=['Simulation', 'Year']
)

```

```
df4 = simulate(df3)
```

```
df4.head()
```

			Mkt
Simulation			Year
0		0	0.0077
		1	0.1633
		2	-0.1479
		3	0.2120
		4	0.3682

Calculate the summary statistics for these new `Balanced` and `Switching` series

```
df4[['Balanced', 'Switching']].agg(['mean', 'std'])
```

	Balanced	Switching
mean	0.0772	0.0774
std	0.1001	0.1487

We see that `Balanced` and `Switching` have the same mean return, but `Balanced` has much lower volatility than `Switching`! A risk-averse investor prefers `Balanced` because it has a higher return/risk ratio than `Switching`. We can quantify this ratio as the Sharpe ratio, which is the mean excess return divided by the volatility of excess returns: $S_i = \frac{r_i - r_f}{\sigma(r_i - r_f)}$. We can calculate excess returns and Sharpe ratios in one code snippet.

```
( df4
  [['Balanced', 'Switching']]
  .sub(df4['RF'], axis=0)
  .agg(lambda x: x.mean() / x.std())
  .to_frame('Sharpe ratio')
)
```

	Sharpe ratio
Balanced	0.4368
Switching	0.2961

We have to do a little more work if we want to combine mean and volatility of *raw* returns with the Sharpe ratio of *excess* returns.

```
pd.concat(
    objs=[
        df4[['Balanced', 'Switching']].agg(['mean', 'std']).transpose(),
        df4[['Balanced', 'Switching']].sub(df4['RF'], axis=0).agg(lambda x: x.mean() / x.std)
    ],
    axis=1
)
```

	mean	std	Sharpe ratio
Balanced	0.0772	0.1001	0.4368
Switching	0.0774	0.1487	0.2961

Which strategy do you prefer?

Why? How sure are you?

We prefer **Balanced** because it has a Sharpe ratio about 50% greater than **Switching**! We can see this in the data above, and here is the theory.

For **Balanced**, $\sigma_p^2 = w_m^2 \sigma_m^2 + w_f^2 \sigma_f^2 + 2w_m w_f \sigma_m \sigma_f \rho_{m,f}$. Because $\sigma_f^2 \approx 0$ and $\rho_{m,f} \approx 0$, $\sigma_p^2 \approx w_m^2 \sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx w_m \sigma_m = \frac{1}{2} \sigma_m$. We see this in the data!

```
0.5 * df3['Mkt'].std()
```

0.0998

```
df4['Balanced'].std()
```

0.1001

For **Switching**, $\sigma_p^2 = w_m\sigma_m^2 + w_f\sigma_f^2 + w_m w_f(\mu_m - \mu_f)^2$. We have a different formula because **Switching** is diversified *over time* instead at a point in time! Because $\sigma_f^2 \approx 0$ and $(\mu_m - \mu_f)^2 \approx 0$, $\sigma_p^2 \approx w_m\sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx \sqrt{w_m}\sigma_m = \sqrt{\frac{1}{2}}\sigma_m$. We see this in the data!

```
np.sqrt(0.5) * df3['Mkt'].std()
```

0.1411

```
df4['Switching'].std()
```

0.1487

The $(\mu_m - \mu_f)^2$ is close to zero but not exactly zero. We can get even closer to the observed **Switching** portfolio volatility if we consider this term!

```
np.sqrt(0.5 * df3['Mkt'].var() + 0.5 * 0.5 * (df3['Mkt'].mean() - df3['RF'].mean())**2)
```

0.1478

Here are two after-class additions to this question.

Can we add a progress bar?

Yes! We have to rewrite the `simulate()` function and replace the list comprehension with a for loop. However, this addition requires the `tqdm` package. To avoid any conflicts, I will not install this package and provide the code as markdown instead of executable code.

```
from tqdm import tqdm

def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced
    and switching portfolio returns.

    Parameters:
    -----
    df : pd.DataFrame
        The input DataFrame.
    cols : list
        List of column names to sample.
    n_iter : int
        Number of iterations for simulation.

    Returns:
    -----
    pd.DataFrame
        A concatenated DataFrame with simulation results.
    """
    all_resamples = []

    # Use tqdm to visualize progress
    for i in tqdm(range(n_iter), desc="Simulating"):
        # Random re-sampling
        resampled = (
            df[cols]
            .sample(frac=1, ignore_index=True, random_state=i)
            .assign(
                Balanced=lambda x: x[cols].mean(axis=1),
                Switching=lambda x: np.where(
                    x.index % 2 == 0,
                    x[cols[0]],
                    x[cols[1]]
                )
            )
        )
        all_resamples.append(resampled)

    # Concatenate final results
    return pd.concat(all_resamples, keys=range(n_iter), names=['Simulation', 'Year'])
```

Ideally, we calculate the mean and volatility of each simulation, then take the average

We can do this easily with either `.groupby()` or `.pivot_table()!`

```
(  
    df4  
    .groupby(level='Simulation')  
    [['Balanced', 'Switching']]  
    .agg(['mean', 'std'])  
    .mean()  
    .unstack()  
)
```

	mean	std
Balanced	0.0772	0.1006
Switching	0.0774	0.1488

These statistics are similar, but not identical!

Portfolio Math Application 2: What are the benefits of diversification?

Use random portfolios of S&P 100 stocks of various portfolio sizes to show that portfolio volatility falls quickly, then slowly, then not at all as we increase portfolio size.

Download daily data for the stocks in the S&P 100

Wikipedia provides tickers for the stocks in the [S&P 100](#). Use a list comprehension to replace
. in tickers with - for compatibility with Yahoo! Finance.

```
wiki = pd.read_html('https://en.wikipedia.org/wiki/S%26P_100')  
tickers = [i.replace('.', '-') for i in wiki[2]['Symbol']]  
data = yf.download(tickers=tickers, auto_adjust=False, progress=False).iloc[:-1]
```

Calculate the past five years of daily returns for these stocks

```
returns = (
    data
    ['Adj Close']
    .dropna(axis=1, how='all')
    .pct_change()
    .iloc[-5*252:]
)
```

Calculate the volatilities of 20 equal-weighted random portfolios of various portfolio sizes

Random portfolios should have portfolio sizes of 1, 2, 4, 6, 8, 10, 20, 30, 40, or 50 stocks each.

You can combine the `.sample(n=?, axis=1, random_state=?), .mean(axis=1),` and `.std()` to calculate the volatilities of equal-weighted portfolios. You can collect these volatilities in a list of lists built with two `for` loops or list comprehensions. Replace the `?`s in `.sample()` with loop counters. The inner loop will calculate a portfolio volatility for each portfolio size, and the outer loop will collect 20 versions of each portfolio. Using the outer loop counter for `random_state=` makes your analysis repeatable!

```
portfolio_size = [1, 2, 4, 6, 8, 10, 20, 30, 40, 50]
portfolio_number = 20

list_of_volatilities = []
for i in range(portfolio_number):
    list_of_volatilities.append([returns.sample(n=j, axis=1, random_state=i).mean(axis=1).std() for j in portfolio_size])

list_of_volatilities[0]
```



```
[0.0153,
 0.0127,
 0.0121,
 0.0118,
 0.0134,
 0.0129,
 0.0128,
 0.0129,
 0.0129,
 0.0129]
```

Combine this list of lists into a data frame


```
volatilities = (
    pd.DataFrame(
        data=list_of_volatilities,
        index=range(1, 1+portfolio_number),
        columns=portfolio_size
    )
    .rename_axis(index='Portfolio Number', columns='Portfolio Size')
)
```

volatilities

Portfolio Size	1	2	4	6	8	10	20	30	40	50
Portfolio Number										
1	0.0153	0.0127	0.0121	0.0118	0.0134	0.0129	0.0128	0.0129	0.0129	0.0129
2	0.0193	0.0144	0.0127	0.0145	0.0144	0.0145	0.0137	0.0136	0.0136	0.0133
3	0.0160	0.0165	0.0161	0.0151	0.0142	0.0144	0.0132	0.0134	0.0133	0.0134
4	0.0193	0.0177	0.0172	0.0171	0.0157	0.0156	0.0141	0.0141	0.0142	0.0139
5	0.0205	0.0199	0.0151	0.0164	0.0141	0.0133	0.0127	0.0131	0.0128	0.0131
6	0.0182	0.0157	0.0123	0.0136	0.0136	0.0139	0.0137	0.0143	0.0137	0.0136
7	0.0199	0.0174	0.0171	0.0160	0.0148	0.0142	0.0124	0.0133	0.0130	0.0128
8	0.0280	0.0178	0.0142	0.0149	0.0144	0.0131	0.0141	0.0138	0.0134	0.0134
9	0.0150	0.0136	0.0115	0.0116	0.0117	0.0119	0.0131	0.0132	0.0131	0.0133
10	0.0141	0.0171	0.0164	0.0150	0.0146	0.0147	0.0138	0.0131	0.0128	0.0126
11	0.0285	0.0210	0.0185	0.0170	0.0163	0.0167	0.0143	0.0142	0.0136	0.0137
12	0.0172	0.0193	0.0148	0.0154	0.0137	0.0140	0.0137	0.0133	0.0125	0.0128
13	0.0150	0.0131	0.0124	0.0126	0.0142	0.0134	0.0138	0.0133	0.0131	0.0135
14	0.0280	0.0198	0.0160	0.0147	0.0147	0.0137	0.0137	0.0131	0.0133	0.0134
15	0.0282	0.0197	0.0152	0.0142	0.0148	0.0142	0.0143	0.0140	0.0133	0.0134
16	0.0188	0.0150	0.0148	0.0136	0.0138	0.0140	0.0138	0.0141	0.0137	0.0136
17	0.0165	0.0132	0.0149	0.0141	0.0133	0.0128	0.0130	0.0133	0.0133	0.0131
18	0.0213	0.0198	0.0177	0.0166	0.0161	0.0156	0.0145	0.0142	0.0135	0.0139
19	0.0132	0.0136	0.0143	0.0140	0.0150	0.0145	0.0139	0.0134	0.0131	0.0127
20	0.0205	0.0209	0.0169	0.0154	0.0151	0.0143	0.0137	0.0134	0.0132	0.0132

Calculate the mean volatility for each portfolio size and replicate the plot above

```
volatilities.mul(100 * np.sqrt(252)).mean().plot()
plt.xlabel('Portfolio Size')
plt.ylabel('Mean of Annualized Volatility (%)')
plt.title('Diminishing Returns to Diversification (Elton and Gruber, 1977)')
plt.show()
```



Herron Topic 1 - Log and Simple Returns, Portfolio Math, and Applications - Sec 04

This notebook covers two topics:

1. Log and simple returns
2. Portfolio returns, plus two applications of portfolio returns

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Log and Simple Returns

We will typically use *simple* returns, calculated as $r_{simple,t} = \frac{p_t + d_t - p_{t-1}}{p_{t-1}}$. This simple return is the return that investors earn on their investments. We can calculate simple returns from Yahoo Finance data with the `.pct_change()` method on the adjusted close column (i.e., `Adj Close`), which adjusts for dividends and splits. The adjusted close column is a reverse-engineered close price (i.e., end-of-trading-day price) that incorporates dividends and splits, making simple return calculations easy.

However, we may see *log* returns elsewhere, which are the (natural) log of one plus simple returns: $r_{log,t} = \log(1 + r_{simple,t})$. Therefore, we calculate log returns as either the log of one plus simple returns or the difference of the logs of the adjusted close column. Log returns are also known as *continuously-compounded* returns.


This section explains the differences between simple and log returns and where each is appropriate.

Simple and Log Returns are Similar for Small Returns

$x \approx \log(1 + x)$ for small values of x , so simple returns and log returns are similar for small returns. Returns are typically small at daily and monthly horizons, so the difference between simple and log returns is small at daily and monthly horizons. The following figure shows that $r_{simple,t} \approx r_{log,t}$ for small values of r .

```
simpler = np.linspace(-0.75, 0.75, 100)
logr = np.log1p(simpler)
```

```
plt.plot(simpler, logr)
plt.plot([-1, 1], [-1, 1])
plt.xlabel('Simple Return')
plt.ylabel('Log Return')
plt.title('Log Versus Simple Returns')
plt.legend(['Actual Relation', 'Relation If Log = Simple'])
plt.show()
```



Simple Return Advantage: Portfolio Calculations

For a portfolio of N assets with portfolio weights w_i , the portfolio return r_p is the weighted average of the returns of its assets: $r_p = \sum_{i=1}^N w_i r_i$. For example, for an equal-weighted portfolio with two stocks, $r_p = 0.5r_1 + 0.5r_2 = \frac{r_1+r_2}{2}$. Therefore, we cannot calculate portfolio returns with log returns because the sum of logs is the log of products. That is $\log(1 + r_i) + \log(1 + r_j) = \log((1 + r_i) \times (1 + r_j))$, which is not what we want to measure! **We cannot perform portfolio calculations with log returns!**

Log Return Advantage: Log Returns are Additive

We compound simple returns with multiplication, *but we compound log returns with addition*. This additive property of log returns makes code simple, computations fast, and proofs easy when we must compound returns.

We compound returns from $t = 0$ to $t = T$ as follows:

$$1 + r_{0,T} = (1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)$$

Next, we take the log of both sides of the previous equation and use the property that the log of products is the sum of logs:

$$\log(1 + r_{0,T}) = \log((1 + r_1) \times (1 + r_2) \times \cdots \times (1 + r_T)) = \log(1 + r_1) + \log(1 + r_2) + \cdots + \log(1 + r_T) = \sum_{t=1}^T \log(1 + r_t)$$

Next, we exponentiate both sides of the previous equation:

$$e^{\log(1 + r_{0,T})} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Next, we use the property that $e^{\log(x)} = x$ to simplify the previous equation:

$$1 + r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)}$$

Finally, we subtract 1 from both sides:

$$r_{0,T} = e^{\sum_{t=0}^T \log(1 + r_t)} - 1$$

So, the return $r_{0,T}$ from $t = 0$ to $t = T$ is the exponentiated sum of log returns. The pandas developers assume users understand the math above and focus on optimizing sums!

```
np.random.seed(42)
df = pd.DataFrame(data={'r': np.exp(np.random.randn(10_000)) - 1})

df.describe()
```

	r
count	10000.0000
mean	0.6529
std	2.1918
min	-0.9802
25%	-0.4896
50%	-0.0026
75%	0.9564
max	49.7158

We can time the calculation of 10-observation rolling returns. We use `.apply()` for the simple return version because `.rolling()` does not have a product method. We find that `.rolling()` is slower with `.apply()` than with `.sum()` by a factor of about 1,000. *We will learn about `.rolling()` and `.apply()` in a few weeks, but they provide the best example of when to use log returns.*

```
%%timeit
df['r10_via_prod'] = (
    df['r']
    .add(1)
    .rolling(10)
    .apply(lambda x: x.prod())
    .sub(1)
)
```

208 ms ± 40.1 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
%%timeit
df['r10_via_sum'] = (
    df['r']
    .add(1)
    .pipe(np.log)
    .rolling(10)
    .sum()
    .pipe(np.exp)
    .sub(1)
)
```

1.49 ms ± 443 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
df.head(15)
```

	r	r10_via_prod	r10_via_sum
0	0.6433	NaN	NaN
1	-0.1291	NaN	NaN
2	0.9111	NaN	NaN
3	3.5861	NaN	NaN
4	-0.2088	NaN	NaN
5	-0.2087	NaN	NaN
6	3.8511	NaN	NaN
7	1.1542	NaN	NaN
8	-0.3747	NaN	NaN
9	0.7204	87.2886	87.2886
10	-0.3709	32.8006	32.8006
11	-0.3723	23.3617	23.3617
12	0.2737	15.2369	15.2369
13	-0.8524	-0.4774	-0.4774
14	-0.8218	-0.8823	-0.8823

```
np.allclose(df['r10_via_prod'], df['r10_via_sum'], equal_nan=True)
```

True

These two approaches calculate the same return series, but the simple-return approach using `.prod()` is about 1,000 times slower than the log-return approach using `.sum()!` **We can use log returns to calculate total returns very quickly!**

Portfolio Math

Portfolio return r_p is the weighted average of its asset returns, so $r_p = \sum_{i=1}^N w_i r_i$. Here N is the number of assets, w_i is the weight on asset i , and $\sum_{i=1}^N w_i = 1$.

The 1/N Portfolio

The $\frac{1}{N}$ portfolio equally weights portfolio assets, so $w_1 = w_2 = \dots = w_N = \frac{1}{N}$. If $w_i = \frac{1}{N}$, then $r_p = \sum_{i=1}^N \frac{1}{N} r_i = \frac{\sum_{i=1}^N r_i}{N} = \bar{r}$. Therefore, we can use `.mean(axis=1)` to calculate $\frac{1}{N}$ portfolio returns!

```
df2 = yf.download(tickers='AAPL AMZN GOOG MSFT NVDA TSLA', auto_adjust=False, progress=False)
returns2 = df2['Adj Close'].pct_change().dropna()

returns2.describe()
```

Ticker	AAPL	AMZN	GOOG	MSFT	NVDA	TSLA
count	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000	3689.0000
mean	0.0011	0.0012	0.0009	0.0010	0.0021	0.0021
std	0.0175	0.0205	0.0173	0.0161	0.0288	0.0362
min	-0.1286	-0.1405	-0.1110	-0.1474	-0.1876	-0.2106
25%	-0.0074	-0.0089	-0.0071	-0.0070	-0.0122	-0.0164
50%	0.0010	0.0010	0.0009	0.0007	0.0017	0.0012
75%	0.0102	0.0119	0.0093	0.0093	0.0161	0.0194
max	0.1198	0.1575	0.1605	0.1422	0.2981	0.2440

```
returns2.mean() # implied axis=0
```

```
Ticker
AAPL    0.0011
AMZN    0.0012
GOOG    0.0009
MSFT    0.0010
NVDA    0.0021
TSLA    0.0021
dtype: float64
```

```
rp2_via_mean = returns2.mean(axis=1)

rp2_via_mean
```

```
Date
2010-06-30   -0.0123
2010-07-01   -0.0107
2010-07-02   -0.0271
2010-07-06   -0.0223
2010-07-07    0.0254
...
2025-02-21   -0.0272
2025-02-24   -0.0127
```

```
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Note that when we apply the same portfolio weights every period, we rebalance at the same frequency as the returns data. If we have daily data, rebalance daily. If we have monthly data, we rebalance monthly, and so on.

A More General Solution

If we combine portfolio weights into vector w and the time series of asset returns into matrix \mathbf{R} , then we can calculate the time series of portfolio returns as $r_p = w^T \mathbf{R}$. The pandas version of this calculation is `R.dot(w)`, where `R` is a data frame of asset returns and `w` is a series or an array of portfolio weights. We can use this approach to calculate $\frac{1}{N}$ portfolio returns, too.

```
returns2.shape
```

```
(3689, 6)
```

```
weights2 = np.ones(returns2.shape[1]) / returns2.shape[1]
```

```
weights2
```

```
array([0.1667, 0.1667, 0.1667, 0.1667, 0.1667, 0.1667])
```

```
rp2_via_dot = returns2.dot(weights2)
```

```
rp2_via_dot
```

Date	
2010-06-30	-0.0123
2010-07-01	-0.0107
2010-07-02	-0.0271
2010-07-06	-0.0223
2010-07-07	0.0254
	...
2025-02-21	-0.0272
2025-02-24	-0.0127

```
2025-02-25    -0.0247
2025-02-26    -0.0055
2025-02-27    -0.0330
Length: 3689, dtype: float64
```

Both approaches give the same answer!

```
np.allclose(rp2_via_mean, rp2_via_dot, equal_nan=True)
```

```
True
```

Portfolio Math Application 1: All stocks half the time or half stocks all the time?

Are you better off investing:

1. 100% in stocks 50% of the time and the riskless asset the other 50% of the time *or*
2. 50% in stocks and 50% in the riskless asset 100% of the time?

Here is a roadmap for convincing yourself with data!

Please see Kritzman (2000, Chapter 5) for a more detailed solution!

Download annual market and risk-free asset returns from Kenneth French's data library

```
pdr.famafrench.get_available_datasets()[:5]
```

```
['F-F_Research_Data_Factors',
 'F-F_Research_Data_Factors_weekly',
 'F-F_Research_Data_Factors_daily',
 'F-F_Research_Data_5_Factors_2x3',
 'F-F_Research_Data_5_Factors_2x3_daily']
```

```
ff = pdr.DataReader(
    name='F-F_Research_Data_Factors',
    data_source='famafrench',
    start='1900'
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_24988\304025689.py:1: FutureWarning: The argument 'ff' is deprecated; use 'pdr.DataReader()' instead.
  ff = pdr.DataReader()
C:\Users\r.herron\AppData\Local\Temp\ipykernel_24988\304025689.py:1: FutureWarning: The argument 'ff' is deprecated; use 'pdr.DataReader()' instead.
  ff = pdr.DataReader()

print(ff['DESCR'])
```

F-F Research Data Factors

```
-----
```

This file was created by CMPT_ME_BEME_RETTS using the 202412 CRSP database. The 1-month TBill

```
0 : (1182 rows x 4 cols)
1 : Annual Factors: January-December (98 rows x 4 cols)
```

Convert these factors to decimal returns and calculate the market return series

```
df3 = ff[1].div(100)
```

```
df3['Mkt'] = df3['Mkt-RF'] + df3['RF']
```

Add a portfolio return series that is half stocks all the time

You might call this portfolio return series Balanced

```
df3['Balanced'] = df3[['Mkt', 'RF']].mean(axis=1)
```

Add a portfolio return series that switches between stocks and bills every year with stocks in odd years

You might call this portfolio return series Switching Stocks Odd

```
df3['Switching Stocks Odd'] = np.where(df3.index.year % 2 == 1, df3['Mkt'], df3['RF'])
```

Add a portfolio return series that switches between stocks and bills every year with stocks in even years

You might call this portfolio return series **Switching Stocks Even**

```
df3['Switching Stocks Even'] = np.where(df3.index.year % 2 == 0, df3['Mkt'], df3['RF'])
```

```
df3.head()
```

Date	Mkt-RF	SMB	HML	RF	Mkt	Balanced	Switching Stocks Odd	Switching Stocks E
1927	0.2947	-0.0204	-0.0454	0.0312	0.3259	0.1785	0.3259	0.0312
1928	0.3539	0.0451	-0.0617	0.0356	0.3895	0.2126	0.0356	0.3895
1929	-0.1954	-0.3070	0.1167	0.0475	-0.1479	-0.0502	-0.1479	0.0475
1930	-0.3123	-0.0517	-0.1154	0.0241	-0.2882	-0.1321	0.0241	-0.2882
1931	-0.4511	0.0370	-0.1395	0.0107	-0.4404	-0.2149	-0.4404	0.0107


Plot the cumulative returns on a \$1 investment and calculate the summary statistics for the Balanced and Switching series

Use the `.describe()` method to calculate summary statistics.

```
portfolios = ['Balanced', 'Switching Stocks Odd', 'Switching Stocks Even']
```

```
import matplotlib.ticker as ticker
```

```
(  
    df3[portfolios]  
    .add(1)  
    .cumprod()  
    .plot()  
)  
  
plt.yscale('log')  
plt.ylabel('Value of $1 Investment')  
plt.title(f'Value of $1 in Investment in Three Strategies\nfrom {df3.index.year[0]} through  
plt.show()
```



Which strategy do you prefer?

Why? How sure are you?

The **Switching Stocks Odd** strategy *seems* best in this example! However, this apparent superiority is largely a matter of **luck**, specifically this combination of the sample period and starting year. In this sample, the number of even and odd years is equal, but the market happened to perform better in odd years. This outcome is entirely driven by **this particular historical sequence**.

To draw broader conclusions, we must break free from the “luck of the draw” tied to this single realization of history. Below, we will explore two approaches to determine which strategy is better:

- Theory:** We will analyze the expected returns and variances mathematically to assess the underlying return-risk tradeoff for each strategy.

2. **Simulation:** To rely on data instead of theory, we will simulate thousands of alternative historical sequences.

By shuffling the 98 years of market data repeatedly, we can generate 10,000 random samples and evaluate the performance of each strategy across these simulations.

This approach allows us to:

- Confirm that **both strategies have the same arithmetic average return.**
- Show that the **balanced strategy consistently delivers a better return-risk tradeoff (Sharpe ratio)** when evaluated across many possible outcomes.

We will do the simulation first.

Use the `simulate()` function to simulate 10,000 different outcomes for the U.S. market

`simulate()` calculates one **Switching** return series because the randomization also randomizes the odd-year and even-year choice.

```
df3[['Mkt', 'RF']].sample(frac=1, ignore_index=True, random_state=42).head()
```

	Mkt	RF
0	0.2886	0.0837
1	0.2870	0.0421
2	0.2361	0.0004
3	0.3871	0.0033
4	-0.3674	0.0160

```
def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced and switching

    Parameters:
    df (pd.DataFrame): The input DataFrame.
    cols (list): List of column names to sample.
    n_iter (int): Number of iterations for simulation.

    Returns:
    pd.DataFrame: A concatenated DataFrame with simulation results.
    """
    return pd.concat(
        objs=[(
```

```

df[cols]
    .sample(frac=1, ignore_index=True, random_state=i)
    .assign(
        Balanced=lambda x: x[cols].mean(axis=1),
        Switching=lambda x: np.where(x.index % 2 == 0, x[cols[0]], x[cols[1]]))
    )
) for i in range(n_iter)],
keys=range(n_iter),
names=['Simulation', 'Year']
)

```

```
df4 = simulate(df3)
```

```
df4.head()
```

			Mkt
Simulation			Year
0	0		0.0077
	1		0.1633
	2		-0.1479
	3		0.2120
	4		0.3682

Calculate the summary statistics for these new `Balanced` and `Switching` series

```
df4[['Balanced', 'Switching']].agg(['mean', 'std'])
```

	Balanced	Switching
mean	0.0772	0.0774
std	0.1001	0.1487

We see that `Balanced` and `Switching` have the same mean return, but `Balanced` has much lower volatility than `Switching`! A risk-averse investor prefers `Balanced` because it has a higher return/risk ratio than `Switching`. We can quantify this ratio as the Sharpe ratio, which is the mean excess return divided by the volatility of excess returns: $S_i = \frac{r_i - r_f}{\sigma(r_i - r_f)}$. We can calculate excess returns and Sharpe ratios in one code snippet.

```
( df4
  [['Balanced', 'Switching']]
  .sub(df4['RF'], axis=0)
  .agg(lambda x: x.mean() / x.std())
  .to_frame('Sharpe ratio')
)
```

	Sharpe ratio
Balanced	0.4368
Switching	0.2961

We have to do a little more work if we want to combine mean and volatility of *raw* returns with the Sharpe ratio of *excess* returns.

```
pd.concat(
    objs=[
        df4[['Balanced', 'Switching']].agg(['mean', 'std']).transpose(),
        df4[['Balanced', 'Switching']].sub(df4['RF'], axis=0).agg(lambda x: x.mean() / x.std)
    ],
    axis=1
)
```

	mean	std	Sharpe ratio
Balanced	0.0772	0.1001	0.4368
Switching	0.0774	0.1487	0.2961

Which strategy do you prefer?

Why? How sure are you?

We prefer **Balanced** because it has a Sharpe ratio about 50% greater than **Switching**! We can see this in the data above, and here is the theory.

For **Balanced**, $\sigma_p^2 = w_m^2 \sigma_m^2 + w_f^2 \sigma_f^2 + 2w_m w_f \sigma_m \sigma_f \rho_{m,f}$. Because $\sigma_f^2 \approx 0$ and $\rho_{m,f} \approx 0$, $\sigma_p^2 \approx w_m^2 \sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx w_m \sigma_m = \frac{1}{2} \sigma_m$. We see this in the data!

```
0.5 * df3['Mkt'].std()
```

0.0998

```
df4['Balanced'].std()
```

0.1001

For **Switching**, $\sigma_p^2 = w_m\sigma_m^2 + w_f\sigma_f^2 + w_m w_f(\mu_m - \mu_f)^2$. We have a different formula because **Switching** is diversified *over time* instead at a point in time! Because $\sigma_f^2 \approx 0$ and $(\mu_m - \mu_f)^2 \approx 0$, $\sigma_p^2 \approx w_m\sigma_m^2$. Therefore, for **Balanced** $\sigma_p \approx \sqrt{w_m}\sigma_m = \sqrt{\frac{1}{2}}\sigma_m$. We see this in the data!

```
np.sqrt(0.5) * df3['Mkt'].std()
```

0.1411

```
df4['Switching'].std()
```

0.1487

The $(\mu_m - \mu_f)^2$ is close to zero but not exactly zero. We can get even closer to the observed **Switching** portfolio volatility if we consider this term!

```
np.sqrt(0.5 * df3['Mkt'].var() + 0.5 * 0.5 * (df3['Mkt'].mean() - df3['RF'].mean())**2)
```

0.1478

Here are two after-class additions to this question.

Can we add a progress bar?

Yes! We have to rewrite the `simulate()` function and replace the list comprehension with a for loop. However, this addition requires the `tqdm` package. To avoid any conflicts, I will not install this package and provide the code as markdown instead of executable code.

```
from tqdm import tqdm

def simulate(df, cols=['Mkt', 'RF'], n_iter=10_000):
    """
    Simulates resampling of the given DataFrame columns and computes balanced
    and switching portfolio returns.

    Parameters:
    -----
    df : pd.DataFrame
        The input DataFrame.
    cols : list
        List of column names to sample.
    n_iter : int
        Number of iterations for simulation.

    Returns:
    -----
    pd.DataFrame
        A concatenated DataFrame with simulation results.
    """
    all_resamples = []

    # Use tqdm to visualize progress
    for i in tqdm(range(n_iter), desc="Simulating"):
        # Random re-sampling
        resampled = (
            df[cols]
            .sample(frac=1, ignore_index=True, random_state=i)
            .assign(
                Balanced=lambda x: x[cols].mean(axis=1),
                Switching=lambda x: np.where(
                    x.index % 2 == 0,
                    x[cols[0]],
                    x[cols[1]]
                )
            )
        )
        all_resamples.append(resampled)

    # Concatenate final results
    return pd.concat(all_resamples, keys=range(n_iter), names=['Simulation', 'Year'])
```

Ideally, we calculate the mean and volatility of each simulation, then take the average

We can do this easily with either `.groupby()` or `.pivot_table()!`

```
(  
    df4  
    .groupby(level='Simulation')  
    [['Balanced', 'Switching']]  
    .agg(['mean', 'std'])  
    .mean()  
    .unstack()  
)
```

	mean	std
Balanced	0.0772	0.1006
Switching	0.0774	0.1488

These statistics are similar, but not identical!

Portfolio Math Application 2: What are the benefits of diversification?

Use random portfolios of S&P 100 stocks of various portfolio sizes to show that portfolio volatility falls quickly, then slowly, then not at all as we increase portfolio size.

Download daily data for the stocks in the S&P 100

Wikipedia provides tickers for the stocks in the [S&P 100](#). Use a list comprehension to replace
. in tickers with - for compatibility with Yahoo! Finance.

```
wiki = pd.read_html('https://en.wikipedia.org/wiki/S%26P_100')  
tickers = [i.replace('.', '-') for i in wiki[2]['Symbol']]  
data = yf.download(tickers=tickers, auto_adjust=False, progress=False).iloc[:-1]
```

Calculate the past five years of daily returns for these stocks

```
returns = (
    data
    ['Adj Close']
    .dropna(axis=1, how='all')
    .pct_change()
    .iloc[-5*252:]
)
```

Calculate the volatilities of 20 equal-weighted random portfolios of various portfolio sizes

Random portfolios should have portfolio sizes of 1, 2, 4, 6, 8, 10, 20, 30, 40, or 50 stocks each.

You can combine the `.sample(n=?, axis=1, random_state=?), .mean(axis=1),` and `.std()` to calculate the volatilities of equal-weighted portfolios. You can collect these volatilities in a list of lists built with two `for` loops or list comprehensions. Replace the `?`s in `.sample()` with loop counters. The inner loop will calculate a portfolio volatility for each portfolio size, and the outer loop will collect 20 versions of each portfolio. Using the outer loop counter for `random_state=` makes your analysis repeatable!

```
portfolio_size = [1, 2, 4, 6, 8, 10, 20, 30, 40, 50]
portfolio_number = 20

list_of_volatilities = []
for i in range(portfolio_number):
    list_of_volatilities.append([returns.sample(n=j, axis=1, random_state=i).mean(axis=1).std() for j in portfolio_size])

list_of_volatilities[0]
```



```
[0.0153,
 0.0127,
 0.0121,
 0.0118,
 0.0134,
 0.0129,
 0.0128,
 0.0129,
 0.0129,
 0.0129]
```

Combine this list of lists into a data frame


```
volatilities = (
    pd.DataFrame(
        data=list_of_volatilities,
        index=range(1, 1+portfolio_number),
        columns=portfolio_size
    )
    .rename_axis(index='Portfolio Number', columns='Portfolio Size')
)
```

volatilities

Portfolio Size	1	2	4	6	8	10	20	30	40	50
Portfolio Number										
1	0.0153	0.0127	0.0121	0.0118	0.0134	0.0129	0.0128	0.0129	0.0129	0.0129
2	0.0193	0.0144	0.0127	0.0145	0.0144	0.0145	0.0137	0.0136	0.0136	0.0133
3	0.0160	0.0165	0.0161	0.0151	0.0142	0.0144	0.0132	0.0134	0.0133	0.0134
4	0.0193	0.0177	0.0172	0.0171	0.0157	0.0156	0.0141	0.0141	0.0142	0.0139
5	0.0205	0.0199	0.0151	0.0164	0.0141	0.0133	0.0127	0.0131	0.0128	0.0131
6	0.0182	0.0157	0.0123	0.0136	0.0136	0.0139	0.0137	0.0143	0.0137	0.0136
7	0.0199	0.0174	0.0171	0.0160	0.0148	0.0142	0.0124	0.0133	0.0130	0.0128
8	0.0280	0.0178	0.0142	0.0149	0.0144	0.0131	0.0141	0.0138	0.0134	0.0134
9	0.0150	0.0136	0.0115	0.0116	0.0117	0.0119	0.0131	0.0132	0.0131	0.0133
10	0.0141	0.0171	0.0164	0.0150	0.0146	0.0147	0.0138	0.0131	0.0128	0.0126
11	0.0285	0.0210	0.0185	0.0170	0.0163	0.0167	0.0143	0.0142	0.0136	0.0137
12	0.0172	0.0193	0.0148	0.0154	0.0137	0.0140	0.0137	0.0133	0.0125	0.0128
13	0.0150	0.0131	0.0124	0.0126	0.0142	0.0134	0.0138	0.0133	0.0131	0.0135
14	0.0280	0.0198	0.0160	0.0147	0.0147	0.0137	0.0137	0.0131	0.0133	0.0134
15	0.0282	0.0197	0.0152	0.0142	0.0148	0.0142	0.0143	0.0140	0.0133	0.0134
16	0.0188	0.0150	0.0148	0.0136	0.0138	0.0140	0.0138	0.0141	0.0137	0.0136
17	0.0165	0.0132	0.0149	0.0141	0.0133	0.0128	0.0130	0.0133	0.0133	0.0131
18	0.0213	0.0198	0.0177	0.0166	0.0161	0.0156	0.0145	0.0142	0.0135	0.0139
19	0.0132	0.0136	0.0143	0.0140	0.0150	0.0145	0.0139	0.0134	0.0131	0.0127
20	0.0205	0.0209	0.0169	0.0154	0.0151	0.0143	0.0137	0.0134	0.0132	0.0132

Calculate the mean volatility for each portfolio size and replicate the plot above

```
volatilities.mul(100 * np.sqrt(252)).mean().plot()
plt.xlabel('Portfolio Size')
plt.ylabel('Mean of Annualized Volatility (%)')
plt.title('Diminishing Returns to Diversification (Elton and Gruber, 1977)')
plt.show()
```



Week 6

McKinney Chapter 10 - Data Aggregation and Group Operations

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

Chapter 10 of McKinney (2022) discusses groupby operations, the pandas equivalent of pivot tables in Excel. Pivot tables calculate statistics (e.g., sums, means, and medians) for one set of variables by groups of other variables (e.g., weekdays and tickers). For example, we could use a pivot table to calculate mean daily stock returns by weekday.

We will focus on:

1. The `.groupby()` method to group by columns and indexes
2. The `.agg()` method to aggregate columns to single values
3. The `.pivot_table()` method as an alternative to `.groupby()`

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

GroupBy Mechanics

“Split-apply-combine” is an excellent way to describe pandas groupby operations.

Hadley Wickham, an author of many popular packages for the R programming language, coined the term split-apply-combine for describing group operations. In the first stage of the process, data contained in a pandas object, whether a Series, DataFrame, or otherwise, is split into groups based on one or more keys that you provide. The splitting is performed on a particular axis of an object. For example, a DataFrame can be grouped on its rows (axis=0) or its columns (axis=1). Once this is done, a function is applied to each group, producing a new value. Finally, the results of all those function applications are combined into a result object. The form of the resulting object will usually depend on what's being done to the data. See Figure 10-1 for a mockup of a simple group aggregation.

[Figure 10-1](#) visualizes a split-apply-combine operation that:

1. Splits by the `key` column (i.e., “groups by `key`”)
2. Applies the sum operation to the `data` column (i.e., “and sums `data`”)
3. Combines the grouped sums (i.e., “combines the output”)

We could describe this operation as “sum the `data` column by groups of the `key` column then combines the output.”

```
np.random.seed(42)
df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                   'key2' : ['one', 'two', 'one', 'two', 'one'],
                   'data1' : np.random.randn(5),
                   'data2' : np.random.randn(5)})
```

df

	key1	key2	data1	data2
0	a	one	0.4967	-0.2341
1	a	two	-0.1383	1.5792
2	b	one	0.6477	0.7674
3	b	two	1.5230	-0.4695
4	a	one	-0.2342	0.5426

Here is the manual way to calculate the means of `data1` by groups of `key1`.

```
df.loc[df['key1'] == 'a', 'data1'].mean()
```

0.0414

```
df.loc[df['key1'] == 'b', 'data1'].mean()
```

1.0854

We can do this calculation more easily!

1. Use the `.groupby()` method to group by `key1`
2. Use the `.mean()` method to calculate the mean of `data1` within each value of `key1`

```
df['data1'].groupby(df['key1']).mean()
```

```
key1
a    0.0414
b    1.0854
Name: data1, dtype: float64
```

We can wrap `data1` with two sets of square brackets if we prefer our result as a data frame instead of a series.

```
df[['data1']].groupby(df['key1']).mean()
```

		data1
		key1
a		0.0414
b		1.0854

We can group by more than one variable!

```
means = df['data1'].groupby([df['key1'], df['key2']]).mean()
means
```

```
key1  key2
a    one    0.1313
      two   -0.1383
b    one    0.6477
      two    1.5230
Name: data1, dtype: float64
```

We can use the `.unstack()` method if we want to use both rows and columns to organize data. Recall that the `.unstack()` method un-stacks the inner index level (i.e., `level = -1`) by default so that `key2` values become the columns.

```
means.unstack()
```

key2	one	two
key1		
a	0.1313	-0.1383
b	0.6477	1.5230

Our grouping variables are typically columns in the data frame we want to group, so the following syntax is more compact and easier to read.

```
df.groupby(['key1', 'key2'])['data1'].mean().unstack()
```

key2	one	two
key1		
a	0.1313	-0.1383
b	0.6477	1.5230

We can wrap long chains in parentheses to insert line breaks and improve readability.

```
(  
    df  
    .groupby(['key1', 'key2'])  
    ['data1']  
    .mean()  
    .unstack()  
)
```

key2	one	two
key1		
a	0.1313	-0.1383
b	0.6477	1.5230

However, we must pass only numerical columns to numerical aggregation methods. Otherwise, pandas will give a type error. For example, in the following code, pandas unsuccessfully tries to calculate the mean of string `key2`.

```
# # TypeError: agg function failed [how->mean,dtype->object]
# df.groupby('key1').mean()
```

We avoid this error by slicing the numerical columns.

```
df.groupby('key1')[['data1', 'data2']].mean()
```

	data1	data2
key1		
a	0.0414	0.6292
b	1.0854	0.1490

Grouping with Functions

We can also group with functions. Below, we group with the `len` function, which calculates the lengths of the labels in the row index.

```
np.random.seed(42)
people = pd.DataFrame(
    data=np.random.randn(5, 5),
    columns=['a', 'b', 'c', 'd', 'e'],
    index=['Joe', 'Steve', 'Wes', 'Jim', 'Travis']
)

people
```

	a	b	c	d	e
Joe	0.4967	-0.1383	0.6477	1.5230	-0.2342
Steve	-0.2341	1.5792	0.7674	-0.4695	0.5426
Wes	-0.4634	-0.4657	0.2420	-1.9133	-1.7249
Jim	-0.5623	-1.0128	0.3142	-0.9080	-1.4123
Travis	1.4656	-0.2258	0.0675	-1.4247	-0.5444

```
people.groupby(len).sum()
```

	a	b	c	d	e
3	-0.5290	-1.6168	1.2039	-1.2983	-3.3714
5	-0.2341	1.5792	0.7674	-0.4695	0.5426
6	1.4656	-0.2258	0.0675	-1.4247	-0.5444

We can mix functions, lists, dictionaries, etc., as arguments to the `.groupby()` method.

```
key_list = ['one', 'one', 'one', 'two', 'two']
people.groupby([len, key_list]).min()
```

		a
3	one	-0.4634
5	two	-0.5623
6	one	-0.2341
	two	1.4656

```
d = {'Joe': 'a', 'Jim': 'b'}
people.groupby([len, d]).min()
```

		a	b
3	a	0.4967	-0
	b	-0.5623	-1

```
d_2 = {'Joe': 'Cool', 'Jim': 'Nerd', 'Travis': 'Cool'}
people.groupby([len, d_2]).min()
```

		a
3	Cool	0.4967
	Nerd	-0.5623
6	Cool	1.4656

Grouping by Index Levels

We can also group by index levels.

```
columns = pd.MultiIndex.from_arrays([['US', 'US', 'US', 'JP', 'JP'],
                                     [1, 3, 5, 1, 3]],
                                     names=['cty', 'tenor'])
hier_df = pd.DataFrame(np.random.randn(4, 5), columns=columns).transpose()
hier_df
```

			0
cty			tenor
US		1	0.1109
		3	-1.1510
		5	0.3757
JP		1	-0.6006
		3	-0.2917

```
hier_df.groupby(level='cty').count()
```

	0	1	2	3
cty				
JP	2	2	2	2
US	3	3	3	3

```
hier_df.groupby(level='tenor').count()
```

	0	1	2	3
tenor				
1	2	2	2	2
3	2	2	2	2
5	1	1	1	1

Data Aggregation

Table 10-1 summarizes the optimized groupby methods:

- **count**: Number of non-NA values in the group
- **sum**: Sum of non-NA values
- **mean**: Mean of non-NA values

- `median`: Arithmetic median of non-NA values
- `std, var`: Unbiased ($n - 1$ denominator) standard deviation and variance
- `min, max`: Minimum and maximum of non-NA values
- `prod`: Product of non-NA values
- `first, last`: First and last non-NA values

These optimized methods are fast and efficient. Still, pandas lets us use non-optimized methods. First, `any` series method is available.

```
df
```

	key1	key2	data1	data2
0	a	one	0.4967	-0.2341
1	a	two	-0.1383	1.5792
2	b	one	0.6477	0.7674
3	b	two	1.5230	-0.4695
4	a	one	-0.2342	0.5426

```
df.groupby('key1')['data1'].quantile(0.9)
```

```
key1
a    0.3697
b    1.4355
Name: data1, dtype: float64
```

```
0.6477 + 0.9 * (1.5230 - 0.6477)
```

1.4355

Second, we can write functions and pass them to the `.agg()` method. These functions should accept an array and return a single value.

```
def max_minus_min(arr):
    return arr.max() - arr.min()
```

```
df.sort_values(by=['key1', 'data1'])
```

	key1	key2	data1	data2
4	a	one	-0.2342	0.5426
1	a	two	-0.1383	1.5792
0	a	one	0.4967	-0.2341
2	b	one	0.6477	0.7674
3	b	two	1.5230	-0.4695

```
df.groupby('key1')['data1'].agg(max_minus_min)
```

```
key1
a    0.7309
b    0.8753
Name: data1, dtype: float64
```

1.5230 - 0.6477

0.8753

Some other methods work, too, even if they do not aggregate an array to a scalar.

```
df.groupby('key1')['data1'].describe()
```

	count	mean	std	min	25%	50%	75%	max
key1								
a	3.0000	0.0414	0.3972	-0.2342	-0.1862	-0.1383	0.1792	0.4967
b	2.0000	1.0854	0.6190	0.6477	0.8665	1.0854	1.3042	1.5230

The .agg() method provides two more handy features:

1. We can pass multiple functions to operate on all columns
2. We can pass specific functions to operate on specific columns

First, here are examples of multiple functions that operate on all columns.

```
df.groupby('key1')['data1'].agg(['mean', 'median', 'min', 'max'])
```

	mean	median	min	max
key1				
a	0.0414	-0.1383	-0.2342	0.4967
b	1.0854	1.0854	0.6477	1.5230

```
df.groupby('key1')[['data1', 'data2']].agg(['mean', 'median', 'min', 'max'])
```

key1	data1				data2			
	mean	median	min	max	mean	median	min	max
	a	0.0414	-0.1383	-0.2342	0.4967	0.6292	0.5426	-0.2341
b	1.0854	1.0854	0.6477	1.5230	0.1490	0.1490	-0.4695	0.7674

Second, here are examples of specific functions that operate on specific columns.

```
df.groupby('key1').agg({'data1': 'mean', 'data2': 'median'})
```

key1	data1	data2
	a	0.0414
b	1.0854	0.1490

We can calculate the mean *and standard deviation* of data1 and the median of data2 by key1.

```
df.groupby('key1').agg({'data1': ['mean', 'std'], 'data2': 'median'})
```

key1	data1		data2
	mean	std	median
	a	0.0414	0.3972
b	1.0854	0.6190	0.1490

Apply: General split-apply-combine

The `.agg()` method aggregates an array to a scalar. We can use the `.apply()` method for more general calculations that do not return a scalar. For example, the following `top()` function selects the top `n` rows in data frame `x` sorted by column `col`. The `.sort_values()` method sorts from low to high by default.

```
def top(x, col, n=1):
    return x.sort_values(col).head(n)
```

`df`

	key1	key2	data1	data2
0	a	one	0.4967	-0.2341
1	a	two	-0.1383	1.5792
2	b	one	0.6477	0.7674
3	b	two	1.5230	-0.4695
4	a	one	-0.2342	0.5426

```
top(
    x=df.loc[df['key1'] == 'a'],
    col='data1',
    n=2
)
```

	key1	key2	data1	data2
4	a	one	-0.2342	0.5426
1	a	two	-0.1383	1.5792

The following code returns the one row with the smallest value of `data1` within each group of `key1`. Note: we include the `include_groups=False` to suppress the `FutureWarning` and adopt the future default behavior now.

The following code returns the *two rows* with the smallest values of `data1` within each group of `key1`.

```
df.groupby('key1').apply(top, col='data1', include_groups=False)
```

	key2	data1	data2
key1			
a	4	one	-0.2342 0.5426
b	2	one	0.6477 0.7674

```
df.groupby('key1').apply(top, col='data1', n=2, include_groups=False)
```

key1	key2	data1	data2
a	4	one	-0.2342
	1	two	-0.1383
b	2	one	0.6477
	3	two	1.5230

We must use the `.reset_index()` method with the `drop=True` argument if we want to drop the index from `df`.

```
( df
    .groupby('key1')
    .apply(top, col='data1', n=2, include_groups=False)
    .reset_index(level=1, drop=True)
)
```

	key2	data1	data2
key1			
a	one	-0.2342	0.5426
a	two	-0.1383	1.5792
b	one	0.6477	0.7674
b	two	1.5230	-0.4695

i Note

The `.agg()` and `.apply()` methods both operate on groups created by the `.groupby()` method. However, they serve different purposes and have distinct use cases.

The `.agg()` method is designed for aggregating data, meaning it applies functions that reduce a group to a single value (e.g., mean, sum, or custom functions that return a

single scalar). This method is useful for summarizing data across groups. In contrast, the `.apply()` method is more general and flexible. The `.apply()` method returns results of varying shapes. The `.agg()` method is limited to scalar outputs for each group, but the `.apply()` method is not.

Pivot Tables and Cross-Tabulation

Above, we manually made pivot tables with the `.groupby()`, `.agg()`, `.apply()` and `.unstack()` methods. pandas provides Excel-style aggregations with the `.pivot_table()` method and the `pandas.pivot_table()` function. It is worthwhile to read the `.pivot_table()` docstring several times.

```
ind = (
    yf.download(
        tickers='^GSPC ^DJI ^IXIC ^FTSE ^N225 ^HSI',
        auto_adjust=False,
        progress=False
    )
    .iloc[:-1]
    .stack(future_stack=True)
)

ind.head()
```

	Price	Adj C
Date	Ticker	
	^DJI	NaN
	^FTSE	NaN
1927-12-30	^GSPC	17.66
	^HSI	NaN
	^IXIC	NaN

The default aggregation function for `.pivot_table()` is `.mean()`. For the remaining examples, we will only consider data from 2015 and later.

```
(  
    ind  
    .loc['2015':]
```

```

    .pivot_table(
        index='Ticker'
    )
)

```

Price Ticker	Adj Close	Close	High	Low	Open	Volume
^DJI	27935.0069	27935.0069	28079.9469	27774.3386	27931.3049	303142509.7886
^FTSE	7163.3634	7163.3634	7203.9626	7121.6113	7162.5127	814461430.5924
^GSPC	3395.6074	3395.6074	3413.1090	3375.7649	3395.1843	4015744083.7901
^HSI	23785.6004	23785.6004	23950.5568	23615.3132	23798.3853	2145570432.1457
^IXIC	9999.4784	9999.4784	10065.3920	9924.6220	9998.5261	3567504600.6265
^N225	25038.6233	25038.6233	25175.5387	24893.2712	25039.7854	98144699.7179

We can specify a different aggregation function with the `aggfunc` argument. We can use `values` to select specific variables, `pd.Grouper()` to sample different date windows, and `aggfunc` to select specific aggregation functions.

```

(
    ind
    .loc['2015':]
    .reset_index()
    .pivot_table(
        values='Close',
        index=pd.Grouper(key='Date', freq='YE'),
        columns='Ticker',
        aggfunc=['min', 'max']
    )
)

```

Ticker Date	min						max		
	^DJI	^FTSE	^GSPC	^HSI	^IXIC	^N225	^DJI	^FTS	
2015-12-31	15666.4404	5874.1001	1867.6100	20556.5996	4506.4902	16795.9609	18312.3906	7104.0	
2016-12-31	15660.1797	5537.0000	1829.0800	18319.5801	4266.8398	14952.0195	19974.6191	7142.7	
2017-12-31	19732.4004	7099.2002	2257.8301	22134.4707	5429.0801	18335.6309	24837.5098	7687.7	
2018-12-31	21792.1992	6584.7002	2351.1001	24585.5293	6192.9199	19155.7402	26828.3906	7877.5	
2019-12-31	22686.2207	6692.7002	2447.8899	25064.3594	6463.5000	19561.9609	28645.2598	7686.6	
2020-12-31	18591.9297	4993.8999	2237.3999	21696.1309	6860.6699	16552.8301	30606.4805	7674.6	

McKinney Chapter 10 - Data Aggregation and Group Operations

Ticker	min	^DJI	^FTSE	^GSPC	^HSI	^IXIC	^N225	max	^DJI	^FTS
Date										
2021-12-31	29982.6191	6407.5000	3700.6499	22744.8594	12609.1602	27013.2500	36488.6289	7420.7		
2022-12-31	28725.5098	6826.2002	3577.0300	14687.0195	10213.2900	24717.5293	36799.6484	7672.3		
2023-12-31	31819.1406	7256.8999	3808.1001	16201.4902	10305.2402	25716.8594	37710.1016	8014.2		
2024-12-31	37266.6719	7446.2998	4688.6802	14961.1797	14510.2998	31458.4199	45014.0391	8445.7		
2025-12-31	41938.4492	8201.5000	5827.0400	18874.1406	18544.4199	38142.3711	44882.1289	8807.4		

McKinney Chapter 10 - Practice - Blank

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

Replicate the following `.pivot_table()` output with `.groupby()`

```
ind = (
    yf.download(
        tickers='^GSPC ^DJI ^IXIC ^FTSE ^N225 ^HSI',
        auto_adjust=False,
        progress=False
    )
    .rename_axis(columns=['Variable', 'Index'])
    .stack(future_stack=True)
)
```

[*****] 2 of 6 completed [*****5]

```
a = (
    ind
    .loc['2015':]
    .reset_index()
    .pivot_table(
        values='Close',
        index=pd.Grouper(key='Date', freq='YE'),
        columns='Index',
        aggfunc=['min', 'max']
    )
)
```

Calculate the mean and standard deviation of returns by ticker for the MATANA (MSFT, AAPL, TSLA, AMZN, NVDA, and GOOG) stocks

Consider only dates with complete returns data. Try this calculation with wide and long data frames, and confirm your results are the same.

```
matana = (
    yf.download(
        tickers='MSFT AAPL TSLA AMZN NVDA GOOG',
        auto_adjust=False,
        progress=False
    )
    .rename_axis(columns=['Variable', 'Ticker'])
)
```

[*****] 2 of 6 completed [*****5

Calculate the mean and standard deviation of returns and the maximum of closing prices by ticker for the MATANA stocks

Calculate monthly means and volatilities for SPY and GOOG returns

Plot the monthly means and volatilities from the previous exercise

Assign the Dow Jones stocks to five portfolios based on the *preceding* month's volatility

Plot the time-series volatilities of these five portfolios

Calculate the *mean* monthly correlation between the Dow Jones stocks

Is market volatility higher during wars?

Here is some guidance:

1. Download the daily factor data from Ken French's website
2. Calculate daily market returns by summing the market risk premium and risk-free rates (Mkt-RF and RF, respectively)
3. Calculate the volatility (standard deviation) of daily returns *every month* by combining pd.Grouper() and .groupby()
4. Multiply by $\sqrt{252}$ to annualize these volatilities of daily returns
5. Plot these annualized volatilities

Is market volatility higher during wars? Consider the following dates:

1. WWII: December 1941 to September 1945
2. Korean War: 1950 to 1953
3. Viet Nam War: 1959 to 1975
4. Gulf War: 1990 to 1991
5. War in Afghanistan: 2001 to 2021

Week 7

McKinney Chapter 11 - Time Series

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

Chapter 11 of McKinney (2022) discusses time series and panel data, which is where pandas *excels!* We will use these time series and panel tools every day for the rest of the course.

We will focus on:

1. Slicing a data frame or series by date or date range
2. Using `.shift()` to create leads and lags of variables
3. Using `.resample()` to change the frequency of variables
4. Using `.rolling()` to aggregate data over moving or rolling windows

Note: Indented block quotes are from McKinney (2022) unless otherwise indicated. The section numbers here differ from McKinney (2022) because we will only discuss some topics.

Time Series Basics

Let us create a time series to play with.

```
from datetime import datetime
dates = [
    datetime(2011, 1, 2),
    datetime(2011, 1, 5),
    datetime(2011, 1, 7),
    datetime(2011, 1, 8),
    datetime(2011, 1, 10),
    datetime(2011, 1, 12)
]
np.random.seed(42)
ts = pd.Series(np.random.randn(6), index=dates)

ts
```

```
2011-01-02    0.4967
2011-01-05   -0.1383
2011-01-07    0.6477
2011-01-08    1.5230
2011-01-10   -0.2342
2011-01-12   -0.2341
dtype: float64
```

Note that pandas converts the `datetime` objects to a pandas `DatetimeIndex` object and a single index value is a `Timestamp` object.

```
ts.index
```

```
DatetimeIndex(['2011-01-02', '2011-01-05', '2011-01-07', '2011-01-08',
                '2011-01-10', '2011-01-12'],
               dtype='datetime64[ns]', freq=None)
```

```
ts.index[0]
```

```
Timestamp('2011-01-02 00:00:00')
```

Recall that pandas automatically aligns objects on indexes.

```
ts
```

```
2011-01-02    0.4967
2011-01-05   -0.1383
2011-01-07    0.6477
2011-01-08    1.5230
2011-01-10   -0.2342
2011-01-12   -0.2341
dtype: float64
```

```
ts.iloc[::2]
```

```
2011-01-02    0.4967
2011-01-07    0.6477
2011-01-10   -0.2342
dtype: float64
```

```
ts + ts.iloc[::2]
```

```
2011-01-02    0.9934
2011-01-05      NaN
2011-01-07    1.2954
2011-01-08      NaN
2011-01-10   -0.4683
2011-01-12      NaN
dtype: float64
```

If we want to assign a default for missing values on the fly, we can use the `.add()` method.

```
ts.add(ts.iloc[::2], fill_value=1_000_000)
```

```
2011-01-02        0.9934
2011-01-05  999999.8617
2011-01-07        1.2954
2011-01-08  1000001.5230
2011-01-10       -0.4683
2011-01-12  999999.7659
dtype: float64
```

Indexing, Selection, Subsetting

pandas uses U.S.-style date strings (e.g., “M/D/Y”) or unambiguous date strings (e.g., “YYYY-MM-DD”) to select data.

```
ts.loc['1/10/2011'] # M/D/YYYY
```

-0.2342

```
ts.loc['2011-01-10'] # YYYY-MM-DD
```

-0.2342

```
ts.loc['20110110'] # YYYYMMDD
```

-0.2342

```
ts.loc['10-Jan-2011'] # D-Mon-YYYY
```

-0.2342

```
ts.loc['Jan-10-2011'] # Mon-D-YYYY
```

-0.2342

pandas *does not* use U.K.-style date strings.

```
# ts.loc['10/1/2011'] # D/M/YYYY # KeyError: '10/1/2011'
```

Let us create a *longer* time series to play with.

```
np.random.seed(42)
longer_ts = pd.Series(
    data=np.random.randn(1000),
    index=pd.date_range('1/1/2000', periods=1000)
)
```

```
longer_ts
```

```
2000-01-01    0.4967
2000-01-02   -0.1383
2000-01-03    0.6477
2000-01-04    1.5230
2000-01-05   -0.2342
...
2002-09-22   -0.2811
2002-09-23    1.7977
2002-09-24    0.6408
2002-09-25   -0.5712
2002-09-26    0.5726
Freq: D, Length: 1000, dtype: float64
```

We can specify a year-month to slice all of the observations in May of 2001.

```
longer_ts.loc['2001-05']
```

```
2001-05-01   -0.6466
2001-05-02   -1.0815
2001-05-03    1.6871
2001-05-04    0.8816
2001-05-05   -0.0080
2001-05-06    1.4799
2001-05-07    0.0774
2001-05-08   -0.8613
2001-05-09    1.5231
2001-05-10    0.5389
2001-05-11   -1.0372
2001-05-12   -0.1903
2001-05-13   -0.8756
2001-05-14   -1.3828
2001-05-15    0.9262
2001-05-16    1.9094
2001-05-17   -1.3986
2001-05-18    0.5630
2001-05-19   -0.6506
2001-05-20   -0.4871
2001-05-21   -0.5924
2001-05-22   -0.8640
```

```
2001-05-23    0.0485
2001-05-24   -0.8310
2001-05-25    0.2705
2001-05-26   -0.0502
2001-05-27   -0.2389
2001-05-28   -0.9076
2001-05-29   -0.5768
2001-05-30    0.7554
2001-05-31    0.5009
Freq: D, dtype: float64
```

We can also specify a year to slice all observations in 2001.

```
longer_ts.loc['2001']
```

```
2001-01-01    0.2241
2001-01-02    0.0126
2001-01-03    0.0977
2001-01-04   -0.7730
2001-01-05    0.0245
...
2001-12-27    0.0184
2001-12-28    0.3476
2001-12-29   -0.5398
2001-12-30   -0.7783
2001-12-31    0.1958
Freq: D, Length: 365, dtype: float64
```

If we sort our data chronologically, we can also slice with a range of date strings.

```
ts.loc['1/6/2011':'1/10/2011']
```

```
2011-01-07    0.6477
2011-01-08    1.5230
2011-01-10   -0.2342
dtype: float64
```

However, we cannot date slice if our data are not sorted chronologically.

```
ts2 = ts.sort_values()
```

```
ts2
```

```
2011-01-10    -0.2342
2011-01-12    -0.2341
2011-01-05    -0.1383
2011-01-02     0.4967
2011-01-07     0.6477
2011-01-08     1.5230
dtype: float64
```

The following date slice fails because `ts2` is not sorted chronologically.

```
# ts2.loc['1/6/2011':'1/11/2011'] # KeyError: 'Value based partial slicing on non-monotonic ...'
```

We can use the `.sort_index()` method first to allow date slices.

```
ts2.sort_index()['1/6/2011':'1/11/2011']
```

```
2011-01-07    0.6477
2011-01-08    1.5230
2011-01-10    -0.2342
dtype: float64
```

As with label slices, date slices are inclusive on both ends.

```
longer_ts.loc['1/6/2001':'1/11/2001']
```

```
2001-01-06    0.4980
2001-01-07    1.4511
2001-01-08    0.9593
2001-01-09    2.1532
2001-01-10   -0.7673
2001-01-11    0.8723
Freq: D, dtype: float64
```

Recall that if we modify a slice, we modify the original series or data frame.

Remember that slicing in this manner produces views on the source time series like slicing NumPy arrays. This means that no data is copied and modifications on the slice will be reflected in the original data.

```
ts3 = ts.copy()
```

```
ts3
```

```
2011-01-02    0.4967
2011-01-05   -0.1383
2011-01-07    0.6477
2011-01-08    1.5230
2011-01-10   -0.2342
2011-01-12   -0.2341
dtype: float64
```

```
ts4 = ts3.iloc[:3]
```

```
ts4
```

```
2011-01-02    0.4967
2011-01-05   -0.1383
2011-01-07    0.6477
dtype: float64
```

```
ts4.iloc[:] = 2001
```

```
ts4
```

```
2011-01-02    2001.0000
2011-01-05    2001.0000
2011-01-07    2001.0000
dtype: float64
```

```
ts3
```

```
2011-01-02    2001.0000
2011-01-05    2001.0000
2011-01-07    2001.0000
```

```
2011-01-08      1.5230
2011-01-10     -0.2342
2011-01-12     -0.2341
dtype: float64
```

Series `ts` is unchanged because `ts3` is a *copy* of `ts`!

```
ts
```

```
2011-01-02      0.4967
2011-01-05     -0.1383
2011-01-07      0.6477
2011-01-08      1.5230
2011-01-10     -0.2342
2011-01-12     -0.2341
dtype: float64
```

Time Series with Duplicate Indices

Most of our data in this course will be well-formed with one observation per date-time for series or one observation per individual per date-time for data frames. However, we may later receive poorly-formed data with duplicate observations. Here, series `dup_ts` has three observations on February 2nd.

```
dates = pd.DatetimeIndex(['1/1/2000', '1/2/2000', '1/2/2000', '1/2/2000', '1/3/2000'])
dup_ts = pd.Series(data=np.arange(5), index=dates)
dup_ts
```

```
2000-01-01    0
2000-01-02    1
2000-01-02    2
2000-01-02    3
2000-01-03    4
dtype: int64
```

The `.is_unique` attribute tells us if an index is unique.

```
dup_ts.index.is_unique
```

```
False
```

```
dup_ts.loc['1/3/2000'] # not duplicated
```

```
np.int64(4)
```

```
dup_ts.loc['1/2/2000'] # duplicated
```

```
2000-01-02    1  
2000-01-02    2  
2000-01-02    3  
dtype: int64
```

The solution to duplicate data depends on the context. For example, we may want the mean of all observations on a given date. The `.groupby()` method can help us here.

```
dup_ts.groupby(level=0).mean()
```

```
2000-01-01    0.0000  
2000-01-02    2.0000  
2000-01-03    4.0000  
dtype: float64
```

Or keep the first value on each date.

```
dup_ts.groupby(level=0).first()
```

```
2000-01-01    0  
2000-01-02    1  
2000-01-03    4  
dtype: int64
```

Date Ranges, Frequencies, and Shifting

Generic time series in pandas are assumed to be irregular; that is, they have no fixed frequency. For many applications this is sufficient. However, it's often desirable to work relative to a fixed frequency, such as daily, monthly, or every 15 minutes, even if that means introducing missing values into a time series. Fortunately pandas has a full suite of standard time series frequencies and tools for resampling, inferring frequencies, and generating fixed-frequency date ranges.

Shifting Data

Shifting is an important feature! Shifting is moving data backward (or forward) through time.

```
np.random.seed(42)
ts = pd.Series(
    data=np.random.randn(4),
    index=pd.date_range('1/1/2000', periods=4, freq='ME')
)
```

```
ts
```

2000-01-31	0.4967
2000-02-29	-0.1383
2000-03-31	0.6477
2000-04-30	1.5230

Freq: ME, dtype: float64

If we specify a *positive integer* N to the `.shift()` method:

1. The date index remains the same
2. Values *shift down* N observations

The `.shift()` method defaults to $N = 1$.

```
ts.shift()
```

2000-01-31	NaN
2000-02-29	0.4967
2000-03-31	-0.1383
2000-04-30	0.6477

Freq: ME, dtype: float64

```
ts.shift(1)
```

2000-01-31	NaN
2000-02-29	0.4967
2000-03-31	-0.1383
2000-04-30	0.6477

Freq: ME, dtype: float64

```
ts.shift(2)
```

2000-01-31	NaN
2000-02-29	NaN
2000-03-31	0.4967
2000-04-30	-0.1383

Freq: ME, dtype: float64

If we specify a *negative integer N* to the `.shift()` method, values *shift up N* observations.

```
ts.shift(-2)
```

2000-01-31	0.6477
2000-02-29	1.5230
2000-03-31	NaN
2000-04-30	NaN

Freq: ME, dtype: float64

i Note

We almost never shift with negative values to prevent a look-ahead bias. That is, assuming chronological sorting, we almost never bring values from the future back to the present. We do not want to assume that financial market participants know the future.

The `.shift()` examples above shift by N observations without considering time stamps. As a result, the time stamps are unchanged, and values shift down for positive `periods` or up for negative `periods`. However, we can specify the `freq` argument to consider time stamps. With the `freq` argument, time stamps shift by `periods` multiples of the `freq` argument.

```
ts
```

2000-01-31	0.4967
2000-02-29	-0.1383
2000-03-31	0.6477
2000-04-30	1.5230

Freq: ME, dtype: float64

```
ts.shift(periods=2, freq='ME')
```

```
2000-03-31    0.4967
2000-04-30   -0.1383
2000-05-31    0.6477
2000-06-30    1.5230
Freq: ME, dtype: float64
```

```
ts.shift(periods=3, freq='D')
```

```
2000-02-03    0.4967
2000-03-03   -0.1383
2000-04-03    0.6477
2000-05-03    1.5230
dtype: float64
```

M is already months, so min is minutes.

```
ts.shift(periods=1, freq='90min')
```

```
2000-01-31 01:30:00    0.4967
2000-02-29 01:30:00   -0.1383
2000-03-31 01:30:00    0.6477
2000-04-30 01:30:00    1.5230
dtype: float64
```

Calculating returns

We can calculate returns in two ways. First, easily with the .pct_change() method.

```
ts.pct_change()
```

```
2000-01-31      NaN
2000-02-29   -1.2784
2000-03-31   -5.6844
2000-04-30    1.3515
Freq: ME, dtype: float64
```

Second, manaully with the .shift() method.

```
(ts - ts.shift()) / ts.shift()
```

```
2000-01-31      NaN
2000-02-29    -1.2784
2000-03-31    -5.6844
2000-04-30     1.3515
Freq: ME, dtype: float64
```

These two return calculations are the same.

```
np.allclose(
    a=ts.pct_change(),
    b=(ts - ts.shift()) / ts.shift(),
    equal_nan=True
)
```

```
True
```

Two observations on these return calculations:

1. The first percent change is `NaN` because there is no previous value to change from
2. The default for `.shift()` and `.pct_change()` is `periods=1`

Shifting dates with offsets

We can also shift time stamps to the beginning or end of a period.

```
from pandas.tseries.offsets import MonthEnd
now = datetime(2011, 11, 17)
```

```
now
```

```
datetime.datetime(2011, 11, 17, 0, 0)
```

`MonthEnd(0)` moves to the end of the month *but does not leave the current month.*

```
now + MonthEnd(0)
```

```
Timestamp('2011-11-30 00:00:00')
```

`MonthEnd(1)` moves to the end of the *current* month. If already at the end of the *current* month, it moves to the end of the *next* month.

```
now + MonthEnd(1)
```

```
Timestamp('2011-11-30 00:00:00')
```

```
now + MonthEnd(1) + MonthEnd(1)
```

```
Timestamp('2011-12-31 00:00:00')
```

```
now + MonthEnd(0) + MonthEnd(1)
```

```
Timestamp('2011-12-31 00:00:00')
```

Be careful! The `MonthEnd()` default is `n=1`!

```
datetime(2021, 10, 31) + MonthEnd(0)
```

```
Timestamp('2021-10-31 00:00:00')
```

```
datetime(2021, 10, 31) + MonthEnd(1)
```

```
Timestamp('2021-11-30 00:00:00')
```

```
datetime(2021, 10, 31) + MonthEnd()
```

```
Timestamp('2021-11-30 00:00:00')
```

Always check your output!

Resampling and Frequency Conversion

Resampling is an important feature!

Resampling refers to the process of converting a time series from one frequency to another. Aggregating higher frequency data to lower frequency is called downsampling, while converting lower frequency to higher frequency is called upsampling. Not all resampling falls into either of these categories; for example, converting W-WED (weekly on Wednesday) to W-FRI is neither upsampling nor downsampling.

We can resample both series and data frames. The `.resample()` method syntax is similar to `.groupby()`.

Downsampling

Aggregating data to a regular, lower frequency is a pretty normal time series task. The data you're aggregating doesn't need to be fixed frequently; the desired frequency defines bin edges that are used to slice the time series into pieces to aggregate. For example, to convert to monthly, 'M' or 'BM', you need to chop up the data into one-month intervals. Each interval is said to be half-open; a data point can only belong to one interval, and the union of the intervals must make up the whole time frame. There are a couple things to think about when using resample to downsample data:

- Which side of each interval is closed
- How to label each aggregated bin, either with the start of the interval or the end

```
rng = pd.date_range(start='2000-01-01', periods=12, freq='min')
ts = pd.Series(np.arange(12), index=rng)
```

```
ts
```

2000-01-01 00:00:00	0
2000-01-01 00:01:00	1
2000-01-01 00:02:00	2
2000-01-01 00:03:00	3
2000-01-01 00:04:00	4
2000-01-01 00:05:00	5
2000-01-01 00:06:00	6
2000-01-01 00:07:00	7
2000-01-01 00:08:00	8

```
2000-01-01 00:09:00    9
2000-01-01 00:10:00   10
2000-01-01 00:11:00   11
Freq: min, dtype: int64
```

We can aggregate the one-minute frequency data above to five-minute frequency data. Resampling requires an aggregation method. Here, we use the `.sum()` method.

```
ts.resample('5min').sum()
```

```
2000-01-01 00:00:00    10
2000-01-01 00:05:00   35
2000-01-01 00:10:00   21
Freq: 5min, dtype: int64
```

When we resample with a minute-frequency:

1. Left edges of the resampling interval are closed (included) and right edges are open (excluded)
2. Labels are by the left edge of the resampling interval by default

In the example above, the first value of 10 at midnight is the sum of values at midnight until 00:05, excluding the value at 00:05. That is, the sums are $10 = 0 + 1 + 2 + 3 + 4$ at 00:00, $35 = 5 + 6 + 7 + 8 + 9$ at 00:05, and so on. We can use the `closed` and `label` arguments to change this behavior.

In finance, we generally prefer `closed='right'` and `label='right'` to avoid a lookahead bias.

```
ts.resample('5min', closed='right', label='right').sum()
```

```
2000-01-01 00:00:00    0
2000-01-01 00:05:00   15
2000-01-01 00:10:00   40
2000-01-01 00:15:00   11
Freq: 5min, dtype: int64
```

These defaults for minute-frequency data may seem odd, but any choice is arbitrary. The defaults for weekly and lower frequencies are `closed='right'` and `label='right'`, which are correct for finance. Still, we should read the docstring and check our output whenever we use the `.resample()` method!

Upsampling and Interpolation

To downsample (i.e., resample from higher to lower frequency), we must aggregate (e.g., `.mean()`, `.sum()`, `.first()`, or `.last()`). To upsample (i.e., resample from lower to higher frequency), we must choose how to fill in the new, higher-frequency observations.

```
np.random.seed(42)
frame = pd.DataFrame(
    data=np.random.randn(2, 4),
    index=pd.date_range('1/1/2000', periods=2, freq='W-WED'),
    columns=['Colorado', 'Texas', 'New York', 'Ohio']
)
```

```
frame
```

	Colorado	Texas	New York	Ohio
2000-01-05	0.4967	-0.1383	0.6477	1.5230
2000-01-12	-0.2342	-0.2341	1.5792	0.7674

We use the `.asfreq()` method to convert to the new frequency and leave the new observations as missing.

```
df_daily = frame.resample('D').asfreq()
```

```
df_daily
```

	Colorado	Texas	New York	Ohio
2000-01-05	0.4967	-0.1383	0.6477	1.5230
2000-01-06	NaN	NaN	NaN	NaN
2000-01-07	NaN	NaN	NaN	NaN
2000-01-08	NaN	NaN	NaN	NaN
2000-01-09	NaN	NaN	NaN	NaN
2000-01-10	NaN	NaN	NaN	NaN
2000-01-11	NaN	NaN	NaN	NaN
2000-01-12	-0.2342	-0.2341	1.5792	0.7674

We use the `.ffill()` method to forward fill values to replace missing values.

```
frame.resample('D').ffill()
```

	Colorado	Texas	New York	Ohio
2000-01-05	0.4967	-0.1383	0.6477	1.5230
2000-01-06	0.4967	-0.1383	0.6477	1.5230
2000-01-07	0.4967	-0.1383	0.6477	1.5230
2000-01-08	0.4967	-0.1383	0.6477	1.5230
2000-01-09	0.4967	-0.1383	0.6477	1.5230
2000-01-10	0.4967	-0.1383	0.6477	1.5230
2000-01-11	0.4967	-0.1383	0.6477	1.5230
2000-01-12	-0.2342	-0.2341	1.5792	0.7674

```
frame.resample('D').ffill(limit=2)
```

	Colorado	Texas	New York	Ohio
2000-01-05	0.4967	-0.1383	0.6477	1.5230
2000-01-06	0.4967	-0.1383	0.6477	1.5230
2000-01-07	0.4967	-0.1383	0.6477	1.5230
2000-01-08	NaN	NaN	NaN	NaN
2000-01-09	NaN	NaN	NaN	NaN
2000-01-10	NaN	NaN	NaN	NaN
2000-01-11	NaN	NaN	NaN	NaN
2000-01-12	-0.2342	-0.2341	1.5792	0.7674

```
frame.resample('W-THU').ffill()
```

	Colorado	Texas	New York	Ohio
2000-01-06	0.4967	-0.1383	0.6477	1.5230
2000-01-13	-0.2342	-0.2341	1.5792	0.7674

Moving Window Functions

Moving or rolling window functions are one of the neatest features of pandas.!


```
df = (
    yf.download(
        tickers=['AAPL', 'MSFT', 'SPY'],
        auto_adjust=False,
        progress=False
    )
    .iloc[:-1]
)
```

```
df
```

Price Ticker Date	Adj Close AAPL	Adj Close MSFT	Adj Close SPY	Close AAPL	Close MSFT	Close SPY	High AAPL	High MSFT	High SPY
1980-12-12	0.0987	NaN	NaN	0.1283	NaN	NaN	0.1289	NaN	NaN
1980-12-15	0.0936	NaN	NaN	0.1217	NaN	NaN	0.1222	NaN	NaN
1980-12-16	0.0867	NaN	NaN	0.1127	NaN	NaN	0.1133	NaN	NaN
1980-12-17	0.0889	NaN	NaN	0.1155	NaN	NaN	0.1161	NaN	NaN
1980-12-18	0.0914	NaN	NaN	0.1189	NaN	NaN	0.1194	NaN	NaN
...
2025-02-21	245.5500	408.2100	599.9400	245.5500	408.2100	599.9400	248.6900	418.0500	610.3000
2025-02-24	247.1000	404.0000	597.2100	247.1000	404.0000	597.2100	248.8600	409.3700	603.0300
2025-02-25	247.0400	397.9000	594.2400	247.0400	397.9000	594.2400	250.0000	401.9200	597.8900
2025-02-26	240.3600	399.7300	594.5400	240.3600	399.7300	594.5400	244.9800	403.6000	599.5800
2025-02-27	237.3000	392.5300	585.0500	237.3000	392.5300	585.0500	242.4600	405.7400	598.0200

The `.rolling()` method accepts a window-width and requires an aggregation method. The following example plots AAPL's observed daily price alongside its 252-trading-day rolling mean, 365-calendar-day rolling mean, and calendar year mean.

```
aapl = df.loc['2012':, ('Adj Close', 'AAPL')]
aapl.plot(label='Observed')
aapl.rolling(window=252).mean().plot(label='252-Trading-Day Mean') # min_periods defaults to
aapl.rolling(window=365D).mean().plot(label='365-Calendar-Day Mean') # min_periods default
aapl.resample('YE').mean().plot(style='.', label='Calendar Year Mean')
plt.legend()
plt.ylabel('AAPL Adjusted Close ($)')
plt.title('Comparison of Rolling and Resampling Means')
plt.show()
```



i Note

If we specify the rolling window width as an integer:

1. Each rolling window is that many observations wide and ignores time stamps
2. Each rolling window must have that many non-missing observations

We can specify `min_periods` to allow incomplete windows. For integer window widths, `min_periods` defaults to the given integer window width. For string date offsets, `min_periods` defaults to 1.

Binary Moving Window Functions


Binary moving window functions accept two inputs. The most common example is the rolling correlation between two return series.

```
returns = df['Adj Close'].pct_change()
```


```
returns
```

Ticker	AAPL	MSFT	SPY
Date			
1980-12-12	NaN	NaN	NaN
1980-12-15	-0.0522	NaN	NaN
1980-12-16	-0.0734	NaN	NaN
1980-12-17	0.0248	NaN	NaN
1980-12-18	0.0290	NaN	NaN
...
2025-02-21	-0.0011	-0.0190	-0.0171
2025-02-24	0.0063	-0.0103	-0.0046
2025-02-25	-0.0002	-0.0151	-0.0050
2025-02-26	-0.0270	0.0046	0.0005
2025-02-27	-0.0127	-0.0180	-0.0160

```
(  
    returns['AAPL']  
    .rolling(126, min_periods=100)  
    .corr(returns['SPY'])  
    .plot()  
)  
plt.ylabel('Correlation between AAPL and SPY')  
plt.title('Rolling Correlation between AAPL and SPY\n (126-Day Window w/ 100-Day Minimum)')  
plt.show()
```



```
(  
    returns[['AAPL', 'MSFT']]  
    .rolling(126, min_periods=100)  
    .corr(returns['SPY'])  
    .plot()  
)  
plt.ylabel('Correlation with SPY')  
plt.title('Rolling Correlation with SPY\n (126-Day Window w/ 100-Day Minimum)')  
plt.show()
```




User-Defined Moving Window Functions

We can define our own moving window functions and use them with the `.apply()` method. However, note that `.apply()` method will be much slower than the optimized methods, like `.mean()` and `.std()`. Here, we will calculate rolling volatility with `.apply()` and `.std()` and compare their speeds.

```

(
    returns['AAPL']
    .rolling(252)
    .apply(np.std)
    .mul(np.sqrt(252) * 100)
    .plot()
)
plt.ylabel('Volatility (%)')
plt.title('Rolling Volatility\n (252-Day Window w/ 252-Day Minimum)')
plt.show()

```



Do not be afraid to use `.apply()`, but realize that `.apply()` is often 1000 times slower than the optimized method!

```
%timeit returns['AAPL'].rolling(252).apply(np.std)
```

2.28 s ± 109 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
%timeit returns['AAPL'].rolling(252).std()
```

613 s ± 80 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

McKinney Chapter 11 - Practice - Blank

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

Download daily returns for ten portfolios formed on book-to-market ratios

Plot cumulative returns for all available data

Calculate total returns for each calendar year

Calculate total returns for all 252-trading-day windows

Calculate total returns for 12-months windows with monthly data

Calculate Sharpe Ratios for each calendar year

Calculate rolling betas

We can calculate CAPM betas as: $\beta_i = \frac{Cov(r_i - r_f, r_M - r_f)}{Var(r_M - r_f)}$

Calculate rolling Sharpe Ratios

Week 8

Project 1

Purpose

I have two goals for this project:

1. Help you learn to compare two different foreign exchange (forex) trading strategies in Python
2. Help you work on a real-world task that an industry partner considers interesting and worthwhile

Assignment

Our industry partner trades forex for a multinational client. This client wants our partner to compare its current active trading approach to a new zero-basis banking approach with passive or automatic trading.

Our assignment: Analyze these two approaches, summarize their advantages and disadvantages, and recommend a plan of action for our partner. Your plan should include an executive summary and supporting figures and tables. Here are ideas to consider:

1. Our partner provided her thought process in `Student Python Research.docx` and plans to visit one section virtually
2. Consider only the USD/CAD pair
3. Consider all transaction costs
 1. Use a variable `spread` to account for transaction costs that increase linearly with trade size (e.g., market orders pay one half of the bid-ask spread or $0.5 * \text{spread} * \text{size}$)
 2. Use a variable `gamma` to account for transaction costs that increase with trade size squared (e.g., large orders have market impacts that increase quickly with size or $\text{gamma} * \text{size} * \text{size}$)
 3. Set `spread` and `gamma` to 0.0001 to approximate the real world
4. At a minimum, compare the final dollar values with both approaches
5. You might also consider the best final dollar value that an active trader could achieve, *plus any other interesting aspects of these data*

Project 1

6. Consider all scenarios that our industry partner provides
7. Title, label, and caption your figures and tables and reference them in your report

Criteria

Table 1 provides the project grading rubric. The project is worth 200 points. The peer reviews are worth 100 points, and students will receive their median score. Almost all students earn perfect peer review scores, so I will factor that into project scores. For example, a project score *without peer review scores* of 77.5% converts to a project score *with perfect peer review scores* of 85% because $\frac{0.775 \times 200 + 1.00 \times 100}{300} = 0.85 = 85\%$.

Table 1: This table provides the project grading rubric

Topic	Points
Clarity, correctness, and completeness of calculations	60
Clarity, correctness, and completeness of visualizations	60
Clarity, correctness, and completeness of discussions	60
Correctness of submission according to the deliverables section	20
Total	200

Deliverables

Upload the following as unzipped files to Canvas by 11:59 PM on 2/28:

1. One Jupyter notebook that contains your report and performs *all* your analysis
 1. Name this file `project_1.ipynb` for me to run your code
 2. Your notebook must run on my computer; I will place the data file in the same folder as your notebook
 3. You may not edit the Word documents after you create them
2. One Quarto-generated Word document *without code* for our partner to review
 1. Name this file `project_1_without_code.docx`
 2. Typing `echo: false` in the first cell of this Jupyter notebook hides code in your Word document
3. One Quarto-generated Word document *with code* for me to grade
 1. Name this file `project_1_with_code.docx`
 2. Typing `echo: true` in the first cell of this Jupyter notebook displays code in your Word document

Here is some additional guidance:

1. Provide a one-page executive summary on your first page; this summary can run up to halfway down the second page to avoid having to render your notebook several times to check lengths
2. Your Word document must not exceed 15 pages in length
3. Do not include your names anywhere in your submission

Data

This project requires two data files.

1. `USDCAD 021345.xlsx` provides recent prices of US dollars in terms of Canadian dollars.
For example, 1.3875 indicates that 1.3875 Canadian dollars buy 1 US dollar
2. `NE University Trades.xlsx` provides a summary of our industry partner's trades

We need to install the `openpyxl` package to read these files. To install it, follow these steps:

1. `conda activate fina6333` (activate your course environment)
2. `conda install openpyxl` (install `openpyxl` into your course environment)
3. `jupyter lab` (open JupyterLab to begin your analysis)

We can use `pd.read_excel()` to read these files. For example, the following code reads the exchange rate file:

```
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import numpy as np
import pandas as pd
import yfinance as yf
```

Edit 2025-02-24: I tweaked this code cell to set `Date` as the index and added comments.

```
df_1 = (
    pd.read_excel( # Load data from an Excel file into a pandas DataFrame
        io='USDCAD 021325.xlsx', # Specify the file path or name of the Excel file to read
        sheet_name='Sheet1', # Select the specific sheet in the Excel file (Sheet1) to load
        names=['Date', 'USDCAD'], # Assign custom column names 'Date' and 'USDCAD' to the data
        parse_dates=['Date'], # Tell pandas to automatically parse the 'Date' column as dates
    )
    .assign( # Add or modify columns in the DataFrame
        Date=lambda x: x['Date'].dt.tz_localize('UTC') # Take the 'Date' column, already a
```

Project 1

```
# 'x' is the DataFrame; .dt accesses datetime properties; .tz_localize('UTC') makes it UTC-aware
)
.set_index('Date') # Set the 'Date' column (now UTC timezone-aware) as the DataFrame's index
)

df_1.head()
```

USDCAD		
Date		
2024-08-01 17:05:00+00:00	1.3877	
2024-08-01 17:10:00+00:00	1.3877	
2024-08-01 17:15:00+00:00	1.3875	
2024-08-01 17:20:00+00:00	1.3874	
2024-08-01 17:25:00+00:00	1.3875	

Edit 2025-02-24: I added this code to parse the trades spreadsheet.

```
df_2 = (
    pd.read_excel( # Load data from an Excel file into a pandas DataFrame
        io='NE University Trades.xlsx', # Specify the file path or name of the Excel file
        sheet_name='Sheet1', # Select the specific sheet in the Excel file to read
        usecols='A:I', # Read only columns A through I (first 9 columns) to focus on relevant data
        skiprows=2, # Skip the first 2 rows (e.g., headers or metadata) to get to the actual data
    )
    .assign( # Add or modify columns in the DataFrame
        Date=lambda x: ( # Create or overwrite the 'Date' column using a lambda function; 'x' refers to the DataFrame
            pd.to_datetime( # Convert a string to a pandas datetime object
                x['Date'].astype(str) + ' ' + # Convert the 'Date' column to strings and add a space separator
                x['Time CST'].apply(lambda t: t.strftime('%H:%M:%S')) # Convert 'Time CST' to UTC
            ) # The result is a combined string like '2025-02-04 15:28:09'
            .dt.tz_localize('America/Chicago') # Assign the 'America/Chicago' timezone (CST)
            .dt.tz_convert('UTC') # Convert the timezone-aware datetime from CST/CDT to UTC
        ) # The 'Date' column now holds UTC timestamps
    )
    .set_index('Date') # Set the 'Date' column (now in UTC) as the DataFrame's index for time-based operations
    .drop(columns='Time CST') # Remove the original 'Time CST' column since it's now part of the 'Date' index
)
```

`df_2`

Date	Interbank Rate	Value Date	Client Rate	Sell CAD	Buy USD	Spread BP
2025-02-04 21:28:09+00:00	1.4329	2025-02-05	1.4374	2M	1391420.50	31
2025-02-12 15:42:37+00:00	1.4295	2025-02-12	1.4341	2M	1394641.79	32
2025-01-29 14:29:55+00:00	1.4453	2025-01-29	1.4496	1M	689840.72	29
2025-01-28 16:18:07+00:00	1.4394	2025-01-29	1.4438	1M	692616.71	30

Quarto

Basics

1. Use [Quarto](#) to generate your Word document from your notebook
2. Use `#` to create a title and `##` to create sections
3. Use `-` or `1.` to create lists
4. Use the first cell in this notebook to hide or display code with `echo=false` or `echo=true`, respectively
5. This first cell must be a `raw` cell instead of a `code` or `markdown` cell
6. Use `quarto render project_1.ipynb` in the same folder as your notebook to render it to a Word document
7. Use the `cd` command in the terminal to change the working directory to the directory with your notebook

Examples

This section provides a sample analysis highlighting how code and formatting work with Quarto. Figure 1 provides a line plot of the value of a \$10,000 investment in SPY, the S&P 500 SPDR ETF. Note that `#| label:` and `#| fig-cap:` comments at the top of the figure cell create the figure reference/link and the figure caption, respectively. You can learn more about cross-referencing figures and tables [here](#).

```
data = (
    yf.download(tickers='SPY', auto_adjust=False, progress=False)
    .assign(
        ret=lambda x: x['Adj Close'].pct_change(),
        value=lambda x: x['ret'].pipe(np.log1p).cumsum().pipe(np.exp).mul(10_000)
    )
)
```

```
data['value'].plot()
plt.ylabel('Value ($)')
plt.suptitle('Value of $10,000 Investment in SPY')
plt.title(f'Investment made at the market close on {data.index[0]: %b %d, %Y}')
plt.gca().yaxis.set_major_formatter(ticker.FuncFormatter(lambda x, p: format(int(x), ',')))

plt.show()
```




Figure 1: This line plot shows the value of a \$10,000 investment in the S&P 500 SPDR ETF at the close of its first day of trading

Artificial Intelligence (AI)

You may use AI (e.g., ChatGPT) to *help* you prepare your analysis and discussion. However:

1. AI will not do very well on this project without significant input from your team

Project 1

2. AI will not be a defense against plagiarism because AI should not *write* your code and slides; If you plagiarize an AI that plagiarizes other sources, you are responsible for plagiarizing the AI and its sources

Week 9

Herron Topic 2 - Trading Strategies Based on Technical Analysis

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

This notebook covers trading strategies based on technical analysis in three parts:

1. What is technical analysis?
2. Why might trading strategies based on technical analysis work (or not work)?
3. Implement a simple moving average (SMA) trading strategy

I based this lecture notebook on Welch (2022, Chapter 12), Lewinson (2020, Chapter 2), and Murphy (1999). If you want to learn technical analysis, Murphy (1999) is the best reference and covers more than we can in a week or semester. The practice notebook will cover several other trading strategies based on technical analysis.

What is technical analysis?

Technical analysis is a methodology that analyzes past market data (e.g., prices and volume, plus open interest in futures and options markets) in an attempt to forecast future price movements. If technical analysis can predict future price movements, the market is not weak-form efficient. Welch (2022, Section 12.2) provides the three degrees of market efficiency:

The Traditional Classification The traditional definition of market efficiency focuses on information. In the traditional classification, market efficiency comes in one of three primary degrees: weak, semi-strong, and strong.

Weak market efficiency says that all information in past prices is reflected in today's stock prices so that technical analysis (trading based solely on historical price patterns) cannot be used to beat the market. Put differently, the market is the best technical analyst.

Semistrong market efficiency says that all public information is reflected in today's stock prices, so that neither fundamental trading (based on underlying firm fundamentals, such as cash flows or discount rates) nor technical analysis can be used to beat the market. Put differently, the market is both the best technical and the best fundamental analyst.

Strong market efficiency says that all information, both public and private, is reflected in today's stock prices, so that nothing — not even private insider information — can be used to beat the market. Put differently, the market is the best analyst and cannot be beat.

In this traditional classification, all finance professors nowadays believe that most U.S. financial markets are not strong-form efficient: Insider trading may be illegal, but it works. However, there are still arguments as to which markets are only semi-strong-form efficient or even only weak-form efficient.

Welch (2022, Section 12.2) goes on to provide his own taxonomy of true, firm, mild, and nonbelievers in market efficiency. Chapter 12 summarizes market efficiency, classical finance, behavioral finance, arbitrage, limits to arbitrage, and their consequences for managers and investors. You can read Chapter 12 [here](#). We will focus on technical analysis in this notebook, but **welch2022** is excellent.

Why might trading strategies based on technical analysis work or not?

...Work?

Technical analysis relies on a few ideas:

1. Market prices and volume reflect all relevant information, so we can focus on past prices and volume instead of fundamentals and news.
2. Market prices move in trends and patterns driven by market participants.
3. These trends and patterns tend to repeat themselves because market participants create them.

...Or Not?

The logic above is reasonable. However, if past market prices reflect all relevant information, they should also reflect any price trends they predict. Therefore, any patterns should be self-defeating, and market prices should follow a [random walk](#). As well, the signal-to-noise ratio in market prices is high! Still, technical analysis provides an opportunity to learn how to implement and back-test trading strategies in Python.

A Random Walk

In a random walk, the price tomorrow equals the price today plus a tiny drift plus noise. In math terms, a random walk is

$$P_t = \rho P_{t-1} + m P_{t-1} + \varepsilon_t$$

where m is a small drift term and $E[\varepsilon] = 0$. If $\rho > 1$, prices would quickly increase, and, if $\rho < 1$, prices would quickly decrease. Let us examine the historical record.

```
ff = (
    pdr.DataReader(
        name='F-F_Research_Data_Factors_daily',
        data_source='famafrench',
        start='1900'
    )
    [0]
    .assign(Mkt=lambda x: x['Mkt-RF'] + x['RF'])
    .div(100)
)
```


```
C:\Users\richa\AppData\Local\Temp\ipykernel_8936\3102596058.py:2: FutureWarning: The argument
```

```
pdr.DataReader()
```

We can compound market returns to impute market prices relative to the last day of June 1926.

```
prices = ff['Mkt'].add(1).cumprod()

prices.plot()
plt.title('Imputed Market Prices\nAssuming $1 on the Last Day of June 1926')
plt.ylabel('Imputed Market Price ($)')
plt.semilogy()
plt.show()
```




We need lagged prices to estimate ρ . We will add 10 lags of P to help us understand the relation between past and future prices.

```
prices_w_lags = (
    pd.concat(
        objs=[prices.shift(t) for t in range(11)],
        keys=[f'Lag {t}' for t in range(11)],
        names=['Price'],
        axis=1,
    )
)
prices_w_lags.tail()
```

Price Date	Lag 0	Lag 1	Lag 2	Lag 3	Lag 4	Lag 5	Lag 6	Lag 7
2024-12-24	15038.5668	14870.9709	14778.3109	14617.9520	14633.0240	15110.9845	15182.7991	15182.7991
2024-12-26	15044.1311	15038.5668	14870.9709	14778.3109	14617.9520	14633.0240	15110.9845	15110.9845
2024-12-27	14870.6722	15044.1311	15038.5668	14870.9709	14778.3109	14617.9520	14633.0240	15110.9845
2024-12-30	14711.1099	14870.6722	15044.1311	15038.5668	14870.9709	14778.3109	14617.9520	14617.9520
2024-12-31	14645.9397	14711.1099	14870.6722	15044.1311	15038.5668	14870.9709	14778.3109	14778.3109

Now we can plot the correlation of price with its lags.

```
(  
    prices_w_lags  
    .dropna()  
    .corr()  
    .loc['Lag 0']  
    .plot(kind='bar')  
)  
plt.title('Correlations between Imputed Market Price and Its Lags')  
plt.ylabel('Correlation')  
plt.show()
```



But these are *pairwise* correlations. If we estimate *conditional* correlations, we see that most of the price information is in the first lag!

```
y = prices_w_lags.dropna()['Lag 0'] # Pt
X = prices_w_lags.dropna().drop('Lag 0', axis=1).pipe(sm.add_constant) # Pt-1, Pt-2, ..., Pt
model = sm.OLS(endog=y, exog=X)
fit = model.fit(cov_type='HAC', cov_kwds={'maxlags': 10})
fit.summary()
```

Dep. Variable:	Lag 0	R-squared:	1.000			
Model:	OLS	Adj. R-squared:	1.000			
Method:	Least Squares	F-statistic:	2.195e+06			
Date:	Fri, 07 Mar 2025	Prob (F-statistic):	0.00			
Time:	17:19:35	Log-Likelihood:	-1.2518e+05			
No. Observations:	25891	AIC:	2.504e+05			
Df Residuals:	25880	BIC:	2.505e+05			
Df Model:	10					
Covariance Type:	HAC					
	coef	std err	z	P> z	[0.025	0.975]
const	-0.0144	0.144	-0.100	0.920	-0.297	0.268
Lag 1	0.9574	0.033	29.042	0.000	0.893	1.022
Lag 2	0.0621	0.058	1.077	0.281	-0.051	0.175
Lag 3	-0.0315	0.037	-0.860	0.390	-0.103	0.040
Lag 4	-0.0220	0.040	-0.553	0.580	-0.100	0.056
Lag 5	0.0290	0.038	0.767	0.443	-0.045	0.103
Lag 6	-0.0498	0.042	-1.181	0.237	-0.132	0.033
Lag 7	0.1215	0.048	2.518	0.012	0.027	0.216
Lag 8	-0.1222	0.053	-2.302	0.021	-0.226	-0.018
Lag 9	0.1331	0.046	2.908	0.004	0.043	0.223
Lag 10	-0.0771	0.030	-2.535	0.011	-0.137	-0.017
Omnibus:	15929.955	Durbin-Watson:	1.994			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4213325.277			
Skew:	-1.820	Prob(JB):	0.00			
Kurtosis:	65.389	Cond. No.	9.90e+03			


Notes:

- [1] Standard Errors are heteroscedasticity and autocorrelation robust (HAC) using 10 lags and without small sample correction
- [2] The condition number is large, 9.9e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

plt.bar(
    x=fit.params.index[1:],
    height=fit.params[1:],
    yerr=2*fit.bse[1:]
)
plt.title('Conditional Correlations between Imputed Market Price and Its Lags\nBars Indicate')
plt.ylabel('Conditional Correlation')
plt.xlabel('Price')
plt.show()

```



Signal-to-Noise Ratio

Recall, we can express a random walk as $P_t = \rho P_{t-1} + mP_{t-1} + \varepsilon_t$. Since $\rho = 1$, we can subtract P_{t-1} from both sides, then divide by P_{t-1} on both sides. This transformation expresses a random walk in terms of returns: $r_{t-1,t} = m + e_t$, where $E[e_t] = 0$ and $SD[e_t] = s$, so $E[r_{t-1,t}] = m$. We can think of the signal-to-noise ratio as $\frac{m}{s}$. How high is this ratio?

```
m, s = ff['Mkt'].mean(), ff['Mkt'].std()
```

Here m is about 4 basis points per day!

```
m
```

```
0.0004
```

However, s is about 108 basis points per day!

```
s
```

```
0.0108
```

Therefore, the signal-to-noise ratio is less than 0.04! We want this ratio above 2 to reject that a drift (or a strategy) is zero.

```
m/s
```

```
0.0397
```

Recall that means grow linearly with time and standard deviations growth with the square-root of time. So, if we want $\sqrt{t} \times \frac{m}{s} \geq 2$, we need $t \geq (2 \times \frac{s}{m})^2$ days! Even with market portfolio noise, which is diversified and low, we need at least a decade! During this decade, the true values of m and s can change!

```
(2 * s / m)**2 / 252
```

```
10.0486
```

Implement a simple moving average (SMA) trading strategy

The over-simplified goal of technical analysis is to “buy low, and sell high.” The n -day SMA reduces noise in market prices, removing market fluctuations and providing estimates of “true” prices. While the market price is *above* the SMA, the SMA *rises*. While the market price is *below* the SMA, the SMA *falls*. So, if we buy the stock as it crosses the SMA from below and sell the stock as it crosses the SMA from above, we mechanically buy low and sell high! Here, we will implement a long-only 20-day SMA strategy with Bitcoin:

1. Buy when the closing price crosses SMA(20) from below
2. Sell when the closing price crosses SMA(20) from above
3. No short-selling

We can simplify this strategy to “long if above SMA(20), otherwise neutral”. First, we will need Bitcoin returns data.

```

btc = (
    yf.download(
        tickers='BTC-USD',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .assign(Return=lambda x: x['Adj Close'].pct_change())
)

```

```

btc.head()

```

Date	Adj Close	Close	High	Low	Open	Volume	Return
2014-09-17	457.3340	457.3340	468.1740	452.4220	465.8640	21056800	NaN
2014-09-18	424.4400	424.4400	456.8600	413.1040	456.8600	34483200	-0.0719
2014-09-19	394.7960	394.7960	427.8350	384.5320	424.1030	37919700	-0.0698
2014-09-20	408.9040	408.9040	423.2960	389.8830	394.6730	36863600	0.0357
2014-09-21	398.8210	398.8210	412.4260	393.1810	408.0850	26580100	-0.0247


Next we:

1. Use `.rolling(20).mean()` to add a `SMA20` column containing `SMA(20)` to our `btc` data frame
2. Use `np.select()` to add a `Position` column containing:
 - 1 (long) when the adjusted close is greater than `SMA(20)`
 - 0 (neutral) when the adjusted close is less than (or equal to) `SMA(20)`
 - We use `.shift()` to compare yesterday's closing prices, avoiding a look-ahead bias**
 - `np.select()` tests multiple conditions and provides a default, making it more flexible framework than `np.where()`
3. Add a `Strategy` column containing:
 1. `Return if Position == 1`
 2. `0 if Position == 0`
 3. We could earn the risk-free rate instead of 0 percent, but earning 0 percent simplifies this example

```
btc = (
    btc
    .assign(
        SMA20=lambda x: x['Adj Close'].rolling(20).mean(),
        Position=lambda x: np.select(
            condlist=[
                x['Adj Close'].shift() > x['SMA20'].shift(),
                x['Adj Close'].shift() <= x['SMA20'].shift()
            ],
            choicelist=[1, 0],
            default=np.nan
        ),
        Strategy=lambda x: x['Position'] * x['Return']
    )
)
```

I find it helpful to plot `Adj Close`, `SMA20`, and `Position` for a sort window with one or more crossings.

```
fig, ax = plt.subplots(2, 1, sharex=True)
df = btc.loc['2023-02-13':'2023-02-28']
df[['Adj Close', 'SMA20']].plot(ax=ax[0], ylabel='BTC-USD ($)')
df[['Position']].plot(ax=ax[1], ylabel='Position', legend=False)
plt.suptitle('Bitcoin SMA(20) Strategy')
plt.show()
```




We can compare the long-run performance of buy-and-hold and SMA(20).

```
df = btc[['Return', 'Strategy']].dropna()

(
    df
    .add(1)
    .cumprod()
    .rename_axis(columns='Strategy')
    .rename(columns={'Return': 'Buy-And-Hold', 'Strategy': 'SMA(20)'})
    .plot()
)
plt.semilogy()
plt.ylabel('Value ($)')
```

Herron Topic 2 - Trading Strategies Based on Technical Analysis

```
plt.title(f'Value of $1 Invested at Close on {df.index[0] - pd.offsets.Day(1)}:{%B %d, %Y}')
plt.show()
```



In the practice notebook, we will dig deeper on this strategy and others.

```
df.add(1).prod()
```

```
Return      260.6042
Strategy    263.2456
dtype: float64
```

Herron Topic 2 - Practice

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf

%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

Implement the SMA(20) strategy with Bitcoin from the lecture notebook

Try to create the `btc` data frame in one code cell with one assignment (i.e., one `=`).

Investigate how SMA(20) generates returns

Consider the following:

1. Does SMA(20) avoid the worst performing days? How many of the worst 20 days does SMA(20) avoid? Try the `.sort_values()` or `.nlargest()` method.
2. Does SMA(20) preferentially avoid low-return days? Try to combine the `.groupby()` method and `pd.qcut()` function.
3. Does SMA(20) preferentially avoid high-volatility days? Try to combine the `.groupby()` method and `pd.qcut()` function.

Implement the SMA(20) strategy with IBM

How often does SMA(20) outperform buy-and-hold with 10-year rolling windows?

Implement a long-only BB(20, 2) strategy with Bitcoin

Bollinger Bands are bands around a trend, typically defined in terms of simple moving averages and volatilities. A long-only BB(20, 2) strategy has upper and lower bands at 2 standard deviations above and below the SMA(20). It invests as follows:

1. Buy when the closing price crosses LB(20) from below
2. Sell when the closing price crosses UB(20) from above
3. No short-selling

The long-only BB(20, 2) is more difficult to implement than the long-only SMA(20) because we need to track buys and sells. For example, if the closing price is between LB(20) and BB(20), we need to know if our last trade was a buy or a sell. Further, if the closing price is below LB(20), we can still be long because we sell when the closing price crosses UB(20) from above.

More on Bollinger Bands [here](#) and [here](#).

Implement a long-short RSI(14) strategy with Bitcoin

From [Fidelity](#):

The Relative Strength Index (RSI), developed by J. Welles Wilder, is a momentum oscillator that measures the speed and change of price movements. The RSI oscillates between zero and 100. Traditionally the RSI is considered overbought when above 70 and oversold when below 30. Signals can be generated by looking for divergences and failure swings. RSI can also be used to identify the general trend.

The RSI formula: $RSI(n) = 100 - \frac{100}{1+RS(n)}$, where $RS(n) = \frac{SMA(U,n)}{SMA(D,n)}$. For “up days”, $U = \Delta\text{Adj Close}$ and $D = 0$. For “down days”, $U = 0$ and $D = -\Delta\text{Adj Close}$.

We will implement a long-short RSI(14) as follows:

1. Buy when the RSI crosses 30 from below, and sell when the RSI crosses 50 from below
2. Short when the RSI crosses 70 from above, and cover when the RSI crosses 50 from above

More about RSI [here](#).

Implement a golden cross with Bitcoin

From Grok:

In technical analysis, a golden cross is a bullish chart pattern that occurs when a short-term moving average (typically the 50-day moving average) crosses above a long-term moving average (typically the 200-day moving average). This crossover is considered a signal that a stock, index, or other asset may be entering a sustained upward trend, suggesting potential buying opportunities for traders and investors.

More [here](#).

Compare all strategies with Bitcoin

Herron Topic 2 - Practice - Sec 02

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import seaborn as sns
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. I am still grading your projects; I hope to finish them by next Tuesday
2. I posted 50 practice problems to prepare for the end-of-course Programming Assessment here: https://northeastern.instructure.com/courses/207607/discussion_topics/2727917
 1. I built 5 autograded notebooks to help you prepare for the assessment, and I will build more in the coming weeks
 2. To run these notebooks, install the `otter-grader` package: In the Anaconda command prompt (or Terminal on Mac) run `conda activate fina6333` then `conda install otter-grader`
 3. See the video at the link above for details

Five-Minute Review

1. Technical analysis (TA) is a method of evaluating trends in trading prices and volume (and open interest in the futures and options markets only)
2. The three tenants of TA are:
 1. Markets discount everything

2. Prices move in trends
3. History repeats itself, so these trends are recurring
3. Academics have not found much evidence that TA generate profits that exceed transaction costs, but we will spend a week on it for three reasons:
 1. You asked for it, and it will help build our data analytics skills
 2. It is a small part of the Chartered Financial Analyst (CFA) curriculum
 3. That TA still receives attention 60 years after the Efficient Markets Hypothesis (EMH) suggests that it has some value that academics have been unable to measure

Practice

Implement the SMA(20) strategy with BTC-USD from the lecture notebook

Try to create the `btc_sma` data frame from the `btc` data frame in one code cell with one assignment (i.e., one `=`).

```
btc = (
    yf.download(
        tickers='BTC-USD',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)
```

After class, I converted our code to a function. We might want to test different moving average parameters, and a function makes these tests easier.

The following `calc_sma()` function accepts:

1. A data frame `df` of daily values from `yfinance.download()`
2. An integer `window` that specifies the number of trading days in the SMA window

And returns the original data frame `df` plus the following columns:

1. `Return` with daily returns
2. SMA for the `window`-trading-day moving average
3. `Position` for the weight on the security each day
4. `Strategy` for the return on the strategy each day

```

def calc_sma(df, window=20):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=window).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['Adj Close'].shift(1) > x['SMA'].shift(1),
                    x['Adj Close'].shift(1) <= x['SMA'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )
)

btc_sma = btc.pipe(calc_sma, window=20)

```

It can be helpful to visualize a few places where the price (here `Adj Close`) crosses the moving average (here `SMA(20)`). In class, we found that there are a few changes in position in the middle of October, 2014. The following code create two subplots in one figure, then puts one plot in each of the subplots.

For the first crossover on October 12, we *buy* at the close on October 12 and earn the security return on October 13. For the second crossover on October 23, we *sell* at the close on October 23 and earn zero return on October 24.


```

columns_sma20 = {
    'Adj Close': 'Price',
    'SMA': 'SMA(20)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)'
}

fig, ax = plt.subplots(2, 1)
df = btc_sma.loc['2014-10-10':'2014-10-25']
df[['Adj Close', 'SMA']].rename(columns=columns_sma20).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('SMA(20) Crossovers and Positions in Mid-October 2014')

```


```
plt.tight_layout()  
plt.show()
```



We can compare the total returns on BTC-USD and our SMA(20) strategy! We .dropna() first because the SMA(20) strategy needs 20 days of data to make its first investing decision.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()
```

```
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```




These total returns, at least today, are similar!

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .prod()  
    .plot(kind='bar')  
)  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
```

```
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies




However, the Sharpe ratio of the SMA(20) is higher, because time out of the market when `Position=0` reduces risk. For simplicity, we will ignore the risk-free rate in the Sharpe ratio formula.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)
```

```
.apply(lambda x: np.sqrt(252) * x.mean() / x.std())
.plot(kind='bar')
)
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies')
plt.ylabel('Sharpe Ratio')
plt.show()
```


Comparison of Reward-to-Risk in BTC-USD Strategies



But how persistent is this outperformance in terms of reward-to-risk ratios? We can quickly modify the code above to calculate Sharpe ratios each year! We find that SMA(20) Sharpe ratio edge is not persistent.

```
df = (
    btc_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .resample('YE')
    .apply(lambda x: np.sqrt(252) * x.mean() / x.std())
)
df.index = df.index.year
df.plot(kind='bar')
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies over Time')
plt.ylabel('Sharpe Ratio')
plt.show()
```

Comparison of Reward-to-Risk in BTC-USD Strategies over Time



We can use our list comprehension skills to easily try several window sizes!

```

def CAGR(x):
    T = x.count()
    return (x.add(1).prod() ** (252 / T)) - 1

def Sharpe(x, ann_fac=np.sqrt(252)):
    return ann_fac * x.mean() / x.std()

columns_sman = {
    'Adj Close': 'Price',
    'SMA': 'SMA(N)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(N)'
}

windows = list(range(5, 55, 5))

btc_smash = (
    pd.concat(
        objs=[
            btc.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

```

btc_smash


Window		Buy
	Statistic	
5	CAGR	0.40
	Sharpe	0.88
10	CAGR	0.40
	Sharpe	0.88

Window	Statistic	Buy
15	CAGR	0.40
	Sharpe	0.88
20	CAGR	0.40
	Sharpe	0.88
25	CAGR	0.40
	Sharpe	0.88
30	CAGR	0.40
	Sharpe	0.88
35	CAGR	0.40
	Sharpe	0.88
40	CAGR	0.40
	Sharpe	0.88
45	CAGR	0.40
	Sharpe	0.88
50	CAGR	0.40
	Sharpe	0.88

```

(
    btc_smas
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of BTC-USD Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()

```



Investigate how SMA(20) generates returns

Consider the following:

1. Does SMA(20) avoid the worst performing days? How many of the worst 20 days does SMA(20) avoid? Try the `.nlargest()` method.
2. Does SMA(20) preferentially avoid low-return days? Try to combine the `.groupby()` method and `pd.qcut()` function.
3. Does SMA(20) preferentially avoid high-volatility days? Try to combine the `.groupby()` method and `pd.qcut()` function.

The SMA(20) does well here because it avoids 17 of the 20 worst days, without avoiding the best days.

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    17
1.0000     3
Name: count, dtype: int64
```

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    10
1.0000    10
Name: count, dtype: int64
```

We can also look at the descriptive statistics by `Postiion` and strategy. Two observations:

1. The `min` column shows that SMA(20) misses the bad days, like we see above
2. The `mean` in `Position=1` is about 5 times higher than in `Position=0`


```
(btc_sma
[[['Position', 'Return', 'Strategy']]]
.dropna()
.groupby('Position')
.describe()
.rename(columns=columns_sma20)
.rename_axis(columns=['Strategy', 'Statistic'])
.stack('Strategy', future_stack=True))
```

Position	Statistic	Strategy
0.0000		Buy-And-Hold SMA(20)
1.0000		Buy-And-Hold SMA(20)

We can also use the seaborn package to visualize `Position` (i.e., the portfolio weight on Bitcoin) during periods of high and low Bitcoin returns and volatility. The SMA(20) strategy is long Bitcoin about 55% of the time, whether Bitcoin returns are high (bin 2) or low (bin 1).

```
(btc_sma
[['Return', 'Position']]
.dropna()
.assign(Return_Bin=lambda x: pd.qcut(x['Return'], q=3))
```


```
.pipe(
    sns.barplot,
    x='Return_Bin',
    y='Position'
)
plt.title('Fraction of Time SMA(20) Strategy is Long by Return Bin')
plt.show()
```



However, the SMA(20) strategy, *for this security, sample, and window*, spends less time in Bitcoin during volatile times.

```
(  
    btc_sma  
    [['Return', 'Position']]  
    .dropna()
```

```
.assign(Volatility_Bin=lambda x: pd.qcut(x['Return'].rolling(20).std(), q=3))
.pipe(
    sns.barplot,
    x='Volatility_Bin',
    y='Position'
)
plt.title('Fraction of Time SMA(20) Strategy is Long by Volatility Bin')
plt.show()
```



Implement the SMA(20) strategy with IBM


How often does SMA(20) outperform buy-and-hold with 10-year rolling windows?

```
ibm = (
    yf.download(
        tickers='IBM',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)

ibm_sma = ibm.pipe(calc_sma, window=20)

(
    ibm_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .add(1)
    .cumprod()
    .plot()
)
plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```


Comparison of One-Dollar Investments in IBM Strategies



Over the full, 60-year sample, SMA(20) underperforms buy-and-hold. What about on rolling ten-year windows? The results look balanced, and neither clearly outperforms.

```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in IBM Strategies



We can quantify how often SMA(20) outperform buy-and-hold. SMA(20) outperforms only 24% of the time!


```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .pipe(lambda x: x['Strategy'] > x['Return'])  
    .mean()  
)
```

0.2382

```
windows = list(range(5, 55, 5))

ibm_smash = (
    pd.concat(
        objs=[
            ibm.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

(
    ibm_smash
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of IBM Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()
```



Implement a long-only BB(20, 2) strategy with Bitcoin

Bollinger Bands are bands around a trend, typically defined in terms of simple moving averages and volatilities. A long-only BB(20, 2) strategy has upper and lower bands at 2 standard deviations above and below the SMA(20). It invests as follows:

1. Buy when the closing price crosses LB(20) from below
2. Sell when the closing price crosses UB(20) from above
3. No short-selling

The long-only BB(20, 2) is more difficult to implement than the long-only SMA(20) because we need to track buys and sells. For example, if the closing price is between LB(20) and BB(20), we need to know if our last trade was a buy or a sell. Further, if the closing price is below LB(20), we can still be long because we sell when the closing price crosses UB(20) from above.

More on Bollinger Bands [here](#) and [here](#).

```
def calc_bb(df, m=20, n=2):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=m).mean(),
            SMV=lambda x: x['Adj Close'].rolling(window=m).std(),
            UB=lambda x: x['SMA'] + n*x['SMV'],
```

```

LB=lambda x: x['SMA'] - n*x['SMV'],
Position_w_nan=lambda x: np.select(
    condlist=[
        (x['Adj Close'].shift(1) > x['LB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)),
        (x['Adj Close'].shift(1) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1))
    ],
    choicelist=[1, 0],
    default=np.nan
),
Position=lambda x: x['Position_w_nan'].ffill(),
Strategy=lambda x: x['Position'] * x['Return']
)
)

btc_bb = btc.pipe(calc_bb)

```

The BB(20, 2) only spends 40% of its time long BTC!

```
btc_bb['Position'].mean()
```

0.4013

And BTC-USD performance is worse when it long than when its neutral!

- Mean daily return is lower
- Volatility of daily returns is higher
- Every percentile of the distribution is worse

```
btc_bb.groupby('Position')['Return'].describe()
```

	count	mean	std	min	25%	50%	75%	max
Position								
0.0000	2268.0000	0.0027	0.0343	-0.1874	-0.0110	0.0019	0.0172	0.2525
1.0000	1520.0000	0.0013	0.0386	-0.3717	-0.0149	0.0007	0.0161	0.2394

```

columns_bb = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'BB(20, 2)'
}

```

```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_bb)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Here are the final values of \$1 investments, which are difficult to read on the log scale above.


```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .rename(columns=columns_bb)  
    .add(1)  
    .prod()  
    .rename_axis('Strategy')  
    .to_frame('Value of One-Dollar Investment')  
)
```

Strategy	Value of One-Dollar Investment
Buy-And-Hold	183.6056
BB(20, 2)	2.1877

We need a more complex plot to better understand what is going on!

```
import matplotlib.ticker as ticker  
  
fig, ax = plt.subplots(nrows=2, ncols=1, sharex=True)  
df = btc_bb.loc['2017-09':'2017-12']  
  
ax[0].plot(df[['Adj Close']], label='Price')  
ax[0].plot(df[['SMA']], label='SMA(20)')  
ax[0].plot(df[['UB']], label='LB and UB', color='green', linestyle=':')  
ax[0].plot(df[['LB']], color='green', linestyle=':')  
ax[0].legend()  
ax[0].set_ylabel('Price ($)')  
  
ax[1].plot(df[['Position']], label='Signal')  
ax[1].legend()  
ax[1].set_ylabel('Position')  
  
ax[0].yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))  
fig.autofmt_xdate()  
  
plt.suptitle('Key Variables in BTC-USD BB(20, 2) Strategy')  
plt.show()
```

Key Variables in BTC-USD BB(20, 2) Strategy



Implement a long-short RSI(14) strategy with Bitcoin

From [Fidelity](#):

The Relative Strength Index (RSI), developed by J. Welles Wilder, is a momentum oscillator that measures the speed and change of price movements. The RSI oscillates between zero and 100. Traditionally the RSI is considered overbought when above 70 and oversold when below 30. Signals can be generated by looking for divergences and failure swings. RSI can also be used to identify the general trend.

The RSI formula: $RSI(n) = 100 - \frac{100}{1+RS(n)}$, where $RS(n) = \frac{SMA(U,n)}{SMA(D,n)}$. For “up days”, $U = \Delta\text{Adj Close}$ and $D = 0$. For “down days”, $U = 0$ and $D = -\Delta\text{Adj Close}$.

We will implement a long-short RSI(14) as follows:

1. Buy when the RSI crosses 30 from below, and sell when the RSI crosses 50 from below

2. Short when the RSI crosses 70 from above, and cover when the RSI crosses 50 from above

More about RSI [here](#).

```
def calc_rsi(df, window=14, lo=30, mid=50, hi=70):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            # This approach with .max() and .min() handles NA values better than the in-class
            Diff=lambda x: x['Adj Close'].diff(),
            Zero=0,
            U=lambda x: x[['Diff', 'Zero']].max(axis=1, skipna=False),
            D=lambda x: -x[['Diff', 'Zero']].min(axis=1, skipna=False),
            SMAU=lambda x: x['U'].rolling(window=window).mean(),
            SMAD=lambda x: x['D'].rolling(window=window).mean(),
            RS=lambda x: x['SMAU'] / x['SMAD'],
            RSI=lambda x: 100 - 100 / (1 + x['RS']),
            Position_w_nan=lambda x: np.select(
                condlist=[
                    (x['RSI'].shift(1) > lo) & (x['RSI'].shift(2) <= lo),
                    (x['RSI'].shift(1) > mid) & (x['RSI'].shift(2) <= mid),
                    (x['RSI'].shift(1) < hi) & (x['RSI'].shift(2) >= hi),
                    (x['RSI'].shift(1) < mid) & (x['RSI'].shift(2) >= mid),
                ],
                choicelist=[1, 0, -1, 0],
                default=np.nan
            ),
            Position=lambda x: x['Position_w_nan'].ffill(),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )

btc_rsi = btc.pipe(calc_rsi)

columns_rsi = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'RSI(14)'
}
```

```
(  
    btc_rsi  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_rsi)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Implement a golden cross with Bitcoin

Someone in Section 04 mentioned two-moving average strategies, so I added this golden cross, where the 50-day SMA crosses the 200-day SMA, to every section.

From Grok:

In technical analysis, a golden cross is a bullish chart pattern that occurs when a short-term moving average (typically the 50-day moving average) crosses above a long-term moving average (typically the 200-day moving average). This crossover is considered a signal that a stock, index, or other asset may be entering a sustained upward trend, suggesting potential buying opportunities for traders and investors.

More [here](#).

```
def calc_cross(df, long=200, short=50):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMAL=lambda x: x['Adj Close'].rolling(window=long).mean(),
            SMAS=lambda x: x['Adj Close'].rolling(window=short).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['SMAS'].shift(1) > x['SMAL'].shift(1),
                    x['SMAS'].shift(1) <= x['SMAL'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )

btc_cross = btc.pipe(calc_cross, long=200, short=50)
```


```
columns_cross = {
    'Adj Close': 'Price',
    'SMAL': 'SMA(200)',
    'SMAS': 'SMA(50)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'Golden Cross(200, 50)'
}
```

```
btc_cross.query('Position == 1')
```

Date	Adj Close	Close	High	Low	Open	Volume	Return	SMAL
2015-07-15	285.8290	285.8290	293.2480	285.3670	288.0450	27486600	-0.0057	247.78
2015-07-16	278.0890	278.0890	291.1830	275.2400	286.0420	49482600	-0.0271	247.58
2015-07-17	279.4720	279.4720	280.2800	272.0430	278.0910	27591400	0.0050	247.42
2015-07-18	274.9010	274.9010	282.5270	274.0750	279.3310	25187100	-0.0164	247.24
2015-07-19	273.6140	273.6140	275.6700	272.5130	274.7670	15332500	-0.0047	247.00
...
2025-03-10	78532.0000	78532.0000	83955.9297	77420.5938	80597.1484	54061099422	-0.0257	83357
2025-03-11	82862.2109	82862.2109	83577.7578	76624.2500	78523.8750	54702837196	0.0551	83451
2025-03-12	83722.3594	83722.3594	84358.5781	80635.2500	82857.3750	40353484454	0.0104	83549
2025-03-13	81066.7031	81066.7031	84301.6953	79931.8516	83724.9219	31412940153	-0.0317	83633
2025-03-14	83969.1016	83969.1016	85263.2891	80797.5625	81066.9922	29588112414	0.0358	83738

```
fig, ax = plt.subplots(2, 1)
df = btc_cross.loc['2015-07':'2015-12']
df[['Adj Close', 'SMAL', 'SMAS']].rename(columns=columns_cross).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('Golden Cross(200, 50) Crossovers and Positions in 2015H2')
plt.tight_layout()
plt.show()
```

Golden Cross(200, 50) Crossovers and Positions in 2015H2



```
(  
    btc_cross  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_cross)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Compare all strategies

I added this comparison of all strategies after class.


```
df = (
    btc_sma[['Return', 'Strategy']]
    .join(btc_bb[['Strategy']], rsuffix='_bb')
    .join(btc_rsi[['Strategy']], rsuffix='_rsi')
    .join(btc_cross[['Strategy']], rsuffix='_cross')
    .dropna()
)

(
    df
```

```
.add(1)
.cumprod()
.rename_axis(columns='Strategy')
.rename(columns={
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)',
    'Strategy_bb': 'BB(20, 2)',
    'Strategy_rsi': 'RSI(14)',
    'Strategy_cross': 'Golden Cross(200, 50)',
})
.plot()
)

plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Herron Topic 2 - Practice - Sec 03

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import seaborn as sns
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. I am still grading your projects; I hope to finish them by next Tuesday
2. I posted 50 practice problems to prepare for the end-of-course Programming Assessment here: https://northeastern.instructure.com/courses/207607/discussion_topics/2727917
 1. I built 5 autograded notebooks to help you prepare for the assessment, and I will build more in the coming weeks
 2. To run these notebooks, install the `otter-grader` package: In the Anaconda command prompt (or Terminal on Mac) run `conda activate fina6333` then `conda install otter-grader`
 3. See the video at the link above for details

Five-Minute Review

1. Technical analysis (TA) is a method of evaluating trends in trading prices and volume (and open interest in the futures and options markets only)
2. The three tenants of TA are:
 1. Markets discount everything

2. Prices move in trends
3. History repeats itself, so these trends are recurring
3. Academics have not found much evidence that TA generate profits that exceed transaction costs, but we will spend a week on it for three reasons:
 1. You asked for it, and it will help build our data analytics skills
 2. It is a small part of the Chartered Financial Analyst (CFA) curriculum
 3. That TA still receives attention 60 years after the Efficient Markets Hypothesis (EMH) suggests that it has some value that academics have been unable to measure

Practice

Implement the SMA(20) strategy with BTC-USD from the lecture notebook

Try to create the `btc_sma` data frame from the `btc` data frame in one code cell with one assignment (i.e., one `=`).

```
btc = (
    yf.download(
        tickers='BTC-USD',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)
```

After class, I converted our code to a function. We might want to test different moving average parameters, and a function makes these tests easier.

The following `calc_sma()` function accepts:

1. A data frame `df` of daily values from `yfinance.download()`
2. An integer `window` that specifies the number of trading days in the SMA window

And returns the original data frame `df` plus the following columns:

1. `Return` with daily returns
2. SMA for the `window`-trading-day moving average
3. `Position` for the weight on the security each day
4. `Strategy` for the return on the strategy each day

```

def calc_sma(df, window=20):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=window).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['Adj Close'].shift(1) > x['SMA'].shift(1),
                    x['Adj Close'].shift(1) <= x['SMA'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )
)

btc_sma = btc.pipe(calc_sma, window=20)

```

It can be helpful to visualize a few places where the price (here `Adj Close`) crosses the moving average (here `SMA(20)`). In class, we found that there are a few changes in position in the middle of October, 2014. The following code create two subplots in one figure, then puts one plot in each of the subplots.

For the first crossover on October 12, we *buy* at the close on October 12 and earn the security return on October 13. For the second crossover on October 23, we *sell* at the close on October 23 and earn zero return on October 24.


```

columns_sma20 = {
    'Adj Close': 'Price',
    'SMA': 'SMA(20)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)'
}

fig, ax = plt.subplots(2, 1)
df = btc_sma.loc['2014-10-10':'2014-10-25']
df[['Adj Close', 'SMA']].rename(columns=columns_sma20).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('SMA(20) Crossovers and Positions in Mid-October 2014')

```


```
plt.tight_layout()  
plt.show()
```



We can compare the total returns on BTC-USD and our SMA(20) strategy! We `.dropna()` first because the SMA(20) strategy needs 20 days of data to make its first investing decision.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()
```

```
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```




These total returns, at least today, are similar!

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .prod()  
    .plot(kind='bar')  
)  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
```

```
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies




However, the Sharpe ratio of the SMA(20) is higher, because time out of the market when `Position=0` reduces risk. For simplicity, we will ignore the risk-free rate in the Sharpe ratio formula.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)
```

```
.apply(lambda x: np.sqrt(252) * x.mean() / x.std())
.plot(kind='bar')
)
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies')
plt.ylabel('Sharpe Ratio')
plt.show()
```


Comparison of Reward-to-Risk in BTC-USD Strategies



But how persistent is this outperformance in terms of reward-to-risk ratios? We can quickly modify the code above to calculate Sharpe ratios each year! We find that SMA(20) Sharpe ratio edge is not persistent.

```
df = (
    btc_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .resample('YE')
    .apply(lambda x: np.sqrt(252) * x.mean() / x.std())
)
df.index = df.index.year
df.plot(kind='bar')
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies over Time')
plt.ylabel('Sharpe Ratio')
plt.show()
```

Comparison of Reward-to-Risk in BTC-USD Strategies over Time



We can use our list comprehension skills to easily try several window sizes!

```

def CAGR(x):
    T = x.count()
    return (x.add(1).prod() ** (252 / T)) - 1

def Sharpe(x, ann_fac=np.sqrt(252)):
    return ann_fac * x.mean() / x.std()

columns_sman = {
    'Adj Close': 'Price',
    'SMA': 'SMA(N)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(N)'
}

windows = list(range(5, 55, 5))

btc_smash = (
    pd.concat(
        objs=[
            btc.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

```

btc_smash


Window		Buy
	Statistic	
5	CAGR	0.40
	Sharpe	0.88
10	CAGR	0.40
	Sharpe	0.88

Window	Statistic	Buy
15	CAGR	0.40
	Sharpe	0.88
20	CAGR	0.40
	Sharpe	0.88
25	CAGR	0.40
	Sharpe	0.88
30	CAGR	0.40
	Sharpe	0.88
35	CAGR	0.40
	Sharpe	0.88
40	CAGR	0.40
	Sharpe	0.88
45	CAGR	0.40
	Sharpe	0.88
50	CAGR	0.40
	Sharpe	0.88

```

(
    btc_smas
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of BTC-USD Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()

```



Investigate how SMA(20) generates returns

Consider the following:

1. Does SMA(20) avoid the worst performing days? How many of the worst 20 days does SMA(20) avoid? Try the `.nlargest()` method.
2. Does SMA(20) preferentially avoid low-return days? Try to combine the `.groupby()` method and `pd.qcut()` function.
3. Does SMA(20) preferentially avoid high-volatility days? Try to combine the `.groupby()` method and `pd.qcut()` function.

The SMA(20) does well here because it avoids 17 of the 20 worst days, without avoiding the best days.

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    17
1.0000     3
Name: count, dtype: int64
```

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    10
1.0000    10
Name: count, dtype: int64
```

We can also look at the descriptive statistics by `Postiion` and strategy. Two observations:

1. The `min` column shows that SMA(20) misses the bad days, like we see above
2. The `mean` in `Position=1` is about 5 times higher than in `Position=0`


```
(btc_sma
[[['Position', 'Return', 'Strategy']]]
.dropna()
.groupby('Position')
.describe()
.rename(columns=columns_sma20)
.rename_axis(columns=['Strategy', 'Statistic'])
.stack('Strategy', future_stack=True))
```

Position	Statistic	Strategy
0.0000		Buy-And-Hold SMA(20)
1.0000		Buy-And-Hold SMA(20)

We can also use the seaborn package to visualize `Position` (i.e., the portfolio weight on Bitcoin) during periods of high and low Bitcoin returns and volatility. The SMA(20) strategy is long Bitcoin about 55% of the time, whether Bitcoin returns are high (bin 2) or low (bin 1).

```
(btc_sma
[['Return', 'Position']]
.dropna()
.assign(Return_Bin=lambda x: pd.qcut(x['Return'], q=3))
```

```
.pipe(
    sns.barplot,
    x='Return_Bin',
    y='Position'
)
plt.title('Fraction of Time SMA(20) Strategy is Long by Return Bin')
plt.show()
```




However, the SMA(20) strategy, *for this security, sample, and window*, spends less time in Bitcoin during volatile times.

```
(  
    btc_sma  
    [['Return', 'Position']]  
    .dropna()
```

```
.assign(Volatility_Bin=lambda x: pd.qcut(x['Return'].rolling(20).std(), q=3))
    .pipe(
        sns.barplot,
        x='Volatility_Bin',
        y='Position'
    )
)

plt.title('Fraction of Time SMA(20) Strategy is Long by Volatility Bin')
plt.show()
```



Implement the SMA(20) strategy with IBM


How often does SMA(20) outperform buy-and-hold with 10-year rolling windows?

```
ibm = (
    yf.download(
        tickers='IBM',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)

ibm_sma = ibm.pipe(calc_sma, window=20)

(
    ibm_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .add(1)
    .cumprod()
    .plot()
)
plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```


Comparison of One-Dollar Investments in IBM Strategies



Over the full, 60-year sample, SMA(20) underperforms buy-and-hold. What about on rolling ten-year windows? The results look balanced, and neither clearly outperforms.

```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in IBM Strategies



We can quantify how often SMA(20) outperform buy-and-hold. SMA(20) outperforms only 24% of the time!


```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .pipe(lambda x: x['Strategy'] > x['Return'])  
    .mean()  
)
```

0.2315

```
windows = list(range(5, 55, 5))

ibm_smash = (
    pd.concat(
        objs=[
            ibm.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

(
    ibm_smash
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of IBM Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()
```



Implement a long-only BB(20, 2) strategy with Bitcoin

Bollinger Bands are bands around a trend, typically defined in terms of simple moving averages and volatilities. A long-only BB(20, 2) strategy has upper and lower bands at 2 standard deviations above and below the SMA(20). It invests as follows:

1. Buy when the closing price crosses LB(20) from below
2. Sell when the closing price crosses UB(20) from above
3. No short-selling

The long-only BB(20, 2) is more difficult to implement than the long-only SMA(20) because we need to track buys and sells. For example, if the closing price is between LB(20) and BB(20), we need to know if our last trade was a buy or a sell. Further, if the closing price is below LB(20), we can still be long because we sell when the closing price crosses UB(20) from above.

More on Bollinger Bands [here](#) and [here](#).

```
def calc_bb(df, m=20, n=2):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=m).mean(),
            SMV=lambda x: x['Adj Close'].rolling(window=m).std(),
            UB=lambda x: x['SMA'] + n*x['SMV'],
            LB=lambda x: x['SMA'] - n*x['SMV']
        )
        .dropna()
    )
```

```

LB=lambda x: x['SMA'] - n*x['SMV'],
Position_w_nan=lambda x: np.select(
    condlist=[
        (x['Adj Close'].shift(1) > x['LB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)),
        (x['Adj Close'].shift(1) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) > x['LB'].shift(1))
    ],
    choicelist=[1, 0],
    default=np.nan
),
Position=lambda x: x['Position_w_nan'].ffill(),
Strategy=lambda x: x['Position'] * x['Return']
)
)

btc_bb = btc.pipe(calc_bb)

```

The BB(20, 2) only spends 40% of its time long BTC!

```
btc_bb['Position'].mean()
```

0.4013

And BTC-USD performance is worse when it long than when its neutral!

- Mean daily return is lower
- Volatility of daily returns is higher
- Every percentile of the distribution is worse

```
btc_bb.groupby('Position')['Return'].describe()
```

	count	mean	std	min	25%	50%	75%	max
Position								
0.0000	2268.0000	0.0027	0.0343	-0.1874	-0.0110	0.0019	0.0172	0.2525
1.0000	1520.0000	0.0013	0.0386	-0.3717	-0.0149	0.0007	0.0161	0.2394

```

columns_bb = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'BB(20, 2)'
}

```

```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_bb)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Here are the final values of \$1 investments, which are difficult to read on the log scale above.


```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .rename(columns=columns_bb)  
    .add(1)  
    .prod()  
    .rename_axis('Strategy')  
    .to_frame('Value of One-Dollar Investment')  
)
```

Strategy	Value of One-Dollar Investment
Buy-And-Hold	183.6056
BB(20, 2)	2.1877

We need a more complex plot to better understand what is going on!

```
import matplotlib.ticker as ticker  
  
fig, ax = plt.subplots(nrows=2, ncols=1, sharex=True)  
df = btc_bb.loc['2017-09':'2017-12']  
  
ax[0].plot(df[['Adj Close']], label='Price')  
ax[0].plot(df[['SMA']], label='SMA(20)')  
ax[0].plot(df[['UB']], label='LB and UB', color='green', linestyle=':')  
ax[0].plot(df[['LB']], color='green', linestyle=':')  
ax[0].legend()  
ax[0].set_ylabel('Price ($)')  
  
ax[1].plot(df[['Position']], label='Signal')  
ax[1].legend()  
ax[1].set_ylabel('Position')  
  
ax[0].yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))  
fig.autofmt_xdate()  
  
plt.suptitle('Key Variables in BTC-USD BB(20, 2) Strategy')  
plt.show()
```

Key Variables in BTC-USD BB(20, 2) Strategy



Implement a long-short RSI(14) strategy with Bitcoin

From [Fidelity](#):

The Relative Strength Index (RSI), developed by J. Welles Wilder, is a momentum oscillator that measures the speed and change of price movements. The RSI oscillates between zero and 100. Traditionally the RSI is considered overbought when above 70 and oversold when below 30. Signals can be generated by looking for divergences and failure swings. RSI can also be used to identify the general trend.

The RSI formula: $RSI(n) = 100 - \frac{100}{1+RS(n)}$, where $RS(n) = \frac{SMA(U,n)}{SMA(D,n)}$. For “up days”, $U = \Delta\text{Adj Close}$ and $D = 0$. For “down days”, $U = 0$ and $D = -\Delta\text{Adj Close}$.

We will implement a long-short RSI(14) as follows:

1. Buy when the RSI crosses 30 from below, and sell when the RSI crosses 50 from below

2. Short when the RSI crosses 70 from above, and cover when the RSI crosses 50 from above

More about RSI [here](#).

```
def calc_rsi(df, window=14, lo=30, mid=50, hi=70):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            # This approach with .max() and .min() handles NA values better than the in-class
            Diff=lambda x: x['Adj Close'].diff(),
            Zero=0,
            U=lambda x: x[['Diff', 'Zero']].max(axis=1, skipna=False),
            D=lambda x: -x[['Diff', 'Zero']].min(axis=1, skipna=False),
            SMAU=lambda x: x['U'].rolling(window=window).mean(),
            SMAD=lambda x: x['D'].rolling(window=window).mean(),
            RS=lambda x: x['SMAU'] / x['SMAD'],
            RSI=lambda x: 100 - 100 / (1 + x['RS']),
            Position_w_nan=lambda x: np.select(
                condlist=[
                    (x['RSI'].shift(1) > lo) & (x['RSI'].shift(2) <= lo),
                    (x['RSI'].shift(1) > mid) & (x['RSI'].shift(2) <= mid),
                    (x['RSI'].shift(1) < hi) & (x['RSI'].shift(2) >= hi),
                    (x['RSI'].shift(1) < mid) & (x['RSI'].shift(2) >= mid),
                ],
                choicelist=[1, 0, -1, 0],
                default=np.nan
            ),
            Position=lambda x: x['Position_w_nan'].ffill(),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )

btc_rsi = btc.pipe(calc_rsi)

columns_rsi = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'RSI(14)'
}
```

```
(  
    btc_rsi  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_rsi)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Implement a golden cross with Bitcoin

Someone in Section 04 mentioned two-moving average strategies, so I added this golden cross, where the 50-day SMA crosses the 200-day SMA, to every section.

From Grok:

In technical analysis, a golden cross is a bullish chart pattern that occurs when a short-term moving average (typically the 50-day moving average) crosses above a long-term moving average (typically the 200-day moving average). This crossover is considered a signal that a stock, index, or other asset may be entering a sustained upward trend, suggesting potential buying opportunities for traders and investors.

More [here](#).

```
def calc_cross(df, long=200, short=50):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMAL=lambda x: x['Adj Close'].rolling(window=long).mean(),
            SMAS=lambda x: x['Adj Close'].rolling(window=short).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['SMAS'].shift(1) > x['SMAL'].shift(1),
                    x['SMAS'].shift(1) <= x['SMAL'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )
```

```
btc_cross = btc.pipe(calc_cross, long=200, short=50)
```


```
columns_cross = {
    'Adj Close': 'Price',
    'SMAL': 'SMA(200)',
    'SMAS': 'SMA(50)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'Golden Cross(200, 50)'
}
```

```
btc_cross.query('Position == 1')
```

Date	Adj Close	Close	High	Low	Open	Volume	Return	SMAL
2015-07-15	285.8290	285.8290	293.2480	285.3670	288.0450	27486600	-0.0057	247.78
2015-07-16	278.0890	278.0890	291.1830	275.2400	286.0420	49482600	-0.0271	247.58
2015-07-17	279.4720	279.4720	280.2800	272.0430	278.0910	27591400	0.0050	247.42
2015-07-18	274.9010	274.9010	282.5270	274.0750	279.3310	25187100	-0.0164	247.24
2015-07-19	273.6140	273.6140	275.6700	272.5130	274.7670	15332500	-0.0047	247.00
...
2025-03-10	78532.0000	78532.0000	83955.9297	77420.5938	80597.1484	54061099422	-0.0257	83357
2025-03-11	82862.2109	82862.2109	83577.7578	76624.2500	78523.8750	54702837196	0.0551	83451
2025-03-12	83722.3594	83722.3594	84358.5781	80635.2500	82857.3750	40353484454	0.0104	83549
2025-03-13	81066.7031	81066.7031	84301.6953	79931.8516	83724.9219	31412940153	-0.0317	83633
2025-03-14	83969.1016	83969.1016	85263.2891	80797.5625	81066.9922	29588112414	0.0358	83738


```
fig, ax = plt.subplots(2, 1)
df = btc_cross.loc['2015-07':'2015-12']
df[['Adj Close', 'SMAL', 'SMAS']].rename(columns=columns_cross).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('Golden Cross(200, 50) Crossovers and Positions in 2015H2')
plt.tight_layout()
plt.show()
```

Golden Cross(200, 50) Crossovers and Positions in 2015H2



```
(  
    btc_cross  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_cross)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Compare all strategies

I added this comparison of all strategies after class.


```
df = (
    btc_sma[['Return', 'Strategy']]
    .join(btc_bb[['Strategy']], rsuffix='_bb')
    .join(btc_rsi[['Strategy']], rsuffix='_rsi')
    .join(btc_cross[['Strategy']], rsuffix='_cross')
    .dropna()
)

(
    df
```

```
.add(1)
.cumprod()
.rename_axis(columns='Strategy')
.rename(columns={
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)',
    'Strategy_bb': 'BB(20, 2)',
    'Strategy_rsi': 'RSI(14)',
    'Strategy_cross': 'Golden Cross(200, 50)',
})
.plot()
)

plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Herron Topic 2 - Practice - Sec 04

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import seaborn as sns
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. I am still grading your projects; I hope to finish them by next Tuesday
2. I posted 50 practice problems to prepare for the end-of-course Programming Assessment here: https://northeastern.instructure.com/courses/207607/discussion_topics/2727917
 1. I built 5 autograded notebooks to help you prepare for the assessment, and I will build more in the coming weeks
 2. To run these notebooks, install the `otter-grader` package: In the Anaconda command prompt (or Terminal on Mac) run `conda activate fina6333` then `conda install otter-grader`
 3. See the video at the link above for details

Five-Minute Review

1. Technical analysis (TA) is a method of evaluating trends in trading prices and volume (and open interest in the futures and options markets only)
2. The three tenants of TA are:
 1. Markets discount everything

2. Prices move in trends
3. History repeats itself, so these trends are recurring
3. Academics have not found much evidence that TA generate profits that exceed transaction costs, but we will spend a week on it for three reasons:
 1. You asked for it, and it will help build our data analytics skills
 2. It is a small part of the Chartered Financial Analyst (CFA) curriculum
 3. That TA still receives attention 60 years after the Efficient Markets Hypothesis (EMH) suggests that it has some value that academics have been unable to measure

Practice

Implement the SMA(20) strategy with BTC-USD from the lecture notebook

Try to create the `btc_sma` data frame from the `btc` data frame in one code cell with one assignment (i.e., one `=`).

```
btc = (
    yf.download(
        tickers='BTC-USD',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)
```

After class, I converted our code to a function. We might want to test different moving average parameters, and a function makes these tests easier.

The following `calc_sma()` function accepts:

1. A data frame `df` of daily values from `yfinance.download()`
2. An integer `window` that specifies the number of trading days in the SMA window

And returns the original data frame `df` plus the following columns:

1. `Return` with daily returns
2. SMA for the `window`-trading-day moving average
3. `Position` for the weight on the security each day
4. `Strategy` for the return on the strategy each day

```

def calc_sma(df, window=20):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=window).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['Adj Close'].shift(1) > x['SMA'].shift(1),
                    x['Adj Close'].shift(1) <= x['SMA'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )
)

btc_sma = btc.pipe(calc_sma, window=20)

```

It can be helpful to visualize a few places where the price (here `Adj Close`) crosses the moving average (here `SMA(20)`). In class, we found that there are a few changes in position in the middle of October, 2014. The following code create two subplots in one figure, then puts one plot in each of the subplots.

For the first crossover on October 12, we *buy* at the close on October 12 and earn the security return on October 13. For the second crossover on October 23, we *sell* at the close on October 23 and earn zero return on October 24.


```

columns_sma20 = {
    'Adj Close': 'Price',
    'SMA': 'SMA(20)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)'
}

fig, ax = plt.subplots(2, 1)
df = btc_sma.loc['2014-10-10':'2014-10-25']
df[['Adj Close', 'SMA']].rename(columns=columns_sma20).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('SMA(20) Crossovers and Positions in Mid-October 2014')

```


```
plt.tight_layout()  
plt.show()
```



We can compare the total returns on BTC-USD and our SMA(20) strategy! We .dropna() first because the SMA(20) strategy needs 20 days of data to make its first investing decision.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()
```

```
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```




These total returns, at least today, are similar!

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .add(1)  
    .prod()  
    .plot(kind='bar')  
)  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
```

```
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies




However, the Sharpe ratio of the SMA(20) is higher, because time out of the market when `Position=0` reduces risk. For simplicity, we will ignore the risk-free rate in the Sharpe ratio formula.

```
(  
    btc_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)
```

```
.apply(lambda x: np.sqrt(252) * x.mean() / x.std())
.plot(kind='bar')
)
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies')
plt.ylabel('Sharpe Ratio')
plt.show()
```


Comparison of Reward-to-Risk in BTC-USD Strategies



But how persistent is this outperformance in terms of reward-to-risk ratios? We can quickly modify the code above to calculate Sharpe ratios each year! We find that SMA(20) Sharpe ratio edge is not persistent.

```
df = (
    btc_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .resample('YE')
    .apply(lambda x: np.sqrt(252) * x.mean() / x.std())
)
df.index = df.index.year
df.plot(kind='bar')
plt.suptitle('Comparison of Reward-to-Risk in BTC-USD Strategies over Time')
plt.ylabel('Sharpe Ratio')
plt.show()
```

Comparison of Reward-to-Risk in BTC-USD Strategies over Time



We can use our list comprehension skills to easily try several window sizes!

```

def CAGR(x):
    T = x.count()
    return (x.add(1).prod() ** (252 / T)) - 1

def Sharpe(x, ann_fac=np.sqrt(252)):
    return ann_fac * x.mean() / x.std()

columns_sman = {
    'Adj Close': 'Price',
    'SMA': 'SMA(N)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(N)'
}

windows = list(range(5, 55, 5))

btc_smash = (
    pd.concat(
        objs=[
            btc.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

```

btc_smash


Window		Buy
	Statistic	
5	CAGR	0.40
	Sharpe	0.88
10	CAGR	0.40
	Sharpe	0.88

Window	Statistic	Buy
15	CAGR	0.40
	Sharpe	0.88
20	CAGR	0.40
	Sharpe	0.88
25	CAGR	0.40
	Sharpe	0.88
30	CAGR	0.40
	Sharpe	0.88
35	CAGR	0.40
	Sharpe	0.88
40	CAGR	0.40
	Sharpe	0.88
45	CAGR	0.40
	Sharpe	0.88
50	CAGR	0.40
	Sharpe	0.88

```

(
    btc_smas
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of BTC-USD Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()

```



Investigate how SMA(20) generates returns

Consider the following:

1. Does SMA(20) avoid the worst performing days? How many of the worst 20 days does SMA(20) avoid? Try the `.nlargest()` method.
2. Does SMA(20) preferentially avoid low-return days? Try to combine the `.groupby()` method and `pd.qcut()` function.
3. Does SMA(20) preferentially avoid high-volatility days? Try to combine the `.groupby()` method and `pd.qcut()` function.

The SMA(20) does well here because it avoids 17 of the 20 worst days, without avoiding the best days.

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    17
1.0000     3
Name: count, dtype: int64
```

```
btc_sma.loc[btc_sma['Return'].nlargest(20).index, ['Position']].value_counts()
```

```
Position
0.0000    10
1.0000    10
Name: count, dtype: int64
```

We can also look at the descriptive statistics by `Postiion` and strategy. Two observations:

1. The `min` column shows that SMA(20) misses the bad days, like we see above
2. The `mean` in `Position=1` is about 5 times higher than in `Position=0`


```
(btc_sma
[[ 'Position', 'Return', 'Strategy']]
.dropna()
.groupby('Position')
.describe()
.rename(columns=columns_sma20)
.rename_axis(columns=['Strategy', 'Statistic'])
.stack('Strategy', future_stack=True))
```

Position	Statistic	Strategy
0.0000		Buy-And-Hold SMA(20)
1.0000		Buy-And-Hold SMA(20)

We can also use the seaborn package to visualize `Position` (i.e., the portfolio weight on Bitcoin) during periods of high and low Bitcoin returns and volatility. The SMA(20) strategy is long Bitcoin about 55% of the time, whether Bitcoin returns are high (bin 2) or low (bin 1).

```
(btc_sma
[['Return', 'Position']]
.dropna()
.assign(Return_Bin=lambda x: pd.qcut(x['Return'], q=3))
```

```
.pipe(
    sns.barplot,
    x='Return_Bin',
    y='Position'
)
plt.title('Fraction of Time SMA(20) Strategy is Long by Return Bin')
plt.show()
```




However, the SMA(20) strategy, *for this security, sample, and window*, spends less time in Bitcoin during volatile times.

```
(  
    btc_sma  
    [['Return', 'Position']]  
    .dropna()
```

```
.assign(Volatility_Bin=lambda x: pd.qcut(x['Return'].rolling(20).std(), q=3))
    .pipe(
        sns.barplot,
        x='Volatility_Bin',
        y='Position'
    )
)

plt.title('Fraction of Time SMA(20) Strategy is Long by Volatility Bin')
plt.show()
```



Implement the SMA(20) strategy with IBM


How often does SMA(20) outperform buy-and-hold with 10-year rolling windows?

```
ibm = (
    yf.download(
        tickers='IBM',
        auto_adjust=False,
        progress=False,
        multi_level_index=False
    )
    .iloc[:-1] # drop incomplete trading day
)

ibm_sma = ibm.pipe(calc_sma, window=20)

(
    ibm_sma
    [['Return', 'Strategy']]
    .dropna()
    .rename(columns=columns_sma20)
    .add(1)
    .cumprod()
    .plot()
)
plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```


Comparison of One-Dollar Investments in IBM Strategies



Over the full, 60-year sample, SMA(20) underperforms buy-and-hold. What about on rolling ten-year windows? The results look balanced, and neither clearly outperforms.

```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_sma20)  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in IBM Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in IBM Strategies



We can quantify how often SMA(20) outperform buy-and-hold. SMA(20) outperforms only 24% of the time!


```
(  
    ibm_sma  
    [['Return', 'Strategy']]  
    .dropna()  
    .pipe(np.log1p)  
    .rolling(window=10*252)  
    .sum()  
    .pipe(np.exp)  
    .pipe(lambda x: x['Strategy'] > x['Return'])  
    .mean()  
)
```

0.2378

```
windows = list(range(5, 55, 5))

ibm_smash = (
    pd.concat(
        objs=[
            ibm.pipe(calc_sma, window=w)[['Return', 'Strategy']].agg([CAGR, Sharpe])
            for w in windows
        ],
        keys=windows,
        names=['Window', 'Statistic']
    )
    .rename(columns=columns_sman)
)

(
    ibm_smash
    .reset_index()
    .melt(
        id_vars=['Window', 'Statistic'],
        value_vars=['Buy-And-Hold', 'SMA(N)'],
        var_name='Strategy',
        value_name='Value'
    )
    .pipe(
        sns.catplot,
        x='Window',
        col='Statistic',
        hue='Strategy',
        y='Value',
        kind='bar',
        sharey=False
    )
)
plt.suptitle('Comparsion of IBM Buy-And-Hold and SMA(N) for Various Ns', y=1.05)
plt.show()
```



Implement a long-only BB(20, 2) strategy with Bitcoin

Bollinger Bands are bands around a trend, typically defined in terms of simple moving averages and volatilities. A long-only BB(20, 2) strategy has upper and lower bands at 2 standard deviations above and below the SMA(20). It invests as follows:

1. Buy when the closing price crosses LB(20) from below
2. Sell when the closing price crosses UB(20) from above
3. No short-selling

The long-only BB(20, 2) is more difficult to implement than the long-only SMA(20) because we need to track buys and sells. For example, if the closing price is between LB(20) and BB(20), we need to know if our last trade was a buy or a sell. Further, if the closing price is below LB(20), we can still be long because we sell when the closing price crosses UB(20) from above.

More on Bollinger Bands [here](#) and [here](#).

```
def calc_bb(df, m=20, n=2):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMA=lambda x: x['Adj Close'].rolling(window=m).mean(),
            SMV=lambda x: x['Adj Close'].rolling(window=m).std(),
            UB=lambda x: x['SMA'] + n*x['SMV'],
            LB=lambda x: x['SMA'] - n*x['SMV'],
            Buy=lambda x: np.where(x['Return'] < x['LB'], 1, 0),
            Sell=lambda x: np.where(x['Return'] > x['UB'], 1, 0),
            Position=lambda x: x['Buy'].cumsum() - x['Sell'].cumsum()
        )
    )
```

```

LB=lambda x: x['SMA'] - n*x['SMV'],
Position_w_nan=lambda x: np.select(
    condlist=[
        (x['Adj Close'].shift(1) > x['LB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) < x['UB'].shift(1)),
        (x['Adj Close'].shift(1) < x['UB'].shift(1)) & (x['Adj Close'].shift(2) > x['LB'].shift(1))
    ],
    choicelist=[1, 0],
    default=np.nan
),
Position=lambda x: x['Position_w_nan'].ffill(),
Strategy=lambda x: x['Position'] * x['Return']
)
)

btc_bb = btc.pipe(calc_bb)

```

The BB(20, 2) only spends 40% of its time long BTC!

```
btc_bb['Position'].mean()
```

0.4013

And BTC-USD performance is worse when it long than when its neutral!

- Mean daily return is lower
- Volatility of daily returns is higher
- Every percentile of the distribution is worse

```
btc_bb.groupby('Position')['Return'].describe()
```

	count	mean	std	min	25%	50%	75%	max
Position								
0.0000	2268.0000	0.0027	0.0343	-0.1874	-0.0110	0.0019	0.0172	0.2525
1.0000	1520.0000	0.0013	0.0386	-0.3717	-0.0149	0.0007	0.0161	0.2394

```

columns_bb = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'BB(20, 2)'
}

```

```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_bb)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Here are the final values of \$1 investments, which are difficult to read on the log scale above.

```
(  
    btc_bb  
    [['Return', 'Strategy']]  
    .rename(columns=columns_bb)  
    .add(1)  
    .prod()  
    .rename_axis('Strategy')  
    .to_frame('Value of One-Dollar Investment')  
)
```

Strategy	Value of One-Dollar Investment
Buy-And-Hold	183.6056
BB(20, 2)	2.1877

We need a more complex plot to better understand what is going on!

```
import matplotlib.ticker as ticker

fig, ax = plt.subplots(nrows=2, ncols=1, sharex=True)
df = btc_bb.loc['2017-09':'2017-12']


ax[0].plot(df[['Adj Close']], label='Price')
ax[0].plot(df[['SMA']], label='SMA(20)')
ax[0].plot(df[['UB']], label='LB and UB', color='green', linestyle=':')
ax[0].plot(df[['LB']], color='green', linestyle=':')
ax[0].legend()
ax[0].set_ylabel('Price ($)')

ax[1].plot(df[['Position']], label='Signal')
ax[1].legend()
ax[1].set_ylabel('Position')

ax[0].yaxis.set_major_formatter(ticker.StrMethodFormatter('{x:,.0f}'))
fig.autofmt_xdate()

plt.suptitle('Key Variables in BTC-USD BB(20, 2) Strategy')
plt.show()
```

Key Variables in BTC-USD BB(20, 2) Strategy



Implement a long-short RSI(14) strategy with Bitcoin

From [Fidelity](#):

The Relative Strength Index (RSI), developed by J. Welles Wilder, is a momentum oscillator that measures the speed and change of price movements. The RSI oscillates between zero and 100. Traditionally the RSI is considered overbought when above 70 and oversold when below 30. Signals can be generated by looking for divergences and failure swings. RSI can also be used to identify the general trend.

The RSI formula: $RSI(n) = 100 - \frac{100}{1+RS(n)}$, where $RS(n) = \frac{SMA(U,n)}{SMA(D,n)}$. For “up days”, $U = \Delta\text{Adj Close}$ and $D = 0$. For “down days”, $U = 0$ and $D = -\Delta\text{Adj Close}$.

We will implement a long-short RSI(14) as follows:

1. Buy when the RSI crosses 30 from below, and sell when the RSI crosses 50 from below

2. Short when the RSI crosses 70 from above, and cover when the RSI crosses 50 from above

More about RSI [here](#).

```
def calc_rsi(df, window=14, lo=30, mid=50, hi=70):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            # This approach with .max() and .min() handles NA values better than the in-class
            Diff=lambda x: x['Adj Close'].diff(),
            Zero=0,
            U=lambda x: x[['Diff', 'Zero']].max(axis=1, skipna=False),
            D=lambda x: -x[['Diff', 'Zero']].min(axis=1, skipna=False),
            SMAU=lambda x: x['U'].rolling(window=window).mean(),
            SMAD=lambda x: x['D'].rolling(window=window).mean(),
            RS=lambda x: x['SMAU'] / x['SMAD'],
            RSI=lambda x: 100 - 100 / (1 + x['RS']),
            Position_w_nan=lambda x: np.select(
                condlist=[
                    (x['RSI'].shift(1) > lo) & (x['RSI'].shift(2) <= lo),
                    (x['RSI'].shift(1) > mid) & (x['RSI'].shift(2) <= mid),
                    (x['RSI'].shift(1) < hi) & (x['RSI'].shift(2) >= hi),
                    (x['RSI'].shift(1) < mid) & (x['RSI'].shift(2) >= mid),
                ],
                choicelist=[1, 0, -1, 0],
                default=np.nan
            ),
            Position=lambda x: x['Position_w_nan'].ffill(),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )

btc_rsi = btc.pipe(calc_rsi)

columns_rsi = {
    'Adj Close': 'Price',
    'Return': 'Buy-And-Hold',
    'Strategy': 'RSI(14)'
}
```

```
(  
    btc_rsi  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_rsi)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Implement a golden cross with Bitcoin

Someone in Section 04 mentioned two-moving average strategies, so I added this golden cross, where the 50-day SMA crosses the 200-day SMA, to every section.

From Grok:

In technical analysis, a golden cross is a bullish chart pattern that occurs when a short-term moving average (typically the 50-day moving average) crosses above a long-term moving average (typically the 200-day moving average). This crossover is considered a signal that a stock, index, or other asset may be entering a sustained upward trend, suggesting potential buying opportunities for traders and investors.

More [here](#).

```
def calc_cross(df, long=200, short=50):
    return (
        df
        .assign(
            Return=lambda x: x['Adj Close'].pct_change(),
            SMAL=lambda x: x['Adj Close'].rolling(window=long).mean(),
            SMAS=lambda x: x['Adj Close'].rolling(window=short).mean(),
            Position=lambda x: np.select(
                condlist=[
                    x['SMAS'].shift(1) > x['SMAL'].shift(1),
                    x['SMAS'].shift(1) <= x['SMAL'].shift(1)
                ],
                choicelist=[1, 0],
                default=np.nan
            ),
            Strategy=lambda x: x['Position'] * x['Return']
        )
    )
```

```
btc_cross = btc.pipe(calc_cross, long=200, short=50)
```


```
columns_cross = {
    'Adj Close': 'Price',
    'SMAL': 'SMA(200)',
    'SMAS': 'SMA(50)',
    'Return': 'Buy-And-Hold',
    'Strategy': 'Golden Cross(200, 50)'
}
```

```
btc_cross.query('Position == 1')
```

Date	Adj Close	Close	High	Low	Open	Volume	Return	SMAL
2015-07-15	285.8290	285.8290	293.2480	285.3670	288.0450	27486600	-0.0057	247.78
2015-07-16	278.0890	278.0890	291.1830	275.2400	286.0420	49482600	-0.0271	247.58
2015-07-17	279.4720	279.4720	280.2800	272.0430	278.0910	27591400	0.0050	247.42
2015-07-18	274.9010	274.9010	282.5270	274.0750	279.3310	25187100	-0.0164	247.24
2015-07-19	273.6140	273.6140	275.6700	272.5130	274.7670	15332500	-0.0047	247.00
...
2025-03-10	78532.0000	78532.0000	83955.9297	77420.5938	80597.1484	54061099422	-0.0257	83357
2025-03-11	82862.2109	82862.2109	83577.7578	76624.2500	78523.8750	54702837196	0.0551	83451
2025-03-12	83722.3594	83722.3594	84358.5781	80635.2500	82857.3750	40353484454	0.0104	83549
2025-03-13	81066.7031	81066.7031	84301.6953	79931.8516	83724.9219	31412940153	-0.0317	83633
2025-03-14	83969.1016	83969.1016	85263.2891	80797.5625	81066.9922	29588112414	0.0358	83738

```
fig, ax = plt.subplots(2, 1)
df = btc_cross.loc['2015-07':'2015-12']
df[['Adj Close', 'SMAL', 'SMAS']].rename(columns=columns_cross).plot(ax=ax[0])
df[['Position']].plot(ax=ax[1])
plt.suptitle('Golden Cross(200, 50) Crossovers and Positions in 2015H2')
plt.tight_layout()
plt.show()
```

Golden Cross(200, 50) Crossovers and Positions in 2015H2



```
(  
    btc_cross  
    [['Return', 'Strategy']]  
    .dropna()  
    .rename(columns=columns_cross)  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')  
plt.ylabel('Value of One-Dollar Investments')  
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Compare all strategies

I added this comparison of all strategies after class.


```
df = (
    btc_sma[['Return', 'Strategy']]
    .join(btc_bb[['Strategy']], rsuffix='_bb')
    .join(btc_rsi[['Strategy']], rsuffix='_rsi')
    .join(btc_cross[['Strategy']], rsuffix='_cross')
    .dropna()
)

(
    df
```

```
.add(1)
.cumprod()
.rename_axis(columns='Strategy')
.rename(columns={
    'Return': 'Buy-And-Hold',
    'Strategy': 'SMA(20)',
    'Strategy_bb': 'BB(20, 2)',
    'Strategy_rsi': 'RSI(14)',
    'Strategy_cross': 'Golden Cross(200, 50)',
})
.plot()
)

plt.semilogy()
plt.suptitle('Comparison of One-Dollar Investments in BTC-USD Strategies')
plt.ylabel('Value of One-Dollar Investments')
plt.show()
```

Comparison of One-Dollar Investments in BTC-USD Strategies



Week 10

Herron Topic 3 - Quantitative Value Investing

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

This notebook adapts the portfolio formation approach from Fama and French (1993), who show that firms with high book-to-market equity ratios (value stocks) tend to outperform those with low ratios (growth stocks), forming the basis of the value factor in their three-factor model. This notebook relies on two large data files. You can find these files in our [shared OneDrive folder](#) or our [Notebook and Syllabus folder on Canvas](#). These two files are:

1. Firm-level market data (e.g., stock prices and share counts) from the [Center for Research in Security Prices \(CRSP\)](#) in `crsp.csv`
2. Firm-level fundamentals (i.e., financial statement information) from [Compustat](#) in `compustat.csv`

This notebook walks you through the steps to test a quantitative value investing strategy based on the book-to-market equity ratio (B/M). I based this lecture notebook on Gray and Carlisle (2012, Chapters 1 and 7), which presents much more on quantitative value investing.

Import the data from CRSP and Compustat

CRSP

The CRSP data provide the following variables:

1. PERMNO is the stock identifier
2. date is the last trading day of the month
3. SHRCD indicates share class, and we can ignore it
4. PRC is the share price in dollars, and negative prices indicate the price is the midpoint of the bid-ask spread instead of a transaction
5. SHROUT is the number of shares outstanding in thousands

Read the CRSP data into a data frame named `crsp`. Be sure to parse dates. The codea A, B, and C indicate missing values.

```
crsp = (
    pd.read_csv(
        filepath_or_buffer='crsp.csv',
        parse_dates=['date'],
        na_values=list('ABC')
    )
    .query('date >= 1965')
    .sort_values(by=['PERMNO', 'date'])
    .assign(
        ME=lambda x: x['PRC'].abs() * x['SHROUT'] / 1000,
        date=lambda x: x['date'] + pd.offsets.MonthEnd(0)
    )
)
```

```
crsp.head()
```

	PERMNO	date	SHRCD	PRC	RET	SHROUT	ME
0	10000	1986-01-31	10	-4.3750	NaN	3680.0000	16.1000
1	10000	1986-02-28	10	-3.2500	-0.2571	3680.0000	11.9600
2	10000	1986-03-31	10	-4.4375	0.3654	3680.0000	16.3300
3	10000	1986-04-30	10	-4.0000	-0.0986	3793.0000	15.1720
4	10000	1986-05-31	10	-3.1094	-0.2227	3793.0000	11.7939

Compustat

The Compustat data provide the following variables:

1. GVKEY is the S&P firm identifier
2. LPERMNO is the link between CRSP and Compustat
3. datadate is the date of the filing, which is the last day of the fiscal year
4. fyear is the fiscal year
5. ceq is the book value of equity
6. ib is income before extraordinary items
7. We can ignore the other variables that Wharton Research Data Services (WRDS) adds to this data set

Read the Compustat data into a data frame named `compustat`. Be sure to parse dates.

```
compustat = (
    pd.read_csv(
        filepath_or_buffer='compustat.csv',
        parse_dates=['datadate']
    )
    .query('datadate >= 1965')
    .sort_values(['LPERMNO', 'fyear', 'datadate'])
    .drop_duplicates(subset=['LPERMNO', 'fyear'], keep='last')
)
```

```
compustat.head()
```

	GVKEY	LPERMNO	datadate	fyear	indfmt	consol	popsrc	datafmt	curcd	ceq
165681	13007	10000	1986-10-31	1986.0000	INDL	C	D	STD	USD	0.418
165586	12994	10001	1986-06-30	1986.0000	INDL	C	D	STD	USD	5.432
165587	12994	10001	1987-06-30	1987.0000	INDL	C	D	STD	USD	5.369
165588	12994	10001	1988-06-30	1988.0000	INDL	C	D	STD	USD	5.512
165589	12994	10001	1989-06-30	1989.0000	INDL	C	D	STD	USD	6.321

Create interim data frames

To test the book/market investing strategy, we need to match the following data:

1. Book value of equity for year $t - 1$
2. Market value of equity from December of year $t - 1$

3. Portfolios for year t based on book and market values of equity in year $t - 1$
4. Returns from July of year t through June of year $t + 1$

We will form portfolios once a year because the book value of equity updates once a year. Fama and French (1993) wait until June of year t to form portfolios, ensuring financial data from year $t - 1$ is public, then track returns from July in year t through June in year $t + 1$. The simplest and clearest way to match these data is to create one data frame for each of these three data sets.

Market value of equity from December of year $t - 1$

Then create a new data frame `mve` that contains the latest market value of equity as of December.

```
mve = (
    crsp
    .query('ME > 0')
    .groupby(['PERMNO', pd.Grouper(key='date', freq='YE-DEC')])
    [['ME']]
    .last()
)
```

```
crsp.query('PERMNO == 10001').head(12)
```

	PERMNO	date	SHRCRD	PRC	RET	SHROUT	ME
18	10001	1986-01-31	11	-6.1250	NaN	985.0000	6.0331
19	10001	1986-02-28	11	-6.2500	0.0204	985.0000	6.1562
20	10001	1986-03-31	11	-6.3125	0.0252	985.0000	6.2178
21	10001	1986-04-30	11	-6.3750	0.0099	985.0000	6.2794
22	10001	1986-05-31	11	-6.3125	-0.0098	985.0000	6.2178
23	10001	1986-06-30	11	-6.1250	-0.0131	985.0000	6.0331
24	10001	1986-07-31	11	-6.0625	-0.0102	985.0000	5.9716
25	10001	1986-08-31	11	-6.5000	0.0722	985.0000	6.4025
26	10001	1986-09-30	11	6.3750	-0.0031	991.0000	6.3176
27	10001	1986-10-31	11	6.6250	0.0392	991.0000	6.5654
28	10001	1986-11-30	11	7.0000	0.0566	991.0000	6.9370
29	10001	1986-12-31	11	7.0000	0.0150	991.0000	6.9370

```
mve.head()
```

PERMNO	date	M
10000	1986-12-31	1.
	1987-12-31	0.
	1986-12-31	6.
10001	1987-12-31	5.
	1988-12-31	6.

Book value of equity for year $t - 1$

Create a new data frame `bve` that contains the latest book value of equity as of December. Because some firms have fiscal years that do not end in December, use `.groupby()`, `pd.Grouper()` with `YE-DEC`, and `.last()` to roll forward earlier fiscal year ends.

```
bve = (
    compustat
    .query('ceq > 0')
    .groupby(by=['LPERMNO', pd.Grouper(key='datadate', freq='YE-DEC')])
    [['ceq']]
    .last()
)
```

```
compustat.query('LPERMNO == 10001').head()
```

	GVKEY	LPERMNO	datadate	fyear	indfmt	consol	popsrc	datafmt	curcd	ceq
165586	12994	10001	1986-06-30	1986.0000	INDL	C	D	STD	USD	5.432
165587	12994	10001	1987-06-30	1987.0000	INDL	C	D	STD	USD	5.369
165588	12994	10001	1988-06-30	1988.0000	INDL	C	D	STD	USD	5.512
165589	12994	10001	1989-06-30	1989.0000	INDL	C	D	STD	USD	6.321
165590	12994	10001	1990-06-30	1990.0000	INDL	C	D	STD	USD	7.179

```
bve.head()
```

LPERMNO	datadate	ce
10000	1986-12-31	0.
	1986-12-31	5.
10001		

LPERMNO	date	date	ce
	1987-12-31	5.	
	1988-12-31	5.	
	1989-12-31	6.	

Portfolios for year t

```
1 + pd.qcut(x=np.arange(10), q=5, labels=False)
```


```
array([1, 1, 2, 2, 3, 3, 4, 4, 5, 5])
```

```
portfolios = (
    bve
    [['ceq']]
    .dropna()
    .rename_axis(index=['PERMNO', 'Date'])
    .join(
        other=(
            mve
            [['ME']]
            .dropna()
            .rename_axis(index=['PERMNO', 'Date'])
        ),
        how='inner'
    )
    .assign(BM=lambda x: x['ceq'] / x['ME'])
    .reset_index()
    .assign(
        Portfolio=lambda x: x.groupby('Date')['BM'].transform(lambda x: 1 + pd.qcut(x=x, q=10))
        Date=lambda x: x['Date'] + pd.DateOffset(months=7)
    )
    [['PERMNO', 'Date', 'Portfolio', 'BM']]
)
```

```
portfolios.head()
```

	PERMNO	Date	Portfolio	BM
0	10000	1987-07-31	2	0.2109
1	10001	1987-07-31	7	0.7830
2	10001	1988-07-31	7	0.9212
3	10001	1989-07-31	7	0.8664
4	10001	1990-07-31	5	0.6109

```
(  
    portfolios  
    .groupby(by=['Portfolio', 'Date'])  
    ['BM']  
    .mean()  
    .groupby(by=['Portfolio'])  
    .mean()  
    .plot(kind='bar')  
)  
plt.ylabel('Mean B/M')  
plt.xlabel('Portfolio (1 = Low B/M, 10 = High B/M)')  
plt.title('Mean B/M for Ten Portfolios Formed on B/M')  
plt.show()
```



Returns from July of year t to June of year $t + 1$

Now we assign returns to each of these ten portfolios.

```
returns = (
    pd.merge_asof(
        left=crsp.sort_values(by=['date', 'PERMNO']),
        right=portfolios.sort_values(by=['Date', 'PERMNO']),
        left_on='date',
        right_on='Date',
        by='PERMNO',
        tolerance=pd.Timedelta('366 days')
    )
    [['PERMNO', 'date', 'RET', 'Portfolio']]
    .dropna()
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int))
```


```
.rename(columns={'date': 'Date'})  
)
```

```
returns
```

	PERMNO	Date	RET	Portfolio
37636	10006	1966-07-31	-0.0149	6
37640	10102	1966-07-31	-0.0545	6
37641	10137	1966-07-31	0.0105	4
37642	10145	1966-07-31	-0.0129	6
37643	10153	1966-07-31	-0.0788	9
...
3457680	93397	2024-12-31	-0.1174	5
3457681	93426	2024-12-31	0.0218	7
3457682	93429	2024-12-31	-0.0947	3
3457683	93434	2024-12-31	0.1333	10
3457684	93436	2024-12-31	0.1700	1


Performance of Portfolios Formed on B/M

```
(  
    returns  
    .groupby(['Portfolio', 'Date'])  
    ['RET']  
    .mean()  
    .groupby(['Portfolio'])  
    .mean()  
    .mul(100)  
    .plot(kind='bar')  
)  
plt.ylabel('Mean Monthly Return (%)')  
plt.xlabel('Portfolio (1 = Low B/M, 10 = High B/M)')  
plt.title('Mean Monthly Returns for Equal-Weighted Portfolios Formed on B/M')  
plt.show()
```



```
(  
    returns  
    .groupby(['Portfolio', 'Date'])  
    ['RET']  
    .mean()  
    .unstack('Portfolio')  
    .add(1)  
    .cumprod()  
    .plot()  
)  
plt.semilogy()  
plt.ylabel('Value of $1 Investment')  
plt.title('Values of $1 Investments in Equal-Weighted Portfolios Formed on B/M')  
plt.show()
```


Values of \$1 Investments in Equal-Weighted Portfolios Formed on B/M



```

(
    returns
    .groupby(['Portfolio', 'Date'])
    ['RET']
    .mean()
    .unstack('Portfolio')
    .assign(Diff=lambda x: x[10] - x[1])
    ['Diff']
    .rolling(12*10)
    .mean()
    .mul(100)
    .plot()
)
plt.ylabel('10-year Rolling Mean of Monthly Spread (%)')
plt.title('10-year Rolling Mean of Monthly Spread\n Between Value and Growth Portfolios')
plt.show()

```



Herron Topic 3 - Practice - Blank

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Review

Practice

Re-implement the value strategy from the lecture notebook

Read the data

Create the interim data frames

Combine the data frames and form portfolios

Backtest the strategy

Re-implement the value strategy from the lecture notebook *with value-weighted portfolios*

Estimate the α s of the equal-weighted and value-weighted portfolios

Implement a momentum strategy

Form deciles on the 11-month returns from months $t - 12$ to month $t - 1$.

Herron Topic 3 - Practice - Sec 02

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Projects:
 1. Project 1:
 1. I am still grading
 2. I plan to finish grading by Friday
 2. Project 2:
 1. Due on Friday, 3/28, at 11:59 PM
 2. We will use class time next week for group work
 3. Ask your questions here: https://northeastern.instructure.com/courses/207607/discussion_topics/2738173
2. Assessments:
 1. ***Both assessments are in class on Tuesday, 4/15***
 2. Programming assessment: Should take 15 minutes and will be based on the questions here:
 3. MSFQ assessment (only for MSFQ students): Should take 45 minutes and based on the 5 required MSFQ courses (standby for a Canvas announcement with more details after I discuss it with the MSFQ program director)

Five-Minute Review

Quantitative value investing is a systematic investment strategy that combines the principles of value investing with data-driven, quantitative techniques. It involves identifying undervalued securities—those trading below their intrinsic value—using predefined, measurable criteria rather than subjective judgment. Common criteria are price/earnings (P/E) and price/book (P/B) ratios. Here, we will use the book-to-market equity ratio (B/M), which is common in the academic literature and quantitative investing.

The key concepts in this topic are creating metrics, forming portfolios, and backtesting the strategy.

Practice

Re-implement the value strategy from the lecture notebook

Read the data

This code reads stock market data from a file called `crsp.csv` and prepares it for analysis. It loads the data into a data frame, makes sure the dates are in the right format, and marks certain letters (A, B, and C) as missing values. Then, it filters the data to include only records from 1965 or later, sorts it by company and date, and adds two new columns: one for market equity (calculated as the number of shares times the stock price, converted to millions), and another to adjust all dates to the end of each month.

```
crsp = (
    pd.read_csv(
        filepath_or_buffer='crsp.csv',      # Read data from crsp.csv file
        parse_dates=['date'],              # Convert 'date' column to datetime format
        na_values=['A', 'B', 'C']         # Treat 'A', 'B', 'C' as missing values (NaN)
    )
    .query('date >= 1965')           # Filter to keep only data from 1965 onward
    .sort_values(['PERMNO', 'date'])   # Sort by company identifier (PERMNO) and date
    .assign(
        ME=lambda x: x['SHROUT'] * x['PRC'].abs() / 1_000, # Calculate market equity (ME) as
        date=lambda x: x['date'] + pd.offsets.MonthEnd(0)     # Adjust dates to the last day of
    )
)

crsp.head()
```

	PERMNO	date	SHRCD	PRC	RET	SHROUT	ME
0	10000	1986-01-31	10	-4.3750	NaN	3680.0000	16.1000
1	10000	1986-02-28	10	-3.2500	-0.2571	3680.0000	11.9600
2	10000	1986-03-31	10	-4.4375	0.3654	3680.0000	16.3300
3	10000	1986-04-30	10	-4.0000	-0.0986	3793.0000	15.1720
4	10000	1986-05-31	10	-3.1094	-0.2227	3793.0000	11.7939

This code brings in financial data from a file named `compustat.csv` and prepares it for analysis. It loads the data into a data frame, formats dates, and filters out any records before 1965. Then, it sorts the data by company, fiscal year, and date. Finally, it removes any duplicate entries for the same company and fiscal year, keeping only the most recent record.

```
compustat = (
    pd.read_csv(
        filepath_or_buffer='compustat.csv', # Read data from compustat.csv file
        parse_dates=['datadate'] # Convert 'datadate' column to datetime format
    )
    .query('datadate >= 1965') # Filter to keep only data from 1965 onward
    .sort_values(['LPERMNO', 'fyear', 'datadate']) # Sort by company identifier (LPERMNO), :
    .drop_duplicates(subset=['LPERMNO', 'fyear'], keep='last') # Keep only the latest record
)
```

`compustat.head()`

	GVKEY	LPERMNO	datadate	fyear	indfmt	consol	popsrc	datafmt	curcd	ceq
165681	13007	10000	1986-10-31	1986.0000	INDL	C	D	STD	USD	0.4180
165586	12994	10001	1986-06-30	1986.0000	INDL	C	D	STD	USD	5.4320
165587	12994	10001	1987-06-30	1987.0000	INDL	C	D	STD	USD	5.3690
165588	12994	10001	1988-06-30	1988.0000	INDL	C	D	STD	USD	5.5120
165589	12994	10001	1989-06-30	1989.0000	INDL	C	D	STD	USD	6.3210

About 64 percent of firm-years have December fiscal-year ends.

```
compustat['datadate'].dt.month.value_counts() / compustat['datadate'].shape[0]
```

```
datadate
12    0.6401
6     0.0723
```

```

9    0.0609
3    0.0518
10   0.0341
1    0.0334
8    0.0203
7    0.0187
5    0.0187
4    0.0179
11   0.0173
2    0.0146
Name: count, dtype: float64

```

Create the interim data frames

This code finds the market value of equity (ME) as of December each year from the `crsp` data frame (to match with book value of equity `ceq` from the `compustat` data frame below). It filters out any rows where market equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last market equity value for each company in each year, giving us the December ME.

```

mve = (
    crsp
    .query('ME > 0')                               # Filter for positive market equity values
    .sort_values(['PERMNO', 'date'])                 # Sort by company identifier (PERMNO) and date
    .groupby(by=['PERMNO', pd.Grouper(key='date', freq='YE-DEC')]) # Group by company and year
    [['ME']]                                         # Select the market equity (ME) column
    .last()                                           # Take the last ME value for each group (December)
)

mve.head()

```

PERMNO	date	M
10000	1986-12-31	1.
	1987-12-31	0.
10001	1986-12-31	6.
	1987-12-31	5.
	1988-12-31	6.

This code finds the book value of equity (`ceq`) as of December each year from the `compustat` data frame. It filters out any rows where common equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last common equity value for each company in each year, giving us the December `ceq`.

```
bve = (
    compustat
    .query('ceq > 0')                                # Filter for positive common equity values
    .sort_values(['LPERMNO', 'datadate'])              # Sort by company identifier (LPERMNO) and date
    .groupby(by=['LPERMNO', pd.Grouper(key='datadate', freq='YE-DEC')]) # Group by company and year
    [['ceq']]                                           # Select the common equity (ceq) column
    .last()                                             # Take the last CEQ value for each group (December)
)

bve.head()
```

LPERMNO	datadate	ceq
10000	1986-12-31	0.
	1986-12-31	5.
10001	1987-12-31	5.
	1988-12-31	5.
	1989-12-31	6.

Combine the data frames and form portfolios

This code shows a simple example of how to form portfolios using the `pd.qcut()` function. It takes a sequence of numbers from 0 to 9 and splits them into 2 equal groups (quantiles), assigning them to either a low or high category. The `labels=False` argument gives us numeric group identifiers (0 or 1), and adding 1 shifts these to 1 or 2, which can represent portfolio numbers.

```
1 + pd.qcut(x=np.arange(10), q=2, labels=False)

array([1, 1, 1, 1, 1, 2, 2, 2, 2])
```

Note

In class, in the following cell, I used multiple cursors to simultaneously type two `.rename_axis(index=['PERMNO', 'Date'])` methods. If you hold down the CTRL key, every mouse click will generate a new cursor. Each cursor will do the same action (e.g., typing, highlighting, and deleting). More [here](#).

This code creates portfolios based on book-to-market (BM) ratios using the `bve` and `mve` data frames. It starts by aligning the company and date indexes of both datasets and merges them, keeping only rows with matching data. Next, it calculates the BM ratio by dividing book value of equity (`ceq`) by market value of equity (`ME`). After resetting the index, it shifts the dates forward by 7 months to time the portfolio formation correctly, giving 6 months between BM calculation and buying stocks at the end of June to receive July returns. Finally, for each date, it sorts the BM ratios into 10 equal groups (deciles) and assigns portfolio numbers from 1 to 10.

```
bm = (
    bve
    .rename_axis(index=['PERMNO', 'Date']) # Standardize index names to PERMNO (company ID)
    .join(
        other=mve.rename_axis(index=['PERMNO', 'Date']), # Combine with market equity data,
        how='inner' # Use inner join to retain only records with matching PERMNO and Date
    )
    .reset_index() # Flatten the index into columns for easier manipulation
    .assign(
        Date=lambda x: x['Date'] + pd.offsets.MonthEnd(7), # Shift dates forward 7 months (from June to July)
        BM=lambda x: x['ceq'] / x['ME'], # Compute book-to-market ratio as ceq / ME
        Portfolio=lambda x: x.groupby('Date')['BM'].transform( # Assign stocks to portfolios
            lambda x: 1 + pd.qcut(x=x, q=10, labels=False) # Divide BM into 10 quantiles (0-10)
        )
    )
)
```

`bm.head()`


	PERMNO	Date	ceq	ME	BM	Portfolio
0	10000	1987-07-31	0.4180	1.9816	0.2109	2
1	10001	1987-07-31	5.4320	6.9370	0.7830	7
2	10001	1988-07-31	5.3690	5.8280	0.9212	7
3	10001	1989-07-31	5.5120	6.3623	0.8664	7
4	10001	1990-07-31	6.3210	10.3477	0.6109	5

This code creates a bar chart showing the average book-to-market (BM) ratio for each portfolio using the `portfolios` data frame. It groups the data by portfolio number (1 to 10), calculates the mean BM ratio for each group, and then plots these averages as bars. As expected, we see that BM rises from portfolio 1 to 10.

```

(
    bm
    .groupby(by=['Portfolio', 'Date']) # Group data by portfolio number and date
    ['BM']                                # Select the book-to-market (BM) column
    .mean()                               # Calculate the mean BM for each portfolio-date combi-
    .groupby(by='Portfolio')               # Group again by portfolio to average across all data
    .mean()                               # Calculate the time-averaged mean BM for each portfo-
    .plot(kind='bar')                     # Plot the results as a bar chart
)
plt.ylabel('Mean B/M for Equal-Weighted Portfolio')
plt.title('Mean B/M for Equal-Weighted B/M Portfolios')
plt.show()

```



This code combines monthly stock returns from the `crsp` data frame with the *most recent* portfolio assignments from the `bm` data frame. The `tolerance` argument tells the merge to only use portfolio assignments within 366 days to avoid stale data.

```
stocks = (
    pd.merge_asof(
        left=crsp.sort_values(['date', 'PERMNO']),
        right=bm.sort_values(['Date', 'PERMNO']),
        left_on='date',
        right_on='Date',
        by='PERMNO',
        tolerance=pd.Timedelta('366d')
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']] # Sort CRSP data by date and company ID
    .dropna() # Sort BM data by date and company ID
    # Match CRSP dates to the closest prior BM date
    # Use BM dates as the reference for the merge
    # Merge on company ID (PERMNO) to link portfolios
    # Limit matches to portfolio assignments within 366 days
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Select key columns: company ID, date, portfolio number, and return
    # Remove rows with missing values (e.g., for companies without assigned portfolios)
    # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Convert 'date' column to 'Date' for consistency
)
```

This code creates equal-weighted portfolios using the `stocks` data frame. It groups the data by portfolio number (`Portfolio`) and date (`Date`), then calculates the mean return (`RET`) for each group. This averaging process gives each stock within a portfolio the same weight, producing equal-weighted portfolio returns.


Note

Here I made `portfolios` a wide data frame (i.e., portfolio numbers in columns) to simplify plots and calculations. This is a difference from class

```
portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading date
    [['RET']] # Isolate the returns column for portfolio-level aggregation
    .mean() # Average returns across stocks within each portfolio-date bin
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's returns
    ['RET'] # Extract the returns data, retaining portfolio numbers
)
```

Finally, we can see that the value portfolio 10 outperforms the growth portfolio 1!

```
portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted B/M Portfolios')
plt.show()
```



Backtest the strategy

We will need the Fama and French (1993) factors for the risk-free rate.

Note

Here I set the Date index to simplify merges/joins below. This is a difference from class

```
ff3 = (
    pd.read_csv(
        name='F-F_Research_Data_Factors',
        data_source='famafrench',
        start='1900'
    )
[0]
.div(100)
.reset_index()
.assign(Date=lambda x: x['Date'].dt.to_timestamp(how='end').dt.normalize()) # convert mon
.set_index('Date')
)
```

C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
 pd.read_csv(
C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
 pd.read_csv(

portfolios_rf = portfolios.sub(ff3['RF'], axis=0)

def CAGR(x, ann_fac=12):
 return (1 + x).prod() ** (ann_fac / x.count()) - 1

def Sharpe(x, ann_fac=np.sqrt(12)):
 return ann_fac * x.mean() / x.std()

def Volatility(x, ann_fac=np.sqrt(12)):
 return ann_fac * x.std()

def Drawdown(x):
 price = x.add(1).cumprod()
 return (price / price.cummax()).min() - 1

The high B/M have higher Sharper ratios and CAGRs than low B/M portfolios. However, the risk relation is not monotonic.

```
import seaborn as sns
```


```
stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)


plt.suptitle('Performance Statistics for Equal-Weighted B/M Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted B/M Portfolios



Also, this outperformance narrows over the past few decades.

```
portfolios.rolling(window=120).apply(CAGR).plot()
plt.ylabel('10-Year CAGR')
plt.title('Rolling 10-Year CAGRs for Equal-Weighted B/M Portfolios')
plt.show()
```



Re-implement the value strategy from the lecture notebook *with value-weighted portfolios*

We need to do two things:

1. Add the beginning of month market values of equity to use as portfolio weights
2. Replace `.mean()` with `np.average()` to calculate weighted means

This code prepares stock data from the `crsp` data frame by adding beginning-of-month market equity (`ME`) and merging it with portfolio assignments from the `bm` data frame. It first selects `PERMNO`, `date`, and `RET` from `crsp`, then merges it with `ME` data shifted forward by 1 month using an inner join to align current returns with next month's market equity. Next, it sorts the CRSP data by `date` and `PERMNO`, and merges it with a subset of `bm` containing only `Date`, `PERMNO`, and `Portfolio`, sorted by `Date` and `PERMNO`, matching each stock's date to the closest prior `Date` in `bm` within a 366-day tolerance. Finally, it keeps `PERMNO`, `date`, `Portfolio`, `RET`, and `ME`, drops missing values, converts `Portfolio` to integers, and renames `date` to `Date` for consistency.

```

# Add beginning-of-month market equity (ME) to CRSP data
crsp_w_bom_me = (
    crsp
    [['PERMNO', 'date', 'RET']] # Select stock ID, date, and returns from CRSP
    .merge(
        right=crsp[['PERMNO', 'date', 'ME']].assign( # Select stock ID, date, and ME, shift
            date=lambda x: x['date'] + pd.offsets.MonthEnd(1)), # Inner join to align current month with next
        how='inner'
    )
)

# Merge CRSP data with book-to-market portfolio assignments
stocks_vw = (
    pd.merge_asof(
        left=crsp_w_bom_me.sort_values(['date', 'PERMNO']), # Sort CRSP with ME by date and stock ID
        right=bm.sort_values(['Date', 'PERMNO'])[['Date', 'PERMNO', 'Portfolio']], # Select BM data
        left_on='date', # Match CRSP dates to the closest BM dates
        right_on='Date', # Use BM Date as the reference
        by='PERMNO', # Join on stock ID (PERMNO) to align dates
        tolerance=pd.Timedelta('366d') # Limit matches to BM data within a year
    )
    [['PERMNO', 'date', 'Portfolio', 'RET', 'ME']] # Select essential columns: stock ID, date, portfolio, return, and market equity
    .dropna() # Remove rows with missing values
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Rename 'date' to 'Date' to align with CRSP
)

```

This code creates value-weighted portfolios using the `stocks` data frame and reshapes the results into a wide format. It groups the data by `Portfolio` (1 to 10) and `Date`, then calculates the value-weighted average return (`RET`) for each group, using market equity (`ME`) as weights. Next, it pivots the data so each portfolio's returns become separate columns, with `Date` as the index.

```


# Calculate value-weighted portfolio returns and reshape into wide format
portfolios_vw = (
    stocks_vw
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and date
    .apply(lambda x: np.average(a=x['RET'], # Compute value-weighted average returns using
                                weights=x['ME']), # Market Equity as weights
          include_groups=False) # Exclude grouping columns (Portfolio, Date)
    .unstack('Portfolio') # Pivot data so each portfolio's returns become columns
)

```

We can briefly compare the equal- and value-weighted portfolios. *We see that the value effect is concentrated in small stocks because the equal-weighted portfolios returns are much larger than the value-weighted portfolio returns.*

```
df = pd.concat(
    objs=[portfolios.add(1).prod(), portfolios_vw.add(1).prod()],
    keys=['Equal', 'Value (ME)'],
    names=['Portfolio Weights'],
    axis=1
)
```

```
df.plot(kind='bar')
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in B/M Portfolios')
plt.show()
```



Estimate the α s of the equal-weighted and value-weighted portfolios

i Note

Here `portfolios` is wide, so I have to `.stack()` it before the join. This is a difference from class

```
portfolios_ff3 = (
    portfolios
    .stack(future_stack=True)
    .to_frame('RET')
    .join(ff3, how='inner')
)
```

```
import statsmodels.formula.api as smf
```

The value portfolio 10 has an α of 82 basis points! Annualized, this is $12 \times 0.0082 = 0.0984$ or almost 10% of returns not explained by market risk.

```
smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query('(Portfolio == 10)').ff3)
```

Dep. Variable:	I(RET - RF)	R-squared:	0.506			
Model:	OLS	Adj. R-squared:	0.505			
Method:	Least Squares	F-statistic:	717.6			
Date:	Sat, 22 Mar 2025	Prob (F-statistic):	2.34e-109			
Time:	06:49:48	Log-Likelihood:	1116.4			
No. Observations:	702	AIC:	-2229.			
Df Residuals:	700	BIC:	-2220.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.0082	0.002	4.370	0.000	0.005	0.012
Q("Mkt-RF")	1.0948	0.041	26.788	0.000	1.015	1.175
Omnibus:	268.363	Durbin-Watson:	1.740			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1787.450			
Skew:	1.560	Prob(JB):	0.00			
Kurtosis:	10.167	Cond. No.	21.9			

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can plot these as (and confidence intervals) with list comprehensions!

```
fits = [
    smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query(f'(Portfolio == {p})')
        for p in range(1, 11)
    ]

df = (
    pd.DataFrame(
        data={
            'Alpha': [f.params['Intercept'] for f in fits],
            'Sterr': [f.bse['Intercept'] for f in fits],
            'Portfolio': range(1, 11)
        }
    )
    .set_index('Portfolio')
)

plt.bar(
    x=df.index,
    height=df['Alpha'],
    yerr=df['Sterr'].mul(1.96)
)
plt.ylabel('Alpha')
plt.title('Monthly Alphas for Equal-Weighted B/M Portfolios\n (black bars indicate 95% confi')
plt.show()
```



Implement a momentum strategy

Form deciles on the 11-month returns from months $t - 12$ to month $t - 2$.

This code calculates 11-month momentum returns (`RET_11`) from 2 to 12 months prior using the `crsp` data frame and assigns portfolios based on those returns. It sets a multi-index with `PERMNO` and `date`, selects the `RET` column, and pivots to a wide format with stocks as columns. After sorting by date, it computes compounded returns over an 11-month rolling window using log transformations (`log1p` and `expm1`), then reshapes back to long format. It drops rows with missing `RET_11` values, creates a `DataFrame`, and resets the index. Finally, it assigns stocks to decile portfolios (1 to 10) based on `RET_11` for each date and shifts the `date` forward by 2 months to reflect the momentum strategy lag. That is, we receive returns in month t based on returns from month $t - 12$ through $t - 2$.

```
# Calculate 2-12 month momentum returns and assign portfolios
mom_02_12 = (
    crsp
```

```

.set_index(['PERMNO', 'date'])          # Set multi-index with stock ID (PERMNO) and date
['RET']                                # Select the monthly returns column
.unstack('PERMNO')                      # Pivot to wide format, with PERMNO as column
.sort_index()                           # Sort dates chronologically for consistent results
.pipe(np.log1p)                         # Convert returns to log scale (log(1 + r))
.rolling(11)                            # Apply an 11-month rolling window to prior returns
.sum()                                  # Sum log returns over the 11-month window
.pipe(np.expm1)                         # Convert summed log returns back to compound returns
.stack(future_stack=True)               # Reshape to long format, aligning with pandas
.dropna()                               # Remove rows with NaN from incomplete windows
.to_frame('RET_11')                     # Create a DataFrame, naming the momentum return
.reset_index()                          # Move PERMNO and date back to columns from index
.assign(
    Portfolio=lambda x: x.groupby('date')['RET_11'].transform( # Assign stocks to decile portfolios
        lambda x: 1 + pd.qcut(x=x, q=10, labels=False)),      # Split RET_11 into 10 quartiles
    date=lambda x: x['date'] + pd.offsets.MonthEnd(2))       # Shift dates forward 2 months
)
)
)

```

```
mom_02_12.head()
```

	date	PERMNO	RET_11	Portfolio
0	1966-01-31	10006	0.1508	5
1	1966-01-31	10014	0.2000	6
2	1966-01-31	10030	0.1489	5
3	1966-01-31	10057	0.4666	8
4	1966-01-31	10102	0.3503	7

```


stocks = (
    pd.merge(
        left=crsp,
        right=mom_02_12,
        on=['PERMNO', 'date'],
        how='inner'
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']]
    .dropna()
    .rename(columns={'date': 'Date'})
)

```

```

portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading da
    [['RET']]
    .mean() # Isolate the returns column for portfolio-level aggregat
    # Average returns across stocks within each portfolio-d
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's
    ['RET'] # Extract the returns data, retaining portfolio numbers
)
portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted Momentum Portfolios')
plt.show()

```



```
portfolios_rf = portfolios.sub(ff3['RF'], axis=0)


stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)

plt.suptitle('Performance Statistics for Equal-Weighted Momentum Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted Momentum Portfolios



Herron Topic 3 - Practice - Sec 03

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Projects:
 1. Project 1:
 1. I am still grading
 2. I plan to finish grading by Friday
 2. Project 2:
 1. Due on Friday, 3/28, at 11:59 PM
 2. We will use class time next week for group work
 3. Ask your questions here: https://northeastern.instructure.com/courses/207607/discussion_topics/2738173
2. Assessments:
 1. ***Both assessments are in class on Tuesday, 4/15***
 2. Programming assessment: Should take 15 minutes and will be based on the questions here:
 3. MSFQ assessment (only for MSFQ students): Should take 45 minutes and based on the 5 required MSFQ courses (standby for a Canvas announcement with more details after I discuss it with the MSFQ program director)

Five-Minute Review

Quantitative value investing is a systematic investment strategy that combines the principles of value investing with data-driven, quantitative techniques. It involves identifying undervalued securities—those trading below their intrinsic value—using predefined, measurable criteria rather than subjective judgment. Common criteria are price/earnings (P/E) and price/book (P/B) ratios. Here, we will use the book-to-market equity ratio (B/M), which is common in the academic literature and quantitative investing.

The key concepts in this topic are creating metrics, forming portfolios, and backtesting the strategy.

Practice

Re-implement the value strategy from the lecture notebook

Read the data

This code reads stock market data from a file called `crsp.csv` and prepares it for analysis. It loads the data into a data frame, makes sure the dates are in the right format, and marks certain letters (A, B, and C) as missing values. Then, it filters the data to include only records from 1965 or later, sorts it by company and date, and adds two new columns: one for market equity (calculated as the number of shares times the stock price, converted to millions), and another to adjust all dates to the end of each month.

```
crsp = (
    pd.read_csv(
        filepath_or_buffer='crsp.csv',      # Read data from crsp.csv file
        parse_dates=['date'],              # Convert 'date' column to datetime format
        na_values=['A', 'B', 'C']         # Treat 'A', 'B', 'C' as missing values (NaN)
    )
    .query('date >= 1965')           # Filter to keep only data from 1965 onward
    .sort_values(['PERMNO', 'date'])   # Sort by company identifier (PERMNO) and date
    .assign(
        ME=lambda x: x['SHROUT'] * x['PRC'].abs() / 1_000, # Calculate market equity (ME) as
        date=lambda x: x['date'] + pd.offsets.MonthEnd(0)     # Adjust dates to the last day of
    )
)

crsp.head()
```

	PERMNO	date	SHRCD	PRC	RET	SHROUT	ME
0	10000	1986-01-31	10	-4.3750	NaN	3680.0000	16.1000
1	10000	1986-02-28	10	-3.2500	-0.2571	3680.0000	11.9600
2	10000	1986-03-31	10	-4.4375	0.3654	3680.0000	16.3300
3	10000	1986-04-30	10	-4.0000	-0.0986	3793.0000	15.1720
4	10000	1986-05-31	10	-3.1094	-0.2227	3793.0000	11.7939

This code brings in financial data from a file named `compustat.csv` and prepares it for analysis. It loads the data into a data frame, formats dates, and filters out any records before 1965. Then, it sorts the data by company, fiscal year, and date. Finally, it removes any duplicate entries for the same company and fiscal year, keeping only the most recent record.

```
compustat = (
    pd.read_csv(
        filepath_or_buffer='compustat.csv', # Read data from compustat.csv file
        parse_dates=['datadate'] # Convert 'datadate' column to datetime format
    )
    .query('datadate >= 1965') # Filter to keep only data from 1965 onward
    .sort_values(['LPERMNO', 'fyear', 'datadate']) # Sort by company identifier (LPERMNO), :
    .drop_duplicates(subset=['LPERMNO', 'fyear'], keep='last') # Keep only the latest record
)
```

`compustat.head()`

	GVKEY	LPERMNO	datadate	fyear	indfmt	consol	popsrc	datafmt	curcd	ceq
165681	13007	10000	1986-10-31	1986.0000	INDL	C	D	STD	USD	0.4180
165586	12994	10001	1986-06-30	1986.0000	INDL	C	D	STD	USD	5.4320
165587	12994	10001	1987-06-30	1987.0000	INDL	C	D	STD	USD	5.3690
165588	12994	10001	1988-06-30	1988.0000	INDL	C	D	STD	USD	5.5120
165589	12994	10001	1989-06-30	1989.0000	INDL	C	D	STD	USD	6.3210

About 64 percent of firm-years have December fiscal-year ends.

```
compustat['datadate'].dt.month.value_counts() / compustat['datadate'].shape[0]
```

```
datadate
12    0.6401
6     0.0723
```

```

9    0.0609
3    0.0518
10   0.0341
1    0.0334
8    0.0203
7    0.0187
5    0.0187
4    0.0179
11   0.0173
2    0.0146
Name: count, dtype: float64

```

Create the interim data frames

This code finds the market value of equity (ME) as of December each year from the `crsp` data frame (to match with book value of equity `ceq` from the `compustat` data frame below). It filters out any rows where market equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last market equity value for each company in each year, giving us the December ME.

```

mve = (
    crsp
    .query('ME > 0')                               # Filter for positive market equity values
    .sort_values(['PERMNO', 'date'])                 # Sort by company identifier (PERMNO) and date
    .groupby(by=['PERMNO', pd.Grouper(key='date', freq='YE-DEC')]) # Group by company and year
    [['ME']]                                         # Select the market equity (ME) column
    .last()                                           # Take the last ME value for each group (December)
)

mve.head()

```

PERMNO	date	M
10000	1986-12-31	1.
	1987-12-31	0.
10001	1986-12-31	6.
	1987-12-31	5.
	1988-12-31	6.

This code finds the book value of equity (`ceq`) as of December each year from the `compustat` data frame. It filters out any rows where common equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last common equity value for each company in each year, giving us the December `ceq`.

```
bve = (
    compustat
    .query('ceq > 0')                                # Filter for positive common equity values
    .sort_values(['LPERMNO', 'datadate'])              # Sort by company identifier (LPERMNO) and date
    .groupby(by=['LPERMNO', pd.Grouper(key='datadate', freq='YE-DEC')]) # Group by company and year
    [['ceq']]                                           # Select the common equity (ceq) column
    .last()                                             # Take the last CEQ value for each group (December)
)

bve.head()
```

LPERMNO	datadate	ceq
10000	1986-12-31	0.
	1986-12-31	5.
10001	1987-12-31	5.
	1988-12-31	5.
	1989-12-31	6.

Combine the data frames and form portfolios

This code shows a simple example of how to form portfolios using the `pd.qcut()` function. It takes a sequence of numbers from 0 to 9 and splits them into 2 equal groups (quantiles), assigning them to either a low or high category. The `labels=False` argument gives us numeric group identifiers (0 or 1), and adding 1 shifts these to 1 or 2, which can represent portfolio numbers.

```
1 + pd.qcut(x=np.arange(10), q=2, labels=False)

array([1, 1, 1, 1, 1, 2, 2, 2, 2])
```

i Note

In class, in the following cell, I used multiple cursors to simultaneously type two `.rename_axis(index=['PERMNO', 'Date'])` methods. If you hold down the CTRL key, every mouse click will generate a new cursor. Each cursor will do the same action (e.g., typing, highlighting, and deleting). More [here](#).

This code creates portfolios based on book-to-market (BM) ratios using the `bve` and `mve` data frames. It starts by aligning the company and date indexes of both datasets and merges them, keeping only rows with matching data. Next, it calculates the BM ratio by dividing book value of equity (`ceq`) by market value of equity (`ME`). After resetting the index, it shifts the dates forward by 7 months to time the portfolio formation correctly, giving 6 months between BM calculation and buying stocks at the end of June to receive July returns. Finally, for each date, it sorts the BM ratios into 10 equal groups (deciles) and assigns portfolio numbers from 1 to 10.


```
bm = (
    bve
    .rename_axis(index=['PERMNO', 'Date']) # Standardize index names to PERMNO (company ID)
    .join(
        other=mve.rename_axis(index=['PERMNO', 'Date']), # Combine with market equity data,
        how='inner' # Use inner join to retain only records with matching PERMNO and Date
    )
    .reset_index() # Flatten the index into columns for easier manipulation
    .assign(
        Date=lambda x: x['Date'] + pd.offsets.MonthEnd(7), # Shift dates forward 7 months (from June to July)
        BM=lambda x: x['ceq'] / x['ME'], # Compute book-to-market ratio as ceq / ME
        Portfolio=lambda x: x.groupby('Date')['BM'].transform( # Assign stocks to portfolios
            lambda x: 1 + pd.qcut(x=x, q=10, labels=False) # Divide BM into 10 quantiles (0-10)
        )
    )
)
```

```
bm.head()
```

	PERMNO	Date	ceq	ME	BM	Portfolio
0	10000	1987-07-31	0.4180	1.9816	0.2109	2
1	10001	1987-07-31	5.4320	6.9370	0.7830	7
2	10001	1988-07-31	5.3690	5.8280	0.9212	7
3	10001	1989-07-31	5.5120	6.3623	0.8664	7
4	10001	1990-07-31	6.3210	10.3477	0.6109	5

This code creates a bar chart showing the average book-to-market (BM) ratio for each portfolio using the `portfolios` data frame. It groups the data by portfolio number (1 to 10), calculates the mean BM ratio for each group, and then plots these averages as bars. As expected, we see that BM rises from portfolio 1 to 10.

```
(  
    bm  
    .groupby(by=['Portfolio', 'Date']) # Group data by portfolio number and date  
    ['BM']  
    .mean() # Select the book-to-market (BM) column  
    .groupby(by='Portfolio') # Calculate the mean BM for each portfolio-date combi  
    .mean()  
    .plot(kind='bar') # Group again by portfolio to average across all date  
    # Calculate the time-averaged mean BM for each portfo  
    # Plot the results as a bar chart  
)  
plt.ylabel('Mean B/M for Equal-Weighted Portfolio')  
plt.title('Mean B/M for Equal-Weighted B/M Portfolios')  
plt.show()
```



This code combines monthly stock returns from the `crsp` data frame with the *most recent* portfolio assignments from the `bm` data frame. The `tolerance` argument tells the merge to only use portfolio assignments within 366 days to avoid stale data.

```
stocks = (
    pd.merge_asof(
        left=crsp.sort_values(['date', 'PERMNO']),
        right=bm.sort_values(['Date', 'PERMNO']),
        left_on='date',
        right_on='Date',
        by='PERMNO',
        tolerance=pd.Timedelta('366d')
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']] # Sort CRSP data by date and company ID
    .dropna() # Sort BM data by date and company ID
    # Match CRSP dates to the closest prior BM date
    # Use BM dates as the reference for the merge
    # Merge on company ID (PERMNO) to link portfolios
    # Limit matches to portfolio assignments within 366 days
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Select key columns: company ID, date, portfolio number, and return
    # Remove rows with missing values (e.g., for companies without recent assignments)
    # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Convert 'date' column to 'Date' for consistency
)
```

This code creates equal-weighted portfolios using the `stocks` data frame. It groups the data by portfolio number (`Portfolio`) and date (`Date`), then calculates the mean return (`RET`) for each group. This averaging process gives each stock within a portfolio the same weight, producing equal-weighted portfolio returns.


Note

Here I made `portfolios` a wide data frame (i.e., portfolio numbers in columns) to simplify plots and calculations. This is a difference from class

```
portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading date
    [['RET']] # Isolate the returns column for portfolio-level aggregation
    .mean() # Average returns across stocks within each portfolio-date bin
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's returns
    ['RET'] # Extract the returns data, retaining portfolio numbers
)
```

Finally, we can see that the value portfolio 10 outperforms the growth portfolio 1!

```
portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted B/M Portfolios')
plt.show()
```



Backtest the strategy

We will need the Fama and French (1993) factors for the risk-free rate.

Note

Here I set the Date index to simplify merges/joins below. This is a difference from class

```
ff3 = (
    pdr.DataReader(
        name='F-F_Research_Data_Factors',
        data_source='famafrench',
        start='1900'
    )
[0]
.div(100)
.reset_index()
.assign(Date=lambda x: x['Date'].dt.to_timestamp(how='end').dt.normalize()) # convert mon
.set_index('Date')
)
```

C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
 pdr.DataReader(
C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
 pdr.DataReader(

portfolios_rf = portfolios.sub(ff3['RF'], axis=0)

def CAGR(x, ann_fac=12):
 return (1 + x).prod() ** (ann_fac / x.count()) - 1

def Sharpe(x, ann_fac=np.sqrt(12)):
 return ann_fac * x.mean() / x.std()

def Volatility(x, ann_fac=np.sqrt(12)):
 return ann_fac * x.std()

def Drawdown(x):
 price = x.add(1).cumprod()
 return (price / price.cummax()).min() - 1

The high B/M have higher Sharper ratios and CAGRs than low B/M portfolios. However, the risk relation is not monotonic.

```
import seaborn as sns
```


```
stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)


plt.suptitle('Performance Statistics for Equal-Weighted B/M Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted B/M Portfolios



Also, this outperformance narrows over the past few decades.

```
portfolios.rolling(window=120).apply(CAGR).plot()
plt.ylabel('10-Year CAGR')
plt.title('Rolling 10-Year CAGRs for Equal-Weighted B/M Portfolios')
plt.show()
```



Re-implement the value strategy from the lecture notebook *with value-weighted portfolios*

We need to do two things:

1. Add the beginning of month market values of equity to use as portfolio weights
2. Replace `.mean()` with `np.average()` to calculate weighted means

This code prepares stock data from the `crsp` data frame by adding beginning-of-month market equity (`ME`) and merging it with portfolio assignments from the `bm` data frame. It first selects `PERMNO`, `date`, and `RET` from `crsp`, then merges it with `ME` data shifted forward by 1 month using an inner join to align current returns with next month's market equity. Next, it sorts the CRSP data by `date` and `PERMNO`, and merges it with a subset of `bm` containing only `Date`, `PERMNO`, and `Portfolio`, sorted by `Date` and `PERMNO`, matching each stock's date to the closest prior `Date` in `bm` within a 366-day tolerance. Finally, it keeps `PERMNO`, `date`, `Portfolio`, `RET`, and `ME`, drops missing values, converts `Portfolio` to integers, and renames `date` to `Date` for consistency.

```

# Add beginning-of-month market equity (ME) to CRSP data
crsp_w_bom_me = (
    crsp
    [['PERMNO', 'date', 'RET']] # Select stock ID, date, and returns from CRSP
    .merge(
        right=crsp[['PERMNO', 'date', 'ME']].assign( # Select stock ID, date, and ME, shift
            date=lambda x: x['date'] + pd.offsets.MonthEnd(1)), # Inner join to align current month with next
        how='inner'
    )
)

# Merge CRSP data with book-to-market portfolio assignments
stocks_vw = (
    pd.merge_asof(
        left=crsp_w_bom_me.sort_values(['date', 'PERMNO']), # Sort CRSP with ME by date and stock ID
        right=bm.sort_values(['Date', 'PERMNO'])[['Date', 'PERMNO', 'Portfolio']], # Select BM data
        left_on='date', # Match CRSP dates to the closest BM dates
        right_on='Date', # Use BM Date as the reference
        by='PERMNO', # Join on stock ID (PERMNO) to align dates
        tolerance=pd.Timedelta('366d') # Limit matches to BM data within a year
    )
    [['PERMNO', 'date', 'Portfolio', 'RET', 'ME']] # Select essential columns: stock ID, date, portfolio, return, and ME
    .dropna() # Remove rows with missing values
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Rename 'date' to 'Date' to align with CRSP
)

```

This code creates value-weighted portfolios using the `stocks` data frame and reshapes the results into a wide format. It groups the data by `Portfolio` (1 to 10) and `Date`, then calculates the value-weighted average return (`RET`) for each group, using market equity (`ME`) as weights. Next, it pivots the data so each portfolio's returns become separate columns, with `Date` as the index.

```


# Calculate value-weighted portfolio returns and reshape into wide format
portfolios_vw = (
    stocks_vw
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and date
    .apply(lambda x: np.average(a=x['RET'], # Compute value-weighted average returns using
                                weights=x['ME']), # Market Equity as weights
          include_groups=False) # Exclude grouping columns (Portfolio, Date)
    .unstack('Portfolio') # Pivot data so each portfolio's returns become columns
)

```

We can briefly compare the equal- and value-weighted portfolios. *We see that the value effect is concentrated in small stocks because the equal-weighted portfolios returns are much larger than the value-weighted portfolio returns.*

```
df = pd.concat(
    objs=[portfolios.add(1).prod(), portfolios_vw.add(1).prod()],
    keys=['Equal', 'Value (ME)'],
    names=['Portfolio Weights'],
    axis=1
)

df.plot(kind='bar')
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in B/M Portfolios')
plt.show()
```



Estimate the α s of the equal-weighted and value-weighted portfolios

i Note

Here `portfolios` is wide, so I have to `.stack()` it before the join. This is a difference from class

```
portfolios_ff3 = (
    portfolios
    .stack(future_stack=True)
    .to_frame('RET')
    .join(ff3, how='inner')
)
```

```
import statsmodels.formula.api as smf
```

The value portfolio 10 has an α of 82 basis points! Annualized, this is $12 \times 0.0082 = 0.0984$ or almost 10% of returns not explained by market risk.

```
smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query('(Portfolio == 10)').ff3)
```

Dep. Variable:	I(RET - RF)	R-squared:	0.506			
Model:	OLS	Adj. R-squared:	0.505			
Method:	Least Squares	F-statistic:	717.6			
Date:	Sat, 22 Mar 2025	Prob (F-statistic):	2.34e-109			
Time:	06:49:48	Log-Likelihood:	1116.4			
No. Observations:	702	AIC:	-2229.			
Df Residuals:	700	BIC:	-2220.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.0082	0.002	4.370	0.000	0.005	0.012
Q("Mkt-RF")	1.0948	0.041	26.788	0.000	1.015	1.175
Omnibus:	268.363	Durbin-Watson:	1.740			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1787.450			
Skew:	1.560	Prob(JB):	0.00			
Kurtosis:	10.167	Cond. No.	21.9			

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can plot these as (and confidence intervals) with list comprehensions!

```
fits = [
    smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query(f'(Portfolio == {p})')
        for p in range(1, 11)
    ]

df = (
    pd.DataFrame(
        data={
            'Alpha': [f.params['Intercept'] for f in fits],
            'Sterr': [f.bse['Intercept'] for f in fits],
            'Portfolio': range(1, 11)
        }
    )
    .set_index('Portfolio')
)

plt.bar(
    x=df.index,
    height=df['Alpha'],
    yerr=df['Sterr'].mul(1.96)
)
plt.ylabel('Alpha')
plt.title('Monthly Alphas for Equal-Weighted B/M Portfolios\n (black bars indicate 95% confi')
plt.show()
```



Implement a momentum strategy

Form deciles on the 11-month returns from months $t - 12$ to month $t - 2$.

This code calculates 11-month momentum returns (`RET_11`) from 2 to 12 months prior using the `crsp` data frame and assigns portfolios based on those returns. It sets a multi-index with `PERMNO` and `date`, selects the `RET` column, and pivots to a wide format with stocks as columns. After sorting by date, it computes compounded returns over an 11-month rolling window using log transformations (`log1p` and `expm1`), then reshapes back to long format. It drops rows with missing `RET_11` values, creates a `DataFrame`, and resets the index. Finally, it assigns stocks to decile portfolios (1 to 10) based on `RET_11` for each date and shifts the `date` forward by 2 months to reflect the momentum strategy lag. That is, we receive returns in month t based on returns from month $t - 12$ through $t - 2$.

```
# Calculate 2-12 month momentum returns and assign portfolios
mom_02_12 = (
    crsp
```

```

.set_index(['PERMNO', 'date'])          # Set multi-index with stock ID (PERMNO) and date
['RET']                                # Select the monthly returns column
.unstack('PERMNO')                      # Pivot to wide format, with PERMNO as column
.sort_index()                           # Sort dates chronologically for consistent results
.pipe(np.log1p)                         # Convert returns to log scale (log(1 + r))
.rolling(11)                            # Apply an 11-month rolling window to prior returns
.sum()                                  # Sum log returns over the 11-month window
.pipe(np.expm1)                         # Convert summed log returns back to compound returns
.stack(future_stack=True)               # Reshape to long format, aligning with pandas
.dropna()                               # Remove rows with NaN from incomplete windows
.to_frame('RET_11')                     # Create a DataFrame, naming the momentum return
.reset_index()                          # Move PERMNO and date back to columns from index
.assign(
    Portfolio=lambda x: x.groupby('date')['RET_11'].transform( # Assign stocks to decile portfolios
        lambda x: 1 + pd.qcut(x=x, q=10, labels=False)),      # Split RET_11 into 10 quartiles
    date=lambda x: x['date'] + pd.offsets.MonthEnd(2))       # Shift dates forward 2 months
)
)
)

```

```
mom_02_12.head()
```

	date	PERMNO	RET_11	Portfolio
0	1966-01-31	10006	0.1508	5
1	1966-01-31	10014	0.2000	6
2	1966-01-31	10030	0.1489	5
3	1966-01-31	10057	0.4666	8
4	1966-01-31	10102	0.3503	7

```


stocks = (
    pd.merge(
        left=crsp,
        right=mom_02_12,
        on=['PERMNO', 'date'],
        how='inner'
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']]
    .dropna()
    .rename(columns={'date': 'Date'})
)

```

```
portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading da
    [['RET']]
    .mean() # Isolate the returns column for portfolio-level aggregat
    # Average returns across stocks within each portfolio-d
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's
    ['RET'] # Extract the returns data, retaining portfolio numbers
)

portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted Momentum Portfolios')
plt.show()
```

Values of \$1 Investments in Equal-Weighted Momentum Portfolios



```
portfolios_rf = portfolios.sub(ff3['RF'], axis=0)


stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)

plt.suptitle('Performance Statistics for Equal-Weighted Momentum Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted Momentum Portfolios



Herron Topic 3 - Practice - Sec 04

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

1. Projects:
 1. Project 1:
 1. I am still grading
 2. I plan to finish grading by Friday
 2. Project 2:
 1. Due on Friday, 3/28, at 11:59 PM
 2. We will use class time next week for group work
 3. Ask your questions here: https://northeastern.instructure.com/courses/207607/discussion_topics/2738173
2. Assessments:
 1. ***Both assessments are in class on Tuesday, 4/15***
 2. Programming assessment: Should take 15 minutes and will be based on the questions here:
 3. MSFQ assessment (only for MSFQ students): Should take 45 minutes and based on the 5 required MSFQ courses (standby for a Canvas announcement with more details after I discuss it with the MSFQ program director)

Five-Minute Review

Quantitative value investing is a systematic investment strategy that combines the principles of value investing with data-driven, quantitative techniques. It involves identifying undervalued securities—those trading below their intrinsic value—using predefined, measurable criteria rather than subjective judgment. Common criteria are price/earnings (P/E) and price/book (P/B) ratios. Here, we will use the book-to-market equity ratio (B/M), which is common in the academic literature and quantitative investing.

The key concepts in this topic are creating metrics, forming portfolios, and backtesting the strategy.

Practice

Re-implement the value strategy from the lecture notebook

Read the data

This code reads stock market data from a file called `crsp.csv` and prepares it for analysis. It loads the data into a data frame, makes sure the dates are in the right format, and marks certain letters (A, B, and C) as missing values. Then, it filters the data to include only records from 1965 or later, sorts it by company and date, and adds two new columns: one for market equity (calculated as the number of shares times the stock price, converted to millions), and another to adjust all dates to the end of each month.

```
crsp = (
    pd.read_csv(
        filepath_or_buffer='crsp.csv',      # Read data from crsp.csv file
        parse_dates=['date'],              # Convert 'date' column to datetime format
        na_values=['A', 'B', 'C']         # Treat 'A', 'B', 'C' as missing values (NaN)
    )
    .query('date >= 1965')           # Filter to keep only data from 1965 onward
    .sort_values(['PERMNO', 'date'])   # Sort by company identifier (PERMNO) and date
    .assign(
        ME=lambda x: x['SHROUT'] * x['PRC'].abs() / 1_000, # Calculate market equity (ME) as
        date=lambda x: x['date'] + pd.offsets.MonthEnd(0)     # Adjust dates to the last day of
    )
)

crsp.head()
```

	PERMNO	date	SHRCD	PRC	RET	SHROUT	ME
0	10000	1986-01-31	10	-4.3750	NaN	3680.0000	16.1000
1	10000	1986-02-28	10	-3.2500	-0.2571	3680.0000	11.9600
2	10000	1986-03-31	10	-4.4375	0.3654	3680.0000	16.3300
3	10000	1986-04-30	10	-4.0000	-0.0986	3793.0000	15.1720
4	10000	1986-05-31	10	-3.1094	-0.2227	3793.0000	11.7939

This code brings in financial data from a file named `compustat.csv` and prepares it for analysis. It loads the data into a data frame, formats dates, and filters out any records before 1965. Then, it sorts the data by company, fiscal year, and date. Finally, it removes any duplicate entries for the same company and fiscal year, keeping only the most recent record.

```
compustat = (
    pd.read_csv(
        filepath_or_buffer='compustat.csv', # Read data from compustat.csv file
        parse_dates=['datadate'] # Convert 'datadate' column to datetime format
    )
    .query('datadate >= 1965') # Filter to keep only data from 1965 onward
    .sort_values(['LPERMNO', 'fyear', 'datadate']) # Sort by company identifier (LPERMNO), :
    .drop_duplicates(subset=['LPERMNO', 'fyear'], keep='last') # Keep only the latest record
)
```

`compustat.head()`

	GVKEY	LPERMNO	datadate	fyear	indfmt	consol	popsrc	datafmt	curcd	ceq
165681	13007	10000	1986-10-31	1986.0000	INDL	C	D	STD	USD	0.4180
165586	12994	10001	1986-06-30	1986.0000	INDL	C	D	STD	USD	5.4320
165587	12994	10001	1987-06-30	1987.0000	INDL	C	D	STD	USD	5.3690
165588	12994	10001	1988-06-30	1988.0000	INDL	C	D	STD	USD	5.5120
165589	12994	10001	1989-06-30	1989.0000	INDL	C	D	STD	USD	6.3210

About 64 percent of firm-years have December fiscal-year ends.

```
compustat['datadate'].dt.month.value_counts() / compustat['datadate'].shape[0]
```

```
datadate
12    0.6401
6     0.0723
```

```

9    0.0609
3    0.0518
10   0.0341
1    0.0334
8    0.0203
7    0.0187
5    0.0187
4    0.0179
11   0.0173
2    0.0146
Name: count, dtype: float64

```

Create the interim data frames

This code finds the market value of equity (ME) as of December each year from the `crsp` data frame (to match with book value of equity `ceq` from the `compustat` data frame below). It filters out any rows where market equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last market equity value for each company in each year, giving us the December ME.

```

mve = (
    crsp
    .query('ME > 0')                               # Filter for positive market equity values
    .sort_values(['PERMNO', 'date'])                 # Sort by company identifier (PERMNO) and date
    .groupby(by=['PERMNO', pd.Grouper(key='date', freq='YE-DEC')]) # Group by company and year
    [['ME']]                                         # Select the market equity (ME) column
    .last()                                           # Take the last ME value for each group (December)
)

mve.head()

```

PERMNO	date	M
10000	1986-12-31	1.
	1987-12-31	0.
10001	1986-12-31	6.
	1987-12-31	5.
	1988-12-31	6.

This code finds the book value of equity (`ceq`) as of December each year from the `compustat` data frame. It filters out any rows where common equity is zero or negative, then sorts the data by company and date. Next, it groups the data by company and year, focusing on December values. Finally, it selects the last common equity value for each company in each year, giving us the December `ceq`.

```
bve = (
    compustat
    .query('ceq > 0')                               # Filter for positive common equity values
    .sort_values(['LPERMNO', 'datadate'])            # Sort by company identifier (LPERMNO) and date
    .groupby(by=['LPERMNO', pd.Grouper(key='datadate', freq='YE-DEC')]) # Group by company and year
    [['ceq']]                                         # Select the common equity (ceq) column
    .last()                                           # Take the last CEQ value for each group (December)
)

bve.head()
```

LPERMNO	datadate	ceq
10000	1986-12-31	0.
	1986-12-31	5.
10001	1987-12-31	5.
	1988-12-31	5.
	1989-12-31	6.

Combine the data frames and form portfolios

This code shows a simple example of how to form portfolios using the `pd.qcut()` function. It takes a sequence of numbers from 0 to 9 and splits them into 2 equal groups (quantiles), assigning them to either a low or high category. The `labels=False` argument gives us numeric group identifiers (0 or 1), and adding 1 shifts these to 1 or 2, which can represent portfolio numbers.

```
1 + pd.qcut(x=np.arange(10), q=2, labels=False)

array([1, 1, 1, 1, 1, 2, 2, 2, 2])
```

Note

In class, in the following cell, I used multiple cursors to simultaneously type two `.rename_axis(index=['PERMNO', 'Date'])` methods. If you hold down the CTRL key, every mouse click will generate a new cursor. Each cursor will do the same action (e.g., typing, highlighting, and deleting). More [here](#).

This code creates portfolios based on book-to-market (BM) ratios using the `bve` and `mve` data frames. It starts by aligning the company and date indexes of both datasets and merges them, keeping only rows with matching data. Next, it calculates the BM ratio by dividing book value of equity (`ceq`) by market value of equity (`ME`). After resetting the index, it shifts the dates forward by 7 months to time the portfolio formation correctly, giving 6 months between BM calculation and buying stocks at the end of June to receive July returns. Finally, for each date, it sorts the BM ratios into 10 equal groups (deciles) and assigns portfolio numbers from 1 to 10.


```
bm = (
    bve
    .rename_axis(index=['PERMNO', 'Date']) # Standardize index names to PERMNO (company ID)
    .join(
        other=mve.rename_axis(index=['PERMNO', 'Date']), # Combine with market equity data,
        how='inner' # Use inner join to retain only records with matching PERMNO and Date
    )
    .reset_index() # Flatten the index into columns for easier manipulation
    .assign(
        Date=lambda x: x['Date'] + pd.offsets.MonthEnd(7), # Shift dates forward 7 months (from June to July)
        BM=lambda x: x['ceq'] / x['ME'], # Compute book-to-market ratio as ceq / ME
        Portfolio=lambda x: x.groupby('Date')['BM'].transform( # Assign stocks to portfolios
            lambda x: 1 + pd.qcut(x=x, q=10, labels=False) # Divide BM into 10 quantiles (0-10)
        )
    )
)
```

`bm.head()`

	PERMNO	Date	ceq	ME	BM	Portfolio
0	10000	1987-07-31	0.4180	1.9816	0.2109	2
1	10001	1987-07-31	5.4320	6.9370	0.7830	7
2	10001	1988-07-31	5.3690	5.8280	0.9212	7
3	10001	1989-07-31	5.5120	6.3623	0.8664	7
4	10001	1990-07-31	6.3210	10.3477	0.6109	5

This code creates a bar chart showing the average book-to-market (BM) ratio for each portfolio using the `portfolios` data frame. It groups the data by portfolio number (1 to 10), calculates the mean BM ratio for each group, and then plots these averages as bars. As expected, we see that BM rises from portfolio 1 to 10.

```
(  
    bm  
    .groupby(by=['Portfolio', 'Date']) # Group data by portfolio number and date  
    ['BM']  
    .mean() # Select the book-to-market (BM) column  
    .groupby(by='Portfolio') # Calculate the mean BM for each portfolio-date combi  
    .mean()  
    .plot(kind='bar') # Group again by portfolio to average across all date  
    # Calculate the time-averaged mean BM for each portfo  
    # Plot the results as a bar chart  
)  
plt.ylabel('Mean B/M for Equal-Weighted Portfolio')  
plt.title('Mean B/M for Equal-Weighted B/M Portfolios')  
plt.show()
```



This code combines monthly stock returns from the `crsp` data frame with the *most recent* portfolio assignments from the `bm` data frame. The `tolerance` argument tells the merge to only use portfolio assignments within 366 days to avoid stale data.

```
stocks = (
    pd.merge_asof(
        left=crsp.sort_values(['date', 'PERMNO']),
        right=bm.sort_values(['Date', 'PERMNO']),
        left_on='date',
        right_on='Date',
        by='PERMNO',
        tolerance=pd.Timedelta('366d')
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']] # Sort CRSP data by date and company ID
    .dropna() # Sort BM data by date and company ID
    # Match CRSP dates to the closest prior BM date
    # Use BM dates as the reference for the merge
    # Merge on company ID (PERMNO) to link portfolios
    # Limit matches to portfolio assignments within 366 days
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Select key columns: company ID, date, portfolio number, and return
    # Remove rows with missing values (e.g., for stocks not assigned to a portfolio)
    # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Convert 'date' column to 'Date' for consistency
)
```

This code creates equal-weighted portfolios using the `stocks` data frame. It groups the data by portfolio number (`Portfolio`) and date (`Date`), then calculates the mean return (`RET`) for each group. This averaging process gives each stock within a portfolio the same weight, producing equal-weighted portfolio returns.


Note

Here I made `portfolios` a wide data frame (i.e., portfolio numbers in columns) to simplify plots and calculations. This is a difference from class

```
portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading date
    [['RET']] # Isolate the returns column for portfolio-level aggregation
    .mean() # Average returns across stocks within each portfolio-date bin
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's returns
    ['RET'] # Extract the returns data, retaining portfolio numbers
)
```

Finally, we can see that the value portfolio 10 outperforms the growth portfolio 1!

```
portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted B/M Portfolios')
plt.show()
```



Backtest the strategy

We will need the Fama and French (1993) factors for the risk-free rate.

Note

Here I set the Date index to simplify merges/joins below. This is a difference from class

```
ff3 = (
    pdr.DataReader(
        name='F-F_Research_Data_Factors',
        data_source='famafrench',
        start='1900'
    )
[0]
.div(100)
.reset_index()
.assign(Date=lambda x: x['Date'].dt.to_timestamp(how='end').dt.normalize()) # convert mon
.set_index('Date')
)
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
```

```
    pdr.DataReader(
```

```
C:\Users\r.herron\AppData\Local\Temp\ipykernel_20756\1889697072.py:2: FutureWarning: The arg
```

```
    pdr.DataReader(
```

```
portfolios_rf = portfolios.sub(ff3['RF'], axis=0)
```

```
def CAGR(x, ann_fac=12):
    return (1 + x).prod() ** (ann_fac / x.count()) - 1
```

```
def Sharpe(x, ann_fac=np.sqrt(12)):
    return ann_fac * x.mean() / x.std()
```

```
def Volatility(x, ann_fac=np.sqrt(12)):
    return ann_fac * x.std()
```

```
def Drawdown(x):
    price = x.add(1).cumprod()
    return (price / price.cummax()).min() - 1
```

The high B/M have higher Sharper ratios and CAGRs than low B/M portfolios. However, the risk relation is not monotonic.

```
import seaborn as sns
```


```
stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)


plt.suptitle('Performance Statistics for Equal-Weighted B/M Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted B/M Portfolios



Also, this outperformance narrows over the past few decades.

```
portfolios.rolling(window=120).apply(CAGR).plot()
plt.ylabel('10-Year CAGR')
plt.title('Rolling 10-Year CAGRs for Equal-Weighted B/M Portfolios')
plt.show()
```



Re-implement the value strategy from the lecture notebook *with value-weighted portfolios*

We need to do two things:

1. Add the beginning of month market values of equity to use as portfolio weights
2. Replace `.mean()` with `np.average()` to calculate weighted means

This code prepares stock data from the `crsp` data frame by adding beginning-of-month market equity (`ME`) and merging it with portfolio assignments from the `bm` data frame. It first selects `PERMNO`, `date`, and `RET` from `crsp`, then merges it with `ME` data shifted forward by 1 month using an inner join to align current returns with next month's market equity. Next, it sorts the CRSP data by `date` and `PERMNO`, and merges it with a subset of `bm` containing only `Date`, `PERMNO`, and `Portfolio`, sorted by `Date` and `PERMNO`, matching each stock's date to the closest prior `Date` in `bm` within a 366-day tolerance. Finally, it keeps `PERMNO`, `date`, `Portfolio`, `RET`, and `ME`, drops missing values, converts `Portfolio` to integers, and renames `date` to `Date` for consistency.

```

# Add beginning-of-month market equity (ME) to CRSP data
crsp_w_bom_me = (
    crsp
    [['PERMNO', 'date', 'RET']] # Select stock ID, date, and returns from CRSP
    .merge(
        right=crsp[['PERMNO', 'date', 'ME']].assign( # Select stock ID, date, and ME, shift
            date=lambda x: x['date'] + pd.offsets.MonthEnd(1)), # Inner join to align current month with next
        how='inner'
    )
)

# Merge CRSP data with book-to-market portfolio assignments
stocks_vw = (
    pd.merge_asof(
        left=crsp_w_bom_me.sort_values(['date', 'PERMNO']), # Sort CRSP with ME by date and stock ID
        right=bm.sort_values(['Date', 'PERMNO'])[['Date', 'PERMNO', 'Portfolio']], # Select BM data
        left_on='date', # Match CRSP dates to the closest BM dates
        right_on='Date', # Use BM Date as the reference
        by='PERMNO', # Join on stock ID (PERMNO) to align dates
        tolerance=pd.Timedelta('366d') # Limit matches to BM data within a year
    )
    [['PERMNO', 'date', 'Portfolio', 'RET', 'ME']] # Select essential columns: stock ID, date, portfolio, return, and ME
    .dropna() # Remove rows with missing values
    .assign(Portfolio=lambda x: x['Portfolio'].astype(int)) # Convert portfolio numbers to integers
    .rename(columns={'date': 'Date'}) # Rename 'date' to 'Date' to align with CRSP
)

```

This code creates value-weighted portfolios using the `stocks` data frame and reshapes the results into a wide format. It groups the data by `Portfolio` (1 to 10) and `Date`, then calculates the value-weighted average return (`RET`) for each group, using market equity (`ME`) as weights. Next, it pivots the data so each portfolio's returns become separate columns, with `Date` as the index.

```


# Calculate value-weighted portfolio returns and reshape into wide format
portfolios_vw = (
    stocks_vw
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and date
    .apply(lambda x: np.average(a=x['RET'], # Compute value-weighted average returns using
                                weights=x['ME']), # Market Equity as weights
          include_groups=False) # Exclude grouping columns (Portfolio, Date)
    .unstack('Portfolio') # Pivot data so each portfolio's returns become columns
)

```

We can briefly compare the equal- and value-weighted portfolios. *We see that the value effect is concentrated in small stocks because the equal-weighted portfolios returns are much larger than the value-weighted portfolio returns.*

```
df = pd.concat(
    objs=[portfolios.add(1).prod(), portfolios_vw.add(1).prod()],
    keys=['Equal', 'Value (ME)'],
    names=['Portfolio Weights'],
    axis=1
)

df.plot(kind='bar')
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in B/M Portfolios')
plt.show()
```



Estimate the α s of the equal-weighted and value-weighted portfolios

i Note

Here `portfolios` is wide, so I have to `.stack()` it before the join. This is a difference from class

```
portfolios_ff3 = (
    portfolios
    .stack(future_stack=True)
    .to_frame('RET')
    .join(ff3, how='inner')
)
```

```
import statsmodels.formula.api as smf
```

The value portfolio 10 has an α of 82 basis points! Annualized, this is $12 \times 0.0082 = 0.0984$ or almost 10% of returns not explained by market risk.

```
smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query('(Portfolio == 10)').ff3)
```

Dep. Variable:	I(RET - RF)	R-squared:	0.506			
Model:	OLS	Adj. R-squared:	0.505			
Method:	Least Squares	F-statistic:	717.6			
Date:	Sat, 22 Mar 2025	Prob (F-statistic):	2.34e-109			
Time:	06:49:48	Log-Likelihood:	1116.4			
No. Observations:	702	AIC:	-2229.			
Df Residuals:	700	BIC:	-2220.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.0082	0.002	4.370	0.000	0.005	0.012
Q("Mkt-RF")	1.0948	0.041	26.788	0.000	1.015	1.175
Omnibus:	268.363	Durbin-Watson:	1.740			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1787.450			
Skew:	1.560	Prob(JB):	0.00			
Kurtosis:	10.167	Cond. No.	21.9			

Notes:


[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We can plot these as (and confidence intervals) with list comprehensions!

```
fits = [
    smf.ols(formula='I(RET-RF) ~ Q("Mkt-RF")', data=portfolios_ff3.query(f'(Portfolio == {p})')
        for p in range(1, 11)
    ]

df = (
    pd.DataFrame(
        data={
            'Alpha': [f.params['Intercept'] for f in fits],
            'Sterr': [f.bse['Intercept'] for f in fits],
            'Portfolio': range(1, 11)
        }
    )
    .set_index('Portfolio')
)

plt.bar(
    x=df.index,
    height=df['Alpha'],
    yerr=df['Sterr'].mul(1.96)
)
plt.ylabel('Alpha')
plt.title('Monthly Alphas for Equal-Weighted B/M Portfolios\n (black bars indicate 95% confi')
plt.show()
```



Implement a momentum strategy

Form deciles on the 11-month returns from months $t - 12$ to month $t - 2$.

This code calculates 11-month momentum returns (`RET_11`) from 2 to 12 months prior using the `crsp` data frame and assigns portfolios based on those returns. It sets a multi-index with `PERMNO` and `date`, selects the `RET` column, and pivots to a wide format with stocks as columns. After sorting by date, it computes compounded returns over an 11-month rolling window using log transformations (`log1p` and `expm1`), then reshapes back to long format. It drops rows with missing `RET_11` values, creates a `DataFrame`, and resets the index. Finally, it assigns stocks to decile portfolios (1 to 10) based on `RET_11` for each date and shifts the `date` forward by 2 months to reflect the momentum strategy lag. That is, we receive returns in month t based on returns from month $t - 12$ through $t - 2$.

```
# Calculate 2-12 month momentum returns and assign portfolios
mom_02_12 = (
    crsp
```

```

.set_index(['PERMNO', 'date'])          # Set multi-index with stock ID (PERMNO) and date
['RET']                                # Select the monthly returns column
.unstack('PERMNO')                      # Pivot to wide format, with PERMNO as column
.sort_index()                           # Sort dates chronologically for consistent results
.pipe(np.log1p)                         # Convert returns to log scale (log(1 + r))
.rolling(11)                            # Apply an 11-month rolling window to prior returns
.sum()                                  # Sum log returns over the 11-month window
.pipe(np.expm1)                         # Convert summed log returns back to compound returns
.stack(future_stack=True)               # Reshape to long format, aligning with pandas
.dropna()                               # Remove rows with NaN from incomplete windows
.to_frame('RET_11')                     # Create a DataFrame, naming the momentum return
.reset_index()                          # Move PERMNO and date back to columns from index
.assign(
    Portfolio=lambda x: x.groupby('date')['RET_11'].transform( # Assign stocks to decile portfolios
        lambda x: 1 + pd.qcut(x=x, q=10, labels=False)),       # Split RET_11 into 10 quartiles
    date=lambda x: x['date'] + pd.offsets.MonthEnd(2))        # Shift dates forward 2 months
)
)
)

```

```
mom_02_12.head()
```

	date	PERMNO	RET_11	Portfolio
0	1966-01-31	10006	0.1508	5
1	1966-01-31	10014	0.2000	6
2	1966-01-31	10030	0.1489	5
3	1966-01-31	10057	0.4666	8
4	1966-01-31	10102	0.3503	7

```

stocks = (
    pd.merge(
        left=crsp,
        right=mom_02_12,
        on=['PERMNO', 'date'],
        how='inner'
    )
    [['PERMNO', 'date', 'Portfolio', 'RET']]
    .dropna()
    .rename(columns={'date': 'Date'})
)


```

```

portfolios = (
    stocks
    .groupby(['Portfolio', 'Date']) # Group stocks by portfolio number (1-10) and trading da
    [['RET']]
    .mean() # Isolate the returns column for portfolio-level aggregat
    # Average returns across stocks within each portfolio-d
    .unstack('Portfolio') # Pivot the data to create columns for each portfolio's
    ['RET'] # Extract the returns data, retaining portfolio numbers
)
portfolios.add(1).cumprod().plot()
plt.semilogy()
plt.ylabel('Value of $1 Investment')
plt.title('Values of $1 Investments in Equal-Weighted Momentum Portfolios')
plt.show()

```

Values of \$1 Investments in Equal-Weighted Momentum Portfolios



```
portfolios_rf = portfolios.sub(ff3['RF'], axis=0)


stats = (
    portfolios_rf
    .agg(Sharpe)
    .to_frame('Sharpe')
    .join(
        portfolios
        .agg([CAGR, Drawdown, Volatility])
        .transpose()
    )
    .rename_axis(columns='Statistic')
)

df = stats.stack().to_frame('Value').reset_index()

sns.catplot(
    data=df,
    x='Portfolio',
    y='Value',
    col='Statistic',
    col_wrap=2,
    height=3,
    kind='bar',
    sharey=False
)

plt.suptitle('Performance Statistics for Equal-Weighted Momentum Portfolios', y=1.05)
plt.show()
```

Performance Statistics for Equal-Weighted Momentum Portfolios



Week 11

Project 2

Purpose

I have two goals for this project:

1. Implement and backtest a 52-week high/low breakout strategy in Python
2. Compare and contrast performance, *explaining any differences*

Assignment

Evaluate a 52-week high/low breakout strategy against a passive buy-and-hold approach. Compare and contrast the strategy on the Dow-Jones Industrial Average (DJIA) ETF, an equal-weighted portfolio of the strategy on the DJIA constituents,¹ and buying-and-holding the DJIA ETF (ticker DIA). *Explain any differences among the three, and recommend one.*

Here is the strategy:

- *Long Position:*
 - *Entry:* Price at the 52-week high and volume more than 1.5 times the 20-day simple moving average
 - *Exit:* Price falls 5% below the entry price or 20 days, whichever comes first
- *Short:*
 - *Entry:* Price at the 52-week low and volume more than 1.5 times the 20-day simple moving average
 - *Exit:* Price rises 5% above the entry price or 20 days, whichever comes first
- *Otherwise:*
 - Hold cash and earn the risk-free rate of return

Here are ideas to consider:

¹Here equal-weighted portfolio means that you find the position for all the DJIA constituents. If you are long stocks 1, 2, and 3, and you are short stock 4. then your portfolio return would be: $r_p = \frac{r_1+r_2+r_3-r_4}{4}$.

Project 2

1. Use the provided dataset for the current DJIA constituents, a DJIA ETF (ticker DIA), and the Fama-French factors
2. At a minimum, compare total returns and Sharpe Ratios
3. You might also explore maximum drawdowns, subsamples, rolling windows, trade frequencies, the volume confirmation, *plus any other compelling insights*
4. Title, label, and caption your figures and tables, referencing them in your summary

Criteria

Table 1 provides the project grading rubric. The project is worth 200 points. The peer reviews are worth 100 points, and students will receive their median score. Almost all students earn perfect peer review scores, so I will factor that into project scores. For example, a project score *without peer review scores* of 77.5% converts to a project score *with perfect peer review scores* of 85% because $\frac{0.775 \times 200 + 1.00 \times 100}{300} = 0.85 = 85\%$.

Table 1: This table provides the project grading rubric

Topic	Points
Clarity, correctness, and completeness of calculations	60
Clarity, correctness, and completeness of visualizations	60
Clarity, correctness, and completeness of discussions	60
Correctness of submission according to the deliverables section	20
Total	200

Deliverables

Upload the following as unzipped files to Canvas by 11:59 PM on 3/28:

1. One Jupyter notebook that contains your report and performs *all* your analysis
 1. Name this file `project_2.ipynb` for me to run your code
 2. Your notebook must run on my computer; I will place the data files in the same folder as your notebook
 3. You may not edit the Word documents after you create them
2. One Quarto-generated Word document *without code* for our partner to review
 1. Name this file `project_2_without_code.docx`
 2. Typing `echo: false` in the first cell of this Jupyter notebook hides code in your Word document

Project 2

3. One Quarto-generated Word document *with code* for me to grade
 1. Name this file `project_2_with_code.docx`
 2. Typing `echo: true` in the first cell of this Jupyter notebook displays code in your Word document

Here is some additional guidance:

1. Provide an up to three-page executive summary at the start of your report, *which is the only writing I will read*
2. Your Word document *without code* must not exceed 15 pages in length
3. Your submission must not include your name

Data

This project requires two data files. Save these data files in the same folder as your `project_2.ipynb` notebook file.

1. `data_djia.csv` provides long-formatted data from Yahoo! Finance for the current DJIA companies, plus a DJIA ETF (ticker DIA), from 1998-01-20 through 2024-12-31
2. `data_ff3` provides factor data from Kenneth French's data library (name `F-F_Research_Data_Factors_da` for the same period

You can read these data files as follows.

Quarto

Basics

1. Use [Quarto](#) to generate your Word document from your notebook
2. Use `#` to create a title and `##` to create sections
3. Use `-` or `1.` to create lists
4. Use the first cell in this notebook to hide or display code with `echo=false` or `echo=true`, respectively
5. This first cell must be a `raw` cell instead of a `code` or `markdown` cell
6. Use `quarto render project_1.ipynb` in the same folder as your notebook to render it to a Word document
7. Use the `cd` command in the terminal to change the working directory to the directory with your notebook

Examples

This section provides a sample analysis highlighting how code and formatting work with Quarto. Figure 1 provides a line plot of the value of a \$10,000 investment in DIA, the DJIA ETF. Note that `#| label:` and `#| fig-cap:` comments at the top of the figure cell create the figure reference/link and the figure caption, respectively. You can learn more about cross-referencing figures and tables [here](#).




Figure 1: This line plot shows the value of a \$10,000 investment in the DJIA ETF at the close of its first day of trading

Artificial Intelligence (AI)

You may use AI (e.g., ChatGPT) to *help* you prepare your analysis and discussion. However:

1. AI will not do very well on this project without significant input from your team

Project 2

2. AI will not be a defense against plagiarism because AI should not *write* your code and slides; If you plagiarize an AI that plagiarizes other sources, you are responsible for plagiarizing the AI and its sources

Week 12

Herron Topic 4 - Portfolio Optimization

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import scipy.optimize as sco # new addition for portfolio optimization
import statsmodels.api as sm
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Introduction

This notebook covers portfolio optimization. In this notebook, we will:

1. Review the $1/N$ portfolio (or equal-weighted portfolio) from the Herron topic 1 notebook
2. Use SciPy's `minimize()` function to:
 1. Find the minimum-variance portfolio
 2. Find the minimum-variance frontier

Ivo Welch provides a clear discussion of the theory and practice of portfolio optimization in Chapter 12 of [his free investments textbook](#).

The $1/N$ Portfolio

We first saw the $1/N$ portfolio (or equal-weighted portfolio) in the Herron topic 1 notebook. The $1/N$ portfolio gives each of N assets an equal portfolio weight of $1/N$. While the $1/N$ strategy seems too simple to be useful, DeMiguel, Garlappi, and Uppal (2009) show the $1/N$ portfolio typically has a higher mean-variance efficiency than advanced portfolio optimizations.

We will investigate the 1/N portfolio and portfolio optimization with the most recent three years of daily returns for the Mag 7 stocks (i.e., GOOGL, AMZN, AAPL, META, MSFT, NVDA, and TSLA).

```
mag7 = (
    yf.download(
        tickers='GOOGL AMZN AAPL META MSFT NVDA TSLA',
        auto_adjust=False,
        progress=False
    )
    .iloc[:-1] # drop incomplete trading day
)

returns = (
    mag7
    ['Adj Close']
    .pct_change()
    .iloc[(-3 * 252):]
)
```

We can manually calculate 1/N portfolio returns three ways.

The first calculation is a literal interpretation of $r_P = \frac{1}{N} \sum_i^N r_i$. The pandas code for this first calculation is `returns.sum(axis=1).div(n)`, which sums across the rows, and then divides by `n`. Recall that with constant weights, we rebalance our portfolio every return period. If we have daily data, we rebalance daily. If we have monthly data, we rebalance monthly, and so on.

```
n = returns.shape[1]
r_p1 = returns.sum(axis=1).div(n)

r_p1.describe()
```

count	756.0000
mean	0.0010
std	0.0196
min	-0.0674
25%	-0.0099
50%	0.0008
75%	0.0128
max	0.0984
dtype:	float64

We learned the last two calculations in the Herron topic 1. The pandas code for the second calculation is `returns.mean(axis=1)`, which calculates the mean of each row.

```
r_p2 = returns.mean(axis=1)
```

```
r_p2.describe()
```

```
count    756.0000
mean      0.0010
std       0.0196
min     -0.0674
25%    -0.0099
50%     0.0008
75%     0.0128
max      0.0984
dtype: float64
```

The pandas code for the third calculations is `returns.dot(weights)`, where `weights` is a NumPy array or a pandas series of portfolio weights. This calculation allows different weights for each asset. We can make $1/N$ portfolio weights with `np.ones(n) / n`.

```
weights = np.ones(n) / n
```

```
weights
```

```
array([0.1429, 0.1429, 0.1429, 0.1429, 0.1429, 0.1429, 0.1429])
```

```
r_p3 = returns.dot(weights)
```

```
r_p3.describe()
```

```
count    756.0000
mean      0.0010
std       0.0196
min     -0.0674
25%    -0.0099
50%     0.0008
75%     0.0128
max      0.0984
dtype: float64
```

We can use `np.allclose()` to show that `r_p1`, `r_p2`, and `r_p3` are similar.

```
np.allclose(r_p1, r_p2)
```

True

```
np.allclose(r_p2, r_p3)
```

True

```
np.allclose(r_p3, r_p1)
```

True

Here is a brief detour to show how the `.dot()` method works.

```
silly_n = 3
silly_w = np.ones(silly_n) / silly_n
silly_r = pd.DataFrame(np.arange(2*silly_n).reshape(2, silly_n))
```

```
print(
    f'silly_n:{silly_n}',
    f'silly_w:{silly_w}',
    f'silly_r:{silly_r}',
    sep='\n\n'
)
```

```
silly_n:
3

silly_w:
[0.3333 0.3333 0.3333]

silly_r:
   0   1   2
0   0   1   2
1   3   4   5
```

```
silly_r.dot(silly_w)
```

```
0    1.0000
1    4.0000
dtype: float64
```

Under the hood, the `.dot()` method efficiently calculates the following:

```
for i, row in silly_r.iterrows():
    print(
        f'Row {i}: ',
        ' + '.join([f'{w}:0.2f} * {y}' for w, y in zip(silly_w, row)]),
        '= ',
        f'{silly_r.dot(silly_w).iloc[i]:0.2f}'
    )
```

```
Row 0:  0.33 * 0 + 0.33 * 1 + 0.33 * 2 =  1.00
Row 1:  0.33 * 3 + 0.33 * 4 + 0.33 * 5 =  4.00
```

SciPy's `minimize()` Function

SciPy's `minimize()` function, from the `optimize` module, finds the input array `x` that minimizes the output of a given function `fun`. The `minimize()` function uses optimization techniques that are beyond the scope of this course. We can consider these optimization techniques as sophisticated trial-and-error (or guess-and-check).

Here are the common arguments for the `minimize()` function:

1. `fun=` is the name of the function who's output we want to minimize
2. `x0=` is our first guess for the input to the function `fun`
3. `args=` is a tuple of additional arguments to the function `fun`
4. `bounds=` is a tuple of tuples with lower and upper bounds for each element in `x`
5. `constraints=` is a tuple of dictionaries that limit `x` by constraining the output of a given function to zero or non-negative values

Here is a simple example that minimizes the function `quadratic()`, which has arguments `x` and `a` and returns $(x - a)^2$.

```
def quadratic(x, a=5):
    return (x - a) ** 2
```

```
quadratic(x=0, a=5)
```

25

```
quadratic(x=5, a=5)
```


0

```
quadratic(x=10, a=5)
```

25

It might be helpful to plot $y = (x - a)^2$.

```
x = np.linspace(-5, 15, 101)
y = quadratic(x=x)
plt.plot(x, y)
plt.xlabel(r'$x$')
plt.ylabel(r'$y$')
plt.title(r'$y = (x - 5)^2$ over $x \in [-5, 15]$')
plt.show()
```



The minimum output of `quadratic()` occurs at $x = 5$ if we do not use bounds or constraints. For this function, we can even start far away from $x = 5$.

```
sco.minimize(
    fun=quadratic,
    x0=np.array([2001])
)
```

```
message: Optimization terminated successfully.
success: True
status: 0
fun: 2.0392713450495178e-16
x: [ 5.000e+00]
nit: 4
jac: [-1.366e-08]
hess_inv: [[ 5.000e-01]]
nfev: 18
```

```
njev: 9
```

The minimum output of `quadratic()` occurs at $x = 6$ if we bound $6 \leq x \leq 10$ with `bounds=((6, 10),)`. `bounds=` requires a tuple of tuples, so we need the trailing comma if we only have one set of bounds.

```
sco.minimize(
    fun=quadratic,
    x0=np.array([2001]),
    bounds=((6, 10),)
)
```

```
message: CONVERGENCE: NORM OF PROJECTED GRADIENT <= PGTOL
success: True
status: 0
fun: 1.0
x: [ 6.000e+00]
nit: 1
jac: [ 2.000e+00]
nfev: 4
njev: 2
hess_inv: <1x1 LbfgsInvHessProduct with dtype=float64>
```

The minimum output of `quadratic()` also occurs at $x = 6$ if we constrain $x - 6 > 0$ with `constraints=({'type': 'ineq', 'fun': lambda x: x - 6})`. For inequality constraints, the given function must evaluate to a non-negative value. We use bounds to limit the search space directly, and we use constraints to limit the search space based on a given function.

```
sco.minimize(
    fun=quadratic,
    x0=np.array([2001]),
    constraints=({'type': 'ineq', 'fun': lambda x: x - 6})
)
```

```
message: Optimization terminated successfully
success: True
status: 0
fun: 1.0000000000000018
x: [ 6.000e+00]
nit: 3
jac: [ 2.000e+00]
```

```
nfev: 6
njev: 3
```

We use the `args=` argument to pass additional arguments to the function `fun`. For example, we change the `a=` argument in `quadratic()` from the default of `a=5` to `a=20` with `args=(20,)`. Like `bounds=`, `args=` expects a tuple, so we need a trailing comma if we have only one argument.

```
sco.minimize(
    fun=quadratic,
    args=(20,),
    x0=np.array([2001]),
)

message: Optimization terminated successfully.
success: True
status: 0
fun: 7.090392030754976e-17
x: [ 2.000e+01]
nit: 4
jac: [-1.940e-09]
hess_inv: [[ 5.000e-01]]
nfev: 18
njev: 9
```

The Minimum-Variance Portfolio

Now we apply `minimize()` to our Mag-7 returns to find the the minimum-variance portfolio. Recall, the `minimize()` function varies an input array `x`, starting from argument `x0=`, to minimize the output of a given function `fun=` subject to the bounds and constraints in `bounds=` and `constraints=`, respectively.

We can define a function `calc_sigmap()` to calculate portfolio volatility. The first argument to `calc_sigmap()` must be the input array `x` that `minimize()` searches over. We call this first argument `w` because it has portfolio weights. Recall that portfolio variance is $\sigma_p^2 = w' \Sigma w$, so portfolio volatility is $\sigma_p = \sqrt{w' \Sigma w}$. Here Σ (and `Sigma`) is the returns covariance matrix, and `w` (and `w`) is the portfolio weights array.

```
def calc_sigmap(w, Sigma, ppy=252):
    return np.sqrt(ppy * w.T @ Sigma @ w)
```

```
Sigma = returns.cov()
```

```
Sigma
```

Ticker	AAPL	AMZN	GOOGL	META	MSFT	NVDA	TSLA
Ticker							
AAPL	0.0003	0.0002	0.0002	0.0002	0.0002	0.0003	0.0003
AMZN	0.0002	0.0005	0.0003	0.0004	0.0003	0.0004	0.0004
GOOGL	0.0002	0.0003	0.0004	0.0003	0.0002	0.0004	0.0003
META	0.0002	0.0004	0.0003	0.0008	0.0003	0.0005	0.0004
MSFT	0.0002	0.0003	0.0002	0.0003	0.0003	0.0004	0.0003
NVDA	0.0003	0.0004	0.0004	0.0005	0.0004	0.0012	0.0006
TSLA	0.0003	0.0004	0.0003	0.0004	0.0003	0.0006	0.0015

We can also define a function `calc_mup()` that calculates a mean portfolio return. Recall that mean portfolio return is $\mu_p = w' \mu$. Here μ (and `mu`) is the mean returns array, and w (and `w`) is the portfolio weights array.

```
def calc_mup(w, mu, ppy=252):
    return ppy * w.T @ mu
```

```
mu = returns.mean()
```

```
mu
```

```
Ticker
AAPL    0.0006
AMZN    0.0006
GOOGL   0.0005
META    0.0018
MSFT    0.0005
NVDA    0.0025
TSLA    0.0006
dtype: float64
```

Finally, we can use `minimize()` to find the portfolio weights that minimize portfolio volatility. Below, the equality constraint is met when `x.sum() - 1` is zero, which is the Python equivalent of $\sum_i w_i = 1$.

```

res_mv = sco.minimize(
    fun=calc_sigmap,
    x0=np.ones(Sigma.shape[0]) / Sigma.shape[0],
    args=(Sigma, 252),
    bounds=[(0,1) for _ in returns],
    constraints=(
        {'type': 'eq', 'fun': lambda x: x.sum() - 1}
    )
)

print(res_mv)

```

```

message: Optimization terminated successfully
success: True
status: 0
fun: 0.24371407736360207
x: [ 4.616e-01  0.000e+00  9.470e-02  0.000e+00  4.437e-01
      0.000e+00  2.212e-17]
nit: 10
jac: [ 2.437e-01  2.575e-01  2.437e-01  2.773e-01  2.437e-01
      3.419e-01  3.062e-01]
nfev: 80
njev: 10

```

What are the attributes of the minimum-variance portfolio? We can write a helper function `print_port_res()` to simplify our lives!

```

def print_port_res(w, mu, Sigma, title, ppy=252):
    width = max(len(title), 32)
    mup = ppy * w.T @ mu
    sigmap = np.sqrt(ppy * w.T @ Sigma @ w)

    return print(
        title,
        '=' * width,
        '',
        'Performance',
        '-' * width,
        'Annualized Mean Return:'.ljust(width - 6) + f'{mup:.4f}',
        'Annualized Volatility:'.ljust(width - 6) + f'{sigmap:.4f}',
        ''
    )

```

```

'Weights',
'-' * width,
'\n'.join([
    f'{_r}'.ljust(width - 6) + f'{_w:0.4f}'
    for _r, _w in zip(Sigma.index, w)
]),
'=' * width,
sep='\n',
)

```

```

print_port_res(
    w=res_mv['x'],
    mu=returns.mean(),
    Sigma=returns.cov(),
    title='minimum-variance portfolio'
)

```

```

minimum-variance portfolio
=====
Performance
-----
Annualized Mean Return: 0.1344
Annualized Volatility: 0.2437

Weights
-----
AAPL:          0.4616
AMZN:          0.0000
GOOGL:         0.0947
META:          0.0000
MSFT:          0.4437
NVDA:          0.0000
TSLA:          0.0000
=====
```

The Minimum-Variance Frontier

We can use `minimize()` to find the minimum-variance portfolio for a *target return*, then use a `for` loop to repeat this calculation for a wide range of target returns. This calculation will map out the minimum-variance frontier. Here are the steps:

1. Create a NumPy array `returns_target` of target returns
2. Create an empty list `res_ef` to hold `minimize()` results
3. Loop over `returns_target`, passing each target return as a constraint to `minimize()`
4. Append each `minimize()` result to `res_ef`

```
returns_target = 252 * np.linspace(
    start=returns.mean().min(),
    stop=returns.mean().max(),
    num=50
)
```

Below, the first equality constraint is met when $x.sum() - 1$ is zero, which is the Python equivalent of $\sum_i w_i = 1$. The second equality constraint is met when `calc_mup(w=x, mu=mu, ppy=252) - r` is zero, which is the Python equivalent of $w'\mu = \text{target return}$.

```
res_ef = []

for r in returns_target:
    _ = sco.minimize(
        fun=calc_sigmap,
        x0=np.ones(Sigma.shape[0]) / Sigma.shape[0],
        args=(Sigma, 252),
        bounds=[(0, 1) for c in Sigma.index],
        constraints=(
            {'type': 'eq', 'fun': lambda x: x.sum() - 1},
            {'type': 'eq', 'fun': lambda x: calc_mup(w=x, mu=mu, ppy=252) - r}
        )
    )
    res_ef.append(_)
```

The list `res_ef` contains the results of all 50 minimum-variance portfolios. For example, `res_ef[0]` is the minimum-variance portfolio for the lowest target return.

```
res_ef[0]

message: Optimization terminated successfully
success: True
status: 0
fun: 0.3240234964602987
x: [ 0.000e+00  1.860e-15  1.000e+00  1.638e-15  7.885e-08
      0.000e+00  0.000e+00]
```

```

nit: 2
jac: [ 1.599e-01  2.423e-01  3.240e-01  2.693e-01  1.777e-01
      2.884e-01  2.554e-01]
nfev: 16
njev: 2

```

We should check that all minimizations succeeded. If a portfolio volatility minimization fails, we should check our function, bounds, and constraints. Portfolio optimization is typically a simple problem for `minimize()`, so we typically have an error if `minimize()` does not succeed.

```

for r in res_ef:
    assert r['success']

```

We can combine the target returns and volatilities into a data frame `mv_frontier`.

```

mv_frontier = pd.DataFrame({
    'returns_target': returns_target,
    'volatility': np.array([r['fun'] if r['success'] else np.nan for r in res_ef])
})

```

Finally, we can plot the minimum-variance frontier.


```

(
    mv_frontier
    .mul(100)
    .plot(x='volatility', y='returns_target', legend=False)
)
plt.ylabel('Annualized Mean Return (%)')
plt.xlabel('Annualized Volatility (%)')

for t in Sigma.index:
    x = 100 * np.sqrt(252 * Sigma.loc[t, t])
    y = 100 * 252 * mu.loc[t]
    plt.annotate(text=t, xy=(x, y))

plt.show()

```



Herron Topic 4 - Practice

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import scipy.optimize as sco # new addition for portfolio optimization
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Recap

Practice

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years

Note that `sco.minimize()` finds *minimums*, so you need to minimize the *negative* Sharpe Ratio.

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow short weights up to 10% on each stock

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow total short weights of up to 30%

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but do not allow any weight to exceed 30% in magnitude

Find the minimum 95% Value at Risk (Var) portfolio of Mag 7 stocks over the last three years

More on VaR [here](#).

Find the minimum draw down portfolio of Mag 7 stocks over the last three years

Find the minimum draw down portfolio for the sample with complete data for the current Dow-Jones Industrial Average (DJIA) stocks

You can find the DJIA tickers on [Wikipedia](#).

Plot the minimum-variance frontier for the sample with complete data for the current the DJIA stocks

Find the maximum Sharpe Ratio portfolio for the sample with complete data for the current the DJIA stocks

Herron Topic 4 - Practice - Sec 02

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import scipy.optimize as sco # new addition for portfolio optimization
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Recap

Practice

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years

Note that `sco.minimize()` finds *minimums*, so you need to minimize the *negative* Sharpe Ratio.

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow short weights up to 10% on each stock

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow total short weights of up to 30%

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but do not allow any weight to exceed 30% in magnitude

Find the minimum 95% Value at Risk (Var) portfolio of Mag 7 stocks over the last three years

More on VaR [here](#).

Find the minimum draw down portfolio of Mag 7 stocks over the last three years

Find the minimum draw down portfolio for the sample with complete data for the current Dow-Jones Industrial Average (DJIA) stocks

You can find the DJIA tickers on [Wikipedia](#).

Plot the minimum-variance frontier for the sample with complete data for the current the DJIA stocks

Find the maximum Sharpe Ratio portfolio for the sample with complete data for the current the DJIA stocks

Herron Topic 4 - Practice - Sec 03

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import scipy.optimize as sco # new addition for portfolio optimization
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Recap

Practice

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years

Note that `sco.minimize()` finds *minimums*, so you need to minimize the *negative* Sharpe Ratio.

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow short weights up to 10% on each stock

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow total short weights of up to 30%

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but do not allow any weight to exceed 30% in magnitude

Find the minimum 95% Value at Risk (Var) portfolio of Mag 7 stocks over the last three years

More on VaR [here](#).

Find the minimum draw down portfolio of Mag 7 stocks over the last three years

Find the minimum draw down portfolio for the sample with complete data for the current Dow-Jones Industrial Average (DJIA) stocks

You can find the DJIA tickers on [Wikipedia](#).

Plot the minimum-variance frontier for the sample with complete data for the current the DJIA stocks

Find the maximum Sharpe Ratio portfolio for the sample with complete data for the current the DJIA stocks

Herron Topic 4 - Practice - Sec 04

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader as pdr
import scipy.optimize as sco # new addition for portfolio optimization
import yfinance as yf
```

```
%precision 4
pd.options.display.float_format = '{:.4f}'.format
# %config InlineBackend.figure_format = 'retina'
```

Announcements

Five-Minute Recap

Practice

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years

Note that `sco.minimize()` finds *minimums*, so you need to minimize the *negative* Sharpe Ratio.

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow short weights up to 10% on each stock

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but allow total short weights of up to 30%

Find the maximum Sharpe Ratio portfolio of Mag 7 stocks over the last three years, but do not allow any weight to exceed 30% in magnitude

Find the minimum 95% Value at Risk (Var) portfolio of Mag 7 stocks over the last three years

More on VaR [here](#).

Find the minimum draw down portfolio of Mag 7 stocks over the last three years

Find the minimum draw down portfolio for the sample with complete data for the current Dow-Jones Industrial Average (DJIA) stocks

You can find the DJIA tickers on [Wikipedia](#).

Plot the minimum-variance frontier for the sample with complete data for the current the DJIA stocks

Find the maximum Sharpe Ratio portfolio for the sample with complete data for the current the DJIA stocks

Week 13

Student's Choice 4

We will vote to determine the “Student’s Choice” topics during the first few weeks of class.

Week 14

MSFQ Assessment Exam

We will take the assessment exam in class on Tuesday of this week.

Week 15

Project 3

I will determine assignment details at least one week before each due date.

References

- DeMiguel, Victor, Lorenzo Garlappi, and Raman Uppal (2009). “Optimal versus naive diversification: How inefficient is the $1/N$ portfolio strategy?” In: *The Review of Financial Studies* 22.5, pp. 1915–1953.
- Fama, Eugene F. and Kenneth R. French (1993). “Common risk factors in the returns on stocks and bonds”. In: *Journal of Financial Economics* 33.1, pp. 3–56.
- Gray, Wesley R. and Tobias E. Carlisle (2012). *Quantitative value: a practitioner’s guide to automating intelligent investment and eliminating behavioral errors*. John Wiley & Sons.
- Kritzman, Mark P. (2000). *Puzzles of Finance: Six Practical Problems and Their Remarkable Solutions*. John Wiley & Sons.
- Lewinson, Eryk (2020). *Python for Finance Cookbook: Over 50 recipes for applying modern Python libraries to financial data analysis*. Packt Publishing Ltd.
- McKinney, Wes (2022). *Python for Data Analysis*. 3rd ed. O’Reilly Media, Inc.
- Murphy, John J. (1999). *Technical analysis of the financial markets: A comprehensive guide to trading methods and applications*. Penguin.
- Welch, Ivo (2022). *Corporate Finance*. 5th ed. Ivo Welch.